

MY451 Introduction to Quantitative Analysis

Jouni Kuha

Department of Methodology

London School of Economics and Political Science

2016-17

Contents

Course information	i
1 Introduction	1
1.1 What is the purpose of this course?	1
1.2 Some basic definitions	3
1.2.1 Subjects and variables	3
1.2.2 Types of variables	6
1.2.3 Description and inference	10
1.2.4 Association and causation	11
1.3 Outline of the course	12
1.4 The use of mathematics and computing	13
1.4.1 Symbolic mathematics and mathematical notation	13
1.4.2 Computing	13
2 Descriptive statistics	15
2.1 Introduction	15
2.2 Example data sets	15
2.3 Single categorical variable	16
2.3.1 Describing the sample distribution	16
2.3.2 Tabular methods: Tables of frequencies	17
2.3.3 Graphical methods: Bar charts	18
2.3.4 Simple descriptive statistics	20
2.4 Two categorical variables	20
2.4.1 Two-way contingency tables	20
2.4.2 Conditional proportions	22
2.4.3 Conditional distributions and associations	23
2.4.4 Describing an association using conditional proportions	24
2.4.5 A measure of association for ordinal variables	25
2.5 Sample distributions of a single continuous variable	29
2.5.1 Tabular methods	29
2.5.2 Graphical methods	30
2.6 Numerical descriptive statistics	36
2.6.1 Measures of central tendency	36
2.6.2 Measures of variation	40
2.7 Associations which involve continuous variables	45
2.8 Presentation of tables and graphs	45
2.9 Appendix: Country data	47
3 Samples and populations	50
3.1 Introduction	50
3.2 Finite populations	50

3.3	Samples from finite populations	51
3.4	Conceptual and infinite populations	53
3.5	Population distributions	55
3.6	Need for statistical inference	56
4	Statistical inference for two-way tables	58
4.1	Introduction	58
4.2	Significance tests	59
4.3	The chi-square test of independence	60
4.3.1	Hypotheses	60
4.3.2	Assumptions of a significance test	62
4.3.3	The test statistic	63
4.3.4	The sampling distribution of the test statistic	64
4.3.5	The P -value	70
4.3.6	Drawing conclusions from a test	71
4.4	Summary of the chi-square test of independence	73
5	Inference for population proportions	75
5.1	Introduction	75
5.2	Examples	76
5.3	Probability distribution of a dichotomous variable	78
5.4	Point estimation of a population probability	79
5.5	Significance test of a single proportion	80
5.5.1	Null and alternative hypotheses	80
5.5.2	The test statistic	81
5.5.3	The sampling distribution of the test statistic and P -values	83
5.5.4	Conclusions from the test	87
5.5.5	Summary of the test	87
5.6	Confidence interval for a single proportion	88
5.6.1	Introduction	88
5.6.2	Calculation of the interval	88
5.6.3	Interpretation of confidence intervals	89
5.6.4	Confidence intervals vs. significance tests	91
5.7	Inference for comparing two proportions	92
6	Continuous variables: Population and sampling distributions	97
6.1	Introduction	97
6.2	Population distributions of continuous variables	98
6.2.1	Population parameters and their point estimates	98
6.3	Probability distributions of continuous variables	100
6.3.1	General comments	100
6.3.2	The normal distribution as a population distribution	101
6.4	The normal distribution as a sampling distribution	106
7	Analysis of population means	112
7.1	Introduction and examples	112
7.2	Descriptive statistics for comparisons of groups	116
7.2.1	Graphical methods of comparing sample distributions	116
7.2.2	Comparing summary statistics	119
7.3	Inference for two means from independent samples	119
7.3.1	Aims of the analysis	119
7.3.2	Significance testing: The two-sample t -test	121

7.3.3	Confidence intervals for a difference of two means	126
7.3.4	Variants of the test and confidence interval	126
7.4	Tests and confidence intervals for a single mean	132
7.5	Inference for dependent samples	138
7.6	Further comments on significance tests	141
7.6.1	Different types of error	141
7.6.2	Power of significance tests	142
7.6.3	Significance vs. importance	143
8	Linear regression models	145
8.1	Introduction	145
8.2	Describing association between two continuous variables	146
8.2.1	Introduction	146
8.2.2	Graphical methods	146
8.2.3	Linear associations	150
8.2.4	Measures of association: covariance and correlation	154
8.3	Simple linear regression models	157
8.3.1	Introduction	157
8.3.2	Definition of the model	158
8.3.3	Interpretation of the model parameters	160
8.3.4	Estimation of the parameters	163
8.3.5	Statistical inference for the regression coefficients	168
8.4	Interlude: Association and causality	171
8.5	Multiple linear regression models	173
8.5.1	Introduction	173
8.5.2	Definition of the model	174
8.5.3	Unchanged elements from simple linear models	175
8.5.4	Interpretation and inference for the regression coefficients	176
8.6	Including categorical explanatory variables	179
8.6.1	Dummy variables	179
8.6.2	A second example	183
8.7	Other issues in linear regression modelling	186
9	Analysis of 3-way contingency tables	189
10	More statistics...	199
Appendix		201
Computer classes		201
General instructions		201
Introduction to SPSS		201
WEEK 2 class: Descriptive statistics for categorical data, and entering data . . .		207
WEEK 3 class		212
WEEK 4 class: Two-way contingency tables		214
WEEK 5 class: Inference for two population means		216
WEEK 7 class: Inference for population proportions		218
WEEK 7 class: Correlation and simple linear regression 1		221
WEEK 8 class: Simple linear regression and 3-way tables		223
WEEK 9 class: Multiple linear regression		225
WEEK 10 class: Review and Multiple linear regression		227
Statistical tables		228
Table of standard normal tail probabilities		229

Table of critical values for t -distributions	230
Table of critical values for χ^2 distributions	231

Course information



Department of
Methodology



MY451: Introduction to Quantitative Analysis

Course Description

This course is intended for those with little or no past training in quantitative methods. The course is an intensive introduction to some of the principles and methods of statistical analysis in social research. Topics covered in MY451 include descriptive statistics, basic ideas of inference and estimation, contingency tables and an introduction to linear regression models. For those with some quantitative training the slightly more advanced course MY452 (Applied Regression Analysis) would be more appropriate, followed by other Department of Methodology and Department of Statistics courses on quantitative methods, such as MY454 (Applied Statistical Computing), MY455 (Multivariate Analysis and Measurement), MY456 (Survey Methodology), MY457 (Causal Inference for Observational and Experimental Studies), MY459 (Quantitative Text Analysis), ST416 (Multilevel Modelling), and ST442 (Longitudinal Data Analysis).

Please note that up to 2010-11 the course was called MI451, so older materials will bear that code. For example, you will find past examination papers on the Library website under MI451.

Course Objectives

This course aims to impart a level of familiarity suitable for a moderately critical understanding of the statistical material in the journals commonly used by students in their work and the ability to use some elementary techniques.

Staff

- **Teachers responsible:** Dominik Hangartner in MT; Jonathan Jackson in LT
- **Other Lecturers:** Benjamin Wilson (MT) and Matthew Mulford (LT)
- **Course Administrator:** Esther Heyhoe, email e.heyhoe@lse.ac.uk
- **Class Teachers:** See the MY451 Moodle site for the list of class teachers.

Teaching

- **Lectures:** Ten two-hour lectures in weeks 1 - 5 and 7 - 11:
 - Michaelmas Term (MY451M): Thu 2.00pm - 4.00pm in the Peacock Theatre
 - Lent Term (MY451L): Wed 2.00pm - 4.00pm in the Hong Kong Theatre (CLM.G.02)
- **Computer classes:** Each student will attend a one-hour class each week, **starting in Week 2**. Please see on-line timetables for the times and places of the class groups. The allocation of students to classes is done through LSE for You or the Moodle system (see below), depending on your status. This will be explained in the first lecture and on the MY451 Moodle page. Please do not change the classes allocated to you without our consent.

Course Materials

- **Coursepack:** This coursepack is the main course text. Printed copies of this coursepack will be available for purchase at the beginning of each term, and instructions on how to obtain a hard copy will be given at the first lecture and on the MY451 Moodle page. Digital copies of this coursepack are available for download for free from the Moodle page.
- **Lecture slides:** Copies of most of the slides displayed during the lectures can be downloaded from the MY451 Moodle page.
- **Recommended course texts:**
 - Alan Agresti and Christine Franklin (2013). *Statistics: The Art and Science of Learning from Data* (Third Ed.). Pearson.
 - Alan Agresti and Barbara Finlay (2013). *Statistical Methods for the Social Sciences* (Fourth Ed.). Pearson

Earlier editions are also suitable. While neither of these books is absolutely required, we recommend that you purchase one if you would like to have additional explanation, examples and exercises to supplement the coursepack. Please note that during the examination you will be allowed to use the coursepack and any other notes or books that you find useful. Of these two, Agresti and Finlay is a bit more advanced. It is also the recommended additional course text for MY452 (which also has a coursepack similar to this one), so you may want to purchase it if you are planning to also take that course.

- **Other text books:** There are hundreds of introductory textbooks on statistics and quantitative methods, many of them covering almost identical material. If you have one which you would like to use, and which looks as if it covers the same material at about the same level as this course, then it is probably suitable as additional reading. If in doubt, please ask one of the lecturers.
 - There are also many books which focus on the SPSS statistical software package used in the computer classes. We do not consider them necessary for this course, or for learning statistics or SPSS. If, however, you feel otherwise, some such books are mentioned in the Appendix.

Homework

There will be homework assignments each week. These are your opportunity to practice interpreting the results of the types of analyses that are introduced in the lectures and then carried out in the computer classes. The homework questions are also broadly similar in style to typical examination questions for MY451. The questions are based on results of the computer class exercises or of further similar analyses. For the homework, you can use output that you produced yourself in the class or afterwards. Alternatively, all the computer outputs that are required to answer the homework questions can be downloaded from the Moodle page.

Homework should be submitted by Monday after its distribution in class, at the latest by 5pm in MT and by 10am in LT. Please see the MY451 Moodle page for instructions on how to submit your homeworks. Homework submitted in time will be corrected and returned at the next class. Please make sure to clearly indicate your name, and the time, class number and class teacher of your computer class on all homework assignments. For the homework assignments, students are encouraged to work together.

MY451 on Moodle

The course materials are all available on Moodle. Go to <http://moodle.lse.ac.uk/> and login using your *username* and *password* (the same as for your LSE e-mail). Then in the *select courses* dialogue box type in MY451, and in *search results* click on MY451. The site contains the structure of the course week by week, the readings, weekly computer class assignments and the associated data sets, coursepack and other materials, as well as a section on news and announcements.

Advisory Sessions

There will be a time each week during the Michaelmas and Lent terms during which a member of the teaching team will be available at the Department of Methodology in Columbia House to answer questions about the course. Information on the times of the advisory hours will be given at the first lecture. These sessions are not intended to be private tutorials and you will be expected to have made a genuine attempt to work through the relevant sections of the coursepack prior to coming to an advisory session. Also, questions addressing material from that week's lecture will be given priority. If you are having difficulty with a topic, come and see us at that time. If you wait until the end of term, when the advisory sessions can be very busy, we cannot guarantee that your questions from earlier sessions will be covered. **There will be no advisory sessions after the end of Lent Term.**

Notes on studying for the course

To learn the material from this course you must do the work every week since it is cumulative; if you miss a week or two (or sleep through it!) there is a chance that you will find yourself hopelessly lost. So this is definitely not a "pick and choose" course! Also bear in mind that most people cannot learn quantitative techniques passively by just turning up to the lectures and reading the occasional chapter in a textbook. To learn statistics you have to do it; there are no shortcuts. Thus in addition to a two-hour weekly lecture there will be one-hour computer classes (in which you do some data analysis and interpretation using SPSS - instructions will be provided) and there will be weekly homework (which will be corrected but not graded by your class teacher). Doing the assignments in the computer classes and the homework demonstrate

whether you have understood and can apply what was covered in the lectures. If you are having any trouble this will reveal what the problem is. Thus the course is designed to have multiple, reinforcing ways of helping you get to grips with this material.

Examinations/assessment

There will be a **two-hour examination in the Summer Term 2016**. This is an **open-book examination**, to which you are allowed to bring in any written material, including the coursepack and any text book with your notes. You are required to bring your own calculators for the examination. Examination papers from previous years are available for revision purposes at the LSE library web site. Students should understand that past examinations should only be used as rough guides to the types of questions that are likely to appear on the examination.

For many of you, MY451 is only one part of a package of methods training with a course code such as MC4M1 or MY4M1. If this is the case, your result for MY451 will contribute only part of the final mark for the package, with a weight determined by the regulations.

Computing

Students must know their Username and Password in time for the first class. This information can be obtained from IT Help Desk (Library, 1st floor). The course uses Microsoft Windows-based software. If you are not familiar with the program, you might want to attend introductory courses in Windows during the first two weeks of the term. The statistical package being used is SPSS, which will be introduced in the first computing class of the course.

Software availability

A personal copy of the program SPSS for Windows can be bought by course participants (maximum 1 copy per person) for the price of £10 from IT services, **not** the Department of Methodology.

Feedback

We would welcome any comments you have on the course. If there are any problems that we can deal with, we will attempt to do so as quickly as possible. Speak to any member of the course team, or to your departmental supervisor if you feel that would be easier for you. Also please let us know if you find any errors or omissions in the coursepack, so that we can correct them for next year.

Acknowledgements

The computer classes were extensively revised for 2012-13. The hard work of designing the new exercises and writing the revised instructions was done by Sally Stares.

This coursepack bears many traces of previous materials and all of their authors, Colm O'Muircheartaigh, Colin Mills, Matt Mulford, Fiona Steele, Paul Mitchell, and Sally Stares. Many thanks to Farimah Daftary, Sue Howard, Jon Jackson, Paul Mitchell, Indraneel Sircar, and many students of previous years for comments and suggestions which are incorporated in the current revision.

Course Programme

Week 1

Lecture Course overview and organisation. Introduction to basic concepts

Class **No class**

Coursepack Chapter 1

Week 2

Lecture Descriptive statistics for categorical variables

Class Introduction to SPSS. Descriptive statistics for categorical variables

Coursepack Sections 2.1–2.4 and 2.8

Week 3

Lecture & Class Descriptive statistics for continuous variables

Coursepack Sections 2.5–2.7

Week 4

Lecture & Class Analysis of two-way contingency tables

Coursepack Chapters 3 and 4

Week 5

Lecture & Class Inference for means in two populations

Coursepack Chapters 6 and 7

Week 6

Reading Week **No lecture, no class.** Revision quiz available on Moodle.

Week 7

Lecture & Class Inference for proportions in one and two populations

Coursepack Chapter 5

Week 8

Lecture & Class Correlation and simple linear regression as descriptive methods

Coursepack Sections 8.1–8.3.4

Week 9

Lectures

Hour 1 Inference for the simple linear regression model

Hour 2 Three-way contingency tables

Class Simple linear regression

Coursepack Section 8.3.5 (Hour 1); Section 8.4 and Chapter 9 (Hour 2)

Week 10

Lecture & Class Multiple linear regression

Coursepack Sections 8.5–8.7

Week 11

Lecture Review and further statistical methods

Class Discussion and additional exercise on multiple linear regression

Coursepack Chapter 10

FAQ: Frequently Asked Questions

Why do we use SPSS? I've heard that SAS/STATA/MINITAB/R/LIMDEP is much better. At this level it does not matter which program you use since we are learning standard procedures that are common to all programs. In favour of SPSS is that it has an easy to learn interface and that it is widely available at sites both within and outside the UK.

Can I get a copy of the SPSS software to use on my home computer? Yes, for a small fee from IT Services helpdesk, not from the Department of Methodology.

I'm taking MY451 because I want to learn how to use SPSS but we don't seem to learn very much about the program. Why is that? MY451 is not a course about learning to use SPSS. We use the program merely to facilitate data analysis and interpretation. Some options for learning more about SPSS will be mentioned in the first lecture.

I'm taking MY451 to help me analyse data for my dissertation. Can I discuss my data and my specific problems with the lecturers? Yes, but not during the course. Staff of the Department of Methodology will be happy to talk to you about problems specific to your dissertation during the weekly sessions of the Methodology Surgery (see the website of the department for more information).

Does the coursepack contain everything I need to know for the exam? Yes. However, you will stand by far the best chance in the exam if you also attend the lectures, where the lecturers emphasise and explain the key parts of the material.

The lecturer introduced some material that was not in the coursepack. Do I need to know that material? This is almost certainly an illusion. The lectures will not introduce any genuinely new material not included in the course pack. However, sometimes the lecturer may of course use different words or a different example to further explain some topic. Copies of the most relevant notes displayed at the lectures will be posted in the MY451 Moodle site. All of the material required for the exam is contained in the coursepack, with the posted lecture notes as additional clarification.

Do I need to bring my coursepack to lectures and classes? Yes, because some of the material displayed during the lectures will be graphs or tables straight from the coursepack. However, an even more important way of using the coursepack is to read the relevant parts of it *before* each lecture. In addition, you will find the instructions for the class exercises at the end of the coursepack.

I can't make it to the advisory hour. Why aren't there more of them? Advisory hours are offered in addition to the normal support arrangements for School courses and rely on the goodwill and enthusiasm of the course team. Given that most team members also offer their time to support advisory hours for other courses there is a limit to how much time they can volunteer. If you genuinely are unable to attend an advisory hour, but desperately need advice, lecturers will always see you in their personal office hours. If these are inconvenient you may be able to set up an appointment at another time by special arrangement.

There was a long queue at advisory and I didn't get much/any time to discuss my particular problem. Isn't there a better way to arrange things? Advisory sessions are meant for quick feedback and clarification of problems on a first-come first served basis. They are not meant to be personal tutorials, seminars or impromptu lectures. They are also not meant to be a substitute for attending the lectures; reading the course-pack; doing the homework; thinking. Priority is always given to problems related to the previous week's lecture material. If several people have the same problem it will usually be possible to talk to everyone together, but there is a finite limit to how many people can be crammed into an office at the same time!

I want to bring my homework to the advisory session because I couldn't understand the class teacher's comments. Is that OK? Yes, but it is unlikely that the person that wrote the comments on your homework will be taking the advisory hour. Usually it will be much better to discuss homework comments directly with your class teacher in class.

Can I work together on the homework with my friends? Yes, we positively encourage you to discuss the homework assignments with your colleagues. If you do this, please submit **one** copy to be marked rather than multiple copies of the same answers.

If I get less than 50% in the exam what happens now? Candidates sit MY451 under different examination regulations depending on which department they are in and what degree they are registered for. For some degrees a pass in MY451 is required, for others the mark from MY451 is just one component of an overall course mark (implying a "pass" mark in MY451 is not required). The regulations for your degree are not set by Methodology, but by your home department. To find out whether a pass in MY451 is required, consult your degree regulations or ask the tutor responsible for your program in your home department. Candidates who fail at the first attempt and whose degree regulations require a pass are, as per the usual School examination rules, entitled, at the next available opportunity, to ONE more attempt. This will be in the Summer Term of the next Session.

Why don't we get our examination scripts returned after they have been marked? Written work is of two sorts, formative and summative. The former, for example the homework exercises, is meant to give you feedback on how you are doing and will, where appropriate, be returned with comments. The latter, for example the examination, is meant to evaluate your performance, not explain to you where you have gone wrong, or for that matter where you have done exceptionally well.

If I don't agree with the mark, how do I appeal? The cultural norm in the UK is that marks are not arrived at by a process of teacher-student negotiation. You can make a formal appeal through the School's appeal process (see the appropriate section of the website for details). NB: appeals cannot be made on grounds of academic substance, only on grounds of procedural irregularities. In other words an appeal will not be allowed if the only grounds you have for appealing is that you/your friend/your personal advisor/your spiritual guru think your script was worth more marks than the examiners did.

The class teachers are always very busy during class helping people with problems and I can't get enough personal attention. Why can't there be more class teachers in the classroom? For MY451 there will usually be two class teachers per classroom i.e. double the normal complement. Even so we realise that you may not be able to get attention at the exact moment you want it. Please be patient and if help is not immediately available you can always try asking the person who is sitting next to you!

I'm not registered at the LSE but at another University of London college. Can I attend this course? Normally yes, but you will have to complete an intercollegiate enrolment form.

I would like to audit the course without taking the exam. Is that OK? Yes, you are welcome to attend the lectures providing you are an LSE/University of London student and there is room for you.

MY451 is not challenging enough for me. Is there a more difficult course? Yes, MY452 and numerous other courses offered by the Department of Methodology and the Statistics department.

Chapter 1

Introduction

1.1 What is the purpose of this course?

The title of any course should be descriptive of its contents. This one is called

MY451: Introduction to Quantitative Analysis

Every part of this tells us something about the nature of the course:

The **M** stands for *Methodology* of social research. Here *research* refers to activities aimed at obtaining new knowledge about the world, in the case of the social sciences the *social* world of people and their institutions and interactions. Here we are concerned solely with *empirical* research, where such knowledge is based on information obtained by *observing* what goes on in that world. There are many different ways (*methods*) of making such observations, some better than others for deriving valid knowledge. “Methodology” refers both to the methods used in particular studies, and the study of research methods in general.

The word **analysis** indicates the area of research methodology that the course is about. In general, any empirical research project will involve at least the following stages:

1. Identifying a research *topic*
2. Formulating *research questions*
3. Deciding what kinds of *information* to collect to try to answer the research questions, and deciding how to collect it and where to collect it from
4. Collecting the information
5. *Analysing* the information in appropriate ways to answer the research questions
6. *Reporting* the findings

The empirical information collected in the research process is often referred to as *data*. This course is mostly about some basic methods for step 5, the *analysis* of such data.

Methods of analysis, however competently used, will not be very useful unless other parts of the research process have also been carried out well. These other parts, which (especially steps 2–4 above) can be broadly termed *research design*, are covered on other courses, such as MY400 (Fundamentals of Social Science Research Design) or comparable courses at your own department. Here we will mostly not consider research design, in effect assuming that we start at a point where we want to analyse some data which have been collected in a sensible way to

answer meaningful research questions. However, you should bear in mind throughout the course that in a real research situation both good design and good analysis are essential for success.

The word **quantitative** in the title of the course indicates that the methods you will learn here are used to analyse quantitative data. This means that the data will enter the analysis in the form of *numbers* of some kind. In social sciences, for example, data obtained from administrative records or from surveys using structured interviews are typically quantitative. An alternative is *qualitative* data, which are not rendered into numbers for the analysis. For example, unstructured interviews, focus groups and ethnography typically produce mostly qualitative data. Both quantitative and qualitative data are important and widely used in social research. For some research questions, one or the other may be clearly more appropriate, but in many if not most cases the research would benefit from collecting both qualitative and quantitative data. This course will concentrate solely on quantitative data analysis, while the collection and analysis of qualitative data are covered on other courses (e.g. MY421, MY426 and MY427), which we hope you will also be taking.

All the methods taught here, and almost all approaches used for quantitative data analysis in the social sciences in general, are *statistical* methods. The defining feature of such methods is that randomness and probability play an essential role in them; some of the ways in which they do so will become apparent later, others need not concern us here. The title of the course could thus also have included the word *statistics*. However, the Department of Methodology courses on statistical methods (e.g. MY451, MY465, MY452, MY455 and MY459) have traditionally been labelled as courses on “quantitative analysis” rather than “statistics”. This is done to indicate that they differ from classical introductory statistics courses in some ways, especially in the presentation being less mathematical.

The course is called an “**Introduction** to Quantitative Analysis” because it is an introductory course which does not assume that you have learned any statistics before. MY451 or a comparable course should be taken before more advanced courses on quantitative methods. Statistics is a cumulative subject where later courses build on material learned on earlier ones. Because MY451 is introductory, it will start with very simple methods, and many of the more advanced (and powerful) ones will only be covered on the later courses. This does not, however, mean that you are wasting your time here even if it is methods from, say, MY452 that you will eventually need most: understanding the material of this course is essential for learning more advanced methods.

Finally, the course has an **MY** code, rather than GV, MC, PS, SO, SP, or whatever is the code of your own department. MY451 is taken by students from many different degrees and departments, and thus cannot be tailored to any one of them specifically. For example, we will use examples from many different social sciences. However, this generality is definitely a good thing: the reason we *can* teach all of you together is that statistical methods (just like the principles of research design or qualitative research) are generic and applicable to the analysis of quantitative data in all fields of social research. There is not, apart from differences in emphases and priorities, one kind of statistics for sociology and another for political science or economics, but one coherent set of principles and methods for all of them (as well as for psychiatry, epidemiology, biology, astrophysics and so on). After this course you will have taken the first steps in learning about all of that.

At the end of the course you should be familiar with certain methods of statistical analysis. This will enable you to be both a user and a consumer of statistics:

- You will be able to use the methods to analyse your own data and to report the results of the analyses.
- Perhaps even more importantly, you will also be able to understand (and possibly criticize)

their use in other people's research. Because interpreting results is typically somewhat easier than carrying out new analyses, and because all statistical methods use the same basic ideas introduced here, you will even have some understanding of many of the techniques not discussed on this course.

Another pair of different but complementary aims of the course is that MY451 is both a self-contained unit and a prerequisite for courses that follow it:

- If this is the last statistics course you will take, it will enable you to understand and use the particular methods covered here. This includes the technique of linear regression modelling (described in Chapter 8), which is arguably the most important and commonly used statistical method of all. This course can, however, introduce only the most important elements of linear regression, while some of the more advanced ones are discussed only on MY452.
- The ideas learned on this course will provide the conceptual foundation for any further courses in quantitative methods that you may take. The basic ideas will then not need to be learned from scratch again, and the other courses can instead concentrate on introducing further, ever more powerful statistical methods for different types of data.

1.2 Some basic definitions

Like any discipline, statistics involves some special terminology which makes it easier to discuss its concepts with sufficient precision. Some of these terms are defined in this section, while others will be introduced later when they are needed.

You should bear in mind that all terminology is arbitrary, so there may be different terms for the same concept. The same is true of notation and symbols (such as n , μ , \bar{Y} , R^2 , and others) which will be introduced later. Some statistical terms and symbols are so well established that they are almost always used in the same way, but for many others there are several versions in common use. While we try to be consistent with the notation and terminology within this coursepack, we cannot absolutely guarantee that we will not occasionally use different terms for the same concept even here. In other textbooks and in research articles you will certainly occasionally encounter alternative terminology for some of these concepts. If you find yourself confused by such differences, please come to the advisory hours or ask your class teacher for clarification.

1.2.1 Subjects and variables

Table 1.1 shows a small set of quantitative data. Once collected, the data are typically arranged and stored in this kind of spreadsheet-type rectangular table, known as a **data matrix**. In the computer classes you will see data in this form in SPSS.

Table 1.1: An example of a small data matrix based on data from the U.S. General Social Survey (GSS), showing measurements of seven variables for 20 respondents in a social survey. The variables are defined as *age*: age in years; *sex*: sex (1=male; 2=female); *educ*: highest year of school completed; *wrkstat*: labour force status (1=working full time; 2=working part time; 3=temporarily not working; 4=unemployed; 5=retired; 6=in education; 7=keeping house; 8=other); *life*: is life exciting or dull? (1=dull; 2=routine; 3=exciting); *income4*: total annual family income (1=\$24,999 or less; 2=\$25,000–\$39,999; 3=\$40,000–\$59,999; 4=\$60,000 or more; 99 indicates a missing value); *pres92*: vote in the 1992 presidential election (0=did not vote or not eligible to vote; 1=Bill Clinton; 2=George H. W. Bush; 3=Ross Perot; 4=Other).

Id	<i>age</i>	<i>sex</i>	<i>educ</i>	<i>wrkstat</i>	<i>life</i>	<i>income4</i>	<i>pres92</i>
1	43	1	11	1	2	3	2
2	44	1	16	1	3	3	1
3	43	2	16	1	3	3	2
4	78	2	17	5	3	4	1
5	83	1	11	5	2	1	1
6	55	2	12	1	2	99	1
7	75	1	12	5	2	1	0
8	31	1	18	1	3	4	2
9	54	2	18	2	3	1	1
10	23	2	15	1	2	3	3
11	63	2	4	5	1	1	1
12	33	2	10	4	3	1	0
13	39	2	8	7	3	1	0
14	55	2	16	1	2	4	1
15	36	2	14	3	2	4	1
16	44	2	18	2	3	4	1
17	45	2	16	1	2	4	1
18	36	2	18	1	2	99	1
19	29	1	16	1	3	3	1
20	30	2	14	1	2	2	1

The rows (moving downwards) and columns (moving left to right) of a data matrix correspond to the first two important terms: the rows to the *subjects* and the columns to the *variables* in the data.

- A **subject** is the smallest unit yielding information in the study. In the example of Table 1.1, the subjects are individual people, as they are in very many social science examples. In other cases they may instead be families, companies, neighbourhoods, countries, or whatever else is relevant in a particular study. There is also much variation in the term itself, so that instead of “subjects”, a study might refer to “units”, “elements”, “respondents” or “participants”, or simply to “persons”, “individuals”, “families” or “countries”, for example. Whatever the term, it is usually clear from the context what the subjects are in a particular analysis.

The subjects in the data of Table 1.1 are uniquely identified only by a number (labelled

“Id”) assigned by the researcher, as in a survey like this their names would not typically be recorded. In situations where the identities of individual subjects are available and of interest (such as when they are countries), their names would typically be included in the data matrix.

- A **variable** is a characteristic which varies between subjects. For example, Table 1.1 contains data on seven variables — age, sex, education, labour force status, attitude to life, family income and vote in a past election — defined and recorded in the particular ways explained in the caption of the table. It can be seen that these are indeed “variable” in that not everyone has the same value of any of them. It is this variation that makes collecting data on many subjects necessary and worthwhile. In contrast, research questions about characteristics which are the same for every subject (i.e. *constants* rather than variables) are rare, usually not particularly interesting, and not very difficult to answer.

The labels of the columns in Table 1.1 (*age*, *wrkstat*, *income4* etc.) are the names by which the variables are uniquely identified in the data file on a computer. Such concise titles are useful for this purpose, but should be avoided when reporting the results of data analyses, where clear English terms can be used instead. In other words, a report should not say something like “The analysis suggests that WRKSTAT of the respondents is...” but instead something like “The analysis suggests that the labour force status of the respondents is...”, with the definition of this variable and its categories also clearly stated.

Collecting quantitative data involves determining the values of a set of variables for a group of subjects and assigning numbers to these values. This is also known as **measuring** the values of the variables. Here the word “measure” is used in a broader sense than in everyday language, so that, for example, we are measuring a person’s sex in this sense when we assign a variable called “Sex” the value 1 if the person is male and 2 if she is female. The value assigned to a variable for a subject is called a **measurement** or an **observation**. Our data thus consist of the measurements of a set of variables for a set of subjects. In the data matrix, each row contains the measurements of all the variables in the data for one subject, and each column contains the measurements of one variable for all of the subjects.

The number of subjects in a set of data is known as the **sample size**, and is typically denoted by n . In a survey, for example, this would be the number of people who responded to the questions in the survey interview. In Table 1.1 we have $n = 20$. This would normally be a very small sample size for a survey, and indeed the real sample size in this one is several thousands. The twenty subjects here were drawn from among them to obtain a small example which fits on a page.

A common problem in many studies is **nonresponse** or **missing data**, which occurs when some measurements are not obtained. For example, some survey respondents may refuse to answer certain questions, so that the values of the variables corresponding to those questions will be missing for them. In Table 1.1, the income variable is missing for subjects 6 and 18, and recorded only as a *missing value code*, here “99”. Missing values create a problem which has to be addressed somehow before or during the statistical analysis. The easiest approach is to simply ignore all the subjects with missing values and use only those with complete data on all the variables needed for a given analysis. For example, any analysis of the data in Table 1.1 which involved the variable *income4* would then exclude all the data for subjects 6 and 18. This method of “complete-case analysis” is usually applied automatically by most statistical software packages, including SPSS. It is, however, not a very good approach. For example, it means that a lot of information will be thrown away if there are many subjects with some observations missing. Statisticians have developed better ways of dealing with missing data, but they are unfortunately beyond the scope of this course.

1.2.2 Types of variables

Information on a variable consists of the observations (measurements) of it for the subjects in our data, recorded in the form of numbers. However, not all numbers are the same. First, a particular way of measuring a variable may or may not provide a good measure of the concept of interest. For example, a measurement of a person's weight from a well-calibrated scale would typically be a good measure of the person's true weight, but an answer to the survey question "How many units of alcohol did you drink in the last seven days?" might be a much less accurate measurement of the person's true alcohol consumption (i.e. it might have *measurement error* for a variety of reasons). So just because you have put a number on a concept does not automatically mean that you have captured that concept in a useful way. Devising good ways of measuring variables is a major part of research design. For example, social scientists are often interested in studying attitudes, beliefs or personality traits, which are very difficult to measure directly. A common approach is to develop *attitude scales*, which combine answers to multiple questions ("items") on the attitude into one number.

Here we will again leave questions of measurement to courses on research design, effectively assuming that the variables we are analysing have been measured well enough for the analysis to be meaningful. Even then we will have to consider some distinctions between different kinds of variables. This is because the type of a variable largely determines which methods of statistical analysis are appropriate for that variable. It will be necessary to consider two related distinctions:

- Between different measurement levels
- Between continuous and discrete variables

Measurement levels

When a numerical value of a particular variable is allocated to a subject, it becomes possible to relate that value to the values assigned to other subjects. The **measurement level** of the variable indicates how much information the number provides for such comparisons. To introduce this concept, consider the variables obtained as answers to the following three questions in the former U.K. General Household Survey:

1

Are you

<i>single, that is, never married?</i>	(coded as 1)
<i>married and living with your husband/wife?</i>	(2)
<i>married and separated from your husband/wife?</i>	(3)
<i>divorced?</i>	(4)
<i>or widowed?</i>	(5)

2

Over the last twelve months, would you say your health has on the whole been good, fairly good, or not good?
 ("Good" is coded as 1, "Fairly Good" as 2, and "Not Good" as 3.)

3

About how many cigarettes A DAY do you usually smoke on weekdays?
(Recorded as the number of cigarettes)

These variables illustrate three of the four possibilities in the most common classification of measurement levels:

- A variable is measured on a **nominal scale** if the numbers are simply labels for different possible values (*levels* or *categories*) of the variable. The only possible comparison is then to identify whether two subjects have the *same* or *different* values of the variable. The marital status variable

1

is measured on a nominal scale. The values of such *nominal-level variables* are not in any order, so we cannot talk about one subject having “more” or “less” of the variable than another subject; even though “divorced” is coded with a larger number (4) than “single” (1), divorced is not more or bigger than single in any relevant sense. We also cannot carry out arithmetical calculations on the values, as if they were numbers in the ordinary sense. For example, if one person is single and another widowed, it is obviously nonsensical to say that they are on average separated (even though $(1 + 5)/2 = 3$).

The only requirement for the codes assigned to the levels of a nominal-level variable is that different levels must receive different codes. Apart from that, the codes are arbitrary, so that we can use any set of numbers for them in any order. Indeed, the codes do not even need to be numbers, so they may instead be displayed in the data matrix as short words (“labels” for the categories). Using successive small whole numbers (1, 2, 3, ...) is just a simple and concise choice for the codes.

Further examples of nominal-level variables are the variables *sex*, *wrkstat*, and *pres92* in Table 1.1.

- A variable is measured on an **ordinal scale** if its values do have a natural ordering. It is then possible to determine not only whether two subjects have the same value, but also whether one or the other has a *higher* value. For example, the self-reported health variable

2

is an ordinal-level variable, as larger values indicate worse states of health. The numbers assigned to the categories now have to be in the correct order, because otherwise information about the true ordering of the categories would be distorted. Apart from the order, the choice of the actual numbers is still arbitrary, and calculations on them are still not strictly speaking meaningful.

Further examples of ordinal-level variables are *life* and *income4* in Table 1.1.

- A variable is measured on an **interval scale** if *differences* in its values are comparable. One example is temperature measured on the Celsius (Centigrade) scale. It is now meaningful to state not only that 20°C is a *different* and *higher* temperature than 5°C, but also that the *difference* between them is 15°C, and that that difference is of the same size as the difference between, say, 40°C and 25°C. Interval-level measurements are “proper” numbers in that calculations such as the average noon temperature in London over a year are meaningful. What we *cannot* do is to compare *ratios* of interval-level variables. Thus 20°C is not four times as warm as 5°C, nor is their real ratio the same as that of 40°C and 10°C. This is because the zero value of the Celsius scale (0°C) is not the lowest possible temperature but an arbitrary point chosen for convenience of definition.

- A variable is measured on a **ratio scale** if it has all the properties of an interval-level variable and also a true zero point. For example, the smoking variable

3

is measured on a ratio level, with zero cigarettes as its point of origin. It is now possible to carry out all the comparisons possible for interval-level variables, and also to compare ratios. For example, it is meaningful to say that someone who smokes 20 cigarettes a day smokes *twice* as many cigarettes as one who smokes 10 cigarettes, and that that ratio is equal to the ratio of 30 and 15 cigarettes.

Further examples of ratio-level variables are *age* and *educ* in Table 1.1.

The distinction between interval-level and ratio-level variables is in practice mostly unimportant, as the same statistical methods can be applied to both. We will thus consider them together throughout this course, and will, for simplicity, refer to variables on either scale as interval level variables. Doing so is logically coherent, because ratio level variables have all the properties of interval level variables, as well the additional property of a true zero point.

Similarly, nominal and ordinal variables can often be analysed with the same methods. When this is the case, we will refer to them together as nominal/ordinal level variables. There are, however, contexts where the difference between them matters, and we will then discuss nominal and ordinal scales separately.

The simplest kind of nominal variable is one with only *two* possible values, for example sex recorded as “male” or “female” or an opinion recorded just as “agree” or “disagree”. Such a variable is said to be **binary** or **dichotomous**. As with any nominal variable, codes for the two levels can be assigned in any way we like (as long as different levels get different codes), for example as 1=Female and 2=Male; later it will turn out that in some analyses it is most convenient to use the values 0 and 1.

The distinction between ordinal-level and interval-level variables is sometimes further blurred in practice. Consider, for example, an attitude scale of the kind mentioned above, let’s say a scale for happiness. Suppose that the possible values of the scale range from 0 (least happy) to 48 (most happy). In most cases it would be most realistic to consider these measurements to be on an ordinal rather than an interval scale. However, statistical methods developed specifically for ordinal-level variables do not cope very well with variables with this many possible values. Thus ordinal variables with many possible values (at least more than ten, say) are typically treated as if they were measured on an interval scale.

Continuous and discrete variables

This distinction is based on the possible values a variable can have:

- A variable is **discrete** if its basic unit of measurement cannot be subdivided. Thus a discrete variable can only have certain values, and the values between these are logically impossible. For example, the marital status variable

1

and the health variable

2

defined under “Measurement Levels” in Section 1.2.2 are discrete, because values like marital status of 2.3 or self-reported health of 1.7 are impossible given the way the variables are defined.

- A variable is **continuous** if it can in principle take infinitely varied fractional values. The idea implies an unbroken scale or continuum of possible values. Age is an example of a continuous variable, as we can in principle measure it to any degree of accuracy we like — years, days, minutes, seconds, micro-seconds. Similarly, distance, weight and even income can be considered to be continuous.

You should note the “in principle” in this definition of continuous variables above. Continuity is here a pragmatic concept, not a philosophical one. Thus we will treat age and income as continuous even though they are in practice measured to the nearest year or the nearest hundred pounds, and not in microseconds or millionths of a penny (nor is the definition inviting you to start musing on quantum mechanics and arguing that nothing is fundamentally continuous). What the distinction between discrete and continuous really amounts to in practice is the difference between variables which in our data tend to take relatively few values (discrete variables) and ones which can take lots of different values (continuous variables). This also implies that we will sometimes treat variables which are undeniably discrete in the strict sense as if they were really continuous. For example, the number of people is clearly discrete when it refers to numbers of registered voters in households (with a limited number of possible values in practice), but effectively continuous when it refers to populations of countries (with very many possible values).

The measurement level of a variable refers to the way a characteristic is recorded in the data, not to some other, perhaps more fundamental version of that characteristic. For example, annual income recorded to the nearest dollar is continuous, but an income variable (c.f. Table 1.1) with values

- if annual income is \$24,999 or less;
- if annual income is \$25,000–\$39,999;
- if annual income is \$40,000–\$59,999;
- if annual income is \$60,000 or more

is discrete. This kind of variable, obtained by grouping ranges of values of an initially continuous measurement, is common in the social sciences, where the exact values of such variables are often not that interesting and may not be very accurately measured.

The term **categorical variable** will be used in this coursepack to refer to a discrete variable which has only a finite (in practice quite small) number of possible values, which are known in advance. For example, a person’s sex is typically coded simply as “Male” or “Female”, with no other values. Similarly, the grouped income variable shown above is categorical, as every income corresponds to one of its four categories (note that it is the “rest” category 4 which guarantees that the variable does indeed cover all possibilities). Categorical variables are of separate interest because they are common and because some statistical methods are designed specifically for them. An example of a non-categorical discrete variable is the population of a country, which does not have a small, fixed set of possible values (unless it is again transformed into a grouped variable as in the income example above).

Relationships between the two distinctions

The distinctions between variables with different measurement levels on one hand, and continuous and discrete variables on the other, are partially related. Essentially all nominal/ordinal-level variables are discrete, and almost all continuous variables are interval-level variables. This leaves one further possibility, namely a discrete interval-level variable; the most common exam-

ple of this is a **count**, such as the number of children in a family or the population of a country. These connections are summarized in Table 1.3.

Table 1.3: Relationships between the types of variables discussed in Section @ref(ss-intro-def-vartypes).

	<i>Measurement level</i>	<i>Measurement level</i>
	Nominal/ordinal	Interval/ratio
Discrete	Many - Always categorical , i.e. having a fixed set of possible values (categories) - If only two categories, variable is binary (dichotomous)	<i>Counts</i> - If many different observed values, often treated as effectively continuous
Continuous	None	Many

In practice the situation may be even simpler than this, in that the most relevant distinction is often between the following two cases:

1. Discrete variables with a small number of observed values. This includes both categorical variables, for which all possible values are known in advance, and variables for which only a small number of values were actually observed even if others might have been possible¹. Such variables can be conveniently summarized in the form of tables and handled by methods appropriate for such tables, as described later in this coursepack. This group also includes all nominal variables, even ones with a relatively large number of categories, since methods for group 2. below are entirely inappropriate for them.
2. Variables with a large number of possible values. This includes all continuous variables and those interval-level or ordinal discrete variables which have so many values that it is pragmatic to treat them as effectively continuous.

Although there are contexts where we need to distinguish between types of variables more carefully than this, for practical purposes this simple distinction is often sufficient.

1.2.3 Description and inference

In the past, the subtitle of this course was “Description and inference”. This is still descriptive of the contents of the course. These words refer to two different although related tasks of statistical analysis. They can be thought of as solutions to what might be called the “too much and not enough” problems with observed data. A set of data is “too much” in that it is very difficult to understand or explain the data, or to draw any conclusions from it, simply by staring at the numbers in a data matrix. Making much sense of even a small data matrix like the one in Table 1.1 is challenging, and the task becomes entirely impossible with bigger ones. There is thus a clear need for methods of statistical description:

- **Description:** summarizing some features of the data in ways that make them easily understandable. Such methods of description may be in the form of numbers or graphs.

¹ESS Round 5: European Social Survey Round 5 Data (2010). Data file edition 2.0. Norwegian Social Science Data Services, Norway Data Archive and distributor of ESS data. The full data can be obtained from <http://ess.nsd.uib.no/ess/round5/>.

The “not enough” problem is that quite often the subjects in the data are treated as representatives of some larger group which is our real object of interest. In statistical terminology, the observed subjects are regarded as a **sample** from a larger **population**. For example, a pre-election opinion poll is not carried out because we are particularly interested in the voting intentions of the particular thousand or so people who answer the questions in the poll (the sample), but because we hope that their answers will help us draw conclusions about the preferences of all of those who intend to vote on election day (the population). The job of statistical inference is to provide methods for generalising from a sample to the population:

- **Inference:** drawing conclusions about characteristics of a population based on the data observed in a sample. The two main tools of statistical inference are **significance tests** and **confidence intervals**.

Some of the methods described on this course are mainly intended for description and others for inference, but many also have a useful role in both.

1.2.4 Association and causation

The simplest methods of analysis described on this course consider questions which involve only one variable at a time. For example, the variable might be the political party a respondent intends to vote for in the next general election. We might then want to know what proportion of voters plan to vote for the Labour party, or which party is likely to receive the most votes.

However, considering variables one at a time is not going to entertain us for very long. This is because most interesting research questions involve associations between variables. One way to define an association is that

- There is an **association** between two variables if knowing the value of one of the variables will help to predict the value of the other variable.

(A more careful definition will be given later.) Other ways of referring to the same concept are that the variables are “related” or that there is a “dependence” between them.

For example, suppose that instead of considering voting intentions overall, we were interested in *comparing* them between two groups of people, homeowners and people who live in rented accommodation. Surveys typically suggest that homeowners are more likely to vote for the Conservatives and less likely to vote for Labour than renters. There is then an association between the two (discrete) variables “type of accommodation” and “voting intention”, and knowing the type of a person’s accommodation would help us better predict who they intend to vote for. Similarly, a study of education and income might find that people with more education (measured by years of education completed) tend to have higher incomes (measured by annual income in pounds), again suggesting an association between these two (continuous) variables.

Sometimes the variables in an association are in some sense on an equal footing. More often, however, they are instead considered asymmetrically in that it is more natural to think of one of them as being used to predict the other. For example, in the examples of the previous paragraph it seems easier to talk about home ownership predicting voting intention than vice versa, and of level of education predicting income than vice versa. The variable used for prediction is then known as an **explanatory variable** and the variable to be predicted as the **response variable** (an alternative convention is to talk about **independent** rather than explanatory variables and **dependent** instead of response variables). The most powerful statistical techniques for analysing associations between explanatory and response variables are known as **regression** methods. They are by far the most important family of methods of quantitative data analysis.

On this course you will learn about the most important member of this family, the method of **linear regression**.

In the many research questions where regression methods are useful, it almost always turns out to be crucially important to be able to consider several different explanatory variables simultaneously for a single response variable. Regression methods allow for this through the techniques of **multiple regression**.

The statistical concept of association is closely related to the stronger concept of **causation**, which is at the heart of very many research questions in the social sciences and elsewhere. The two concepts are not the same. In particular, association is not *sufficient* evidence for causation, i.e. finding that two variables are statistically associated does not prove that either variable has a causal effect on the other. On the other hand, association is almost always *necessary* for causation: if there is no association between two variables, it is very unlikely that there is a direct causal effect between them. This means that analysis of associations is a necessary part, but not the only part, of the analysis of causal effects from quantitative data. Furthermore, statistical analysis of associations is carried out in essentially the same way whether or not it is intended as part of a causal argument. On this course we will mostly focus on associations. The kinds of additional arguments that are needed to support causal conclusions are based on information on the research design and the nature of the variables. They are discussed only briefly on this course, and at greater length on courses of research design such as MY400 (and the more advanced MY457, which considers design and analysis for causal inference together).

1.3 Outline of the course

We have now defined three separate distinctions between different problems for statistical analysis, according to (1) the types of variables involved, (2) whether description or inference is required, and (3) whether we are examining one variable only or associations between several variables. Different combinations of these elements require different methods of statistical analysis. They also provide the structure for the course, as follows:

- **Chapter 2:** Description for single variables of any type, and for associations between categorical variables.
- **Chapter 3:** Some general concepts of statistical inference.
- **Chapter 4:** Inference for associations between categorical variables.
- **Chapter 5:** Inference for single dichotomous variables, and for associations between a dichotomous explanatory variable and a dichotomous response variable.
- **Chapter 6:** More general concepts of statistical inference.
- **Chapter 7:** Description and inference for associations between a dichotomous explanatory variable and a continuous response variable, and inference for single continuous variables.
- **Chapter 8:** Description and inference for associations between any kinds of explanatory variables and a continuous response variable.
- **Chapter 9:** Some additional comments on analyses which involve three or more categorical variables.

As well as in Chapters 3 and 6, general concepts of statistical inference are also gradually introduced in Chapters 4, 5 and 7, initially in the context of the specific analyses considered in these chapters.

1.4 The use of mathematics and computing

Many of you will approach this course with some reluctance and uncertainty, even anxiety. Often this is because of fears about mathematics, which may be something you never liked or never learned that well. Statistics does indeed involve a lot of mathematics in both its algebraic (symbolical) and arithmetic (numerical) senses. However, the understanding and use of statistical concepts and methods can be usefully taught and learned even without most of that mathematics, and that is what we hope to do on this course. It is perfectly possible to do well on the course without being at all good at mathematics of the secondary school kind.

1.4.1 Symbolic mathematics and mathematical notation

Statistics *is* a mathematical subject in that its concepts and methods are expressed using mathematical formalism, and grounded in a branch of mathematics known as probability theory. As a result, heavy use of mathematics is essential for those who develop these methods (i.e. statisticians). However, those who only *use* them (i.e. you) can ignore most of it and still gain a solid and non-trivialised understanding of the methods. We will thus be able to omit most of the mathematical details. In particular, we will not show you how the methods are derived or prove theorems about them, nor do we expect you to do anything like that.

We will, however, use mathematical notation whenever necessary to state the main results and to define the methods used. This is because mathematics is the language in which many of these results are easiest to express clearly and accurately, and trying to avoid all mathematical notation would be contrived and unhelpful. Most of the notation is fairly simple and will be explained in detail. We will also interpret such formulas in English as well to draw attention to their most important features.

Another way of explaining statistical methods is through applied examples. These will be used throughout the course. Most of them are drawn from real data from research in a range of social sciences. If you wish to find further examples of how these methods are used in your own discipline, a good place to start is in relevant books and research journals.

1.4.2 Computing

Statistical analysis involves also a lot of mathematics of the numerical kind, i.e. various calculations on the numbers in the data. Doing such calculations by hand or with a pocket calculator would be tedious and unenlightening, and in any case impossible for all but the smallest samples and simplest methods. We will mostly avoid doing that by leaving the drudgery of calculation to computers, where the methods are implemented in statistical software packages. This also means that you can carry out the analyses without understanding all the numerical details of the calculations. Instead, we can focus on trying to understand when and why certain methods of analysis are used, and learning to interpret their results.

A simple pocket calculator is still more convenient than a computer for some very simple calculations. You will also need one for this purpose in the examination, where computers are not allowed. Any such calculations required in the examination will be extremely simple to do (assuming you know what you are trying to do, of course). For more complex analyses, the exam questions will involve interpreting computer output rather than carrying out the calculations. The homework questions that follow the computer classes contain examples of both of these types of questions.

The software package used in the computer classes of this course is called SPSS. There are other comparable packages, for example SAS, Minitab, Stata and R. Any one of them could be used for the analyses on this course, and the exact choice does not matter very much. SPSS is convenient for our purposes, because it is widely used, has a reasonably user-friendly menu interface, and is available on a cheap licence even for the personal computers of LSE students.

Sometimes you may see a phrase such as “SPSS course” used apparently as a synonym for “Statistics course”. This makes as little sense as treating an introduction to Microsoft Word as a course on how to write good English. It is not possible to learn quantitative data analysis well by just sitting down in front of SPSS or any other statistics package and trying to figure out what all those menus are for. On the other hand, using SPSS to apply statistical methods to analyse real data is an effective way of strengthening the understanding of those methods *after* they have first been introduced in lectures. That is why this course has weekly computer classes.

The software-specific questions on how to carry out statistical analyses are typically of a lesser order of difficulty once the methods themselves are reasonably well understood. In other words, once you have a clear idea of what you want to do, finding out how to do it in SPSS tends not to be that difficult. For example, in the next chapter we will discuss the mean, one simple tool of descriptive statistics. Suppose that you then want to calculate the mean of a variable called *Age* in a data set. Learning how to do this in SPSS is then a matter of (1) finding the menu item where SPSS can be told to calculate a mean, (2) finding out which part of that menu is used to tell SPSS that you want the mean of *Age* specifically, and (3) finding the part of the SPSS output where the calculated mean of *Age* is reported. Instructions for steps like this for techniques covered on this course are given in the descriptions of the corresponding computer classes.

There are, however, some tasks which have more to do with specific software packages than with statistics in general. For example, the fact that SPSS has a menu interface, and the general style of those menus, need to be understood first. You also need to learn how to get data into SPSS in the first place, how to manipulate the data in various ways, and how to export output from the analyses to other packages. Some instructions on how to do such things are given in the first computer class. The introduction to the computer classes also includes details of some SPSS guidebooks and other sources of information which you may find useful if you want to know more about the program.

Chapter 2

Descriptive statistics

2.1 Introduction

This chapter introduces some common descriptive statistical methods. It is organised around two dichotomies:

- Methods that are used only for variables with small numbers of values, vs. methods that are used also or only for variables with many values (see the end of Section 1.2.2 for more on this distinction). The former include, in particular, descriptive methods for categorical variables, and the latter the methods for continuous variables.
- **Univariate** descriptive methods which consider only one variable at a time, vs. **bivariate** methods which aim to describe the association between *two* variables.

Section 2.3 describes univariate methods for categorical variables and Section 2.4 bivariate methods for cases where both variables are categorical. Sections 2.5 and 2.6 cover univariate methods which are mostly used for continuous variables. Section 2.7 lists some bivariate methods where at least one variable is continuous; these methods are discussed in detail elsewhere in the coursepack. The chapter concludes with some general guidelines for presentation of descriptive tables and graphs in Section 2.8.

2.2 Example data sets

Two examples are used to illustrate the methods throughout this chapter:

Example: Country data

Consider data for 155 countries on three variables:

- The **region** where the country is located, coded as 1=Africa, 2=Asia, 3=Europe, 4=Latin America, 5=Northern America, 6=Oceania.
- A measure of the level of **democracy** in the country, measured on an 11-point scale from 0 (lowest level of democracy) to 10 (highest).
- Gross Domestic Product (**GDP**) per capita, in thousands of U.S. dollars.

Further information on the variables is given in the appendix to this chapter (Section 2.9), together with the whole data set, shown in Table 2.14.

Region is clearly a discrete (and categorical), nominal-level variable, and GDP a continuous, interval-level variable. The democracy index is discrete; it is most realistic to consider its measurement level to be ordinal, and it is regarded as such in this chapter. However, it is the kind of variable which might in many analyses be treated instead as an effectively continuous, interval-level variable.

Example: Survey data on attitudes towards income redistribution

The data for the second example come from Round 5 of the European Social Survey (ESS), which was carried out in 2010.¹ The survey was fielded in 28 countries, but here we use only data from 2344 respondents in the UK. Two variables are considered:

- **Sex** of the respondent, coded as 1=Male, 2=Female.
- Answer to the following survey question:
“The government should take measures to reduce differences in income levels”,
 with five response options coded as “Agree strongly”=1, “Agree”=2, “Neither agree nor disagree”=3, “Disagree”=4, and “Disagree strongly”=5. This is a measure of the respondent’s **attitude** towards income redistribution.

Both of these are discrete, categorical variables. Sex is binary and attitude is ordinal.

Attitudes towards income redistribution are an example of the broader topic of public opinion on welfare state policies. This is a large topic of classic and current interest in the social sciences, and questions on it have been included in many public opinion surveys.² Of key interest is to explore the how people’s attitudes are associated with their individual characteristics (including such factors as age, sex, education and income) and the contexts in which they live (for example the type of welfare regime adopted in their country). In section 2.4 below we use descriptive statistics to examine such associations between sex and attitude in this sample.

2.3 Single categorical variable

2.3.1 Describing the sample distribution

The term *distribution* is very important in statistics. In this section we consider the distribution of a single variable in the observed data, i.e. its *sample distribution*:

- The **sample distribution** of a variable consists of a list of the values of the variable which occur in a sample, together with the number of times each value occurs.

Later we will discuss other kinds of distributions, such as population, probability and sampling distributions, but they will all be variants of the same concept.

The task of descriptive statistics for a single variable is to summarize the sample distribution or some features of it. This can be done in the form of tables, graphs or single numbers.

¹ESS Round 5: European Social Survey Round 5 Data (2010). Data file edition 2.0. Norwegian Social Science Data Services, Norway Data Archive and distributor of ESS data. The full data can be obtained from <http://ess.nsd.uib.no/ess/round5/>.

²The data were obtained from <http://data.london.gov.uk/datastore/package/london-borough-profiles>.

If you download the “Profiles in Excel” workbook, you will find that one of the pages contains a map of the boroughs, and a tool for visualising the data on that map. A regular map of the boroughs can be found at for example at

<http://www.londoncouncils.gov.uk/londonfacts/londonlocalgovernment/londonmapandlinks/default.htm>.

2.3.2 Tabular methods: Tables of frequencies

When a variable has only a limited number of distinct values, its sample distribution can be summarized directly from the definition given above. In other words, we simply count and display the number of times each of the values appears in the data. One way to do the display is as a table, like the ones for region and the democracy index in the country data, and attitude in the survey example, which are shown in Tables 2.1, 2.2 and 2.3 respectively.

Table 2.1: Frequency distribution of the region variable in the country data.

Region	Frequency	Proportion	%
Africa	48	0.310	31.0
Asia	44	0.284	28.4
Europe	34	0.219	21.9
Latin America	23	0.148	14.8
Northern America	2	0.013	1.3
Oceania	4	0.026	2.6
Total	155	1.000	100.0

Table 2.2: Frequency distribution of the democracy index in the country data.

Democracy score	Frequency	Proportion	%	Cumulative %
0	35	0.226	22.6	22.6
1	12	0.077	7.7	30.3
2	4	0.026	2.6	32.9
3	6	0.039	3.9	36.8
4	5	0.032	3.2	40.0
5	5	0.032	3.2	43.2
6	12	0.077	7.7	50.9
7	13	0.084	8.4	59.3
8	16	0.103	10.3	69.6
9	15	0.097	9.7	79.3
10	32	0.206	20.6	99.9
Total	155	0.999	99.9	

Table 2.3: Frequency distribution of responses to a question on attitude towards income redistribution in the survey example.

Response	Frequency	Proportion	%	Cumulative %
Agree strongly (1)	366	0.156	15.6	15.6
Agree (2)	1090	0.465	46.5	62.1
Neither agree nor disagree (3)	426	0.182	18.2	80.3
Disagree (4)	387	0.165	16.5	96.8
Disagree strongly (5)	75	0.032	3.2	100.0
Total	2344	1.00	100.0	

Each row of such a table corresponds to one possible value of a variable, and the second column shows the number of units with that value in the data. Thus there are 48 countries from Africa and 44 from Asia in the country data set and 32 countries with the highest democracy score 10, and so on. Similarly, 366 respondents in the survey sample strongly agreed with the attitude question, and 75 strongly disagreed with it. These counts are also called **frequencies**, a distribution like this is a **frequency distribution**, and the table is also known as a **frequency table**. The sum of the frequencies, given on the line labelled “Total” in the tables, is the sample size n , here 155 for the country data and 2344 for the survey data.

It is sometimes more convenient to consider relative values of the frequencies instead of the frequencies themselves. The **relative frequency** or **proportion** of a category of a variable is its frequency divided by the sample size. For example, the proportion of countries from Africa in the country data is $48/155 = 0.310$ (rounded to three decimal places). A close relative of the proportion is the **percentage**, which is simply proportion multiplied by a hundred; for example, 31% of the countries in the sample are from Africa. The sum of the proportions is one, and the sum of the percentages is one hundred (because of rounding error, the sum in a reported table may be very slightly different, as it is in Table 2.2).

2.3.3 Graphical methods: Bar charts

Graphical methods of describing data (*statistical graphics*) make use of our ability to process and interpret even very large amounts of visual information. The basic graph for summarising the sample distribution of a discrete variable is a **bar chart**. It is the graphical equivalent of a one-way table of frequencies.

Figures 2.1, 2.2 and 2.3 show the bar charts for region, democracy index and attitude, corresponding to the frequencies in Tables 2.1, 2.2 and 2.3. Each bar corresponds to one category of the variable, and the height of the bar is proportional to the frequency of observations in that category. This visual cue allows us to make quick comparisons between the frequencies of different categories by comparing the heights of the bars.

Some guidelines for drawing bar charts are:

- The heights of the bars may represent frequencies, proportions or percentages. This only changes the units on the vertical axis but not the relative heights of the bars. The shape of the graph will be the same in each case. In Figure 2.1, the units are frequencies, while in Figures 2.2 and 2.3 they are percentages.
- The bars do not touch each other, to highlight the discrete nature of the variable.
- The bars *must* start at zero. If they do not, visual comparisons between their heights are distorted and the graph becomes useless.
- If the variable is ordinal, the bars must be in the natural order of the categories, as in Figures 2.2 and 2.3. If the variable is nominal, the order is arbitrary. Often it makes sense to order the categories from largest (i.e. the one with the largest frequency) to the smallest, possibly leaving any “Others” category last. In Figure 2.1, the frequency ordering would swap Northern America and Oceania, but it seems more natural to keep Northern and Latin America next to each other.

A bar chart is a relatively unexciting statistical graphic in that it does not convey very much visual information. For nominal variables, in particular, the corresponding table is often just as easy to understand and takes less space. For ordinal variables, the bar chart has the additional advantage that its shape shows how the frequencies vary across the ordering of the categories.




Figure 2.1: Bar chart of regions in the country data.




Figure 2.2: Bar chart of the democracy index in the country data.




Figure 2.3: Bar chart of the attitude variable in the survey data example. Agreement with statement: “The government should take measures to reduce differences in income levels”. European Social Survey, Round 5 (2010), UK respondents only.

For example, Figure 2.2 quite effectively conveys the information that the most common values of the democracy index are the extreme scores 0 and 10.

Sometimes you may see graphs which look like bar charts of this kind, but which actually show the values of a single variable for some units rather than frequencies or percentages. For example, a report on the economies of East Asia might show a chart of GDP per capita for Japan, China, South Korea and North Korea, with one bar for each country, and their heights proportional to 28.2, 5.0, 17.8 and 1.3 respectively (c.f. the data in Table 2.14). The basic idea of such graphs is the same as that of standard bar charts. However, they are not particularly useful as descriptive statistics, since they simply display values in the original data without any summarization or simplification.

2.3.4 Simple descriptive statistics

Instead of the whole sample distribution, we may want to summarise only some individual aspects of it, such as its central tendency or variation. Descriptive statistics that are used for this purpose are broadly similar for both discrete and continuous variables, so they will be discussed together for both in Section 2.6.

2.4 Two categorical variables

2.4.1 Two-way contingency tables

The next task we consider is how to describe the sample distributions of two categorical variables together, and in so doing also summarise the association between these variables. The key tool is a table which shows the **crosstabulation** of the frequencies of the variables. This is also known as a **contingency table**. Table 2.4 shows such a table for the respondents’ sex and

attitude in our survey example. We use it to introduce the basic structure and terminology of contingency tables:

Table 2.4: “*The government should take measures to reduce differences in income levels*”: Two-way table of frequencies of respondents in the survey example, by sex and attitude towards income redistribution. Data: European Social Survey, Round 5, 2010, UK respondents only.

Sex	Agree strongly	Agree	Neither agree nor disagree	Disagree	Disagree strongly	Total
Male	160	439	187	200	41	1027
Female	206	651	239	187	34	1317
Total	366	1090	426	387	75	2344

- Because a table like 2.4 summarizes the values of two variables, it is known as a **two-way** contingency table. Similarly, the tables of single variables introduced in Section 2.3.2 are *one-way* tables. It is also possible to construct tables involving more than two variables, i.e. three-way tables, four-way tables, and so on. These are discussed in Chapter 9.
- The variables in a contingency table may ordinal or nominal (including dichotomous). Often an ordinal variable is derived by grouping an originally continuous, interval-level variable, a practice which is discussed further in Section 2.5.
- The horizontal divisions of a table (e.g. the lines corresponding to the two sexes in Table 2.4) are its **rows**, and the vertical divisions (e.g. the survey responses in Table 2.4) are its **columns**.
- The size of a contingency table is stated in terms of the numbers of its rows and columns. For example, Table 2.4 is a 2×5 (pronounced “two-by-five”) table, because it has two rows and five columns. This notation may also be used symbolically, so that we may refer generically to $R \times C$ tables which have some (unspecified) number of R rows and C columns. The smallest two-way table is thus a 2×2 table, where both variables are dichotomous.
- The intersection of a row and a column is a **cell** of the table. The basic two-way contingency table shows in each cell the number (frequency) of units in the data set with the corresponding values of the row variable and the column variable. For example, Table 2.4 shows that there were 160 male respondents who strongly agreed with the statement, and 239 female respondents who neither agreed nor disagreed with it. These frequencies are also known as **cell counts**.
- The row and column labelled “Total” in Table 2.4 are known as the **margins** of the table. They show the frequencies of the values of the row and the column variable separately, summing the frequencies over the categories of the other variable. For example, the table shows that there were overall 1027 ($= 160 + 439 + 187 + 200 + 41$) male respondents, and that overall 75 ($= 41 + 34$) respondents strongly disagreed with the statement. In other words, the margins are *one-way* tables of the frequencies of each of the two variables, so for example the frequencies on the margin for attitude in Table 2.4 are the same as the ones in the one-way table for this variable shown in Table 2.3. The distributions described by the margins are known as the **marginal distributions** of the row and column variables. In contrast, the frequencies in the internal cells of the table, which

show how many units have each possible *combination* of the row and column variables, describe the **joint distribution** of the two variables.

- The number in the bottom right-hand corner of the table is the sum of all of the frequencies, i.e. the total sample size n .

In addition to frequencies, it is often convenient to display proportions or percentages. Dividing the frequencies by the sample size gives overall proportions and (multiplying by a hundred) percentages. This is illustrated in Table 2.5, which shows the overall proportions, obtained by dividing the frequencies in Table 2.4 by $n = 2344$. For example, out of all these respondents, the proportion of 0.102 ($= 239/2344$) were women who neither agreed nor disagreed with the statement. The proportions are also shown for the marginal distributions: for example, 15.6% (i.e. the proportion $0.156 = 366/2344$) of the respondents strongly agreed with the statement. The sum of the proportions over all the cells is 1, as shown in the bottom right corner of the table.

Table 2.5: “*The government should take measures to reduce differences in income levels*”: Two-way table of joint proportions of respondents in the survey example, with each combination of sex and attitude towards income redistribution. Data: European Social Survey, Round 5, 2010, UK respondents only.

Sex	Agree strongly	Agree	Neither agree nor disagree	Disagree	Disagree strongly	Total
Male	0.068	0.187	0.080	0.085	0.017	0.438
Female	0.088	0.278	0.102	0.080	0.015	0.562
Total	0.156	0.465	0.182	0.165	0.032	1.000

2.4.2 Conditional proportions

A two-way contingency table is symmetric in that it does not distinguish between explanatory and response variables. In many applications, however, this distinction is useful for interpretation. In our example, for instance, it is natural to treat sex as the explanatory variable and attitude towards income redistribution as the response response, and so to focus the interpretation on how attitude may depend on sex.

The overall proportions are in such cases not the most relevant quantities for interpretation of a table. Instead, we typically calculate proportions within each category of the row variable or the column variable, i.e. the **conditional proportions** of one variable given the other. The numbers in brackets in Table 2.6 show these proportions calculated for each *row* of Table 2.4 (Table 2.6 also includes the actual frequencies; it is advisable to include them even when conditional proportions are of most interest, to show the numbers on which the proportions are based). In other words, these are the conditional proportions of attitude towards income redistribution given sex, i.e. separately for men and women. For example, the number 0.156 in the top left-hand corner of Table 2.6 is obtained by dividing the number of male respondents who agreed strongly with the statement (160) by the total number of male respondents (1027). Thus 15.6% of the men strongly agreed, and for example 2.6% of women strongly disagreed with the statement. The (1.0) in the last column of the table indicate that the proportions sum to 1 along each row, to remind us that the conditional proportions have been calculated within the rows. The bracketed proportions in the ‘Total’ row are the proportions of the *marginal* distribution of the attitude variable, so they are the same as the proportions in the ‘Total’ row

of Table 2.5.

Table 2.6: “The government should take measures to reduce differences in income levels”: Two-way table of frequencies of respondents in the survey example, by sex and attitude towards income redistribution. The numbers in brackets are proportions within the rows, i.e. conditional proportions of attitude given sex. Data: European Social Survey, Round 5, 2010, UK respondents only.

Sex	Agree strongly	Agree	Neither agree nor disagree	Disagree	Disagree strongly	Total
Male	160 (0.156)	439 (0.428)	187 (0.182)	200 (0.195)	41 (0.040)	1027 (1.0)
Female	206 (0.156)	651 (0.494)	239 (0.182)	187 (0.142)	34 (0.026)	1317 (1.0)
Total	366 (0.156)	1090 (0.465)	426 (0.182)	387 (0.165)	75 (0.032)	2344 (1.0)

We could also have calculated conditional proportions within the *columns*, i.e. for sex given attitude. For example, the proportion $0.563 = 206/366$ of all respondents who strongly agreed with the statement are women. These, however, seem less interesting, because it seems more natural to examine how attitude varies by sex rather than how sex varies by attitude. In general, for any two-way table we can calculate conditional proportions for both the rows and the columns, but typically only one of them is used for interpretation.

2.4.3 Conditional distributions and associations

Suppose that we regard one variable in a two-way table as the explanatory variable (let us denote it by X) and the other variable as the response variable (Y). In our survey example, sex is thus X and attitude is Y . Here the dichotomous X divides the full sample into two groups, identified by the observed value of X — men and women. We may then think of these two groups as two separate samples, and consider statistical quantities separately for each of them. In particular, in Table 2.6 we calculated conditional proportions for Y given X , i.e. for attitude given sex. These proportions describe two distinct sample distributions of Y , one for men and one for women. They are examples of *conditional distributions*:

- The **conditional distribution** of a variable Y given another variable X is the distribution of Y among those units which have a particular value of X .

This concept is not limited to two-way tables but extends also to other kinds of variables and distributions that are discussed later in this coursepack. Both the response variable Y and the explanatory variable X may be continuous as well as discrete, and can have any number of values. In all such cases there is a separate conditional distribution for Y for each possible value of X . A particular one of these distributions is sometimes referred to more explicitly as the conditional distribution of Y given $X = x$, where the “ $X = x$ ” indicates that X is considered at a particular value x (as in “the distribution of Y given $X = 2$ ”, say).

Conditional distributions of one variable given another allow us to define and describe associations between the variables. The informal definition in Section 1.2.4 stated that there is an association between two variables if knowing the value of one of them will help to predict the value of the other. We can now give a more precise definition:

- There is an **association** between variables X and Y if the conditional distribution of Y given X is different for different values of X .

This definition coincides with the more informal one. If the conditional distribution of Y varies with X and if we know X , it is best to predict Y from its conditional distribution given the known value of X . This will indeed work better than predicting Y without using information on X , i.e. from the marginal distribution of Y . Prediction based on the conditional distribution would still be subject to error, because in most cases X does not predict Y perfectly. In other words, the definition of an association considered here is *statistical* (or *probabilistic*) rather than *deterministic*. In our example a deterministic association would mean that there is one response given by all the men and one response (possibly different from the men's) given by all the women. This is of course not the case here nor in most other applications in the social sciences. It is thus crucially important that we have the tools also to analyse statistical associations.

In our example, sex and attitude are associated if men and women differ in their attitudes toward income redistribution. Previous studies suggest that such an association exists, and that it takes the form that women tend to have higher levels of support than men for redistribution.³ As possible explanations for this pattern, both structural reasons (women tend to have lower incomes than men and to rely more on welfare state support) and cultural or psychological ones (women are more likely than men to adopt social values of equality and caring) have been suggested.

2.4.4 Describing an association using conditional proportions

Two variables presented in a contingency table are associated in the sample if the conditional distributions of one of them vary across the values of the other. This is the case in our data set: for example, 4.0% of men but 2.6% of women strongly disagree with the statement. There is thus some association between sex and attitude in this sample. This much is easy to conclude. What requires a little more work is a more detailed description of the pattern and strength of the association, i.e. how and where the conditional distributions differ from each other.

The most general way of summarising associations in a contingency table is by comparing the conditional proportions of the same level of the response given different levels of the explanatory variable. There is no simple formula for how this should be done, so you should use your common sense to present comparisons which give a good summary of the patterns across the table. Unless both variables in the table are dichotomous, several different comparisons may be needed, and may not all display similar patterns. For example, in Table 2.6 the same proportion (0.156, or 15.6%) of both men and women strongly agree with the statement, whereas the proportion who respond “Agree” is higher for women (49.4%) than for men (42.8%).

When the response variable is ordinal, it is often more illuminating to focus on comparisons of *cumulative* proportions which add up conditional proportions over two or more adjacent categories. For instance, the combined proportion of respondents who either strongly agree or agree with the statement is a useful summary of the general level of agreement among the respondents. In our example this is 58.4% ($= 15.5\% + 42.8\%$) for men but 65.0% for women.

A comparison between two proportions may be further distilled into a single number by reporting the *difference* or *ratio* between them. For example, for the proportions of agreeing or strongly agreeing above, the difference is $0.650 - 0.584 = 0.066$, so the proportion is 0.066 (i.e. 6.6 percentage points) higher for women than for men. The ratio of these proportions is $0.650/0.584 = 1.11$, so the proportion for women is 1.11 times the proportion for men (i.e. 11%

³The data can be obtained from <http://www3.norc.ox.ac.uk/gss+website/>, which gives further information on the survey, including the full text of the questionnaires.

higher). Both of these indicate that in this sample women were more likely to agree or strongly agree with the statement than were men. In a particular application we might report a difference or a ratio like this, depending on which of them was considered more relevant or easily understandable. Other summaries are also possible; for example, on MY452 we will discuss a measure called the *odds ratio*, which turns out to be convenient for more general methods of analysing associations involving categorical variables.

The broad conclusion in the example is that there is an association between sex and attitude in these data from the European Social Survey, and that it is of the kind suggested by existing literature. A larger proportion of women than of men indicate agreement with the statement that the government should take measures to reduce income differences, and conversely larger proportion of men disagree with it (e.g. 23.5% of men but only 16.8% of women disagree or strongly disagree). Thus in this sample women do indeed demonstrate somewhat higher levels of support for income redistribution. Whether these differences also warrant a generalisation of the conclusions to people outside the sample is a question which we will take up in Chapters 3 and 4.

2.4.5 A measure of association for ordinal variables

In the previous example the explanatory variable (sex) had 2 categories and the response variable (attitude) had 5. A full examination of the individual conditional distributions of attitude given sex then involved comparisons of five pairs of proportions, one for each level of the attitude variable. This number gets larger still if the explanatory variable also has several levels, as in the following example:

Example: Importance of short-term gains for investors

Information on the behaviour and expectations of individual investors was collected by sending a questionnaire to a sample of customers of a U.S. brokerage house.⁴ One of the questions asked the respondents to state how much importance they placed on quick profits (short-term gains) as an objective when they invested money. The responses were recorded in four categories as “Irrelevant”, “Slightly important”, “Important” or “Very important”. Table 2.7 shows the crosstabulation of this variable with the age of the respondent in four age groups.

Table 2.7: Importance of short-term gains: Frequencies of respondents in the investment example, by age group and attitude towards short-term gains as investment goal. Conditional proportions of attitude given age group are shown in brackets. The value of the γ measure of association is -0.377 .

Age group	Irrelevant	Slightly important	Important	Very important	Total
Under 45	37 (0.253)	45 (0.308)	38 (0.260)	26 (0.178)	146 (1.00)
45–54	111 (0.394)	77 (0.273)	57 (0.202)	37 (0.131)	282 (1.00)
55–64	153 (0.605)	49 (0.194)	31 (0.123)	20 (0.079)	253 (1.00)

⁴ESS Round 5: European Social Survey Round 5 Data (2010). Data file edition 2.0. Norwegian Social Science Data Services, Norway Data Archive and distributor of ESS data. The full data can be obtained from <http://ess.nsd.uib.no/ess/round5/>.

65 and over	193	64	19	15	291
	(0.663)	(0.220)	(0.065)	(0.052)	(1.00)
Total	494	235	145	98	972

Here there are four conditional distributions, one for each age group, and each of them is described by four proportions of different levels of attitude. There are then many possible comparisons of the kind discussed above. For example, we might want to compare the proportions of respondents who consider short-term gains irrelevant between the oldest and the youngest age group, the proportions for whom such gains are very important between these two groups, or, in general, the proportions in any category of the response variable between any two age groups.

Although pairwise comparisons like this are important and informative, they can clearly become cumbersome when the number of possible comparisons is large. A potentially attractive alternative is then to try to summarise the strength of the association between the variables in a single number, a **measure of association** of some kind. There are many such measures for two-way contingency tables, labelled with a range of Greek and Roman letters (e.g. ϕ , λ , γ , ρ , τ , V , Q , U and d). The most useful of them are designed for tables where both of the variables are measured at the ordinal level, as is the case in Table 2.7. The ordering of the categories can then be exploited to capture the strength of the association in a single measure. This is not possible when at least one of the variables is measured at the nominal level, as any attempt to reduce the patterns of the conditional probabilities into one number will then inevitably obscure much of the information in the table. It is better to avoid measures of association defined for nominal variables, and to describe their associations only through comparisons of conditional probabilities as described in the previous section.

Here we will discuss only one measure of association for two-way tables of ordinal variables. It is known as γ (“gamma”). It characterises one possible general pattern of association between two ordinal variables, namely the extent to which high values of one variable tend to be associated with high or low values of the other variable. Here speaking of “low” and “high” values, or of “increasing” or “decreasing” them, is meaningful when the variables are ordinal. For example, in Table 2.7 the categories corresponding to the bottom rows and right-most columns are in an obvious sense “high” values of age and importance respectively.

Consider the conditional proportions of importance given age group shown in Table 2.7. It is clear that, for example, the proportion of respondents for whom short-term gains are very important is highest in the youngest, and lowest in the oldest age group. Similarly, the proportion of respondents for whom such gains are irrelevant increases consistently from the youngest to the oldest group. In other words, respondents with *high* values of the explanatory variable (age group) tend to have *low* values the response variable (importance of short-term gains). Such an association is said to be *negative*. A *positive* association would be seen in a table where high values of one variable were associated with high values of the other.

Measures of association for summarising such patterns are typically based on the numbers of concordant and discordant pairs of observations. Suppose we compare two units classified according to the two variables in the table. These units form a *concordant pair* if one of them has a higher value of both variables than the other. For example, consider two respondents in Table 2.7, one with values (Under 45; Irrelevant) and the other with (45–54; Important). This is a concordant pair, because the second respondent has both a higher value of age group (45–54 vs. Under 45) and a higher value of the importance variable (Important vs. Irrelevant) than the first respondent. In contrast, in a *discordant pair* one unit has a higher value of one variable but a lower value of the other variable than the other unit. For example, a pair of respondents

with values (45–54; Very important) and (55–64; Irrelevant) is discordant, because the latter has a higher value of age group but a lower value of the importance variable than the former. Pairs of units with the same value of one or both of the variables are known as *tied* pairs. They are not used in the calculations discussed below.

The γ measure of association is defined as

$$\gamma = \frac{C - D}{C + D} \quad (2.1)$$

where C is the total number of concordant pairs in the table, and D is the number of discordant pairs. For Table 2.7, the value of this is $\gamma = -0.377$.

Calculation of C and D is straightforward but tedious and uninteresting, and can be left to a computer. Remembering the exact form of (2.1) is also not crucial. More important than the formula of γ (or any other measure of association) is its interpretation. This can be considered on several levels of specificity, which are discussed separately below. The discussion is relatively detailed, as these considerations are relevant and useful not only for γ , but also for all other measures of association in statistics.

The sign of the statistic: It can be seen from (2.1) that γ is positive (greater than zero) when there are more concordant pairs than discordant ones (i.e. $C > D$), and negative when there are more discordant than concordant pairs ($C < D$). This also implies that γ will be positive when the association is positive in the sense discussed above, and negative when the association is negative. A value of $\gamma = 0$ indicates a complete lack of association of this kind. In Table 2.7 we have $\gamma = -0.377$, indicating a negative association. This agrees with the conclusion obtained informally above.

The extreme values of the statistic: Clearly $\gamma = 1$ if there are no discordant pairs ($D = 0$), and $\gamma = -1$ if there are no concordant pairs ($C = 0$). The values $\gamma = -1$ and $\gamma = 1$ are the smallest and largest possible values of γ , and indicate the strongest possible levels of negative and positive association respectively. More generally, the closer γ is to -1 or 1 , the stronger is the (negative or positive) association.

The formal interpretation of the statistic: This refers to any way of interpreting the value more understandably than just vaguely as a measure of “strength of association”. Most often, such an interpretation is expressed as a *proportion* of some kind. For γ , this is done using a principle known as **Proportional reduction of error** (PRE). Because the PRE idea is also used to interpret many other measures of association in statistics, we will first describe it in general terms which are not limited to γ .

Suppose we consider an explanatory variable X and a response variable Y , and want to make predictions of the values of Y in a data set. This is done twice, first in a way which makes no use of X , and then in a way which predicts the value of Y for each unit using information on the corresponding value of X and on the strength and direction of the association between X and Y . Recalling the connection between association and prediction, it is clear that the second approach should result in better predictions if the two variables are associated. The comparison also reflects the *strength* of the association: the stronger it is, the bigger is the improvement in prediction gained by utilising information on X .

A PRE measure describes the size of this improvement. Suppose that the magnitude or number of errors made in predicting the values of Y in a data set using the first scheme, i.e. ignoring information on X , is somehow measured by a single number E_1 , and that E_2 is the same measure of errors for the second prediction scheme which makes use of X . The difference $E_1 - E_2$ is thus

the improvement in prediction achieved by the second scheme over the first. A PRE measure of association is the ratio

$$\text{PRE} = \frac{E_1 - E_2}{E_1}, \quad (2.2)$$

i.e. the improvement in predictions as a *proportion* of the number of errors E_1 under the first scheme. This formulation is convenient for interpretation, because a proportion is easily understandable even if E_1 and E_2 themselves are expressed in some unfamiliar units. The smallest possible value of (2.2) is clearly 0, obtained when $E_2 = E_1$, i.e. when using information on X gives no improvement in predictions. The largest possible value of PRE is 1, obtained when $E_2 = 0$, i.e. when Y can be predicted perfectly from X . The values 0 and 1 indicate no association and perfect association respectively.

The γ statistic is a PRE measure, although with a somewhat convoluted explanation. Suppose that we consider a pair of observations which is known to be either concordant or discordant (the PRE interpretation of γ ignores tied observations). One of the two observations thus has a higher value of X than the other. For example, suppose that we consider two respondents in Table 2.7 from different age groups. We are then asked to predict the *order* of the values of Y , i.e. which of the two units has the higher value of Y . In the example of Table 2.7, this means predicting whether the older respondent places a higher or lower level of importance on short-term gains than the younger respondent. Two sets of predictions are again compared. The first approach makes the prediction at random and with equal probabilities, essentially tossing a coin to guess whether the observation with the higher value of X has the higher or lower value of Y . The second prediction makes use of information on the direction of the association between X and Y . If the association is known to be negative (i.e. there are more discordant than concordant pairs), every pair is predicted to be discordant; if it is positive, every pair is predicted to be concordant. For example, in Table 2.7 the association is negative, so we would always predict that the older of two respondents places a lower value of importance on short-term gains.

If these predictions are repeated for every non-tied pair in the table, the expected number of incorrect predictions under the first scheme is $E_1 = (C + D)/2$. Under the second scheme it is $E_2 = D$ if the association is positive and $E_2 = C$ if it is negative. Substituting these into the general formula (2.2) shows that the γ statistic (2.1) is of the PRE form when γ is positive; when it is negative, the absolute value of γ (i.e. its value with the minus sign omitted) is a PRE measure, and the negative sign of γ indicates that the association is in the negative direction. In our example $\gamma = -0.377$, so age and attitude are negatively associated. Its absolute value 0.377 shows that we will make 37.7% fewer errors if we predict for every non-tied pair that the older respondent places less importance on short-term gains, compared to predictions made by tossing a coin for each pair.

The final property of interest is the **substantive interpretation** of the strength of association indicated by γ for a particular table. For example, should $\gamma = -0.377$ for Table 2.7 be regarded as evidence of weak, moderate or strong negative association between age and attitude? Although this is usually the most (or only) interesting part of the interpretation, it is also the most difficult, and one to which a statistician's response is likely to be a firm "it depends". This is because the strength of associations we may expect to observe depends on the variables under consideration: a γ of 0.5, say, might be commonplace for some types of variables but never observed for others. Considerations of the magnitude of γ are most useful in comparisons of associations between the same two variables in different samples or groups. For example, in Chapter 9 we will calculate γ for the variables in Table 2.7 separately for men and women (see

Table 9.5). These turn out to be very similar, so the strength of the association appears to be roughly similar in these two groups.

Three further observations complete our discussion of γ :

- Since “high” values of a variable were defined as ones towards the bottom and right of a table, reversing the order in which the categories are listed will also reverse the interpretation of “high” and “low” and of a “negative” or “positive” association. Such a reversal for one variable will change the sign of γ but not its absolute value. For example, in Table 2.7 we could have listed the age groups from the oldest to the youngest, in which case we would have obtained $\gamma = 0.377$ instead of $\gamma = -0.377$. Reversing the ordering of both of the variables will give the same value of γ as when neither is reversed. The nature and interpretation of the association remain unchanged in each case.
- γ can also be used when one or both of the variables are dichotomous, but not when either is nominal and has more than two categories. If, for example, the table includes a nominal variable with four categories, there are 24 different and equally acceptable ways of ordering the categories, each giving a different value of γ (or rather 12 different positive values and their negatives). An interpretation of the value obtained for any particular ordering is then entirely meaningless.
- γ can also be treated as an estimate of the corresponding measure of association in a population from which the observed table is a sample. To emphasise this, the symbol $\hat{\gamma}$ is sometimes used for the sample statistic we have discussed here, reserving γ for the population parameter. It is then also possible to define significance tests and confidence intervals for the population γ . These are given, for example, in SPSS output for two-way tables. Here, however, we will not discuss them, but will treat γ purely as a descriptive measure of association. Statistical inference on associations for two-way tables will be considered only in the context of a different test, introduced in Chapter 4.

2.5 Sample distributions of a single continuous variable

2.5.1 Tabular methods

A table of frequencies and proportions or percentages is a concise and easily understandable summary of the sample distribution of a categorical variable or any variable for which only a small number of different values have been observed. On the other hand, applying the same idea to a continuous variable or a discrete variable with many different values is likely to be less useful, because all of the individual frequencies may be small. For example, in this section we illustrate the methods using the GDP variable in the country data introduced at the beginning of Section 2.2. This has 99 different values among the 155 countries, 66 of these values appear only once, and the largest frequency (for 0.8) is five. A frequency table of these values would be entirely unenlightening.

Table 2.8: Frequency distribution of GDP per capita in the country data.

GDP (thousands of dollars)	Frequency	%
less than 2.0	49	31.6
2.0–4.9	32	20.6
5.0–9.9	29	18.7

10.0–19.9	21	13.5
20.0–29.9	19	12.3
30.0 or more	5	3.2
Total	155	99.9

Instead, we can count the frequencies for some *intervals* of values. Table 2.8 shows an example of this for the GDP variable. The frequency on its first line shows that there are 49 countries with GDP per capita of less than \$2000, the second line that there are 32 countries with the GDP per capita between \$2000 and \$4900 (these values included), and so on. We have thus in effect first created an ordinal categorical variable by grouping the original continuous GDP variable, and then drawn a frequency table of the grouped variable in the same way as we do for categorical variables. Some information about the distribution of the original, ungrouped variable will be lost in doing so, in that the exact values of the observations within each interval are obscured. This, however, is a minor loss compared to the benefit of obtaining a useful summary of the main features of the distribution.

The intervals must be *mutually exclusive*, so that no value belongs to more than one interval, and *exhaustive*, so that all values in the data belong to some interval. Otherwise the choice is arbitrary, in that we can choose the intervals in any way which is sensible and informative. Often this is a question of finding the right balance between too few categories (losing too much of the original information) and too many categories (making the table harder to read).

2.5.2 Graphical methods

Histograms

A **histogram** is the graphical version of a frequency table for a grouped variable, like that in Table 2.8. Figure 2.4 shows a histogram for the GDP variable (the histogram consists of the bars; the lines belong to a different graph, the frequency polygon explained below). The basic idea of a histogram is very similar to that of the bar chart, except that now the bars touch each other to emphasise the fact that the original (ungrouped) variable is considered continuous. Because the grouped variable is ordinal, the bars of a histogram must be in the correct order.

A good choice of the grouping intervals of the variable and thus the number of bars in the histogram is important for the usefulness of the graph. If there are too few bars, too much information is obscured; if too many, the shape of the histogram may become confusingly irregular. Often the number of intervals used for a histogram will be larger than what would be sensible for a table like 2.8. Furthermore, intervals like those in Table 2.8 are not even allowed in a histogram, because they are of different widths (of 2, 3, 5, 10 and 10 units for the first five, and unbounded for the last one). The intervals in a histogram must be of equal widths, because otherwise the visual information in it becomes distorted (at least unless the histogram is modified in ways not discussed here). For example, the intervals in Figure 2.4 (less than 2.5, 2.5–less than 5.0, 5.0–less than 7.5 etc.) are all 2.5 units wide. The exact choice can usually be left to computer packages such as SPSS which use automatic rules for choosing sensible intervals.

Frequency polygons

Figure 2.4 also shows a **frequency polygon** of the GDP variable. This is obtained by drawing lines to connect the mid-points of the tops of the bars in a histogram. At each end of the




Figure 2.4: Histogram of GDP per capita in the country data, together with the corresponding frequency polygon.

histogram the lines are further connected to zero, as if the histogram had additional bars of zero height to the left and right of the smallest and largest observed categories. The result is a curve with a similar shape as the corresponding histogram, and its interpretation is similar to that of the histogram.

A histogram is usually preferable to a frequency polygon for presenting a single distribution, especially since histograms are typically much easier to produce in standard software such as SPSS. However, frequency polygons will later be useful for making comparisons between several distributions.

Stem and leaf plots

A **stem and leaf plot** is a close relative of the histogram, and is used for much the same purposes, mostly in small data sets. It is easiest to explain through an example, so let us consider the GDP variable again. The stem and leaf plot for it is shown in Figure 2.9. First, note that the values of the variable in the sample (from \$500 to \$37800, recorded as 0.5 to 37.8 thousands of dollars) have at most three significant digits. If the observations have too many digits to be convenient for a stem and leaf plot, they can be rounded first; for example, if the GDP figures had actually been recorded down to the last dollar, we would have rounded them to the nearest hundred dollars (as in Table 2.14) for the plot. The last digit (here hundreds of dollars) will determine the *leaves* for the plot, while other digits (here round thousands of dollars) will define the *stem*.

Table 2.9: Stem and leaf plot of GDP per capita in the country data (Stem=thousands of dollars, Leaf=hundreds of dollars).

0	5566677778888899
1	0001112233334445566677788899999
2	1122234556799
3	02334579
4	00013567889
5	014588
6	0013334779
7	002466
8	9
9	000159
10	267
11	12448
12	38
13	139
14	
15	7
16	9
17	8
18	0
19	0028
20	0
21	56
22	0
23	247
24	
25	
26	78
27	4667
28	26
29	0168
30	0
31	1
32	7
33	
34	
35	
36	
37	88

The left-hand column in 2.9 lists the stem values in the data, from smallest (0) to the largest (37). Each data value with the same stem is represented on the same line by its leaf, i.e. its last digit. Thus the smallest value, 0.5 for Sierra Leone, is shown as a leaf “5” on the “0” stem, East Timor (another 0.5) as another “5” next to it, and so on up to the largest value 37.8 for Norway, shown as an “8” leaf on the “37” stem.

The stem and leaf plot is very similar to a histogram (try turning Figure 2.9 on its side, and compare to Figure 2.4). It has the additional advantage that it also shows the actual numerical

values of the observations. In some rather special cases this can reveal additional features of the data. Consider, for example, the plot shown in Figure 2.10. The variable here is the number of hours 86 respondents in a social survey (a small subset of all the respondents, drawn purely for this illustration) reported their *spouse* worked in the previous week. An obvious feature of the plot is the prevalence of zeroes as the leaves, especially the many observations with 40 reported hours. This suggests that most respondents probably did not carefully recall and add up the exact hours their spouses worked the previous week; instead, a round “40” is likely to be effectively a synonym for “my spouse has a regular nine-to-five job”. Such *digit preference* is quite common for many variables in surveys, and serves as a reminder that our measurements are not always as precise as they may appear.

Table 2.10: Stem and leaf plot of the reported hours worked last week by the spouses of respondents in a social survey (the data are a sample from data from the U.S. General Survey; observations with less than 12 reported hours have been excluded). The stems and leaves indicate tens of hours and single hours respectively. The main disadvantage of a stem and leaf plot is that since every data value is shown separately, the plot can only be used when the sample size is relatively small. In such cases it is, however, a very useful and user-friendly graph. Also, “small” does not mean “tiny”. For example, the country data set has as many as $n = 155$ observations, yet Figure 2.9 is still quite readable and fits on a single page.

1	55
2	0000000555
3	00002222556889
4	0000000000000000000000000000255556888
5	000000355
6	000000555
7	022

Box plots

A **box plot** differs from the graphs discussed so far in that it does not attempt to display the whole distribution, but only certain characteristics of it. The quantities included in a box plot are some of the summary statistics defined in Section 2.6. To introduce the idea, one box plot is shown in Figure 2.5. The variable considered here is again GDP per capita. The vertical axis shows possible values of the variable, and the plot itself contains the following elements:

- The line inside the central box is the **median** of the variable. Here it is 4.7.
- The end points of the **box** are the **first and third quartile** of the variable, here 1.7 and 11.4 respectively. The length of the box is thus the interquartile range (IQR), here $\text{IQR} = 11.4 - 1.7 = 9.7$. The range of values covered by the box contains the middle 50% of the observations. Half of the countries in this sample have GDPs between \$1700 and \$11400.
- The two lines extending from the box on either side are known as the **whiskers**. Their length is determined as follows:




Figure 2.5: An annotated box plot of GDP per capita in the country data ($n = 155$).

- Calculate the value of 1.5 times the IQR. This is the maximum length of each whisker. Here this is $1.5 \times 9.7 = 14.6$
- The lower whisker extends to the smallest value (**minimum**) of the variable in the sample, or to the smallest value which is at most $1.5 \times \text{IQR}$ units below the first quartile, whichever is larger. Here the minimum is 0.5, which is less than 14.6 units below the first quartile of 1.7, so the lower whisker ends at 0.5.
- The upper whisker extends to the largest value (**maximum**) in the sample, or to the largest value which is at most $1.5 \times \text{IQR}$ units above the third quartile, whichever is smaller. Here the maximum is 37.8, which is further than the maximum distance of 14.6 above the third quartile of 11.4 allowed for a whisker. Thus the upper whisker could be drawn at most to $11.4 + 14.6 = 26$. In this sample there are actually no observations of exactly 26, so the whisker ends at the next smallest observed value, which is 23.7.
- If the minimum is further than $1.5 \times \text{IQR}$ below the first quartile, or maximum further than $1.5 \times \text{IQR}$ above the third quartile, there are still observations which are not in the range spanned by the box and the whiskers. Such extreme observations are considered **outliers** in the plot. The values for each outlier are plotted separately as points. Here there are 15 different outlying values, all with large values of the variable (because in two cases two countries have the same value, these 15 points actually represent 17 countries).

A box plot thus shows some of the main features of a distribution with the following visual cues:

- The central line shows a central value (the median) of the distribution.
- The box shows the location of the central bulk (middle 50%) of the observations
- The whiskers show the range of the regular (non-outlying) observations.
- Very extreme values (outliers), if any, are shown individually.

This can be quite effective for summarizing a distribution. For example, a box plot where the median line is not roughly in the middle of the box, or where one of the whiskers is much longer than the other, indicates that the sample distribution is skewed in the direction of the longer half of the box and the longer whisker. Here the distribution of GDP per capita is clearly positively skewed, as we have already observed. However, for a single distribution all such information and more can also be obtained from a histogram. It is instead for *comparisons* of distributions between two or more samples that box plots are particularly convenient, because it is easy to place two or more of them side by side. This will be illustrated later in Section 7.2.1.

Other graphs for single variables

Other types of graphs that are not described here are also sometimes used for displaying distributions. One of them is a **pie chart**, which shows the proportions of the levels of a categorical (or grouped continuous) variable as sectors of a circle. The relative area of a sector indicates the proportion of the category. We will not discuss pie charts further here, because we do not find them particularly useful (the same information can usually be presented more clearly in a table, bar chart or histogram). That, however, is partly a matter of taste, and there is nothing inherently wrong with (clearly and properly presented) pie charts.

2.6 Numerical descriptive statistics

The tabular and graphical methods discussed above aim to display the whole sample distribution of a variable in an understandable form. The methods introduced in this section have a different purpose. Each of them is used to summarize some important single feature of the distribution in one number. In general, any such number calculated from the data is called a **statistic**. When it is used for data description, it is a **descriptive statistic**, also known as a **summary statistic**. This creates some terminological confusion, as the phrase “descriptive statistics” can mean either all statistical methods used for description or those statistics (i.e. numbers calculated from the data) with a descriptive purpose. The difference is usually unimportant or clear from the context.

The two salient features of a distribution for which we will define descriptive statistics are its *central tendency* and its *variation*.

2.6.1 Measures of central tendency

If you were allowed to know only one feature of the sample distribution of a variable, chances are you would ask for something like its most typical value, the middle value, or the average value — in short, you would be interested in a measure of *central tendency*. We will discuss three such measures below: the mode, the median and the mean (corresponding, respectively, to the phrases “most typical”, “middle” and “average” used above).

The mode

The **mode** is the value of the variable which occurs most often in the data, i.e. the one with the highest frequency. For example, Tables 2.1 and 2.2 show that the mode of the region variable in the country data is “Africa” and the mode of the democracy score is 0. The GDP variable has two modes, 0.8 and 1.9, which both appear five times (a distribution can have several modes; one with two modes is said to be *bimodal*).

The mode can be used for variables of any measurement level. For *nominal* variables it is the only available measure of central tendency, as the median and the mean are not appropriate for such variables.

The mode does not need to be a *central* value in the sense that it can even be the largest or smallest value of the variable, if this occurs most often. This is the case for the democracy index in our example.

The mode is most useful for categorical variables, where the number of possible values is small, and the most common value thus has a high frequency. With continuous variables (like GDP) and discrete variables with many different values, the mode may be unstable and misleading. For example, it is perfectly possible that all but one value appear once each in a sample, and the mode is the value which happens to occur twice.

The median

Suppose that the values of a variable in a sample are first ordered from the smallest to the largest. For example, in Table 2.14 the countries are ordered in this way according to their GDP (starting from the bottom of the table). The **median** is the value which falls in the middle of this ordering, so that it divides the observed values into two halves. Because this

requires a meaningful ordering of the values, the median is appropriate only for ordinal and interval-level variables, but not for nominal ones.

More specifically, suppose that there are n observations, indexed from 1 for the smallest to n for the largest. The index of the middle observation is then $(n + 1)/2$. If n is an odd number, the median is simply the observation in the ordered sample with this index. If n is even, $(n + 1)/2$ falls between two whole numbers, and the median is the mean (of the kind defined below) of the observations with these two indices. For example, in the country data set $n = 155$ (an odd number), and $(n + 1)/2 = 78$, so the median is the value of the 78th observation in the ordered sample; if instead there had been $n = 156$ countries, $(n + 1)/2 = 78.5$, so the median would have been the mean of the 78th and 79th observations.

In the country data set the median of the democracy score is 6, and the median GDP is \$4700 (the 78th observation in GDP order is Paraguay). In practice these are of course found using a computer package like SPSS. For an ordinal categorical variable like the democracy score the median can also be found easily from the frequency table by considering the *cumulative percentages* (or proportions) of the categories. These are obtained by adding up the percentages up to and including each category, as shown in the last column of Table 2.2. The median is then the category in which the cumulative percentage reaches or passes 50%. For the democracy index this happens for the score of 6, which has a cumulative percentage of 50.9%.

The mean

The **mean** is the best-known and most widely used measure of central tendency. It is also known as the **average**. To define the mean, we need to introduce our first pieces of mathematical notation. Suppose first that the variable of interest is denoted by Y . In practice the variable is of course called something else, like GDP or Age or Income, but in the formulas below it is much more convenient to refer to any such variable generically by one letter (note also that the choice of the letter itself is arbitrary; for example, you may often see X used instead of Y when the mean is defined). Individual observations of Y are denoted generically by Y_i , where the subscript i identifies a single subject. The values of i range from 1 to n , so all of the observations in the sample are Y_1, Y_2, \dots, Y_n , e.g. in the country example (with $n = 155$) Y_1, Y_2, \dots, Y_{155} . The ordering of the observations is arbitrary here, so it might for example be the order in which they are listed in your SPSS data file. The mean \bar{Y} (“Y-bar”) of the observations of Y in the sample is defined as

$$\bar{Y} = \frac{\sum Y_i}{n}. \quad (2.3)$$

Here n is again the sample size. The symbol Σ (upper-case Greek letter “Sigma”) is a *summation symbol*, which indicates that we calculate the sum of all Y_i (often this is stated more explicitly by the notation $\sum_i Y_i$ or $\sum_{i=1}^n Y_i$ to make it clear that the summation is over all the values of i). In other words, (2.3) is a concise expression for

$$\bar{Y} = \frac{Y_1 + Y_2 + \dots + Y_n}{n}$$

or, in English, “calculate the sum of all the observations of the variable Y in the sample, and divide this sum by the number of observations to obtain the mean of Y in the sample”. For example, for GDP per capita this calculation gives

$$\bar{Y} = \frac{37.8 + 37.8 + 32.7 + \dots + 0.6 + 0.5 + 0.5}{155} = \frac{1335.1}{155} = 8.6$$

(rounded to one decimal place), i.e. mean GDP among these countries is about \$8600.

Because the mean requires arithmetical calculations (summation and division) on the observations Y_i , it is strictly speaking only appropriate for interval level variables, but not for ordinal ones, for which the numbers of the categories are ordered labels rather than real numbers. However, it is common to see this instruction ignored and means calculated for ordinal variables, especially when they have a large number of categories (see also the discussion under “Measurement Levels” in Section 1.2.2). For example, the mean democracy score in our sample (using the codes 0–10 as its values) is 5.4. This may be used as a summary of the central tendency of the variable, but it should not be overinterpreted as its meaning is not quite as clear as that of, say, mean GDP.

For interval level variables the mean is by far the most commonly used measure of central tendency. It does, however, have one arguably undesirable feature. This is illustrated by the statistics for the GDP variable, as shown in Table 2.11. Its mean (\$8600) is clearly much larger than the median (\$4700). This is due to the shape of the distribution of GDP, as revealed by Figure 2.4 or even more clearly by the stem and leaf plot of Figure 2.9. While most of the countries are concentrated around a fairly narrow range of GDPs, there are also a number of countries with much larger GDPs. The ranges of values in the small and large ends of the values in a distribution are (for fairly obvious visual reasons) called the **tails** of the distribution. A distribution with a (much) longer tail in one end than the other is said to be **skewed**. A distribution like that of GDP in Figure 2.4, with its long tail towards the large values, is said to be **skewed to the right** or **positively skewed**. A distribution shown in panel A of Figure 2.6 is **skewed to the left** (**negatively skewed**): while the examination marks of most students are relatively high, there are some students with very low marks. A distribution which has no clear skewness in either direction, like the distribution of typical weekly working hours in panel B of Figure 2.6 is (approximately) **symmetric**.

Table 2.11: Summary statistics for the three variables in the country data. IQR=interquartile range; s.d.=standard deviation; *: inappropriate for a nominal variable; †: if the democracy index is treated as an interval-level variable.

	Measures	of central	tendency	Measures	of	variation
Variable	Mode	Median	Mean	Range	IQR	s.d.
Region	Africa	*	*	*	*	*
Democracy index	0	6	5.4 [†]	10 [†]	8 [†]	3.9 [†]
GDP per capita	\$800 and \$1900	\$4700	\$8600	\$37300	\$9700	\$9450

The mean is almost always further in the direction of skewness than the median. That is why the mean of the positively skewed GDP variable is larger than its median. In general, a comparison between the two statistics will reveal the direction of any skewness, and give an indication of its magnitude. When the difference is large, as it is here, it is typically sensible to report both the mean and the median.

The mean is sensitive even to individual observations far in the tails of the distribution. Such observations, which are very different (much larger or smaller) from the rest of the data, are known as **outliers**. Even a single outlier can, if it is extreme enough, pull the mean far towards itself, even beyond the range of all the other observations, as in the following example:




Figure 2.6: Examples of a negatively skewed and an approximately symmetric sample distribution. Panel A shows the distribution of examination marks for MY451 (2003; $n = 419$), and B shows the distribution of the number of hours a person usually works in their main job in the 3 per cent Individual Sample of Anonymized Records from the 2001 U.K. Census ($n = 867,016$, respondents with hours 0 or not applicable omitted) Source of the data for panel B: Cathie Marsh Centre for Census and Survey Research, University of Manchester, <http://www.ccsr.ac.uk/sars/>.

Example: A sample with an outlier

Suppose that an M.Sc. student, preparing her dissertation on elements of social capital in Canada, is examining various measures of community activities in a sample of forty municipalities in the province of Manitoba⁵. As part of an initial description of these communities, she wants to summarize their populations, which are

5, 79, 143, 226, 303, 317, 384, 417, 448, 505, 524, 525, 538, 619, 621, 629, 637, 760, 801, 906, 955, 959, 964, 1047, 1111, 1152, 1457, 1491, 1722, 1907, 2079, 2405, 2723, 3950, 4012, 4032, 4183, 4427, 12602, 619544.

The outlier in this case is the city of Winnipeg, whose population of nearly 620,000 is 49 times as large as that of the next largest municipality in the sample. With it included in the sample, the mean population of the 40 municipalities is about 17000; without it, the mean for the other 39 is 1600. The two numbers give rather different pictures of the size of an “average” community in the data (similar differences would probably be observed for other variables too, so the large city would be an outlier in many respects in a study like this). The median, on the other hand, is 906 for the 39 smaller communities, and 930.5 with Winnipeg included. It is thus essentially unaffected by the outlier, basically because it is only influenced by the fact that 619,554 is bigger than the mid-point of the data, but not by how much bigger it is.

2.6.2 Measures of variation

A measure of central tendency is not a complete summary of a distribution, in that there can be distributions which have the same central tendency but which are different in some other respect. To illustrate this with a hypothetical example, suppose we are studying the students in three classrooms of the same grade at a local school. Each class has 14 students, and all students have just taken the same test, graded 1 (low) to 10 (high). The marks of the students are found to be as shown in Table 2.12.

Both the mean and the median of the marks are 6 in every class. However, the classes are otherwise clearly not similar. In particular, the **variation** (or **dispersion**) of the marks is very different. There is no variation at all in Class 1 where everyone has the same score, and quite a lot of variation in Class 3, while Class 2 seems to fall between the two. To capture this, some **measure of variation** will be needed. Three such measures are described here. All of them strictly speaking require the variable to be measured at an interval level, because they involve calculations of differences between its values. Using them on an ordinal variable is thus subject to similar cautions as for the mean above. These measures of variation are entirely inappropriate for nominal-level variables. There are some measures which can be used for such variables, but they are not described here.

Table 2.12: A hypothetical examples of test marks of students in three classes.

Class 1:	6 6 6 6 6 6 6 6 6 6 6 6 6 6
Class 2:	4 4 5 5 5 6 6 6 6 7 7 7 8 8
Class 3:	1 2 2 3 4 4 4 8 8 9 9 10 10 10

⁵Strictly speaking, the analysis should incorporate sampling weights (variable *DWEIGHT*) to adjust for different sampling probabilities for different types of respondents. Here the weights are ignored. Using them would not change the main conclusions for these variables.

Range

The **range** of a variable is simply the difference between its largest and smallest observed values (the **maximum** and **minimum** in statistical terminology). In the class example above,

Class 1: Range = $6 - 6 = 0$

Class 2: Range = $8 - 4 = 4$

Class 3: Range = $10 - 1 = 9$

The measure is largest for Class 3 and smallest for Class 1, so it seems to capture the differences in variation suggested by an initial look at the numbers themselves. For Class 1 the range is 0, because all of the observations are the same. In general, any sensible measure of variation should be zero when there is no variation (all observations are identical), and all of the measures described here have that property.

In the country data, the range of GDP is $\$37800 - \$500 = \$37300$, and the range of the democracy score (if we cautiously treat it as an interval-level variable) is $10 - 0 = 10$.

Interquartile range

The range is often not a particularly useful measure of variation, because it depends *only* on the two extremes of the data. It is thus very sensitive to outliers. If, for example, there is one large outlier, the range will be large even if all of the other observations are very similar to each other.

One way to reduce the effects of outliers is to ignore the tails of the distribution and consider the variation only among the central range of the data. This idea is expressed in the **Interquartile range**. First we have to define the quartiles:

- **The first quartile** is the value such that 25% (one quarter) of the observations are smaller than (or equal to) it, and 75% (three quarters) bigger than (or equal to) it.
- **The third quartile** is the value such that 75% of the observations are smaller than (or equal to) it, and 25% bigger than (or equal to) it.

The quartiles are thus similar in spirit to the median. Just as the median divides the observations into two equal halves (those below and those above the median), the quartiles divide them into two groups at different points. For example, the first quartile divides the observations into the smallest 25% and the remaining largest 75%. (The median can thus also be described as the *second quartile*, and all of these statistics are special cases of a larger class of similar statistics known as *percentiles*.)

The interquartile range (IQR) is the difference between the third and the first quartile. It is the range of the middle 50% of the observations, leaving out the smallest 25% and the largest 25%. This effectively eliminates the effects of any outliers, so IQR is a useful measure of variation (often used together with the median as measure of central tendency) when there are serious outliers or when the distribution is very skewed.

For the class example the interquartile ranges are

Class 1: IQR = $6 - 6 = 0$

Class 2: IQR = $7 - 5 = 2$

Class 3: IQR = $9.25 - 2.75 = 6.5$

These are again in the expected order.⁶ For the country data, the first and third quartiles for GDP are 1.7 and 11.4 respectively, and $IQR=11.4-1.7=9.7$. For the democracy score the quartiles are 1 and 9, and $IQR=8$.

Standard deviation

The most commonly used measure of variation is based on the **deviations**

$$Y_i - \bar{Y}$$

where Y_i again denotes an individual observation of a variable, and \bar{Y} is its mean. A deviation is the difference between an individual observation and the average value in the sample. Table 2.13 shows the deviations for Class 3 in the class example, together with the other calculations discussed below. Here a negative deviation indicates that an observation is smaller than the mean of 6 (e.g. $1 - 6 = -5$), and a positive deviation that an observation is larger than the mean (e.g. $10 - 6 = +4$).

Table 2.13: Calculating the standard deviation of test marks for Class 3 in the class example at the beginning of Section 2.6.2.

Student	Y_i	$Y_i - \bar{Y}$	$(Y_i - \bar{Y})^2$
1	1	-5	25
2	2	-4	16
3	2	-4	16
4	3	-3	9
5	4	-2	4
6	4	-2	4
7	4	-2	4
8	8	+2	4
9	8	+2	4
10	9	+3	9
11	9	+3	9
12	10	+4	16
13	10	+4	16
14 = n	10	+4	16
Sum	$\sum Y_i = 84$	$\sum (Y_i - \bar{Y}) = 0$	$\sum (Y_i - \bar{Y})^2 = 152$
	$\bar{Y} = 84/14 = 6$	$\sum (Y_i - \bar{Y})/n = 0$	$s^2 = 152/13 = 11.69$
			$s = \sqrt{11.69} = 3.4$

The deviations are clearly related to variation, as a sample with little variation will have small deviations (most observations are close to the mean) and one with a lot of variation will have many large deviations (many observations are far from the mean). All that remains is to aggregate them in some sensible way into a single number.

An inappropriate summary of the deviations is their mean, i.e. $\sum (Y_i - \bar{Y})/n$. In the class example this turns out to be zero (see the second column of Table 2.13), and not by coincidence.

⁶The data can be obtained from <http://bes2009-10.org/>, which gives further information on the survey, including the full text of the questionnaires. The data analysed in this class and homework are from the BES Campaign Internet Panel Survey, which has been divided into two data sets corresponding to two time periods leading up to the General Election.

It can be shown that the mean of the deviations is in fact zero for any set of numbers. This happens because positive and negative deviations will always exactly cancel out each other in the sum. This is clearly not what we want, because a negative deviation of, say, -2 (an observation two units below the mean) should be equally strong evidence of variation as a positive deviation of $+2$ (an observation two units above the mean). The signs of the deviations thus need to be eliminated somehow. Just dropping the negative signs (so that -2 becomes 2) means calculating the *absolute values* of the deviations, denoted $|Y_i - \bar{Y}|$. Taking the mean of these gives the **mean absolute deviation** or MAD, defined as

$$\text{MAD} = \frac{\sum |Y_i - \bar{Y}|}{n}.$$

This is a perfectly sensible measure of variation, but it is not very commonly used. This is largely because absolute values are mathematically rather difficult to work with, and this would make MAD very inconvenient for more sophisticated analyses, where measures of variation will also be needed.⁷ Instead, we eliminate the signs of the deviations by using their squares $(Y_i - \bar{Y})^2$, i.e. by multiplying each deviation by itself (c.f. the third column of Table 2.13 for an illustration). These are used to calculate the **variance**, denoted s^2 and defined as

$$s^2 = \frac{\sum (Y_i - \bar{Y})^2}{n - 1}. \quad (2.4)$$

This is (apart from the $n - 1$ rather than n as the divisor) essentially the mean of the squared deviations. Its units of measurement are also squares of the units of the original measurements. For example, the variance of the GDP variable, which is itself measured in (thousands of) dollars, is expressed in dollars squared. This is rather inconvenient for any meaningful interpretation. To obtain a measure of variation expressed in the original units, we can take the square root (indicated below by $\sqrt{}$) of the variance. This statistic is the **standard deviation**, often abbreviated as S.D., denoted by s and defined as

$$s = \sqrt{\frac{\sum (Y_i - \bar{Y})^2}{n - 1}}. \quad (2.5)$$

For the class example, this is 0 for Class 1, 1.3 for Class 2, and 3.4 for class 3. In the country data, the standard deviation of GDP is \$9450 and that of the democracy score (if it is treated as an interval-level variable) is 3.9, as shown in Table 2.11.

Like the mean, the standard deviation is sensitive to outliers and skewness of the distribution, so sometimes other measures of variation (e.g. IQR or MAD) should be reported instead of, or in addition to it. Nevertheless, the standard deviation is by far the most commonly used measure of variation. One reason for this is that it is very important not just as a descriptive statistic but also as an element in several forms of statistical inference. For description it is typically less immediately interpretable than measures of central tendency. Often the most revealing descriptive uses of the standard deviation are in comparisons between samples, like in the class example above. The following is a real example of this kind, where variation was in fact of more interest than central tendency:

Example: Variation in rates of economic growth

In an article titled “Dancing in step” on November 13th 2004, *The Economist* discussed a set of data (collected by the J. P. Morgan Chase bank) on the annual growth rates (in percentage points) of the Gross Domestic Products (GDP) of 30 countries for each year since 1971. Measures

⁷Official results obtained from www.olympic.org/london-2012-summer-olympics.

of central tendency, such as average growth rates for each country and each year, are clearly interesting in this case. However, most of the discussion in the article concerned *variation* in growth rates, measured by their standard deviation across countries for each year, and especially changes in this variation over time. The standard deviation of growth rates was around 3–5 percentage points for every year until the early 1990s, had fallen to about 2 percentage points in 2003, and was forecast to decline further in subsequent years. There had thus previously been a fair amount of variation in rates of economic growth (with some economies growing faster and some slower, some perhaps being in recession), whereas recently the growth rates had become more similar across countries. The article summarized this in its subtitle as “The world’s economies are more synchronised than ever before”, and went on to discuss the implications of this development for global economy.

The formula (2.5) for the standard deviation involves the divisor $n - 1$, where the discussion leading up to it might make you expect n instead. The reasons for this will be discussed briefly in Section 6.2.1. The definition is not entirely consistent in that some textbooks do use n instead of $n - 1$. The difference is of no great importance, and using either n or $n - 1$ would be fine for our purposes. Whenever n is even moderately large, the difference between n and $n - 1$ is in any case small, and both definitions of standard deviation give very similar values.

Finally, measures of central tendency and measures of variation, even though they summarise the two most important features of a sample distribution of a variable, may still miss some important features of the distribution. Consider, for example, the class marks in Classes 2 and 3 in our hypothetical example. These are summarized by the bar charts of Figure 2.7. The distribution for Class 2 is symmetric and concentrated around the mean value of 6. The most noticeable feature of the marks in Class 3, on the other hand, is that there appear to be two distinct groups of students, one with very low scores and one with high scores. A similar feature was also noted in the distribution of the democracy index in the country data (c.f. Figure 2.2). This property would not be revealed by measures of central tendency or variation, so it is an illustration of why it is always sensible to also examine the whole distribution of a variable using frequency tables or graphical methods.



Figure 2.7: Bar charts of the test marks in the class example at the beginning of Section 2.6.2 for Classes 2 (on the left) and 3 (on the right).

2.7 Associations which involve continuous variables

Bivariate descriptive methods which are designed for situations where at least one of the two variables is continuous are not described here but in later sections:

- Explanatory variable is categorical and response variable continuous: Parallel histograms, frequency polygons and box plots (Section 7.2).
- Both explanatory and response variables are continuous: Scatter plots and line plots (Section 8.2.2).

We do not discuss the remaining possibility, where the explanatory variable is continuous and the response is categorical. The simplest and usually quite sufficient way to give an initial description of the associations in this case is to group the explanatory variable into a categorical variable and then apply the methods of Section 2.4.

2.8 Presentation of tables and graphs

The purpose of statistical tables and graphs is to communicate information correctly, clearly and effectively. If they do not do that, that is, if they leave the reader misled, confused or uninformed, they have failed and should not have been shown at all. Creating good tables and graphics is not only a matter of understanding the technical details described above. It also involves general principles of design and presentation. Most of these should be simple common sense but clearly are not, judging by the many entirely unhelpful tables and graphs appearing in all kinds of publications. This section discusses very briefly some principles of good practice in presenting descriptive statistics in tables and graphs. Much of the section is based on two books, *The Visual Display of Quantitative Information* by Edward R. Tufte (Graphics Press, 1983) and *Visual Revelations* by Howard Wainer (Copernicus, 1997). These can be consulted for further information and examples of both good and bad practice.

First, a reader of a table or graph should be able to understand what it is about:

- The variables should be labelled clearly. In particular, the names used in computer data files should not be used unless they are also understandable words. So even if a variable is called ATTDFOXH in your SPSS file, it should still be labelled “Attitude to foxhunting” or something similar in presentation. Similarly, the categories of variables should be labelled in words wherever appropriate.
- Items such as the columns of a table or the vertical axis of a bar chart should also be labelled clearly (e.g. whether they are for frequencies or percentages).
- More generally, a table or figure and its caption should be (within reason) as self-contained as possible, in that the reader should be able to understand them with little reference to the rest of the text for explanation (remember that tables and figures often float, i.e. they may appear on a different page from where they are referred to in the main text). This may also include giving the source of the data in a note or caption to the table or figure.

Some guidelines for constructing tables are

- A table produced by software such as SPSS, although it contains the necessary numbers, is rarely suitable for presentation directly. Tables included in research reports should be retyped and reformatted.
- The categories of the variable should be in a sensible order. For ordinal variables (including those obtained by grouping a continuous one), this should obviously be the natural

ordering of the categories. For a nominal variable, the order can be chosen in whichever way is most useful for presentation. Often it makes sense to order categories from the largest to the smallest, typically leaving any “Others” category last.

- If only proportions or percentages are shown, the sample size n should also be reported, perhaps in a note or caption to the table. This will allow the reader to judge how informative the table is. A percentage of 20% is clearly richer information when it corresponds to a frequency of 2,000 in a sample 10,000 than when it means 2 out of 10 observations. When n is very small, proportions and percentages should be avoided altogether: reporting 1 out of 7 as 14.3% is simply nonsensical.
- Proportions and percentages can and should be rounded. It is rarely necessary to see percentages with more than one decimal place, if even that.

With graphs, it is always useful to bear in mind Wainer’s principle:

The aim of good data graphics is to display data accurately and clearly

The way to produce *bad* graphs is thus to break some part of this, for example by (1) not showing much data, (2) showing much that is not data, (3) showing the data inaccurately, or (4) obscuring the data. Graphs with these characteristics are a form of visual lying, distorting the graphical cues in a plot in ways which make it difficult or impossible to obtain accurate information from it.

One example of a lying graph already mentioned is the “cut” bar chart where the bars do not begin at zero. Another is the pseudo third dimension, an example of which is shown in Figure 2.8. The information presented in this graph is the same as that of Figure 2.1, i.e. frequencies of different regions. These are represented by the heights of the bars. The additional information conveyed by the apparent thickness of the bars, represented in perspective to give an illusion of three-dimensional bars, is then — exactly nothing. The fake third dimension represents no data, and serves only to distort the real data that are being shown.

We can thus give a simple instruction: using a fake third dimension like the one in Figure 2.8 is always wrong and not acceptable under any circumstances. This is true irrespective of the fact that such graphs are often seen and easily (often almost automatically) produced by software packages like Microsoft Excel. All this proves is that the programmers of those packages have little graphical sense, or perhaps that their companies have discovered that their customers are willing to pay for such “features” as colourful but pointless graphs. Indeed, many if not most of the graph styles provided by, say, Excel (exploding pie charts, doughnuts, cones, pyramids and so on) are entirely useless for accurate presentation of data.

An objection sometimes offered to such a severe rule is that bad graphs “look good”. This can be answered in two ways. First, a statistical graphic is not a decoration, but a tool for presenting information. Authors who confuse the two often end up displaying pretty colours and strange shapes to hide the lack of actual information in a graph. Second, even in an aesthetic sense a useful and accurate graph is preferable to a bad one, in the way that any well-designed object with a function tends to be more attractive than a badly designed one.

What, then, is the recipe for good graphics? Mostly this is just a matter of using basic graph types in a sensible and restrained manner, focusing on presenting information and avoiding all distracting decoration. Some such examples have been given earlier in this chapter. Other types of graphs are used to illustrate associations between variables, which we have not yet discussed. To anticipate that a little, Figure 2.9 shows one (good but not in any way exceptional) example of such graphs. It is a reproduction of a graph originally published in a survey of Spain in *The Economist*, and shows changes in average house prices in Spain, Germany and Britain between




Figure 2.8: An example of an unacceptable graph: a bar chart with a pseudo three-dimensional effect. The data are the same as in Figure 2.1.

1993 and 2003. Even without an introductory statistics course, the main message of Figure 2.9 is immediately clear: increases in Spanish house prices over the period have been comparable to those in Britain, with prices more than doubling in both countries, and very unlike those in Germany, where the prices have remained unchanged. Note also that the graph distinguishes between the lines for different countries by using different types of line. Different colours can of course be used instead, but their differences will become obscured if the graph is photocopied or printed in black and white.

In addition to such modest but sensible and useful basic graphs, you may sometimes encounter inspired examples of special graphs which manage to describe particular data sets in exceptionally vivid and informative ways. Some such examples are shown at www.datavis.ca/gallery/index.php, on the web page maintained by Michael Friendly at York University in Canada (unfortunately, however, the electronic images do not always do justice to the originals; crisper versions can be found in the books mentioned above). For example, the page shows what Edward Tufte has described as possibly “the best statistical graphic ever drawn”. This is Charles Joseph Minard’s graphical memorial, drawn in 1861, to the fate of Napoleon I’s army in their invasion of Russia in 1812. For contrast, the page also shows a number of examples of visual lying and other terrible graphs, including a mould-breaking re-interpretation of the idea of a pie chart by Fox News, and a colourful effort that Tufte has called possibly “the worst graphic ever to find its way into print”. Clearly not all pictures tell us as much as a thousand words.

2.9 Appendix: Country data

The data used for illustration throughout this chapter are given in Table 2.14. The variables are defined as follows:

- **region** indicates the macro region where the country is located, coded as 1=Africa, 2=Asia, 3=Europe, 4=Latin America, 5=Northern America, 6=Oceania. The list of regions and the assignment of countries to regions are those used by the UN Statistics



Figure 2.9: An example of an informative graph: house prices in three countries between 1993 and 2003, indexed to 100 in 1993. Source: *The Economist*, June 26th, 2004. The numbers were estimated from the graph in the magazine, so they are approximate.

Division (see unstats.un.org/unsd/methods/m49/m49.htm).

- **democracy** is a measure of institutionalised democracy by the Polity IV project⁸. The values refer to each country's classification in 2002. The variable has an 11-point scale from 0 (lowest level of democracy) to 10 (highest). Countries coded as being in the state of "interruption" or "interregnum" have been omitted.
- **GDP** is the country's Gross Domestic Product per capita (in thousands of U.S. dollars), adjusted for purchasing power parity. The data were obtained from CIA's *The World Factbook 2004* (<https://www.cia.gov/library/publications/resources/the-world-factbook/>). The figures refer to slightly different years for different countries.

The data set contains those 155 countries for which recent data on all of the three variables were available at the time the example created.

Table 2.14:

Country	R	D	GDP	Country	R	D	GDP	Country	R	D	GDP
Norway	3	10	37.8	Bulgaria	3	9	7.6	Pakistan	2	0	2.1
USA	5	10	37.8	Thailand	2	9	7.4	Angola	1	1	1.9
Switzerland	3	10	32.7	Namibia	1	6	7.2	Bangladesh	2	6	1.9
Denmark	3	10	31.1	Iran	2	4	7.0	Cambodia	2	3	1.9
Austria	3	10	30.0	Romania	3	8	7.0	Sudan	1	0	1.9
Canada	5	10	29.8	Tunisia	1	1	6.9	Zimbabwe	1	0	1.9
Ireland	3	10	29.6	Macedonia	3	9	6.7	Burma	2	0	1.8
Belgium	3	10	29.1	Turkey	2	8	6.7	Cameroon	1	1	1.8
Australia	6	10	29.0	Libya	1	0	6.4	Mauritania	1	0	1.8
Netherlands	3	10	28.6	Colombia	4	7	6.3	Moldova	3	8	1.8
Japan	2	10	28.2	Kazakhstan	2	0	6.3	Mongolia	2	10	1.8

⁸The data can be obtained from www3.norc.ox.ac.uk/GSS+Website/, which gives further information on the survey, including the full text of the questionnaires.

Country	R	D	GDP	Country	R	D	GDP	Country	R	D	GDP
UK	3	10	27.7	Panama	4	9	6.3	Laos	2	0	1.7
France	3	9	27.6	Belarus	3	0	6.1	Gambia	1	0	1.7
Germany	3	10	27.6	Algeria	1	1	6.0	Uzbekistan	2	0	1.7
Finland	3	10	27.4	Dominican R.	4	8	6.0	Haiti	4	1	1.6
Sweden	3	10	26.8	Fiji	6	6	5.8	Kyrgyzstan	2	1	1.6
Italy	3	10	26.7	Turkmenistan	2	0	5.8	Senegal	1	8	1.6
Singapore	2	2	23.7	Gabon	1	0	5.5	Iraq	2	0	1.5
Taiwan	2	9	23.4	Ukraine	3	7	5.4	Togo	1	1	1.5
UAE	2	0	23.2	Peru	4	9	5.1	Cote d'Ivoire	1	5	1.4
Spain	3	10	22.0	China	2	0	5.0	Nepal	2	1	1.4
NZ	6	10	21.6	Swaziland	1	0	4.9	Uganda	1	0	1.4
Qatar	2	0	21.5	El Salvador	4	7	4.8	Bhutan	2	0	1.3
Greece	3	10	20.0	Venezuela	4	6	4.8	Djibouti	1	3	1.3
Israel	2	10	19.8	Paraguay	4	7	4.7	N. Korea	2	0	1.3
Cyprus	2	10	19.2	Philippines	2	8	4.6	Rwanda	1	0	1.3
Kuwait	2	0	19.0	Albania	3	7	4.5	Chad	1	1	1.2
Slovenia	3	10	19.0	Jordan	2	2	4.3	Mozambique	1	6	1.2
Portugal	3	10	18.0	Guatemala	4	8	4.1	Benin	1	6	1.1
S. Korea	2	8	17.8	Egypt	1	0	4.0	Burkina Faso	1	2	1.1
Bahrain	2	0	16.9	Guyana	4	6	4.0	C. Afr. R.	1	5	1.1
Czech R.	3	10	15.7	Morocco	1	0	4.0	Kenya	1	8	1.0
Hungary	3	10	13.9	Jamaica	4	9	3.9	Liberia	1	3	1.0
Slovakia	3	9	13.3	Sri Lanka	2	7	3.7	Tajikistan	2	2	1.0
Oman	2	0	13.1	Armenia	2	6	3.5	Mali	1	6	.9
Uruguay	4	10	12.8	Azerbaijan	2	0	3.4	Nigeria	1	4	.9
Estonia	3	7	12.3	Ecuador	4	6	3.3	Guinea-Bissau	1	5	.8
Saudi Ar.	2	0	11.8	Syria	2	0	3.3	Madagascar	1	7	.8
Lithuania	3	10	11.4	Indonesia	2	8	3.2	Niger	1	4	.8
Mauritius	1	10	11.4	Lesotho	1	8	3.0	Yemen	2	1	.8
Argentina	4	8	11.2	Cuba	4	0	2.9	Zambia	1	3	.8
Poland	3	9	11.1	India	2	9	2.9	Comoros	1	4	.7
S. Africa	1	9	10.7	Equatorial G.	1	0	2.7	Eritrea	1	0	.7
Croatia	3	7	10.6	Honduras	4	7	2.6	Ethiopia	1	3	.7
Latvia	3	8	10.2	Georgia	2	5	2.5	Congo (Br.)	1	0	.7
Trinidad	4	10	9.5	Vietnam	2	0	2.5	Burundi	1	1	.6
Costa Rica	4	10	9.1	Bolivia	4	9	2.4	Malawi	1	6	.6
Botswana	1	9	9.0	Nicaragua	4	8	2.3	Tanzania	1	3	.6
Malaysia	2	4	9.0	Ghana	1	7	2.2	East Timor	2	6	.5
Mexico	4	8	9.0	PNG	6	10	2.2	Sierra Leone	1	5	.5
Russia	3	7	8.9	Serbia	3	7	2.2				
Brazil	4	8	7.6	Guinea	1	1	2.1				

Chapter 3

Samples and populations

3.1 Introduction

So far we have discussed statistical description, which is concerned with summarizing features of a sample of observed data. From now on, most of the attention will be on statistical inference. As noted in Section 1.2.3, the purpose of inference is to draw conclusions about the characteristics of some larger population based on what is observed in a sample. In this chapter we will first give more careful definitions of the concepts of populations and samples, and of the connections between them. In Section 3.5 we then consider the idea of a population distribution, which is the target of statistical inference. The discussion of statistical inference will continue in Chapters 4–7 where we gradually introduce the basic elements of inference in the contexts of different types of analyses.

3.2 Finite populations

In many cases the population of interest is a particular group of real people or other units. Consider, for example, the European Social Survey (ESS) which we used in Chapter 2 (see early in Section 2.2)¹. The ESS is a cross-national survey carried out biennially in around 30 European countries. It is an academically-driven social survey which is designed to measure a wide range attitudes, beliefs and behaviour patterns among the European population, especially for purposes for cross-national comparisons.

The target population of ESS is explicitly stated as being “all persons aged 15 and over resident within private households, regardless of their nationality, citizenship, language or legal status” in each of the participating countries. This is, once “private household” has been defined carefully, and notwithstanding the inevitable ambiguity in that the precise number and composition of households are constantly changing, a well-defined, existing group. It is also a large group: in the UK, for example, there are around 50 million such people. Nevertheless, we have no conceptual difficulty with imagining this collection of individuals. We will call any such population a *finite population*.

The main problem with studying a large finite population is that it is usually not feasible to collect data on all of its members. A **census** is a study where some variables *are* in fact measured for the entire population. The best-known example is the Census of Population, which

¹ESS Round 5: European Social Survey Round 5 Data (2010). Data file edition 2.0. Norwegian Social Science Data Services, Norway Data Archive and distributor of ESS data. The full data can be obtained from <http://ess.nsd.uib.no/ess/round5/>.

at least aims to be a complete evaluation of all persons living in a country on a particular date with respect to basic demographic data. Similarly, we have the Census of Production, Census of Distribution etc. For most research, however, a census is not feasible. Even when one is attempted, it is rarely truly comprehensive. For example, all population censuses which involve collecting the data from the people themselves end up missing a substantial (and non-random) proportion of the population. For most purposes a well-executed sample of the kind described below is actually preferable to an unsuccessful census.

3.3 Samples from finite populations

When a census is not possible, information on the population is obtained by observing only a subset of units from it, i.e. a sample. This is meant to be *representative* of the population, so that we can *generalise* findings from the sample to the population. To be representative in a sense appropriate for statistical inference, a sample from a finite population must be a *probability sample*, obtained using

- **probability sampling:** a sampling method where every unit in the population has a **known, non-zero** probability of being selected to the sample.

Probability sampling requires first a **sampling frame**, essentially one or more lists of units or collections of units which make it possible to select and contact members of the sample. For example, the first stage of sampling for many UK surveys uses the Postcode Address File, a list of postal addresses in the country. A **sampling design** is then created in such a way that it assigns a **sampling probability** for each unit, and the sample is drawn so that each unit's probability of being selected into the sample is given by their sampling probability. The selection of the specific set of units actually included in the sample thus involves *randomness*, usually implemented with the help of random number generators on computers.

The simplest form of probability sampling is

- **simple random sampling**, where every unit in the population has the *same* probability of selection.

This requirement of equal selection probabilities is by no means essential. Other probability sampling methods which relax it include

- **stratified sampling**, where the selection probabilities are set separately for different groups (*strata*) in the population, for example separately for men and women, different ethnic groups or people living in different regions.
- **cluster sampling**, where the units of interest are not sampled individually but in groups (*clusters*). For example, a school survey might involve sampling entire classes and then interviewing every pupil in each selected class.
- **multistage sampling**, which employs a sequence of steps, often with a combination of stratification, clustering and simple random sampling. For example, many social surveys use a *multistage area sampling* design which begins with one or more stages of sampling areas, then households (addresses) within selected small areas, and finally individuals within selected households.

These more complex sampling methods are in fact used for most large-scale social surveys to improve their accuracy and/or cost-efficiency compared to simple random sampling. For example, the UK component of the European Social Survey uses a design of three stages: (1) a stratified sample of postcode sectors, stratified by region, level of deprivation, percentage

of privately rented households, and percentage of pensioners; (2) simple random sample of addresses within the selected sectors; and (3) simple random sample of one adult from each selected address.

Some analyses of such data require the use of *survey weights* to adjust for the fact that some units were more likely than others to end up in the sample. The questions of how and when the weights should be used are, however, beyond the scope of this course. Here we will omit the weights even in examples where they might normally be used.²

Not all sampling methods satisfy the requirements of probability sampling. Such techniques of **non-probability sampling** include

- *purposive sampling*, where the investigator uses his or her own “expert” judgement to select units considered to be representative of the population. It is very difficult to do this well, and very easy to introduce conscious or unconscious biases into the selection. In general, it is better to leave the task to the random processes of probability sampling.
- *haphazard* or *convenience* sampling, as when a researcher simply uses the first n passers-by who happen to be available and willing to answer questions. One version of this is *volunteer* sampling, familiar from call-in “polls” carried out by morning television shows and newspapers on various topics of current interest. All we learn from such exercises are the opinions of those readers or viewers who felt strongly enough about the issue to send in their response, but these tell us essentially nothing about the average attitudes of the general population.
- *quota sampling*, where interviewers are required to select a certain number (quota) of respondents in each of a set of categories (defined, for example, by sex, age group and income group). The selection of specific respondents within each group is left to the interviewer, and is usually done using some (unstated) form of purposive or convenience sampling. Quota sampling is quite common, especially in market research, and can sometimes give reasonable results. However, it is easy to introduce biases in the selection stage, and almost impossible to know whether the resulting sample is a representative one.

A famous example of the dangers of non-probability sampling is the survey by the *Literary Digest* magazine to predict the results of the 1936 U.S. presidential election. The magazine sent out about 10 million questionnaires on post cards to potential respondents, and based its conclusions on those that were returned. This introduced biases in at least two ways. First, the list of those who were sent the questionnaire was based on registers such as the subscribers to the magazine, and of people with telephones, cars and various club memberships. In 1936 these were mainly wealthier people who were more likely to be Republican voters, and the typically poorer people not on the source lists had no chance of being included. Second, only about 25% of the questionnaires were actually returned, effectively rendering the sample into a volunteer sample. The magazine predicted that the Republican candidate Alf Landon would receive 57% of the vote, when in fact his Democratic opponent F. D. Roosevelt gained an overwhelming victory with 62% of the vote. The outcome of the election was predicted correctly by a much smaller probability sample collected by George Gallup.

A more recent example is the “GM Nation” public consultation exercise on attitudes to genetically modified (GM) agricultural products, carried out in the U.K. in 2002–3³. This involved

²The data were obtained from <http://data.london.gov.uk/datastore/package/london-borough-profiles>.

If you download the “Profiles in Excel” workbook, you will find that one of the pages contains a map of the boroughs, and a tool for visualising the data on that map. A regular map of the boroughs can be found at for example at

<http://www.londoncouncils.gov.uk/londonfacts/londonlocalgovernment/londonmapandlinks/default.htm>.

³The data can be obtained from <http://www3.norc.org/gss+website/>, which gives further information on

various activities, including national, regional and local events where interested members of the public were invited to take part in discussions on GM foods. At all such events the participants also completed a questionnaire, which was also available on the GM Nation website. In all, around 37000 people completed the questionnaire, and around 90% of those expressed opposition to GM foods. While the authors of the final report of the consultation drew some attention to the unrepresentative nature of this sample, this fact had certainly been lost by the time the results were reported in the national newspapers as “5 to 1 against GM crops in biggest ever public survey”. At the same time, probability samples suggested that the British public is actually about evenly split between supporters and opponents of GM foods.

3.4 Conceptual and infinite populations

Even a cursory inspection of academic journals in the social sciences will reveal that a finite population of the kind discussed above is not always clearly defined, nor is there often any reference to probability sampling. Instead, the study designs may for example resemble the following two examples:

Example: A psychological experiment

Fifty-nine undergraduate students from a large U.S. university took part in a psychological experiment⁴, either as part of a class project or for extra credit on a psychology course. The participants were randomly assigned to listen to one of two songs, one with clearly violent lyrics and one with no violent content. One of the variables of interest was a measure (from a 35-item attitude scale) of state hostility (i.e. temporary hostile feelings), obtained after the participants had listened to a song, and the researchers were interested in comparing levels of hostility between the two groups.

Example: Voting in a congressional election

A political-science article⁵ considered the U.S. congressional election which took place between June 1862 and November 1863, i.e. during a crucial period in the American Civil War. The units of analysis were the districts in the House of Representatives. One part of the analysis examined whether the likelihood of the candidate of the Republican Party (the party of the sitting president Abraham Lincoln) being elected from a district was associated with such explanatory variables as whether the Republican was the incumbent, a measure of the quality of the other main candidate, number of military casualties for the district, and the timing of the election in the district (especially in relation to the Union armies’ changing fortunes over the period).

There is no reference here to the kinds of finite populations and probability samples discussed Sections 3.2 and 3.3. In the experiment, the participants were a convenience sample of respondents easily available to the researcher, while in the election study the units represent (nearly) all the districts in a single (and historically unique) election. Yet both articles contain plenty of statistical inference, so the language and concepts of samples and populations are clearly being used. How is this to be justified?

In the example of the psychological experiment the subjects will clearly not be representative of a general (non-student) population in many respects, e.g. in age and education level. However,

the survey, including the full text of the questionnaires.

⁴ESS Round 5: European Social Survey Round 5 Data (2010). Data file edition 2.0. Norwegian Social Science Data Services, Norway Data Archive and distributor of ESS data. The full data can be obtained from <http://ess.nsd.uib.no/ess/round5/>.

⁵Strictly speaking, the analysis should incorporate sampling weights (variable *DWEIGHT*) to adjust for different sampling probabilities for different types of respondents. Here the weights are ignored. Using them would not change the main conclusions for these variables.

it is not really such characteristics that the study is concerned with, nor is the population of interest really a population of people. Instead, the implicit “population” being considered is that of possible values of level of hostility after a person has listened to one of the songs in the experiment. In this extended framework, these possible values include not just the levels of hostility possibly obtained for different people, but also those that a single person might have after listening to the song at different times or in different moods etc. The generalisation from the observed data in the experiment is to this hypothetical population of possible reactions.

In the political science example the population is also a hypothetical one, namely those election results that *could* have been obtained if something had happened differently, i.e. if different people turned up to vote, if some voters had made different decisions, and so on (or if we considered a different election in the same conditions, although that is less realistic in this example, since other elections have not taken place in the middle of a civil war). In other words, votes that actually took place are treated as a sample from the population of votes that could conceivably have taken place.

In both cases the “population” is in some sense a hypothetical or conceptual one, a population of possible realisations of events, and the data actually observed are a sample from that population. Sometimes it is useful to apply similar thinking even to samples from ostensibly quite finite populations. Any such population, say the residents of a country, is exactly fixed at one moment only, and was and will be slightly different at any other time, or would be even now if any one of a myriad of small events had happened slightly differently in the past. We could thus view the finite population itself at a single moment as a sample from a conceptual population of possible realisations. This is known in survey literature as a *superpopulation*. The data actually observed are then also a sample from the superpopulation. With this extension, it is possible to regard almost any set of data as a sample from some conceptual superpopulation.

The highly hypothetical notion of a conceptual population of possible events is clearly going to be less easy both to justify and to understand than the concept of a large but finite population of real subjects defined in Section 3.2. If you find the whole idea distracting, you can focus in your mind on the more understandable latter case, at least if you are willing to believe that the idea of a conceptual population is also meaningful. Its main justification is that much of the time it works, in the sense that useful decision rules and methods of analysis are obtained based on the idea. Most of the motivation and ideas of statistical inference are essentially the same for both kinds of populations.

Even when the idea of a conceptual population is invoked, questions of representativeness of and generalisability to real, finite populations will still need to be kept in mind in most applications. For example, the assumption behind the psychological experiment described above is that the findings about how hearing a violent song affects levels of hostility are generalisable to some larger population, beyond the 59 participants in the experiment and beyond the body of students in a particular university. This may well be the case at least to some extent, but it is still open to questioning. For this reason findings from studies like this only become really convincing when they are *replicated* in comparable experiments among different kinds of participants.

Because the kinds of populations discussed in this section are hypothetical, there is no sense of them having a particular fixed number of members. Instead, they are considered to be *infinite* in size. This also implies (although it may not be obvious why) that we can essentially always treat samples from such populations as if they were obtained using simple random sampling.

3.5 Population distributions

We will introduce the idea of a population distribution first for finite populations, before extending it to infinite ones. The discussion in this section focuses on categorical variables, because the concepts are easiest to explain in that context; generalisations to continuous variables are discussed in Chapter 7.

Suppose that we have drawn a sample of n units from a finite population and determined the values of some variables for them. The units that are not in the sample also possess values of the variables, even though these are not observed. We can thus easily imagine how any of the methods which were in Chapter 2 used to describe a sample could also be applied in the same way to the whole population, if only we knew all the values in it. In particular, we can, paralleling the sample distribution of a variable, define the **population distribution** as the set of values of the variable which appear in the population, together with the frequencies of each value.

For illustration, consider again the example introduced early in Section 2.2. The two variables there are a person's sex and his or her attitude toward income redistribution. We have observed them for a sample $n = 2344$ people drawn from the population of all UK residents aged 15 or over. The sample distributions are summarised by Table 2.3.

Table 3.1: “*The government should take measures to reduce differences in income levels*”: Attitude towards income redistribution by sex, in a hypothetical population of 50 million people. The numbers in the table are frequencies in millions of people, row percentages (in parentheses) and overall percentages in square brackets.

Sex	Agree strongly	Agree	Neither agree nor disagree	Disagree	Disagree strongly	Total
Male	3.84 (16.00)	10.08 (42.00)	4.56 (19.00)	4.32 (18.00)	1.20 (5.00)	24.00 (100)
	7.68	20.16	9.12	8.64	2.40	48.00
Female	4.16 (16.00)	13.00 (50.00)	4.68 (18.00)	3.38 (13.00)	0.78 (3.00)	26.00 (100)
	8.32	26.00	9.36	6.76	1.56	52.00
Total	8.00 (16.00)	23.08 (46.16)	9.24 (18.48)	7.70 (15.40)	1.98 (3.96)	50 (100)

Imagine now that the full population consisted of 50 million people, and that the values of the two variables for them were as shown in Table 3.1. The frequencies in this table describe the population distribution of the variables in this hypothetical population, with the joint

distribution of sex and attitude shown by the internal cells of the table and the marginal distributions by its margins. So there are for example 3.84 million men and 4.16 million women in the population who strongly agree with the attitude statement, and 1.98 million people overall who strongly disagree with it.

Rather than the frequencies, it is more helpful to discuss population distributions in terms of proportions. Table 3.1 shows two sets of them, the overall proportions in square brackets out of the total population size, and the two rows of conditional proportions of attitude given sex (in parentheses). Either of these can be used to introduce the ideas of population distributions, but we focus on the conditional proportions because they will be more convenient for the discussion in later chapters. In this population we observe, for example, that the conditional proportion of “Strongly disagree” given that a person is a woman is 0.03, i.e. 3% of women strongly disagree with the statement, while among men the corresponding conditional proportion is 0.05.

Instead of “proportions”, when we discuss population distributions we will usually talk of “probabilities”. The two terms are equivalent when the population is finite and the variables are categorical, as in Table 3.1, but the language of probabilities is more appropriate in other cases. We can then say that Table 3.1 shows two sets of **conditional probabilities** in the population, which define two conditional **probability distributions** for attitude given sex.

The notion of a probability distribution creates a conceptual connection between population distributions and sampling from them. This is that the probabilities of the population distribution can also be thought of as sampling probabilities in (simple random) sampling from the population. For example, here the conditional probability of “Strongly disagree” among men is 0.05, while the probability of “Strongly agree” is 0.16. The sampling interpretation of this is that if we sample a man at random from the population, the probability is 0.05 that he strongly disagrees and 0.16 that he strongly agrees with the attitude statement.

The view of population distributions as probability distributions works also in other cases than the kind that is illustrated by Table 3.1. First, it applies also for continuous variables, where proportions of individual values are less useful (this is discussed further in Chapter 7). Second, it is also appropriate when the population is regarded as an infinite superpopulation, in which case the idea of population *frequencies* is not meaningful. With this device we have thus reached a formulation of a population distribution which is flexible enough to cover all the situations where we will need it.

3.6 Need for statistical inference

We have now introduced the first key concepts that are involved in statistical inference:

- The population, which may be regarded as finite or infinite. Distributions of variables in the population are the population distributions, which are formulated as probability distributions of the possible values of the variables.
- Random samples from the population, and sample distributions of variables in the sample.

Substantive research questions are most often questions about population distributions. This raises the fundamental challenge of inference: what we are interested in — the population — is not fully observed, while what we do observe — the sample — is not of main interest for itself. The sample is, however, what information we do have to draw on for conclusions about the population. Here a second challenge arises: because of random variation in the sampling, sample distributions will not be identical to population distributions, so inference will not be as simple

as concluding that whatever is true of the sample is also true of the population. Something cleverer is needed to weigh the evidence in the sample, and that something is statistical inference.

The next three chapters are mostly about statistical inference. Each of them discusses a particular type of analysis and inferential and descriptive statistical methods for it. These methods are some of the most commonly used in basic statistical analyses of empirical data. In addition, we will also use them as contexts in which to introduce the general concepts of statistical inference. This will be done gradually, with each chapter both building on previous concepts and introducing new ones, as follows:

- Chapter 4: Associations in two-way contingency tables (significance testing, sampling distributions of statistics).
- Chapter 5: Single proportions and comparisons of proportions (probability distributions, parameters, point estimation, confidence intervals).
- Chapter 7: Means of continuous variables (probability distributions of continuous variables, and inference for such variables).

Chapter 4

Statistical inference for two-way tables

4.1 Introduction

In this section we continue the discussion of methods of analysis for two-way contingency tables that was begun in Section 2.4.1. We will use again the example from the European Social Survey that was introduced early in Section 2.2. The two variables in the example are a person's sex and his or her attitude toward income redistribution measured as an ordinal variable with five levels. The two-way table of these variables in the sample is shown again for convenience in Table 4.1, including both the frequencies and the conditional proportions for attitude given sex.

Table 4.1: “*The government should take measures to reduce differences in income levels*”: Frequencies of respondents in the survey example, by sex and attitude towards income redistribution. The numbers in parentheses are conditional proportions of attitude given sex. Data: European Social Survey, Round 5, 2010, UK respondents only.

Sex	Agree strongly	Agree	Neither agree nor disagree	Disagree	Disagree strongly	Total
Male	160 (0.156)	439 (0.428)	187 (0.182)	200 (0.195)	41 (0.040)	1027 (1.0)
Female	206 (0.156)	651 (0.494)	239 (0.182)	187 (0.142)	34 (0.026)	1317 (1.0)
Total	366 (0.156)	1090 (0.465)	426 (0.182)	387 (0.165)	75 (0.032)	2344 (1.0)

Unlike in Section 2.4.1, we will now go beyond description of sample distributions and into statistical inference. The observed data are thus treated as a sample from a population, and we wish to draw conclusions about the population distributions of the variables. In particular, we want to examine whether the sample provides evidence that the two variables in the table are associated in the population — in the example, whether attitude depends on sex in the population. This is done using a statistical significance test known as χ^2 test of independence.

We will use it also as a vehicle for introducing the basic ideas of significance testing in general.

This initial explanation of significance tests is be lengthy and detailed, because it is important to gain a good understanding of these fundamental concepts from the beginning. From then on, the same ideas will be used repeatedly throughout the rest of the course, and in practically all statistical methods that you may encounter in the future. You will then be able to draw on what you will have learned in this chapter, and that learning will also be reinforced through repeated appearances of the same concepts in different contexts. It will then not be necessary to restate the basic ideas of the tools of inference in similar detail. A short summary of the χ^2 test considered in this chapter is given again at the end of the chapter, in Section 4.4.

4.2 Significance tests

A **significance test** is a method of statistical inference that is used to assess the plausibility of *hypotheses* about a population. A hypothesis is a question about population distributions, formulated as a *claim* about those distributions. For the test considered in this chapter, the question is whether or not the two variables in a contingency table are associated in the population. In the example we want to know whether men and women have the same distribution of attitudes towards income redistribution in the population. For significance testing, this question is expressed as the claim “The distribution of attitudes towards income redistribution *is* the same for men and women”, to which we want to identify the correct response, either “Yes, it is” or “No, it isn’t”.

In trying to answer such questions, we are faced with the complication that we only have information from a sample. For example, in Table 4.1 the conditional distributions of attitude are certainly not identical for men and women. According to the definition in Section 2.4.3, this shows that sex and attitude are associated *in the sample*. This, however, does not prove that they are also associated *in the population*. Because of sampling variation, the two conditional distributions are very unlikely to be exactly identical in a sample even if they are the same in the population. In other words, the hypothesis will not be exactly true in a sample even if it is true in the population.

On the other hand, some sample values differ from the values claimed by the hypothesis by so much that it would be difficult to explain them as a result of sampling variation alone. For example, if we had observed a sample where 99% of the men but only 1% of the women disagreed with the attitude statement, it would seem obvious that this should be evidence against the claim that the corresponding probabilities were nevertheless equal in the population. It would certainly be stronger evidence against such a claim than the difference of 19.5% vs. 14.2% that was actually observed in our sample, which in turn would be stronger evidence than, say, 19.5% vs. 19.4%. But how are we to decide where to draw the line, i.e. when to conclude that a particular sample value is or is not evidence against a hypothesis? The task of statistical significance testing is to provide explicit and transparent rules for making such decisions.

A significance test uses a statistic calculated from the sample data (a *test statistic*) which has the property that its values will be large if the sample provides evidence against the hypothesis that is being tested (the *null hypothesis*) and small otherwise. From a description (a *sampling distribution*) of what kinds of values the test statistic might have had if the null hypothesis was actually true in the population, we derive a measure (the *P-value*) that summarises in one number the strength of evidence against the null hypothesis that the sample provides. Based on this summary, we may then use conventional decision rules (*significance levels*) to make a discrete decision about the null hypothesis about the population. This decision will be either

to *fail to reject* or *reject* the null hypothesis, in other words to conclude that the observed data are or are not consistent with the claim about the population stated by the null hypothesis.

It only remains to put these general ideas into practice by defining precisely the steps of statistical significance tests. This is done in the sections below. Since some of the ideas are somewhat abstract and perhaps initially counterintuitive, we will introduce them slowly, discussing one at a time the following basic elements of significance tests:

- The hypotheses being tested
- Assumptions of a test
- Test statistics and their sampling distributions
- P -values
- Drawing and stating conclusions from tests

The significance test considered in this chapter is known as the χ^2 **test of independence** (χ^2 is pronounced “chi-squared”). It is also known as “Pearson’s χ^2 test”, after Karl Pearson who first proposed it in 1900.¹ We use this test to explain the elements of significance testing. These principles are, however, not restricted to this case, but are entirely general. This means that all of the significance tests you will learn on this course or elsewhere have the same basic structure, and differ only in their details.

4.3 The chi-square test of independence

4.3.1 Hypotheses

The null hypothesis and the alternative hypothesis

The technical term for the hypothesis that is tested in statistical significance testing is the **null hypothesis**. It is often denoted H_0 . The null hypothesis is a specific claim about population distributions. The χ^2 test of independence concerns the association between two categorical variables, and its null hypothesis is that there is no such association in the population.

In the context of this test, it is conventional to use alternative terminology where the variables are said to be **statistically independent** when there is no association between them, and **statistically dependent** when they are associated. Often the word “statistically” is omitted, and we talk simply of variables being independent or dependent. In this language, the null hypothesis of the χ^2 test of independence is that

$$H_0 : \text{The variables are statistically independent in the population.} \quad (4.1)$$

In our example the null hypothesis is thus that a person’s sex and his or her attitude toward income redistribution are independent in the population of adults in the UK.

The null hypothesis (4.1) and the χ^2 test itself are symmetric in that there is no need to designate one of the variables as explanatory and the other as the response variable. The hypothesis can, however, also be expressed in a form which does make use of this distinction. This links it more clearly with the definition of associations in terms of conditional distributions. In this form, the

¹ESS Round 5: European Social Survey Round 5 Data (2010). Data file edition 2.0. Norwegian Social Science Data Services, Norway Data Archive and distributor of ESS data. The full data can be obtained from <http://ess.nsd.uib.no/ess/round5/>.

null hypothesis (4.1) can also be stated as the claim that the conditional distributions of the response variable are the same at all levels of the explanatory variable, i.e. in our example as

H_0 : The conditional distribution of attitude is the same for men as for women.

The hypothesis could also be expressed for the conditional distributions the other way round, i.e. here that the distribution of sex is the same at all levels of the attitude. All three versions of the null hypothesis mean the same thing for the purposes of the significance test. Describing the hypothesis in particular terms is useful purely for easy interpretation of the test and its conclusions in specific examples.

As well as the null hypothesis, a significance test usually involves an **alternative hypothesis**, often denoted H_a . This is in some sense the opposite of the null hypothesis, which indicates the kinds of observations that will be taken as evidence against H_0 . For the χ^2 test of independence this is simply the logical opposite of (4.1), i.e.

H_a : The variables are not statistically independent in the population. (4.2)

In terms of conditional distributions, H_a is that the conditional distributions of one variable given the other are not all identical, i.e. that for at least one pair of levels of the explanatory variable the conditional probabilities of at least one category of the response variable are not the same.

Statistical hypotheses and research hypotheses

The word “hypothesis” appears also in research design and philosophy of science. There a **research hypothesis** means a specific claim or prediction about observable quantities, derived from subject-matter theory. The prediction is then compared to empirical observations. If the two are in reasonable agreement, the hypothesis and corresponding theory gain support or *corroboration*; if observations disagree with the predictions, the hypothesis is *falsified* and the theory must eventually be modified or abandoned. This role of research hypotheses is, especially in the philosophy of science originally associated with Karl Popper, at the heart of the scientific method. A theory which does not produce empirically falsifiable hypotheses, or fails to be modified even if its hypotheses are convincingly falsified, cannot be considered scientific.

Research hypotheses of this kind are closely related to the kinds of **statistical hypotheses** discussed above. When empirical data are quantitative, decisions about research hypotheses are in practice usually made, at least in part, as decisions about statistical hypotheses implemented through significance tests. The falsification and corroboration of research hypotheses are then paralleled by rejection and non-rejection of statistical hypotheses. The connection is not, however, entirely straightforward, as there are several differences between research hypotheses and statistical hypotheses:

- Statistical significance tests are also often used for testing hypotheses which do not correspond to any theoretical research hypotheses. Sometimes the purpose of the test is just to identify those observed differences and regularities which are large enough to deserve further discussion. Sometimes claims stated as null hypotheses are interesting for reasons which have nothing to do with theoretical predictions but rather with, say, normative or policy goals.
- Research hypotheses are typically stated as predictions about theoretical concepts. Translating them into testable statistical hypotheses requires further operationalisation of these concepts. First, we need to decide how the concepts are to be measured. Second, any

test involves also assumptions which are imposed not by substantive theory but by constraints of statistical methodology. Their appropriateness for the data at hand needs to be assessed separately.

- The conceptual connection is clearest when the research hypothesis matches the null hypothesis of a test in general form. Then the research hypothesis remains unfalsified as long as the null hypothesis remains not rejected, and gets falsified when the null hypothesis is rejected. Very often, however, the statistical hypotheses are for technical reasons defined the other way round. In particular, for significance tests that are about associations between variables, a research hypothesis is typically that there *is* an association between particular variables, whereas the null hypothesis is that there is *no* association (i.e. “null” association). This leads to the rather confusing situation where the research hypothesis is supported when the null hypothesis is rejected, and possibly falsified when the null hypothesis is not rejected.

4.3.2 Assumptions of a significance test

In the following discussion we will sometimes refer to Figure 4.1, which shows SPSS output for the χ^2 test of independence for the data in Table 4.1. Output for the test is shown on the line labelled “Pearson Chi-Square”, and “N of valid cases” gives the sample size n . The other entries in the table are output for other tests that are not discussed here, so they can be ignored.

Chi-Square Tests			
	Value	df	Asymp. Sig. (2-sided)
Pearson Chi-Square	18.862 ^a	4	.001
Likelihood Ratio	18.777	4	.001
Linear-by-Linear Association	11.842	1	.001
N of Valid Cases	2344		

a. 0 cells (.0%) have expected count less than 5. The minimum expected count is 32.86.

Figure 4.1: SPSS output of the χ^2 test of independence (here labelled “Pearson Chi-square”) for the data in Table 4.1.

When we apply any significance test, we need to be aware of its **assumptions**. These are conditions on the data which are not themselves being tested, but which need to be approximately satisfied for the conclusions from the test to be valid. Two broad types of such assumptions are particularly common. The first kind are assumptions about the measurement levels and population distributions of the variables. For the χ^2 test of independence these are relatively mild. The two variables must be categorical variables. They can have any measurement level, although in most cases this will be either nominal or ordinal. The test makes no use of the ordering of the categories, so it effectively treats all variables as if they were nominal.

The second common class of assumptions are conditions on the sample size. Many significance tests are appropriate only if this is sufficiently large. For the χ^2 test, the expected frequencies f_e (which will be defined below) need to be large enough in *every cell* of the table. A common rule of thumb is that the test can be safely used if all expected frequencies are at least 5. Another, slightly more lenient rule requires only that no more than 20% of the expected frequencies are less than 5, and that none are less than 1. These conditions can easily be checked with the help

of SPSS output for the χ^2 test, as shown in Figure 4.1. This gives information on the number and proportion of expected frequencies (referred to as “expected counts”) less than five, and also the size of the smallest of them. In our example the smallest expected frequency is about 33, so the sample size condition is easily satisfied.

When the expected frequencies do not satisfy these conditions, the χ^2 test is not fully valid, and the results should be treated with caution (the reasons for this will be discussed below). There are alternative tests which do not rely on these large-sample assumptions, but they are beyond the scope of this course.

In general, the hypotheses of a test define the questions it can answer, and its assumptions indicate the types of data it is appropriate for. Different tests have different hypotheses and assumptions, which need to be considered in deciding which test is appropriate for a given analysis. We will introduce a number of different significance tests in this coursepack, and give guidelines for choosing between them.

4.3.3 The test statistic

A **test statistic** is a number calculated from the sample (i.e. a statistic in the sense defined at the beginning of Section 2.6) which is used to test a null hypothesis. We will describe the calculation of the χ^2 test statistic step by step, using the data in Table 4.1 for illustration. All of the elements of the test statistic for this example are shown in Table 4.2. These elements are

- The **observed frequencies**, denoted f_o , one for each cell of the table. These are simply the observed cell counts (compare the f_o column of Table 4.2 to the counts in Table 4.1).
- The **expected frequencies** f_e , also one for each cell. These are cell counts in a hypothetical table which would show no association between the variables. In other words, they represent a table for a sample which would exactly agree with the null hypothesis of independence in the population. To explain how the expected frequencies are calculated, consider the cell in Table 4.1 for Male respondents who strongly agree with the statement. As discussed above, if the null hypothesis of independence is true in the population, then the conditional probability of strongly agreeing is the same for both men and women. This also implies that it must then be equal to the overall (marginal) probability of strongly agreeing. The sample version of this is that the proportion who strongly agree should be the same for men as among all respondents overall. This overall proportion in Table 4.1 is $366/2344 = 0.156$. If this proportion applied also to the 1027 male respondents, the number of of them who strongly agreed would be

$$f_e = \left(\frac{366}{2344} \right) \times 1027 = \frac{366 \times 1027}{2344} = 160.4.$$

Here 2344 is the total sample size, and 366 and 1027 are the marginal frequencies of strongly agreeers and male respondents respectively, i.e. the two marginal totals corresponding to the cell (Male, Strongly agree). The same rule applies also in general: the expected frequency for any cell in this or any other table is calculated as the product of the row and column totals corresponding to the cell, divided by the total sample size.

- The difference $f_o - f_e$ between observed and expected frequencies for each cell. Since f_e are the cell counts in a table which exactly agrees with the null hypothesis, the differences indicate how closely the counts f_o actually observed agree with H_0 . If the differences are small, the observed data are consistent with the null hypothesis, whereas large differences indicate evidence against it. The test statistic will be obtained by aggregating information about these differences across all the cells of the table. This cannot, however, be done by

adding up the differences themselves, because positive (f_o is larger than f_e) and negative (f_o is smaller than f_e) differences will always exactly cancel each other out (c.f. their sum on the last row of Table 4.2). Instead, we consider...

- ...the squared differences $(f_o - f_e)^2$. This removes the signs from the differences, so that the squares of positive and negative differences which are equally far from zero will be treated as equally strong evidence against the null hypothesis.
- Dividing the squared differences by the expected frequencies, i.e. $(f_o - f_e)^2/f_e$. This is an essential but not particularly interesting scaling exercise, which expresses the sizes of the squared differences relative to the sizes of f_e themselves.
- Finally, aggregating these quantities to get the χ^2 test statistic

$$\chi^2 = \sum \frac{(f_o - f_e)^2}{f_e}. \quad (4.3)$$

Here the summation sign Σ indicates that χ^2 is obtained by adding up the quantities $(f_o - f_e)^2/f_e$ across all the cells of the table.

Table 4.2: Calculating the χ^2 test statistic for Table 4.1. In the second column, SA, A, 0, D, and SD are abbreviations for Strongly agree, Agree, Neither agree nor disagree, Disagree and Strongly disagree respectively.

Cell	Cell						
tabularnewline Sex	Attitude	f_o	f_e	$f_o - f_e$	$(f_o - f_e)^2$	$(f_o - f_e)^2/f_e$	
Male	SA	160	160.4	-0.4	0.16	0.001	
Male	A	439	477.6	-38.6	1489.96	3.120	
Male	0	187	186.6	0.4	0.16	0.001	
Male	D	200	169.6	30.4	924.16	5.449	
Male	SD	41	32.9	8.1	65.61	1.994	
Female	SA	206	205.6	0.4	0.16	0.001	
Female	A	651	612.4	38.6	1489.96	2.433	
Female	0	239	239.4	-0.4	0.16	0.001	
Female	D	187	217.4	-30.4	924.16	4.251	
Female	SD	34	42.1	-8.1	65.61	1.558	
	Sum	2344	2344	0	4960.1	$\chi^2 = 18.81$	

The calculations can be done even by hand, but we will usually leave them to a computer. The last column of Table 4.2 shows that for Table 4.1 the test statistic is $\chi^2 = 18.81$ (which includes some rounding error, the correct value is 18.862). In the SPSS output in Figure 4.1, it is given in the “Value” column of the “Pearson Chi-Square” row.

4.3.4 The sampling distribution of the test statistic

We now know that the value of the χ^2 test statistic in the example is 18.86. But what does that mean? Why is the test statistic defined as (4.3) and not in some other form? And what does the number mean? Is 18.86 small or large, weak or strong evidence against the null hypothesis that sex and attitude are independent in the population?

In general, a test statistic for any null hypothesis should satisfy two requirements:

1. The value of the test statistic should be small when evidence against the null hypothesis is weak, and large when this evidence is strong.
2. The sampling distribution of the test statistic should be known and of convenient form when the null hypothesis is true.

Taking the first requirement first, consider the form of (4.3). The important part of this are the squared differences $(f_o - f_e)^2$ for each cell of the table. Here the expected frequencies f_e reveal what the table would look like if the sample was in perfect agreement with the claim of independence in the population, while the observed frequencies f_o show what the observed table actually does look like. If f_o in a cell is close to f_e , the squared difference is small and the cell contributes only a small addition to the test statistic. If f_o is very different from f_e — either much smaller or much larger than it — the squared difference and hence the cell's contribution to the test statistic are large.

Summing the contributions over all the cells, this implies that the overall value of the test statistic is small when the observed frequencies are close to the expected frequencies under the null hypothesis, and large when at least some of the observed frequencies are far from the expected ones. (Note also that the smallest possible value of the statistic is 0, obtained when the observed and the expected frequency are exactly equal in each cell.) It is thus *large* values of χ^2 which should be regarded as evidence *against* the null hypothesis, just as required by condition 1 above.

Turning then to condition 2, we first need to explain what is meant by “sampling distribution of the test statistic ... when the null hypothesis is true”. This is really the conceptual crux of significance testing. Because it is both so important and relatively abstract, we will introduce the concept of a sampling distribution in some detail, starting with a general definition and then focusing on the case of test statistics in general and the χ^2 test in particular.

Sampling distribution of statistic: General definition

The χ^2 test statistic (4.3) is a *statistic* as defined at the beginning of Section 2.6, that is a number calculated from data in a sample. Once we have observed a sample, the value of a statistic in that sample is known, such as the 18.862 for χ^2 in our example.

However, we also realise that this value would have been different if the sample had been different, and also that the sample could indeed have been different because the sampling is a process that involves randomness. For example, in the actually observed sample in Table 4.1 we had 200 men who disagreed with the statement and 41 who strongly disagreed with it. It is easily imaginable that another random sample of 2344 respondents from the same population could have given us frequencies of, say, 195 and 46 for these cells instead. If that had happened, the value of the χ^2 statistic would have been 19.75 instead of 18.86. Furthermore, it also seems intuitively plausible that not all such alternative values are equally likely for samples from a given population. For example, it seems quite improbable that the population from which the sample in Table 4.1 was drawn would instead produce a sample which also had 1027 men and 1317 women but where all the men strongly disagreed with the statement (which would yield $\chi^2 = 2210.3$).

The ideas that different possible samples would give different values of a sample statistic, and that some such values are more likely than others, are formalised in the concept of a sampling distribution:

- The **sampling distribution of a statistic** is the distribution of the statistic (i.e. its possible values and the proportions with which they occur) in the set of all possible

random samples of the same size from the population.

To observe a sampling distribution of a statistic, we would thus need to draw samples from the population over and over again, and calculate the value of the statistic for each such sample, until we had a good idea of the proportions with which different values of the statistic appeared in the samples. This is clearly an entirely hypothetical exercise in most real examples where we have just one sample of actual data, whereas the number of possible samples of that size is essentially or actually infinite. Despite this, statisticians can find out what sampling distributions would look like, under specific assumptions about the population. One way to do so is through mathematical derivations. Another is a *computer simulation* where we use a computer program to draw a large number of samples from an artificial population, calculate the value of a statistic for each of them, and examine the distribution of the statistic across these repeated samples. We will make use of both of these approaches below.

Sampling distribution of a test statistic under the null hypothesis

The sampling distribution of any statistic depends primarily on what the population is like. For test statistics, note that requirement 2 above mentioned only the situation where the null hypothesis is true. This is in fact the central conceptual ingredient of significance testing. The basic logic of drawing conclusions from such tests is that we consider what we would expect to see if the null hypothesis was in fact true in the population, and compare that to what was actually observed in our sample. The null hypothesis should then be rejected if the observed data would be surprising (i.e. unlikely) if the null hypothesis was actually true, and not rejected if the observed data would not be surprising under the null hypothesis.

We have already seen that the χ^2 test statistic is in effect a measure of the discrepancy between what is expected under the null hypothesis and what is observed in the sample. All test statistics for any hypotheses have this property in one way or another. What then remains to be determined is exactly how surprising or otherwise the observed data are relative to the null hypothesis. A measure of this is derived from the sampling distribution of the test statistic *under the null hypothesis*. It is the only sampling distribution that is needed for carrying out a significance test.

Sampling distribution of the χ^2 test statistic under independence

For the χ^2 test, we need the sampling distribution of the test statistic (4.3) under the independence null hypothesis (4.1). To make these ideas a little more concrete, the upper part of Table 4.4 shows the crosstabulation of sex and attitude in our example for a finite population where the null hypothesis holds. We can see that it does because the two conditional distributions for attitude, among men and among women, are the same (this is the only aspect of the distributions that matters for this demonstration; the exact values of the probabilities are otherwise irrelevant). These are of course hypothetical population distributions, as we do not know the true ones. We also do not claim that this hypothetical population is even close to the true one. The whole point of this step of hypothesis testing is to set up a population where the null hypothesis holds as a fixed point of comparison, to see what samples from such a population would look like and how they compare with the real sample that we have actually observed.

Population (frequencies are in millions of people):

Table 4.3: “The government should take measures to reduce differences in income levels”: Attitude towards income redistribution by sex (with row proportions in parentheses), in a hypothetical population of 50 million people where sex and attitude are independent, and in one random sample from this population.

Sex	Agree strongly	Agree	Neither agree nor disagree	Disagree	Disagree strongly	Total
Male	3.744 (0.156)	11.160 (0.465)	4.368 (0.182)	3.960 (0.165)	0.768 (0.032)	24.00 (1.0)
Female	4.056 (0.156)	12.090 (0.465)	4.732 (0.182)	4.290 (0.165)	0.832 (0.032)	26.00 (1.0)
Total	7.800 (0.156)	23.250 (0.465)	9.100 (0.182)	8.250 (0.165)	1.600 (0.032)	50 (1.0)

Sample:

Table 4.4: $\chi^2 = 2.8445$

Sex	Agree strongly	Agree	Neither agree nor disagree	Disagree	Disagree strongly	Total
Male	181 (0.161)	505 (0.450)	191 (0.170)	203 (0.181)	41 (0.037)	1121 (1.0)
Female	183 (0.150)	569 (0.465)	229 (0.187)	202 (0.165)	40 (0.033)	1223 (1.0)
Total	364 (0.155)	1074 (0.458)	420 (0.179)	405 (0.173)	81 (0.035)	2344 (1.0)

In the example we have a sample of 2344 observations, so to match that we want to identify the sampling distribution of the χ^2 statistic in random samples of size 2344 from the population like the one in the upper part of Table 4.4. The lower part of that table shows one such sample. Even though it comes from a population where the two variables are independent, the same is not exactly true in the sample: we can see that the conditional sample distributions are not the same for men and women. The value of the χ^2 test statistic for this simulated sample is 2.8445.

Before we proceed with the discussion of the sampling distribution of the χ^2 statistic, we should note that it will be a *continuous* probability distribution. In other words, the number of distinct values that the test statistic can have in different samples is so large that their distribution is clearly effectively continuous. This is true even though the two *variables* in the contingency table are themselves categorical. The two distributions, the population distribution of the variables and the sampling distribution of a test statistic, are quite separate entities and need not resemble each other. We will consider the nature of continuous probability distributions in more detail in Chapter 7. In this chapter we will discuss them relatively superficially and only to the extent that is absolutely necessary.

Figure 4.2 shows what we observe if do a computer simulation to draw many more samples from the population in Table 4.4. The figure shows the histogram of the values of the χ^2 test statistic

calculated from 100,000 such samples. We can see, for example, that χ^2 is between 0 and 10 for most of the samples, and larger than that for only a small proportion of them. In particular, we note already that the value $\chi^2 = 18.8$ that was actually observed in the real sample occurs very rarely if samples are drawn from a population where the null hypothesis of independence holds.




Figure 4.2: Example of the sampling distribution of the χ^2 test statistic for independence. The plot shows a histogram of the values of the statistic in 100,000 simulated samples of size $n = 2344$ drawn from the population distribution in the upper part of Table 4.4. Superimposed on the histogram is the curve of the approximate sampling distribution, which is the χ^2 distribution with 4 degrees of freedom.

The form of the sampling distribution can also be derived through mathematical arguments. These show that for any two-way contingency table, the approximate sampling distribution of the χ^2 statistic is a member of a class of continuous probability distributions known as the **χ^2 distributions** (the same symbol χ^2 is rather confusingly used to refer both to the test statistic and its sampling distribution). The χ^2 distributions are a family of individual distributions, each of which is identified by a number known as the **degrees of freedom** of the distribution. Figure 4.3 shows the probability curves of some χ^2 distributions (what such curves mean is explained in more detail below, and in Chapter 7). All of the distributions are skewed to the right, and the shape of a particular curve depends on its degrees of freedom. All of the curves give non-zero probabilities only for positive values of the variable on the horizontal axis, indicating that the value of a χ^2 -distributed variable can never be negative. This is appropriate for the χ^2 test statistic (4.3), which is also always non-negative.

For the χ^2 test statistic of independence we have the following result:

- When the null hypothesis (4.1) is true in the population, the sampling distribution of the test statistic (4.3) calculated for a two-way table with R rows and C columns is approximately the χ^2 distribution with $df = (R - 1)(C - 1)$ degrees of freedom.

The degrees of freedom are thus given by the number of rows in the table minus one, multiplied by the number of columns minus one. Table 4.1, for example, has $R = 2$ rows and $C = 5$ columns, so its degrees of freedom are $df = (2 - 1) \times (5 - 1) = 4$ (as indicated by the “df”



Figure 4.3: Probability curves of some χ^2 distributions with different degrees of freedom (df).

column of the SPSS output of Figure 4.1). Figure 4.2 shows the curve of the χ^2 distribution with $df = 4$ superimposed on the histogram of the sampling distribution obtained from the computer simulation. The two are in essentially perfect agreement, as mathematical theory indicates they should be.

These degrees of freedom can be given a further interpretation which relates to the structure of the table². We can, however, ignore this and treat df simply as a number which identifies the appropriate χ^2 distribution to be used for the χ^2 test for a particular table. Often it is convenient to use the notation χ^2_{df} to refer to a specific distribution, e.g. χ^2_4 for the χ^2 distribution with 4 degrees of freedom.

The χ^2 sampling distribution is “approximate” in that it is an *asymptotic approximation* which is exactly correct only if the sample size is infinite and approximately correct when it is sufficiently large. This is the reason for the conditions for the sizes of the expected frequencies that were discussed in Section 4.3.2. When these conditions are satisfied, the approximation is accurate enough for all practical purposes and we use the appropriate χ^2 distribution as the sampling distribution.

In Section 4.3.4, under requirement 2 for a good test statistic, we mentioned that its sampling distribution under the null hypothesis should be “known” and “of convenient form”. We now know that for the χ^2 test it is a χ^2 distribution. The “convenient form” means that the sampling distribution should not depend on too many specific features of the data at hand. For the χ^2 test, the approximate sampling distribution depends (through the degrees of freedom) only on the size of the table but not on the sample size or the marginal distributions of the two variables. This is convenient in the right way, because it means that we can use the same χ^2 distribution for any table with a given number of rows and columns, as long as the sample size is large

²The data were obtained from <http://data.london.gov.uk/datastore/package/london-borough-profiles>.

If you download the “Profiles in Excel” workbook, you will find that one of the pages contains a map of the boroughs, and a tool for visualising the data on that map. A regular map of the boroughs can be found at for example at

<http://www.londoncouncils.gov.uk/londonfacts/londonlocalgovernment/londonmapandlinks/default.htm>.

enough for the conditions in Section 4.3.2 to be satisfied.

4.3.5 The P -value

The last key building block of significance testing operationalises the comparison between the observed value of a test statistic and its sampling distribution under the null hypothesis. In essence, it provides a way to determine whether the test statistic in the sample should be regarded as “large” or “not large”, and with this the measure of evidence against the null hypothesis that is the end product of the test:

- The **P-value** is the probability, if the null hypothesis was true in the population, of obtaining a value of the test statistic which provides as strong or stronger evidence against the null hypothesis, and in the direction of the alternative hypothesis, as the the value of the test statistic in the sample actually observed.

The relevance of the phrase “in the direction of the alternative hypothesis” is not apparent for the χ^2 test, so we can ignore it for the moment. As argued above, for this test it is large values of the test statistic which indicate evidence against the null hypothesis of independence, so the values that correspond to “as strong or stronger evidence” against it are the ones that are as large or larger than the observed statistic. Their probability is evaluated from the χ^2 sampling distribution defined above.

Figure 4.4 illustrates this calculation. It shows the curve of the χ^2_4 distribution, which is the relevant sampling distribution for the test for the 2×5 table in our example. Suppose first, hypothetically, that we had actually observed the sample in the lower part of Table 4.4, for which the value of the test statistic is $\chi^2 = 2.84$. The P -value of the test for this sample would then be the probability of values of 2.84 or larger, evaluated from the χ^2_4 distribution.

For a probability curve like the one in Figure 4.4, areas under the curve correspond to probabilities. For example, the area under the whole curve from 0 to infinity is 1, because a variable which follows the χ^2_4 distribution is certain to have one of these values. Similarly, the probability that we need for the P -value for $\chi^2 = 2.84$ is the area under the curve to the right of the value 2.84, which is shown in grey in Figure 4.4. This is $P = 0.585$.

The test statistic for the real sample in Table 4.1 was $\chi^2 = 18.86$, so the P -value is the combined probability of this and all larger values. This is also shown in Figure 4.4. However, this area is not really visible in the plot because 18.86 is far into the tail of the distribution where the probabilities are low. The P -value is then also low, specifically $P = 0.0008$.

In practice the P -value is usually calculated by a computer. In the SPSS output of Figure 4.1 is shown in the column labelled “Asymp. Sig. (2-sided)” which is short for “Asymptotic significance level” (you can ignore the “2-sided” for this test). The value is listed as 0.001. SPSS reports, by default, P -values rounded to three decimal places. Sometimes even the smallest of these is zero, in which case the value is displayed as “.000”. This is bad practice, as the P -value for most significance tests is never *exactly* zero. P -values given by SPSS as “.000” should be reported instead as “ $P < 0.001$ ”.

Before the widespread availability of statistical software, P -values had to be obtained approximately using tables of distributions. Since you may still see this approach described in many text books, it is briefly explained here. You may also need to use the table method in the examination, where computers are not allowed. Otherwise, however, this approach is now of little interest: if the P -value is given in the computer output, there is no need to refer to distributional tables.



Figure 4.4: Illustration of the P -value for a χ^2 test statistic with 4 degrees of freedom and with values $\chi^2 = 2.84$ (area of the grey region under the curve) and $\chi^2 = 18.86$.

All introductory statistical text books include a table of χ^2 distributions, although its format may vary slightly from book to book. Such a table is also included in the Appendix of this coursepack. An extract from the table is shown in Table 4.5. Each row of the table corresponds to a χ^2 distribution with the degrees of freedom given in the first column. The other columns show so-called “critical values” for the probability levels given on the first row. Consider, for example, the row for 4 degrees of freedom. The figure 7.78 in the column for probability level 0.100 indicates that the probability of a value of 7.78 or larger is exactly 0.100 for this distribution. The 9.49 in the next column shows that the probability of 9.49 or larger is 0.050. Another way of saying this is that if the appropriate degrees of freedom were 4, and the test statistic was 7.78, the P -value would be exactly 0.100, and if the statistic was 9.49, P would be 0.050.

Table 4.5: An extract from a table of critical values for χ^2 distributions. Row 2-5 show the right-hand tail probability

df	0.100	0.050	0.010	0.001
1	2.71	3.84	6.63	10.83
2	4.61	5.99	9.21	13.82
3	6.25	7.81	11.34	16.27
4	7.78	9.49	13.28	18.47
\vdots	\vdots			\vdots

The values in the table also provide bounds for other values that are not shown. For instance, in the hypothetical sample in Table 4.4 we had $\chi^2 = 2.84$, which is smaller than 7.78. This implies that the corresponding P -value must be larger than 0.100, which (of course) agrees with the precise value of $P = 0.585$ (see also Figure 4.4). Similarly, $\chi^2 = 18.86$ for the real data in Table 4.1, which is larger than the 18.47 in the “0.001” column of the table for the χ^2_4 distribution. Thus the corresponding P -value must be smaller than 0.001, again agreeing with the correct value of $P = 0.0008$.

4.3.6 Drawing conclusions from a test

The P -value is the end product of any significance test, in that it is a complete quantitative summary of the strength of evidence against the null hypothesis provided by the data in the

When the P -value is *smaller* than a conventional level of significance (i.e. the test statistic is *larger* than the corresponding critical value), it is said that the null hypothesis is **rejected** at that level of significance, or that the results (i.e. evidence against the null hypothesis) are **statistically significant** at that level. In our example the P -value was smaller than 0.001. The null hypothesis is thus “rejected at the 0.1 % level of significance”, i.e. the evidence that the variables are not independent in the population is “statistically significant at the 0.1% level” (as well as the 10%, 5% and 1% levels of course, but it is enough to state only the strongest level).

The strict decision formulation of significance testing is much overused and misused. It is in fact quite rare that the statistical analysis will immediately be followed by some practical action which absolutely requires a decision about whether to act on the basis of the null hypothesis or the alternative hypothesis. Typically the analysis which a test is part of aims to examine some research question, and the results of the test simply contribute new information to add support for one or the other side of the argument about the question. The P -value is the key measure of the strength and direction of that evidence, so it should *always* be reported. The standard significance levels used for rejecting or not rejecting null hypotheses, on the other hand, are merely useful conventional reference points for structuring the reporting of test results, and their importance should not be overemphasised. Clearly P -values of, say, 0.049 and 0.051 (i.e. ones either side of the most common conventional significance level 0.05) indicate very similar levels of evidence against a null hypothesis, and acting as if one was somehow qualitatively more decisive is simply misleading.

How to state the conclusions

The final step of a significance test is describing its conclusions in a research report. This should be done with appropriate care:

- The report should make clear which test was used. For example, this might be stated as something like “The χ^2 test of independence was used to test the null hypothesis that in the population the attitude toward income redistribution was independent of sex in the population”. There is usually no need to give literature references for the standard tests described on this course.
- The numerical value of the P -value should be reported, rounded to two or three decimal places (e.g. $P = 0.108$ or $P = 0.11$). It can also be reported in an approximate way as, for example, “ $P < 0.05$ ” (or the same in symbols to save space, e.g. for $P < 0.1$, ** for $P < 0.05$, and so on). Very small P -values can always be reported as something like “ $P < 0.001$ ”.
- When (cautiously) discussing the results in terms of discrete decisions, the most common practice is to say that the null hypothesis was either *not rejected* or *rejected* at a given significance level. It is *not* acceptable to say that the null hypothesis was “accepted” as an alternative to “not rejected”. Failing to reject the hypothesis that two variables are independent in the population is not the same as proving that they actually *are* independent.
- A common mistake is to describe the P -value as the probability that the null hypothesis is true. This is understandably tempting, as such a claim would seem more natural and convenient than the correct but convoluted interpretation of the P -value as “the probability of obtaining a test statistic as or more extreme as the one observed in the data if the test was repeated many times for different samples from a population where the null hypothesis was true”. Unfortunately, however, the P -value is *not* the probability

of the null hypothesis being true. Such a probability does not in fact have any real meaning at all in the statistical framework considered here³.

- The results of significance tests should be stated using the names and values of the variables involved, and not just in terms of “null” and “alternative” hypotheses. This also forces you to recall what the hypotheses actually were, so that you do not accidentally describe the result the wrong way round (e.g. that the data support a claim when they do just the opposite). There are no compulsory phrases for stating the conclusions, so it can be done in a number of ways. For example, a fairly complete and careful statement in our example would be
 - “There is strong evidence that the distributions of attitudes toward income redistribution are not the same for men and women in the population ($P < 0.001$).”

Other possibilities are

- “The association between sex and attitude toward income redistribution in the sample is statistically significant ($P < 0.001$).”
- “The analysis suggests that there is an association between sex and attitude toward income redistribution in the population ($P < 0.001$).”

The last version is slightly less clear than the other statements in that it relies on the reader recognizing that the inclusion of the P -value implies that the word “differs” refers to a statistical claim rather a statement of absolute fact about the population. In many contexts it would be better to say this more explicitly.

Finally, if the null hypothesis of independence is rejected, the test should not usually be the only statistical analysis that is reported for a two-way table. Instead, we would then go on to describe *how* the two variables appear to be associated, using the of descriptive methods discussed in Section 2.4.

4.4 Summary of the chi-square test of independence

We have now described the elements of a significance test in some detail. Since it is easy to lose sight of the practical steps of a test in such a lengthy discussion, they are briefly repeated here for the χ^2 test of independence. The test of the association between sex and attitude in the survey example is again used for illustration:

1. Data: observations of two categorical variables, here sex and attitude towards income redistribution for $n = 2344$ respondents, presented in the two-way, 2×5 contingency table 4.1.
2. Assumptions: the variables can have any measurement level, but the expected frequencies f_e must be large enough. A common rule of thumb is that f_e should be at least 5 for every cell of the table. Here the smallest expected frequency is 32.9, so the requirement is comfortably satisfied.
3. Hypotheses: null hypothesis H_0 that the two variables are statistically independent (i.e. not associated) in the population, against the alternative hypothesis that they are dependent.

³The data can be obtained from <http://www3.norc.ox.ac.uk/gss+website/>, which gives further information on the survey, including the full text of the questionnaires.

4. The test statistic: the χ^2 statistic

$$\chi^2 = \sum \frac{(f_o - f_e)^2}{f_e}$$

where f_o denotes observed frequencies in the cells of the table and f_e the corresponding expected frequencies under the null hypothesis, and the summation is over all of the cells. For Table 4.1, $\chi^2 = 18.86$.

5. The sampling distribution of the test statistic when H_0 is true: a χ^2 distribution with $(R - 1) \times (C - 1) = 1 \times 4 = 4$ degrees of freedom, where R ($= 2$) and C ($= 5$) denote the numbers of rows and columns in the table respectively.
6. The P -value: the probability that a randomly selected value from the χ^2_4 distribution is at least 18.86. This is $P = 0.0008$, which may also be reported as $P < 0.001$.
7. Conclusion: the null hypothesis of independence is strongly rejected. The χ^2 test indicates very strong evidence that sex and attitude towards income redistribution are associated in the population ($P < 0.001$).

When the association is judged to be statistically significant, its nature and magnitude can be further explored using the descriptive methods for two way tables discussed in Section 2.4.

Chapter 5

Inference for population proportions

5.1 Introduction

In this chapter we still consider statistical analyses which involve only discrete, categorical variables. In fact, we now focus on the simplest case of all, that of **dichotomous** (binary) variables which have only two possible values. Four examples which will be used for illustration throughout this chapter are introduced in Section 5.2. In the first two of them we consider a binary variable in a single population, while in the last two examples the question of interest involves a comparison of the distributions of the variable between two populations (*groups*).

The data for such analyses can be summarised in simple tables, the one-group case with a one-way table of two cells, and the two-group case with a 2×2 contingency table. Here, however, we formulate the questions of interest slightly differently, with primary emphasis on the *probability* of one of the two values of the variable of interest. In the one-group case the questions of interest are then about the population value of a single probability, and in the two-group case about the comparison of the values of this probability between the two groups.

While we describe specific methods of inference for these cases, we also use them to introduce some further general elements of statistical inference:

- Population **parameters** of probability distributions.
- **Point estimation** of population parameters.
- Hypotheses about the parameters, and significance tests for them.
- **Confidence intervals** for population parameters.

The comparisons in the two-group analyses again address questions about associations, now between the group and the dichotomous variable of interest. Here it will be useful to employ the terminology introduced in Section 1.2.4, which distinguishes between the **explanatory variable** and the **response variable** in the association. Following a common convention, we will denote the explanatory variable by X and the response variable by Y . In the two-group cases of this chapter, X will be the group (which is itself also binary) and Y the binary variable whose probabilities we are interested in. We will use Y to denote this binary variable also in the one-group examples.

5.2 Examples

The following four examples will be discussed in this chapter. Examples 5.1 and 5.2 concern only one group, while in Examples 5.3 and 5.4 two groups are to be compared. Table 5.1 shows basic sample statistics for the examples, together with the results of the significance tests and confidence intervals described later.

Example 5.1: An EU referendum

A referendum about joining the European Union was held in Finland on the 16th of October, 1994. Suppose that in an opinion poll conducted around October 4th (just before the beginning of postal voting), 330 of the $n = 702$ respondents (47.0%) indicated that they would definitely vote Yes to joining the EU, 190 (27.1%) said they would definitely vote No, and 182 (25.9%) were still undecided¹. Here we will consider the dichotomous variable with the values of Yes (330 respondents) versus No or Undecided (372 respondents, or 53.0%). The proportion of voters who definitely intend to vote Yes provides a lower bound for the proportion of Yes-votes in the referendum, even if all of the currently undecided voters eventually decided to vote No.

Example 5.2: Evidence of possible racial bias in jury selection

As part of an official inquiry into the extent of racial and gender bias in the justice system in the U.S. state of Pennsylvania², an investigation was made of whether people from minority racial groups were underrepresented in trial juries. One part of the assessment was a survey administered to all those called for the jury panel for criminal trials (from which the juries for actual trials will be selected) in Allegheny County, Pennsylvania (the city of Pittsburgh and its surrounding areas) between May and October, 2001. We will consider the dichotomous variable of whether a respondent to the survey identified his or her own race as Black (African American) or some other race category. Of the $n = 4950$ respondents, 226 (4.57%) identified themselves as black. This will be compared to the the percentage of blacks in the whole population of people aged 18 and over (those eligible for jury service) in the county, which is 12.4% (this is a census estimate which will here be treated as a known population quantity, ignoring any possible census error in it).

¹ESS Round 5: European Social Survey Round 5 Data (2010). Data file edition 2.0. Norwegian Social Science Data Services, Norway Data Archive and distributor of ESS data. The full data can be obtained from <http://ess.nsd.uib.no/ess/round5/>.

²The data were obtained from <http://data.london.gov.uk/datastore/package/london-borough-profiles>.

If you download the “Profiles in Excel” workbook, you will find that one of the pages contains a map of the boroughs, and a tool for visualising the data on that map. A regular map of the boroughs can be found at for example at

<http://www.londoncouncils.gov.uk/londonfacts/londonlocalgovernment/londonmapandlinks/default.htm>.

Table 5.1: Examples of analyses of population proportions used in Chapter 5. In addition to sample sizes n and proportions $\hat{\pi}$, the table shows for the one-sample examples 5.1 and 5.2 the z -test statistic for the hypothesis $H_0 : \pi = \pi_0$, its P -value against a two-sided alternative, and a 95% confidence interval for π . For the two-sample examples 5.3 and 5.4, the table shows the estimated between-group difference of proportions $\hat{\Delta}$, the z -test statistic for the hypothesis $H_0 : \Delta = 0$, its P -value against a two-sided alternative, and a 95% confidence interval for Δ .

One sample	strut	strut	strut	strut	strut	strut	strut
Example 5.1: Voting intention in an EU referendum	n	Yes	$\hat{\pi}$	π_0	z	P	95% CI
	702	330	0.470	0.5	-1.59	0.112	(0.433; 0.507)
Example 5.2: Race of members of jury panel	n	Black	$\hat{\pi}$	π_0	z	P	95% CI
	4950	226	0.0457	0.124	-16.71	< 0.001	(0.040; 0.052)
Two Independent samples							
Example 5.3: Polio diagnoses in a vaccine trial	n	Yes	$\hat{\pi}$	Diff. ($\hat{\Delta}$)	z	P	95% CI
Control group (placebo)	201,229	142	0.000706				
Treatment group (vaccine)	200,745	57	0.000284	0.000422	-6.01	< 0.001	(-0.000560; -0.000284)
Example 5.4: Optimistic about young people's future	n	Yes	$\hat{\pi}$	Diff. ($\hat{\Delta}$)	z	P	95% CI
Negative question	921	257	0.279				
Positive question	929	338	0.364	0.085	3.92	< 0.001	(0.043; 0.127)

Example 5.3: The Salk polio vaccine field trial of 1954

The first large-scale field trials of the “killed virus” polio vaccination developed by Dr. Jonas Salk were carried out in the U.S. in 1954³. In the randomized, double-blind placebo-control part

³The data can be obtained from <http://www3.norc.ox.ac.uk/gss+website/>, which gives further information on the survey, including the full text of the questionnaires.

of the trial, a sample of schoolchildren were randomly assigned to receive either three injections of the polio vaccine, or three injections of a placebo, inert saltwater which was externally indistinguishable from the real vaccine. The explanatory variable X is thus the group (vaccine or “treatment” group vs. placebo or “control” group). The response variable Y is whether the child was diagnosed with polio during the trial period (yes or no). There were $n_1 = 201,229$ children in the control group, and 142 of them were diagnosed with polio; in the treatment group, there were 57 new polio cases among $n_2 = 200,745$ children (in both cases only those children who received all three injections are included here). The proportions of cases of polio were thus 0.000706 in the control group and 0.000284 in the vaccinated group (i.e. 7.06 and 2.84 cases per 10,000 subjects, respectively).

Example 5.4: Split-ballot experiment on acquiescence bias

Survey questions often ask whether respondents agree or disagree with given statements on opinions or attitudes. *Acquiescence bias* means the tendency of respondents to agree with such statements, regardless of their contents. If it is present, we will overestimate the proportion of people holding the opinion corresponding to agreement with the statement. The data used in this example come from a study which examined acquiescence bias through a randomized experiment⁴. In a survey carried out in Kazakhstan, the respondents were presented with a number of attitude statements, with four response categories: “Fully agree”, “Somewhat agree”, “Somewhat disagree”, and “Fully disagree”. Here we combine the first two and the last two, and consider the resulting dichotomous variable, with values labelled “Agree” and “Disagree”.

We consider one item from the survey, concerning the respondents’ opinions on the expectations of today’s young people. There were two forms of the question:

- “A young person today can expect little of the future”
- “A young person today can expect much of the future”

We will call these the “Negative” and “Positive” question respectively. Around half of the respondents were randomly assigned to receive the positive question, and the rest got the negative question. The explanatory variable X indicates the type of question, with Negative and Positive questions coded here as 1 and 2 respectively. The dichotomous response variable Y is whether the respondent gave a response which was optimistic about the future (i.e. agreed with the positive or disagreed with the negative question) or a pessimistic response. The sample sizes and proportions of optimistic responses in the two groups are reported in Table 5.1. The proportion is higher when the question was worded positively, as we would expect if there was acquiescence bias. Whether this difference is statistically significant remains to be determined.

5.3 Probability distribution of a dichotomous variable

The response variables Y considered in this section have only two possible values. It is common to code them as 0 and 1. In our examples, we will define the values of the variable of interest as follows:

- Example 5.1: 1 if a person says that he or she will definitely vote Yes, and 0 if the respondent will vote No or is undecided
- Example 5.2: 1 for black respondents and 0 for all others

⁴ESS Round 5: European Social Survey Round 5 Data (2010). Data file edition 2.0. Norwegian Social Science Data Services, Norway Data Archive and distributor of ESS data. The full data can be obtained from <http://ess.nsd.uib.no/ess/round5/>.

- Example 5.3: 1 if a child developed polio, 0 if not
- Example 5.4: 1 if the respondent gave an optimistic response, 0 if not

The population distribution of such a variable is completely specified by one number, the **probability** that a randomly selected member of the population will have the value $Y = 1$ rather than 0. It can also be thought of as the **proportion** of units in the population with $Y = 1$; we will use the two terms interchangeably. This probability is denoted here π (the lower-case Greek letter “pi”⁵). The value of π is between 0 (no-one in the population has $Y = 1$) and 1 (everyone has $Y = 1$). Because Y can have only two possible values, and the sum of probabilities must be one, the population probability of $Y = 0$ is $1 - \pi$.

The probability distribution which corresponds to this kind of population distribution is the *Binomial distribution*. For later use, we note already here that the mean of this distribution is π and its variance is $\pi(1 - \pi)$.

In Example 5.1, the population is that of eligible voters at the time of the opinion poll, and π is the probability that a randomly selected eligible voter definitely intended to vote Yes. In Example 5.2, π is the probability that a black person living in the county will be selected to the jury panel. In Example 5.3, π is the probability (possibly different in the vaccinated and unvaccinated groups) that a child will develop polio, and in Example 5.4 it is the probability (which possibly depends on how the question was asked) that a respondent will give an optimistic answer to the survey question.

The probability π is the **parameter** of the binomial distribution. In general, the parameters of a probability distribution are one or more numbers which fully determine the distribution. For example, in the analyses of Chapter 4 we considered conditional distributions of a one variable in a contingency table given the other variable. Although we did not make use of this terminology there, these distributions also have their parameters, which are the probabilities of (all but one of) the categories of the response variable. Another case will be introduced in Chapter 7, where we consider a probability distribution for a continuous variable, and its parameters.

5.4 Point estimation of a population probability

Questions and hypotheses about population distributions are usually most conveniently formulated in terms of the parameters of the distributions. For a binary variable Y , this means that statistical inference will be focused on the probability π .

The most obvious question about a parameter is “what is our best guess of the value of the parameter in the population?” The answer will be based on the information in the sample, using some sample statistic as the best guess or **estimate** of the population parameter. Specifically, this is a **point estimate**, because it is expressed as a single value or a “point”, to distinguish it from *interval* estimates defined later.

We denote a point estimate of π by $\hat{\pi}$. The “^” or “hat” is often used to denote an estimate of a parameter indicated by the symbol under the hat; $\hat{\pi}$ is read as “pi-hat”. As π for a binomial distribution is the population proportion of $Y = 1$, the obvious choice for a point estimate of it is the *sample* proportion of units with $Y = 1$. If we denote the *number* of such units by m , the proportion is thus $\hat{\pi} = m/n$, i.e. m divided by the sample size n . In Example 5.1, $m = 330$ and

⁵Strictly speaking, the analysis should incorporate sampling weights (variable *DWEIGHT*) to adjust for different sampling probabilities for different types of respondents. Here the weights are ignored. Using them would not change the main conclusions for these variables.

$n = 702$, and $\hat{\pi} = 330/702 = 0.47$. This and the estimated proportions in the other examples are shown in Table 5.1, in the two-sample examples 5.3 and 5.4 separately for the two groups.

When Y is coded with values 0 and 1, $\hat{\pi}$ is also equal to the sample mean of Y , since

$$\bar{Y} = \frac{Y_1 + Y_2 + \cdots + Y_n}{n} = \frac{0 + 0 + \cdots + 0 + \overbrace{1 + 1 + \cdots + 1}^{m \text{ ones}}}{n} = \frac{m}{n} = \hat{\pi}. \quad (5.1)$$

5.5 Significance test of a single proportion

5.5.1 Null and alternative hypotheses

A null hypothesis about a single population probability π is of the form

$$H_0 : \pi = \pi_0 \quad (5.2)$$

where π_0 is a given number which is either of specific interest or in some other sense a suitable benchmark in a given application. For example, in the voting example 5.1 we could consider $\pi_0 = 0.5$, i.e. that the referendum was too close to call. In the jury example 5.2 the value of interest would be $\pi_0 = 0.124$, the proportion of blacks in the general adult population of the county.

An alternative but equivalent form of (5.2) is expressed in terms of the difference

$$\Delta = \pi - \pi_0 \quad (5.3)$$

(Δ is the upper-case Greek letter “Delta”). Then (5.2) can also be written as

$$H_0 : \Delta = 0, \quad (5.4)$$

i.e. that there is no difference between the true population probability and the hypothesised value π_0 . This version of the notation allows us later to draw attention to the similarities between different analyses in this chapter and in Chapter 7. In all of these cases the quantities of interest turn out to be differences of some kind, and the formulas for test statistics and confidence intervals will be of essentially the same form.

The alternative hypothesis to the null hypothesis (5.4) requires some further comments, because there are some new possibilities that did not arise for the χ^2 test of independence in Chapter 4. For the difference Δ , we may consider two basic kinds of alternative hypotheses. The first is a **two-sided alternative hypothesis**

$$H_a : \Delta \neq 0 \quad (5.5)$$

(where “ \neq ” means “not equal to”). This claims that the true value of the population difference Δ is some unknown value which is *not* 0 as claimed by the null hypothesis. With a two-sided H_a , sample evidence that the true difference differs from 0 will be regarded as evidence against the null hypothesis, irrespective of whether it suggests that Δ is actually smaller or larger than 0 (hence the word “two-sided”). When $\Delta = \pi - \pi_0$, this means that we are trying to

assess whether the true probability π is different from the claimed value π_0 , but without any expectations about whether π might be smaller or larger than π_0 .

The second main possibility is one of the two **one-sided alternative hypotheses**

$$H_a : \Delta > 0 \quad (5.6)$$

or

$$H_a : \Delta < 0 \quad (5.7)$$

Such a hypothesis is only interested in values of Δ to one side of 0, either larger or smaller than it. For example, hypothesis (5.6) in the referendum example 5.1, with $\pi_0 = 0.5$, is $H_a : \pi > 0.5$, i.e. that the proportion who intend to vote Yes is *greater* than one half. Similarly, in the jury example 5.2, with $\pi_0 = 0.124$, (5.7) is the hypothesis $H_a : \pi < 0.124$, i.e. that the probability that an eligible black person is selected to a jury panel is *smaller* than the proportion of blacks in the general population.

Whether we choose to consider a one-sided or a two-sided alternative hypothesis depends largely on the research questions. In general, a one-sided hypothesis would be used when deviations from the null hypothesis only in one direction would be interesting and/or surprising. This draws on background information about the variables. A two-sided alternative hypothesis is neutral in this respect. Partly for this reason, two-sided hypotheses are in practice used more often than one-sided ones. Choosing a two-sided alternative hypothesis is not wrong even when a one-sided one could also be considered; this will simply lead to a more cautious (conservative) approach in that it takes stronger evidence to reject the null hypothesis when the alternative is two-sided than when it is one-sided. Such conservatism is typically regarded as a desirable feature in statistical inference (this will be discussed further in Section 7.6.1).

The two-sided alternative hypothesis (5.5) is clearly the logical opposite of the null hypothesis (5.4): if Δ is not equal to 0, it must be “not equal” to 0. So a two-sided alternative hypothesis must correspond to a “point” null hypothesis (5.4). For a one-sided alternative hypothesis, the same logic would seem to imply that the null hypothesis should also be one-sided: for example, $H_0 : \Delta \leq 0$ and $H_a : \Delta > 0$ would form such a logical pair. Often such “one-sided” null hypothesis is also closest to our research questions: for example, it would seem more interesting to try to test the hypothesis that the proportion of Yes-voters is less than or equal to 0.5 than that it is exactly 0.5. It turns out, however, that when the alternative hypothesis is, say, $H_a : \Delta > 0$, the test will be the same when the null hypothesis is $H_0 : \Delta \leq 0$ as when it is $H_0 : \Delta = 0$, and rejecting or not rejecting one of them is equivalent to rejecting or not rejecting the other. We can thus here always take the null hypothesis to be technically of the form (5.4), even if we are really interested in a corresponding “one-sided” null hypothesis. It is then only the alternative hypothesis which is explicitly either two-sided or one-sided.

5.5.2 The test statistic

The test statistic used to test hypotheses of the form (5.2) is the **z-test statistic**

$$z = \frac{\hat{\Delta}}{\hat{\sigma}_{\hat{\Delta}}} = \frac{\text{Estimate of the population difference } \Delta}{\text{Estimated standard error of the estimate of } \Delta}. \quad (5.8)$$

The statistic is introduced first in this form in order to draw attention to its generality. Null hypotheses in many ostensibly different situations can be formulated as hypotheses of the form (5.4) about population differences of some kind, and each can be tested with the test statistic (5.8). For example, all of the test statistics discussed in Chapters 5, 7 and 8 of this course pack will be of this type (but the χ^2 test statistic of Chapter 4 is not). The principles of the use and interpretation of the test that are introduced in this section apply almost unchanged also in these other contexts, and only the exact formulas for calculating $\hat{\Delta}$ and $\hat{\sigma}_{\hat{\Delta}}$ will need to be defined separately for each of them. In some applications considered in Chapter 7 the test statistic is typically called the **t-test statistic** instead of the *z*-test statistic, but its basic idea is still the same.

In (5.8), $\hat{\Delta}$ denotes a sample estimate of Δ . For a test of a single proportion, this is

$$\hat{\Delta} = \hat{\pi} - \pi_0, \quad (5.9)$$

i.e. the difference between the sample proportion and π_0 . This is the core of the test statistic. Although the forms of the two statistics seem rather different, (5.9) contains the comparison of the observed and expected sample values that was also at the heart of the χ^2 test statistic (see formula at end of Section 4.3.3) in Chapter 4. Here the “observed value” is the sample estimate $\hat{\pi}$ of the probability parameter, “expected value” is the value π_0 claimed for it by the null hypothesis, and $\hat{\Delta} = \hat{\pi} - \pi_0$ is their difference. (Equivalently, we could also say that the expected value of $\Delta = \pi - \pi_0$ under the null hypothesis (5.4) is 0, its observed value is $\hat{\Delta}$, and $\hat{\Delta} = \hat{\Delta} - 0$ is their difference.)

If the null hypothesis was true, we would expect the observed difference $\hat{\Delta}$ to be close to 0. If, on the other hand, the true π was different from π_0 , we would expect the same to be true of $\hat{\pi}$ and thus $\hat{\Delta}$ to be different from 0. In other words, the difference $\hat{\Delta} = \hat{\pi} - \pi_0$ tends to be small (close to zero) when the null hypothesis is true, and large (far from zero) when it is not true, thus satisfying one of the requirements for a good test statistic that were stated at the beginning of Section 4.3.4. (Whether in this we count as “large” both large positive and large negative values, or just one or the other, depends on the form of the alternative hypothesis, as explained in the next section.)

The $\hat{\sigma}_{\hat{\Delta}}$ in (5.8) denotes an estimate of the standard deviation of the sampling distribution of $\hat{\Delta}$, which is also known as the estimated **standard error** of $\hat{\Delta}$. For the test statistic (5.8), it is evaluated under the null hypothesis. The concept of a standard error of an estimate will be discussed in more detail in Section 6.4. Its role in the test statistic is to provide an interpretable scale for the size of $\hat{\Delta}$, so that the sampling distribution discussed in the next section will be of a convenient form.

For a test of the hypothesis (5.2) about a single proportion, the estimated standard error under the null hypothesis is

$$\hat{\sigma}_{\hat{\Delta}} = \sqrt{\frac{\pi_0(1 - \pi_0)}{n}}, \quad (5.10)$$

and the specific formula of the test statistic (5.8) is then

$$z = \frac{\hat{\pi} - \pi_0}{\sqrt{\pi_0(1 - \pi_0)/n}}. \quad (5.11)$$

This is the **one-sample *z*-test statistic for a population proportion**.

In Example 5.1 we have $\hat{\pi} = 0.47$, $\pi_0 = 0.5$, and $n = 702$, so

$$z = \frac{\hat{\pi} - \pi_0}{\sqrt{\pi_0(1 - \pi_0)/n}} = \frac{0.47 - 0.50}{\sqrt{0.50 \times (1 - 0.50)/702}} = -1.59.$$

Similarly, in Example 5.2 we have $\hat{\pi} = 0.0457$, $\pi_0 = 0.124$, $n = 4950$, and

$$z = \frac{0.0457 - 0.124}{\sqrt{0.124 \times (1 - 0.124)/4950}} = \frac{-0.0783}{\sqrt{0.10862/4950}} = -16.71.$$

Strangely, SPSS does not provide a direct way of calculating this value. However, since the formula (5.11) is very simple, we can easily calculate it with a pocket calculator, after first using SPSS to find out $\hat{\pi}$. This approach will be used in the computer classes.

5.5.3 The sampling distribution of the test statistic and P -values

Like the χ^2 test of Chapter 4, the z -test for a population proportion requires some conditions on the sample size in order for the approximate sampling distribution of the test statistic to be appropriate. These depend also on the value of π , which we can estimate by $\hat{\pi}$. One rule of thumb is that n should be larger than 10 divided by π or $1 - \pi$, whichever is smaller. When π is not very small or very large, e.g. if it is between 0.3 and 0.7, this essentially amounts to the condition that n should be at least 30. In the voting example 5.1, where $\hat{\pi} = 0.47$, the sample size of $n = 702$ is clearly large enough. In the jury example 5.2, $\hat{\pi} = 0.0457$ is much closer to zero, but since $10/0.0457$ is a little over 200, a sample of $n = 4950$ is again sufficient.

When the sample size is large enough, the sampling distribution of z defined by (5.11) is approximately the **standard normal distribution**. The probability curve of this distribution is shown in Figure 5.1. For now we just take it as given, and postpone a general discussion of the normal distribution until Chapter 7.

The P -value of the test is calculated from this distribution using the general principles introduced in Section 4.3.5. In other words, the P -value is the probability that the test statistic z has a value that is as or more extreme than the value of z in the observed sample. Now, however, the details of this calculation depend also on the alternative hypothesis, so some additional explanation is needed.

Consider first the more common case of a two-sided alternative hypothesis (5.5), that $\Delta \neq 0$. As discussed in the previous section, it is *large* values of the test statistic which indicate evidence against the null hypothesis, because a large z is obtained when the sample difference $\hat{\Delta} = \hat{\pi} - \pi_0$ is very different from the zero difference claimed by the null hypothesis. When the alternative is two-sided, “large” is taken to mean any value of z far from zero, i.e. either large positive or large negative values, because both indicate that the sample difference is far from 0. If z is large and positive, $\hat{\Delta}$ is much *larger* than 0. In example 5.1 this would indicate that a much larger proportion than 0.5 of the sample say they intend to vote Yes. If z is large and negative, $\hat{\Delta}$ is much *smaller* than 0, indicating a much smaller sample proportion than 0.5. Both of these cases would count as evidence against H_0 when the alternative hypothesis is two-sided.

The observed value of the z -test statistic in Example 5.1 was actually $z = -1.59$. Evidence would thus be “as strong” against H_0 as the observed z if we obtained a z -test statistic of -1.59 or 1.59 , the value exactly as far from 0 as the observed z but above rather than below 0. Similarly, evidence against the null would be even stronger if z was further from zero than 1.59 , i.e. larger than 1.59 or smaller than -1.59 . To obtain the P -value, we thus need to calculate the probability of observing a z -test statistic which is at most -1.59 or at least 1.59 when the null hypothesis is true in the population. In general, the P -value for testing the null hypothesis




Figure 5.1: Illustration of the calculation of P -values from the standard normal distribution. Here the value of the z -test statistic is $z = -1.59$ (as in the referendum example 5.1). The areas in grey indicate the two-sided P -values, i.e. the probabilities of values at least as far from 0 as the observed value of z .

against a two-sided alternative is the probability of obtaining a value at least z or at most $-z$ (when z is positive, vice versa when it is negative), where z here denotes the value of the test statistic in the sample. Such probabilities are calculated from the approximately standard normal sampling distribution of the test statistic under H_0 .

This calculation of the P -value is illustrated graphically in Figure 5.1. The curve in the plot is that of the standard normal distribution. Two areas are shown in grey under the curve, one on each tail of the distribution. The one on the left corresponds to values of -1.59 and smaller, and the one on the right to values of 1.59 or larger. Each of these areas is about 0.056 , and the P -value for a test against a two-sided alternative is their combined area, i.e. $P = 0.056 + 0.056 = 0.112$. This means that even if the true population proportion of Yes-voters was actually exactly 0.5 , there would be a probability of 0.112 of obtaining a test statistic as or more extreme than the $z = -1.59$ that was actually observed in Example 5.1.

In example 5.2 the observed test statistic was $z = -16.71$. The two-sided P -value is then the probability of values that are at most -16.71 or at least 16.71 . These areas are not shown in Figure 5.1 because they would not be visible in it. The horizontal axis of the figures runs from -4 to $+4$, so -16.71 is clearly far in the tail of the distribution and the corresponding probability is very small; we would report it as $P < 0.001$.

Consider now the case of a one-sided alternative hypothesis. For example, in the referendum example we might have decided beforehand to focus only on the possibility that the proportion of people who intend to vote Yes is smaller than 0.5 , and hence consider the alternative hypothesis that $\Delta < 0$. Two situations might then arise. First, suppose that the observed value of the sample difference is in the direction indicated by the alternative hypothesis. This is the case in the example, where the sample difference $\hat{\Delta} = -0.03$ is indeed smaller than zero, and the

test statistic $t = -1.59$ is negative. The possible values of z contributing to the P -value are then those of -1.59 or smaller. Values of 1.59 and larger are now not included, because positive values of the test statistic (corresponding to sample differences greater than 0) would not be regarded as evidence in favour of the claim that Δ is smaller than 0. The P -value is thus only the probability corresponding to the area on the left tail of the curve in Figure 5.1, and the corresponding area on the right tail is not included. Since both areas have the same size, the one-sided P -value is half the two-sided value, i.e. 0.056 instead of 0.112. In general, the one-sided P -value for a z -test of a proportion and other similar tests is always obtained by dividing the two-sided value by 2, if the sample evidence is in the direction of the one-sided alternative hypothesis.

The second case occurs when the sample difference is not in the direction indicated by a one-sided alternative hypothesis. For example, suppose that the sample proportion of Yes-voters had actually been 0.53, i.e. 0.03 larger than 0.5, so that we had obtained $z = +1.59$ instead. The possible values of the test statistic which contributed to the P -value would then be $z = 1.59$ and all smaller values. These are “as strong or stronger evidence against the null hypothesis and in the direction of the alternative hypothesis” as required by the definition at the beginning of Section 4.3.5, since they agree with the alternative hypothesis (negative values of z) or at least disagree with it less than the observed z (positive values from 0 to 1.59). In Figure 5.1, these values would correspond to the area under the whole curve, apart from the region to the right of 1.59 on the right tail. Since the probability of the latter is 0.056 and the total probability under the curve is 1, the required probability is $P = 1 - 0.056 = 0.944$. However, calculating the P -value so precisely is hardly necessary in this case, as it is clearly going to be closer to 1 than to 0. The conclusion from such a large P -value will always be that the null hypothesis should not be rejected. This is also intuitively obvious, as a sample difference in the opposite direction from the one claimed by the alternative hypothesis is clearly not to be regarded as evidence in favour of that alternative hypothesis. In short, if the sample difference is in a different direction than a one-sided alternative hypothesis, the P -value can be reported simply as $P > 0.5$ without further calculations.

If a statistical software package is used to carry out the test, it will also report the P -value and no further calculations are needed (except dividing a two-sided P -value by 2, if a one-sided value is needed and only a two-sided one is reported). However, since SPSS does not currently provide a procedure for this test, and for exam purposes, we will briefly outline how an approximate P -value is obtained using critical values from a table. This is done in a very similar way as for the χ^2 test in Section 4.3.5.

The first part of Table 5.2 shows a table of critical values for the standard normal distribution. These values are also shown in the Appendix at the end of this course pack, on the last row of a larger table (the other parts of this table will be explained later, in Section 7.3.4). A version of this table is included in all introductory text books on statistics, although its format may be slightly different in different books.

Table 5.2: A table of critical values for the standard normal distribution. The upper part of the table shows the critical values in one row, as in standard statistical tables (see the last row of the table in the Appendix). The lower part of the table includes the same numbers rearranged to show critical values for conventional significance levels for one- and two-sided tests.

	0.100	0.050	0.025	0.01	0.005	0.001	0.0005
Critical value	1.282	1.645	1.960	2.326	2.576	3.090	3.291

Alternative hypothesis	Significance levels			
	0.10	0.05	0.01	0.001
Two-sided	1.65	1.96	2.58	3.29
One-sided	1.28	1.65	2.33	3.09

The columns of the first part of Table 5.2 are labelled “Right-hand tail probabilities”, with separate columns for some values from 0.100 to 0.0005. This means that the probability that a value from the standard normal distribution is at least as large as the value given in a particular column is the number given at the top of that column. For example, the value in the column labelled “0.025” is 1.960, indicating that the probability of obtaining a value equal to or greater than 1.960 from the standard normal distribution is 0.025. Because the distribution is symmetric, the probability of values of at most -1.960 is also 0.025, and the total probability that a value is at least 1.960 units from zero is $0.025 + 0.025 = 0.05$.

These values can be used to obtain bounds for P -values, expressed in terms of conventional significance levels of 0.10, 0.05, 0.01 and 0.001. The values at which these tail probabilities are obtained are the corresponding critical values for the test statistic. They are shown in the lower part of Table 5.2, slightly rearranged for clarity of presentation and rounded to two decimal places (which is accurate enough for practical purposes). The basic idea of using the critical values is that if the observed (absolute value of) the z -test statistic is *larger* than a critical value (for the required kind of alternative hypothesis) shown in the lower part of Table 5.2, the P -value is *smaller* than the significance level corresponding to that critical value.

The table shows only positive critical values. If the observed test statistic is actually negative, its negative ($-$) sign is omitted and the resulting positive value (i.e. the absolute value of the statistic) is compared to the critical values. Note also that the critical value for a given significance level depends on whether the alternative hypothesis is two-sided or one-sided. In the one-sided case, the test statistic is compared to the critical values only if it is actually in the direction of the alternative hypothesis; if not, we can simply report $P > 0.5$ as discussed above.

The P -value obtained from the table is reported as being smaller than the smallest conventional significance level for which the corresponding critical value is exceeded by the observed test statistic. For instance, in the jury example 5.2 we have $z = -16.71$. Considering a two-sided alternative hypothesis, 16.71 is larger than the critical values 1.65, 1.96, 2.58 and 3.29 for all the standard significance levels, so we can report that $P < 0.001$. For Example 5.1, in contrast, $z = -1.59$, the absolute value of which is smaller than even the critical value 1.65 for the 10% significance level. For this example, we would report $P > 0.1$.

The intuitive idea of the critical values and their connection to the P -values is illustrated for Example 5.1 by Figure 5.1. Here the observed test statistic is $t = -1.59$, so the two-sided P -value is the probability of values at least 1.59 or at most -1.59 , which correspond to the two gray areas in the tails of the distribution. Also shown in the plot is one of the critical values for two-sided tests, the 1.96 for significance level 0.05. By definition of the critical values, the combined tail probability of values at least 1.96 from 0, i.e. the probability of values at least 1.96 or at most -1.96 , is 0.05. It is clear from the plot that since 1.59 is smaller than 1.96, these areas are smaller than the tail areas corresponding to 1.59 and -1.59 , and the combined area of the latter must be more than 0.05, i.e. it must be that $P > 0.05$. Similar argument for the 10% critical value of 1.65 shows that P is here also larger than 0.1.

5.5.4 Conclusions from the test

The general principles of drawing and stating conclusions from a significance test have already been explained in Section 4.3.6, so they need not be repeated here. Considering two-sided alternative hypotheses, the conclusions in our two examples are as follows:

- In the referendum example 5.1, $P = 0.112$ for the null hypothesis that $\pi = 0.5$ in the population of eligible voters. The null hypothesis is not rejected at conventional levels of significance. There is not enough evidence to conclude that the proportion of voters who definitely intend to vote Yes differs from one half.
- In the jury example 5.2, $P < 0.001$ for the null hypothesis that $\pi = 0.124$. The null hypothesis is thus overwhelmingly rejected at any conventional level of significance. There is very strong evidence that the probability of a black person being selected to the jury pool differs from the proportion of blacks in the population of the county.

5.5.5 Summary of the test

As a summary, let us again repeat the main steps of the test described in this section in a concise form, using the voting variable of Example 5.1 for illustration:

1. Data: a sample of size $n = 702$ of a dichotomous variable Y with values 1 (Yes) and 0 (No or undecided), with the sample proportion of ones $\hat{\pi} = 0.47$.
2. Assumptions: the observations are a random sample from a population distribution with some population proportion (probability) π , and the sample size n is large enough for the test to be valid (for example, $n \geq 30$ when π_0 is between about 0.3 and 0.7, as it is here).
3. Hypotheses: null hypothesis $H_0 : \pi = \pi_0$ against the alternative hypothesis $H_a : \pi \neq \pi_0$, where $\pi_0 = 0.5$.
4. The test statistic: the z -statistic

$$z = \frac{\hat{\pi} - \pi_0}{\sqrt{\pi_0(1 - \pi_0)/n}} = \frac{0.47 - 0.50}{\sqrt{0.50 \times (1 - 0.50)/702}} = -1.59.$$

5. The sampling distribution of the test statistic when H_0 is true: a standard normal distribution.
6. The P -value: the probability that a randomly selected value from the the standard normal distribution is at most -1.59 or at least 1.59 , which is $P = 0.112$.
 - If the precise P -value is not available, we can observe that 1.59 is smaller than the two-sided critical value 1.65 for the 10% level of significance. Thus it must be that $P > 0.1$.
7. Conclusion: The null hypothesis is not rejected ($P = 0.112$). There is not enough evidence to conclude that the proportion of eligible voters who definitely intend to vote Yes differs from one half. Based on this opinion poll, the referendum remains too close to call.

5.6 Confidence interval for a single proportion

5.6.1 Introduction

A significance test assesses whether it is plausible, given the evidence in the observed data, that a population parameter or parameters have a specific set of values claimed by the null hypothesis. For example, in Section 5.5 we asked such a question about the probability parameter of a binary variable in a single population.

In many ways a more natural approach would be try to identify all of those values of a parameter which *are* plausible given the data. This leads to a form of statistical inference known as **interval estimation**, which aims to present not only a single best guess (i.e. a point estimate) of a population parameter, but also a range of plausible values (an **interval estimate**) for it. Such an interval is known as a **confidence interval**. This section introduces the idea of confidence intervals, and shows how to construct them for a population probability. In later sections, the same principles will be used to calculate confidence intervals for other kinds of population parameters.

Interval estimation is an often underused part of statistical inference, while significance testing is arguably overused or at least often misused. In most contexts it would be useful to report confidence intervals in addition to, or instead of, results of significance tests. This is not done often enough in research publications in the social sciences.

5.6.2 Calculation of the interval

Our aim is again to draw inference on the difference $\Delta = \pi - \pi_0$ or, equivalently, the population probability π . The point estimate of Δ is $\hat{\Delta} = \hat{\pi} - \pi_0$ where $\hat{\pi}$ is the sample proportion corresponding to π . Suppose that the conditions on the sample size n that were discussed in Section 5.5.3 are again satisfied.

Consider now Figure @ref(fig:f-pval-prob)}. One of the results illustrated by it is that if π_0 is the true value of the population probability π , so that $\Delta = \pi - \pi_0 = 0$, there is a probability of 0.95 that for a randomly drawn sample from the population the z -test statistic $z = \hat{\Delta}/\hat{\sigma}_{\hat{\Delta}}$ is between -1.96 and $+1.96$. This also implies that the probability is 0.95 that in such a sample the observed value of $\hat{\Delta}$ will be between $\Delta - 1.96\hat{\sigma}_{\hat{\Delta}}$ and $\Delta + 1.96\hat{\sigma}_{\hat{\Delta}}$. Furthermore, it is clear from the figure that all of the values within this interval are more likely to occur than any of the values outside the interval (i.e. those in the two tails of the sampling distribution). The interval thus seems like a sensible summary of the “most likely” values that the estimate $\hat{\Delta}$ may have in random samples.

A confidence interval essentially turns this around, into a statement about the unknown true value of Δ in the population, even in cases where Δ is not 0. This is done by substituting $\hat{\Delta}$ for Δ above, to create the interval

$$\text{from } \hat{\Delta} - 1.96 \times \hat{\sigma}_{\hat{\Delta}} \text{ to } \hat{\Delta} + 1.96 \times \hat{\sigma}_{\hat{\Delta}}. \quad (5.12)$$

This is the **95 % confidence interval** for the population difference Δ . It is usually written more concisely as

$$\hat{\Delta} \pm 1.96 \hat{\sigma}_{\hat{\Delta}} \quad (5.13)$$

where the “plusminus” symbol \pm indicates that we calculate the two endpoints of the interval as in (5.12), one below and one above $\hat{\Delta}$.

Expression (5.13) is general in the sense that many different quantities can take the role of Δ in it. Here we consider for now the case of $\Delta = \pi - \pi_0$. The estimated standard error $\hat{\sigma}_{\hat{\Delta}}$ is analogous to (5.10) used for the z -test, but not the same. This is because the confidence interval is not calculated under the null hypothesis $H_0 : \pi = \pi_0$, so we cannot use π_0 for π in the standard error. Instead, π is estimated by the sample proportion $\hat{\pi}$, giving the estimated standard error

$$\hat{\sigma}_{\hat{\Delta}} = \sqrt{\frac{\hat{\pi}(1 - \hat{\pi})}{n}} \quad (5.14)$$

and thus the 95% confidence interval

$$(\hat{\pi} - \pi_0) \pm 1.96 \sqrt{\frac{\hat{\pi}(1 - \hat{\pi})}{n}}$$

for $\Delta = \pi - \pi_0$. Alternatively, a confidence interval for π itself is given by

$$\hat{\pi} \pm 1.96 \sqrt{\frac{\hat{\pi}(1 - \hat{\pi})}{n}}. \quad (5.15)$$

This is typically the most useful interval for use in practice. For instance, in the referendum example 5.1 this gives a 95% confidence interval of

$$0.470 \pm 1.96 \times \sqrt{\frac{0.470 \times (1 - 0.470)}{702}} = 0.470 \pm 0.0369 = (0.433, 0.507)$$

for the proportion of definite Yes-voters in the population. Similarly, in Example 5.2 the 95% confidence interval for the probability of a black person being selected for the jury pool is (0.040, 0.052). These intervals are also shown in Table 5.1.

5.6.3 Interpretation of confidence intervals

As with the P -value of a significance test, the precise interpretation of a confidence interval refers to probabilities calculated from a sampling distribution, i.e. probabilities evaluated from a hypothetical exercise of repeated sampling:

- If we obtained many samples from the population and calculated the confidence interval for each such sample using the formula (5.15), approximately 95% of these intervals would contain the true value of the population proportion π .

This is undeniably convoluted, even more so than the precise interpretation of a P -value. In practise a confidence interval would not usually be described in exactly these words. Instead, a research report might, for example, write that (in the referendum example) “the 95 % confidence interval for the proportion of eligible voters in the population who definitely intend to vote Yes is (0.433, 0.507)”, or that “we are 95 % confident that the proportion of eligible voters in the population who definitely intend to vote Yes is between 43.3% and 50.7%”. Such a statement in effect assumes that the readers will be familiar enough with the idea of confidence intervals to understand the claim. It is nevertheless useful to be aware of the more careful interpretation of a confidence interval, if only to avoid misunderstandings. The most common error is to claim

that “there is a 95% probability that the proportion in the population is between 0.433 and 0.507”. Although the difference to the interpretation given above may seem small, the latter statement is not really true, or strictly speaking even meaningful, in the statistical framework considered here.

In place of the 1.96 in (5.13), we may also use other numbers. To allow for this in the notation, we can also write

$$\hat{\Delta} \pm z_{\alpha/2} \hat{\sigma}_{\hat{\Delta}}. \quad (5.16)$$

where the multiplier $z_{\alpha/2}$ is a number which depends on two things. One of them is the sampling distribution of $\hat{\Delta}$, which is here assumed to be the normal distribution (another possibility is discussed in Section 7.3.4). The second is the **confidence level** which we have chosen for the confidence interval. For example, the probability of 0.95 in the interpretation of a 95% confidence interval discussed above is the confidence level of that interval. Conventionally the 0.95 level is most commonly used, while other standard choices are 0.90 and 0.99, i.e. 90% and 99% confidence intervals.

In the symbol $z_{\alpha/2}$, α is a number such that $1 - \alpha$ equals the required confidence level. In other words, $\alpha = 0.1$, 0.05 , and 0.01 for confidence levels of $1 - \alpha = 0.90$, 0.95 and 0.99 respectively. The values that are required for the conventional levels are $z_{0.10/2} = z_{0.05} = 1.64$, $z_{0.05/2} = z_{0.025} = 1.96$, and $z_{0.01/2} = z_{0.005} = 2.58$, which correspond to intervals at the confidence levels of 90%, 95% and 99% respectively. These values are also shown in Table 5.4.

Table 5.4: Multipliers $z_{\alpha/2}$ used to obtain confidence intervals based on the normal distribution, for three standard confidence levels. These values are substituted for $z_{\alpha/2}$ in formula (5.16) to obtain the confidence interval.

	Confidence levels		
	90%	95%	99%
Multiplier $z_{\alpha/2}$	1.64	1.96	2.58

A confidence interval contains, loosely speaking, those numbers which are considered plausible values for the unknown population difference Δ in the light of the evidence in the data. The *width* of the interval thus reflects our uncertainty about the exact value of Δ , which in turn is related to the amount of information the data provide about Δ . If the interval is wide, many values are consistent with the observed data, so there is still a large amount of uncertainty; if the interval is narrow, we have much information about Δ and thus little uncertainty. Another way of stating this is that when the confidence interval is narrow, estimates of Δ are very *precise*.

The width of the interval (5.16) is $2 \times z_{\alpha/2} \times \hat{\sigma}_{\hat{\Delta}}$. This depends on

- The confidence level: the higher the level, the wider the interval. Thus a 99% confidence interval is always wider than a 95% interval for the same data, and wider still than a 90% interval. This is logically inevitable: if we want to state with high level of confidence that a parameter is within a certain interval, we must allow the interval to contain a wide range of values. It also explains why we do not consider a 100% confidence interval: this would contain all possible values of Δ and exclude none, making no use of the data at all. Instead, we aim for a high but not perfect level of confidence, obtaining an interval which contains some but not all possible values, for the price of a small chance of incorrect conclusions.

- The standard error $\hat{\sigma}_{\hat{\Delta}}$, which in the case of a single proportion is (5.14). This in turn depends on
 - the sample size n : the larger this is, the narrower the interval. Increasing the sample size thus results (other things being equal) in reduced uncertainty and higher precision.
 - the true population proportion π : the closer this is to 0.5, the wider the interval. Unlike the sample size, this determinant of the estimation uncertainty is not in our control.

Opinion polls of the kind illustrated by the referendum example are probably where non-academic audiences are most likely to encounter confidence intervals, although not under that label. Media reports of such polls typically include a *margin of error* for the results. For example, in the referendum example it might be reported that 47% of the respondents said that they would definitely vote Yes, and that “the study has a margin of error of plus or minus four percentage points”. In most cases the phrase “margin of error” refers to a 95% confidence interval. Unless otherwise mentioned, we can thus take a statement like the one above to mean that the 95% confidence interval for the proportion of interest is approximately 47 ± 4 percentage points. For a realistic interpretation of the implications of the results, the width of this interval is at least as important as the point estimate of the proportion. This is often neglected in media reports of opinion polls, where the point estimate tends to be headline news, while the margin of error is typically mentioned only in passing or omitted altogether.

5.6.4 Confidence intervals vs. significance tests

There are some obvious similarities between the conclusions from significance tests and confidence intervals. For example, a z -test in the referendum example 5.1 showed that the null hypothesis that the population proportion π was 0.5 was not rejected ($P = 0.112$). Thus 0.5 is a plausible value for π in light of the observed data. The 95% confidence interval for π showed that, at this level of confidence, plausible values for π are those between 0.433 and 0.507. In particular, these include 0.5, so the confidence interval also indicates that a proportion of 0.5 is plausible. This connection between the test and the confidence interval is in fact exact:

- If the hypothesis $H_0 : \Delta = 0$ about a population quantity Δ is rejected at the 5% level of significance using the z -test against a *two-sided* alternative hypothesis, the 95 % confidence interval for Δ will not contain 0, and vice versa. Similarly, if H_0 is not rejected, the confidence interval will contain 0, and vice versa.

The same is true for other matching pairs of levels of significance and confidence, e.g. for a test with a 1% level of significance and a 99% (i.e. (100-1)%) confidence interval. In short, the significance test and the confidence interval will in these cases always give the same answer about whether or not a parameter value is plausible (consistent with the data) at a given level of significance/confidence.

These pairs of a test and an interval are exactly comparable in that they concern the same population parameter, estimate all parameters in the same way, use the same sampling distribution for inference, and use the same level of significance/confidence. Not all tests and confidence intervals have exact pairs in this way. Also, some tests are for hypotheses about more than one parameter at once, so there is no corresponding single confidence interval. Nevertheless, the connection stated above is useful for understanding the ideas of both tests and confidence intervals.

These results also illustrate how confidence intervals are inherently more informative than significance tests. For instance, in the jury example 5.2, both the test and the confidence interval agree on the implausibility of the claim that the population probability of being selected to the jury panel is the same as the proportion (0.124) of black people in the population, since the claim that $\pi = 0.124$ is rejected by the test (with $P < 0.001$) and outside the interval (0.040; 0.052). Unlike the test, however, the confidence interval summarizes the plausibility of *all* possible values of π and not just $\pi_0 = 0.124$. One way to describe this is to consider what would have happened if we had carried out a series of significance tests of null hypotheses of the form $H_0 : \pi = \pi_0$ for a range of values of π_0 . The confidence interval contains all those values π_0 which would not have been rejected by the test, while all the values outside the interval would have been rejected. Here $H_0 : \pi = \pi_0$ would thus not have been rejected at the 5% level if π_0 had been between 0.040 and 0.052, and rejected otherwise. This, of course, is not how significance tests are actually conducted, but it provides a useful additional interpretation of confidence intervals.

A confidence interval is particularly useful when the parameter of interest is measured in familiar units, such as the proportions considered so far. We may then try to judge, in substantive terms, how wide the interval is and how far it is from particular values of interest. In the jury example the 95% confidence interval ranges from 4.0% to 5.2%, which suggests that the population probability is estimated fairly precisely by this survey. The interval also reveals that even its upper bound is less than half of the figure of 12.4% which would correspond to proportional representation of blacks in the jury pool, a result which suggests quite substantial underrepresentation in the pool.

5.7 Inference for comparing two proportions

In Examples 5.3 and 5.4, the aim is to compare the proportion of a dichotomous response variable Y between two groups of a dichotomous explanatory variable X :

- Example 5.3: compare the proportion of polio cases among the unvaccinated (π_1) and vaccinated (π_2) children.
- Example 5.4: compare the proportion of optimistic responses to a negative (π_1) vs. positive wording of the question (π_2).

The quantity of interest is then the population difference

$$\Delta = \pi_2 - \pi_1. \quad (5.17)$$

For a significance test of this, the null hypothesis will again be $H_0 : \Delta = 0$, which is in this case equivalent to the hypothesis of equal proportions

$$H_0 : \pi_1 = \pi_2. \quad (5.18)$$

The null hypothesis thus claims that there is no association between the group variable X and the dichotomous response variable Y , while the alternative hypothesis (e.g. the two-sided one $H_a : \pi_1 \neq \pi_2$, i.e. $H_a : \Delta \neq 0$) implies that there is an association.

The obvious estimates of π_1 and π_2 are the corresponding sample proportions $\hat{\pi}_1$ and $\hat{\pi}_2$, calculated from samples of sizes n_1 and n_2 respectively, and the estimate of Δ is then

$$\hat{\Delta} = \hat{\pi}_2 - \hat{\pi}_1. \quad (5.19)$$

This gives $\hat{\Delta} = 0.000284 - 0.000706 = -0.000422$ in Example 5.3 and $\hat{\Delta} = 0.364 - 0.279 = 0.085$ in Example 5.4. In the samples, the proportion of polio cases is thus lower in the vaccinated group, and the proportion of optimistic answers is higher in response to a positively worded question. Note also that although the inference discussed below focuses on the difference of the proportions, for purely descriptive purposes we might prefer to use some other statistic, such as the ratio of the proportions. For example, the difference of 0.000422 in polio incidence between vaccine and control groups may seem small, because the proportions in both groups are small. A better idea of the the magnitude of the contrast is given by their ratio of $0.000706/0.000284 = 2.49$ (this is known as the *risk ratio*). In other words, the rate of polio infection in the unvaccinated group was two and a half times the rate in the vaccinated group.

The tests and confidence intervals discussed below are again based on the assumption that the relevant sampling distributions are approximately normal, which is true when the sample sizes n_1 and n_2 are large enough. The conditions for this are not very demanding: one rule of thumb states that the methods described in this section are reasonably valid if in both groups the number of observations with Y having the value 1, and of ones with the value 0, are both more than 5. This condition is satisfied in both of the examples considered here.

The validity of the test, as well as the amount of information the data provide about π_1 and π_2 in general, thus depends not just on the overall sample sizes but on having enough observations of both values of Y . The critical quantity is then the number of observations in the rarer category of Y . In Example 5.3 this means the numbers of children diagnosed with polio, because the probability of polio was low in the study population. The numbers of eventual polio cases were 142 and 57 in the control and treatment groups respectively, so the rule of thumb stated above was satisfied. With such low probabilities of polio incidence, sufficient numbers of cases were achieved only by making the overall sample sizes n_1 and n_2 large enough. That is why the trial had to be very large, involving hundreds of thousands of participants.

The standard error of $\hat{\Delta}$ is

$$\sigma_{\hat{\Delta}} = \sqrt{\frac{\pi_2(1 - \pi_2)}{n_2} + \frac{\pi_1(1 - \pi_1)}{n_1}}. \quad (5.20)$$

As in the one-sample case above, the best way to estimate this is different for a significance test than for a confidence interval. For a test, the standard error can be estimated under assumption that the null hypothesis (5.18) is true, in which case the population proportion is the same in both groups. A good estimate of this common proportion, denoted below by $\hat{\pi}$, is the proportion of observations with value 1 for Y in the total sample of $n_1 + n_2$ observations, pooling observations from both groups together; expressed in terms of the group-specific estimates, this is

$$\hat{\pi} = \frac{n_1 \hat{\pi}_1 + n_2 \hat{\pi}_2}{n_1 + n_2}. \quad (5.21)$$

Using this for both π_1 and π_2 in (5.20) gives the estimated standard error

$$\hat{\sigma}_{\hat{\Delta}} = \sqrt{\hat{\pi}(1 - \hat{\pi}) \left(\frac{1}{n_2} + \frac{1}{n_1} \right)}, \quad (5.22)$$

and using (5.19) and (5.22) in the general formula (5.8) gives the **two-sample z -test statistic for proportions**

$$z = \frac{\hat{\pi}_2 - \hat{\pi}_1}{\sqrt{\hat{\pi}(1 - \hat{\pi})(1/n_2 + 1/n_1)}} \quad (5.23)$$

where $\hat{\pi}$ is given by (5.21). When the null hypothesis is true, the sampling distribution of this test statistic is approximately standard normal when the sample sizes are large enough.

For a confidence interval, the calculation of the estimated standard error cannot assume that (5.18) is true. Instead, we use the estimate

$$\hat{\sigma}_{\hat{\Delta}} = \sqrt{\frac{\hat{\pi}_2(1 - \hat{\pi}_2)}{n_2} + \frac{\hat{\pi}_1(1 - \hat{\pi}_1)}{n_1}} \quad (5.24)$$

and, substituting this to the general formula (5.16), we get

$$(\hat{\pi}_2 - \hat{\pi}_1) \pm z_{\alpha/2} \sqrt{\frac{\hat{\pi}_2(1 - \hat{\pi}_2)}{n_2} + \frac{\hat{\pi}_1(1 - \hat{\pi}_1)}{n_1}} \quad (5.25)$$

as the confidence interval for $\Delta = \pi_2 - \pi_1$, with confidence level $1 - \alpha$.

For an illustration of the calculations, consider Example 5.4. Denoting the group of respondents answering the negatively worded question by 1 and those with the positive question by 2, the basic quantities are $n_1 = 921$, $\hat{\pi}_1 = 0.279$, $n_2 = 929$ and $\hat{\pi}_2 = 0.364$. The estimated difference in the proportions of respondents giving an optimistic answer is thus

$$\hat{\Delta} = \hat{\pi}_2 - \hat{\pi}_1 = 0.364 - 0.279 = 0.085.$$

For a significance test, the estimated standard error of $\hat{\Delta}$ uses the pooled estimate (5.21) of the population proportion, which is given by

$$\hat{\pi} = \frac{921 \times 0.279 + 929 \times 0.364}{921 + 929} = \frac{257 + 338}{921 + 929} = 0.322.$$

The standard error from (5.22) is then

$$\hat{\sigma}_{\hat{\Delta}} = \sqrt{0.322 \times (1 - 0.322) \times \left(\frac{1}{929} + \frac{1}{921} \right)} = \sqrt{\frac{0.2182}{462.5}} = 0.0217,$$

and the test statistic (5.23) is

$$z = \frac{0.085}{0.0217} = 3.92.$$

For the confidence interval, the standard error of $\hat{\Delta}$ is estimated from (5.24) as

$$\begin{aligned} \hat{\sigma}_{\hat{\Delta}} &= \sqrt{\frac{0.364 \times (1 - 0.364)}{929} + \frac{0.279 \times (1 - 0.279)}{921}} \\ &= \sqrt{\frac{0.2315}{929} + \frac{0.2012}{921}} = 0.0216 \end{aligned}$$

and a 95% confidence interval from (5.25) is

$$0.085 \pm 1.96 \times 0.0216 = 0.085 \pm 0.042 = (0.043; 0.127).$$

The P -value for the test statistic is clearly very low (in fact about 0.00009), so the null hypothesis of equal proportions is convincingly rejected. There is very strong evidence that the probability that a respondent will give an answer indicating optimism for the future is different for the two differently worded questions. The confidence interval indicates that we are 95% confident that the proportion of optimistic answers is between 4.3 and 12.7 percentage points higher when the question is worded positively than when it is worded negatively. This suggests quite a substantial acquiescence bias arising from changing just one word in the survey question, as described in the introduction to Example 5.4 at the beginning of this chapter.

In Example 5.3, the estimated difference is $\hat{\Delta} = -0.000422$ (see Table 5.1, i.e. 422 fewer polio cases per million children in the vaccinated group than in the unvaccinated group. Similar calculations as above show that the value of the test statistic is $z = -6.01$, so the P -value is again very small (in fact about 0.000000001) and the null hypothesis of equal probabilities is strongly rejected. There is thus overwhelming evidence that the proportion of polio cases was different among the vaccinated children than among the unvaccinated ones. The 95% confidence interval for the difference shows that we are 95% confident that this difference was a reduction of between 284 and 560 polio cases per million children⁶. This was acknowledged as a convincing demonstration that the Salk vaccine worked (see Figure 5.2), and it (and later other types of polio vaccination) was soon put to widespread use. The resulting dramatic decline in the incidence of polio is one of the great success stories of modern medicine. Compared to the 199 children with polio in 1954 among the less than half a million participants of the vaccine trial alone, in 2014 there were 414 confirmed cases of polio in the whole world (see www.polioeradication.org/Dataandmonitoring/Poliothisweek.aspx). There is hope that that number will reach 0 in a not-too-distant future, so that the once devastating disease will one day be entirely eradicated.

⁶The data can be obtained from <http://bes2009-10.org/>, which gives further information on the survey, including the full text of the questionnaires. The data analysed in this class and homework are from the BES Campaign Internet Panel Survey, which has been divided into two data sets corresponding to two time periods leading up to the General Election.



Figure 5.2: Public reaction to statistical inference.

Chapter 6

Continuous variables: Population and sampling distributions

6.1 Introduction

This chapter serves both as an explanation of some topics that were skimmed over previously, and as preparation for later chapters. Its central theme is probability distributions of continuous variables. These may appear in two distinct roles:

- As population distributions of continuous variables, for instance blood pressure in the illustrative example of this chapter. This contrasts with the kinds of discrete variables that were considered in Chapters 4 and 5. Methods of inference for continuous variables will be introduced in Chapters 7 and 8.
- As sampling distributions of sample statistics. These are typically continuous even in analyses of discrete variables, such as in Chapter 5 where the variable of interest Y was binary but the sampling distributions of a sample proportion $\hat{\pi}$ and the z -test statistic for population probability π were nevertheless continuous. We have already encountered two continuous distributions in this role, the χ^2 distributions in Chapter 4 and the standard normal distribution in Chapter 5. Their origins are explained in more detail below.

To illustrate the concepts, we use data from the Health Survey for England 2002 (HES).¹ One part of the survey was a short physical examination by a nurse. Figure 6.1 shows a histogram and frequency polygon of diastolic blood pressure (in mm Hg) for 4489 respondents, measured by the mean of the last two of three measurements taken during the examination. Data from respondents for whom the measurements were not obtained or were considered invalid have been excluded. Respondents aged under 25 have also been excluded for simplicity, because this age group was oversampled in the 2002 HES.

The respondents whose blood pressures are summarized in Figure 6.1 are in reality a sample from a larger population in the sense of Sections 3.2 and 3.3. However, for illustrative purposes we will here pretend that they are actually an entire finite population of 4489 people (the adults in a small town, say). The values summarised in Figure 6.1 then form the population distribution of blood pressure in this population. It is clear that blood pressure is best treated as a continuous variable.

¹ESS Round 5: European Social Survey Round 5 Data (2010). Data file edition 2.0. Norwegian Social Science Data Services, Norway Data Archive and distributor of ESS data. The full data can be obtained from <http://ess.nsd.uib.no/ess/round5/>.



Figure 6.1: Histogram of diastolic blood pressure, with the corresponding frequency polygon, from Health Survey for England 2002 (respondents aged 25 or over, $n = 4489$).

6.2 Population distributions of continuous variables

6.2.1 Population parameters and their point estimates

If we knew all of its values, we could summarise a finite population distribution by, say, a histogram like Figure 6.1. We can also consider specific characteristics of the distribution, i.e. its *parameters* in the sense introduced in Section 5.3. For the distribution of a continuous variable, the most important parameters are the **population mean**

$$\mu = \frac{\sum Y_i}{N} \quad (6.1)$$

and the **population variance**

$$\sigma^2 = \frac{\sum (Y_i - \mu)^2}{N} \quad (6.2)$$

or, instead of the variance, the **population standard deviation**

$$\sigma = \sqrt{\frac{\sum (Y_i - \mu)^2}{N}}. \quad (6.3)$$

Here μ and σ are the lower-case Greek letters “mu” and “sigma” respectively, and σ^2 is read “sigma squared”. It is common to use Greek letters for population parameters, as we did also for the probability parameter π in Chapter 5.

In (6.1)–(6.3), N is the number of units in a finite population and the sums indicated by Σ are over all of these N units. If we treat the data in Figure 6.1 as a population, $N = 4489$ and these population parameters are $\mu = 74.2$, $\sigma^2 = 127.87$ and $\sigma = 11.3$.

Because the formulas (6.1)–(6.3) involve the population size N , they apply in this exact form only to finite populations like the one in this example (and as discussed more generally in Section 3.2) but not to infinite ones of the kind discussed in Section 3.4. However, the definitions of μ , σ^2 , σ and other parameters can be extended to apply also to infinite populations. These definitions, which will be omitted here, involve the concept of continuous probability distributions that is discussed in the next section. The interpretations of the population parameters turn out to be intuitively similar for both the finite and infinite-population cases, and the same methods of analysis apply to both, so we can here ignore the distinction without further comment.

The population formulas (6.1)–(6.3) clearly resemble those of some sample statistics introduced in Chapter 2, specifically the *sample* mean, variance and standard deviation

$$\bar{Y} = \frac{\sum Y_i}{n}, \quad (6.4)$$

$$s^2 = \frac{\sum (Y_i - \bar{Y})^2}{n - 1} \quad (6.5)$$

and

$$s = \sqrt{\frac{\sum (Y_i - \bar{Y})^2}{n - 1}} \quad (6.6)$$

where the sums are now over the n observations in a sample. These can be used as descriptions of the sample distribution as discussed in Chapter 2, but also as *point estimates* of the corresponding population parameters in the sense defined in Section 5.4. We may thus use the sample mean \bar{Y} as a point estimate of the population mean μ , and the sample variance s^2 and sample standard deviation s as point estimates of population variance σ^2 and standard deviation σ respectively. These same estimates can be used for both finite and infinite population distributions.

For further illustration of the connection between population and sample quantities, we have also drawn a simple random sample of $n = 50$ observations from the finite population of $N = 4489$ observations in Figure 6.1. Table 6.1 shows the summary statistics (6.4)–(6.6) in this sample and the corresponding parameters (6.1)–(6.3) in the population.

Table 6.1: Summary statistics for diastolic blood pressure in the population and a sample from it in the example used for illustration in Sections 6.2–6.4.

	Size	Mean	Standard Deviation	Variance
Population	$N = 4489$	$\mu = 74.2$	$\sigma = 11.3$	$\sigma^2 = 127.87$
Sample	$n = 50$	$\bar{Y} = 72.6$	$s = 12.7$	$s^2 = 161.19$

You may have noticed that the formulas of the sample variance (6.5) and sample standard deviation (6.6) involve the divisor $n - 1$ rather than the n which might seem more natural, while the population formulas (6.2) and (6.3) do use N rather than $N - 1$. The reason for this

is that using $n - 1$ gives the estimators certain mathematically desirable properties (s^2 is an *unbiased* estimate of σ^2 , but $\hat{\sigma}^2$ below is not). This detail need not concern us here. In fact, the statistics which use n instead, i.e.

$$\hat{\sigma}^2 = \frac{\sum (Y_i - \bar{Y})^2}{n} \quad (6.7)$$

for σ^2 and $\hat{\sigma} = \sqrt{\hat{\sigma}^2}$ for σ , are also sensible estimates and very similar to s^2 and s unless n is very small. In general, there are often several possible sample statistics which could be used as estimates for the same population parameter.

6.3 Probability distributions of continuous variables

6.3.1 General comments

Thinking about population distributions of continuous distributions using, say, histograms as in Figure 6.1 would present difficulties for statistical inference, for at least two reasons. First, samples cannot in practice give us enough information to make reliable inferences on all the details of a population distribution, such as the small kinks and bumps of Figure 6.1. Such details would typically not even be particularly interesting compared to major features like the central tendency and variation of the population distribution. Second, this way of thinking about the population distribution is not appropriate when the population is regarded as infinite.

Addressing both of these problems requires one more conceptual leap. This is to make the assumption that the population distribution is well-represented by a continuous *probability distribution*, and focus on inference on the parameters of that distribution.

We have already introduced the concept of probability distributions in Section 3.5, and considered instances of it in Chapters 4 and 5. There, however, the term was not emphasised because it added no crucial insight into the methods of inference. This was because for discrete variables a probability distribution is specified simply by listing the probabilities of all the categories of the variable. The additional terminology of probability distributions and their parameters seems almost redundant in that context.

The situation is very different for continuous variables. This is illustrated by Figure 6.2, which shows the same frequency polygon as in Figure 6.1, now supplemented by a smooth curve. This curve (“a probability density function”) describes a particular probability distribution. It can be thought of as a smoothed version of the shape of the frequency polygon. What we will do in the future is to use some such probability distribution to represent the population distribution. This means effectively arguing that we believe that the shape of the true population distribution is sufficiently regular to be well described by a smooth curve such as the one in Figure 6.2.

In Figure 6.2 the curve and the frequency polygon have reasonably similar shapes, so the assumption that the former is a good representation of the latter does not seem far-fetched. However, the two are clearly not exactly the same, nor do we expect that even the blood pressures of all English adults exactly match this curve or any other simple probability distribution. All we require is that a population distribution is close enough to a specified probability distribution for the results from analyses based on this assumption to be meaningful and not misleading.

Such a simplifying assumption about the population distribution is known as a **statistical model** for the population. The reason for working with a model is that it leads to much simpler methods of analysis than would otherwise be required. For example, the shape of



Figure 6.2: The frequency polygon of Figure 6.1, together with a normal curve with the same mean and variance.

the distribution shown in Figure 6.2 is entirely determined by just two parameters, its mean and variance. Under this model, all questions about the population distribution can thus be reduced to questions about these two population parameters, and inference can focus on tests and confidence intervals for them.

The potential cost of choosing a specific probability distribution as the statistical model for a particular application is that the assumption may be inappropriate for the data at hand, and if it is, conclusions about population parameters derived from analyses based on this assumption may be misleading. The distribution should thus be chosen carefully, usually based on both substantive considerations and initial descriptive examination of the observed data.

For example, the particular probability distribution shown in Figure 6.2, which is known as the normal distribution, is by definition symmetric around its mean. While it is an adequate approximation of many approximately symmetric population distributions of continuous variables, such as that of blood pressure, many other population distributions are not even roughly symmetric. It would be unrealistic to assume the population distributions of such variables to be normal. Instead, we might consider other continuous probability distributions which can be skewed. Examples of these are the *Exponential*, *Gamma*, *Weibull*, and *Beta* distributions. *Discrete* distributions, of course, will require quite different probability distributions, such as the *Binomial* distribution discussed in Chapter 5, or the *Multinomial* and *Poisson* distributions. On this course, however, we will not include further discussion of these various possibilities.

6.3.2 The normal distribution as a population distribution

The particular probability distribution that is included in Figure 6.2 is a **normal distribution**, also known as the *Gaussian* distribution, after the great German mathematician Karl Friedrich Gauss who was one of the first to derive it in 1809. Figure 6.3 shows a portrait of Gauss from the former German 10-DM banknote, together with pictures of the university town of Göttingen and of the normal curve (even the mathematical formula of the curve is engraved on the note).

The curve of the normal distribution is also known as the “bell curve” because of its shape.



Figure 6.3: A portrait of Gauss and the normal curve on a former German 10-DM banknote.

The normal distribution is by far the most important probability distribution in statistics. The main reason for this is its use as a sampling distribution in a wide range of contexts, for reasons that are explained in Section 6.4. However, the normal distribution is also useful for describing many approximately symmetric population distributions, and it is in this context that we introduce its properties first.

A normal distribution is completely specified by two numbers, its mean (or “expected value”) μ and variance σ^2 . This is sometimes expressed in notation as $Y \sim N(\mu, \sigma^2)$, which is read as “ Y is normally distributed with mean μ and variance σ^2 ”. Different values for μ and σ^2 give different distributions. For example, the curve in Figure 6.2 is that of the $N(74.2, 127.87)$ distribution, where the mean $\mu = 74.2$ and variance $\sigma^2 = 127.87$ are the same as the mean and variance calculated from formulas (6.1) and (6.2) for the 4489 observations of blood pressure. This ensures that this particular normal curve best matches the frequency polygon in Figure 6.2.

The mean μ describes the central tendency of the distribution, and the variance σ^2 its variability. This is illustrated by Figure 6.4, which shows the curves for three different normal distributions. The mean of a normal distribution is also equal to both its median and its mode. Thus μ is the central value in the sense that it divides the distribution into two equal halves, and it also indicates the peak of the curve (the highest probability, as discussed below). In Figure 6.4, the curves for $N(0, 1)$ and $N(0, 9)$ are both centered around $\mu = 0$; the mean of the $N(5, 1)$ distribution is $\mu = 5$, so the whole curve is shifted to the right and centered around 5.

The variance σ^2 determines how widely spread the curve is. In Figure 6.4, the curves for $N(0, 1)$ and $N(5, 1)$ have the same variance $\sigma^2 = 1$, so they have the same shape in terms of their spread. The curve for $N(0, 9)$, on the other hand, is more spread out, because it has a higher variance of $\sigma^2 = 9$. As before, it is often more convenient to describe variability in terms of the standard deviation σ , which is the square root of the variance. Thus we may also say that the $N(0, 9)$ distribution has the standard deviation $\sigma = \sqrt{9} = 3$ (for $\sigma^2 = 1$ the two numbers are the same, since $\sqrt{1} = 1$).

In the histogram in Figure 6.1, the heights of the bars correspond to the proportions of different ranges of blood pressure among the 4489 people in the data set. Another way of stating this is that if we were to sample a person from this group at random, the heights of the bars indicate the **probabilities** that the selected person’s blood pressure would be in a particular range.




Figure 6.4: Three normal distributions with different means and/or variances.

Some values are clearly more likely than others. For example, for blood pressures in the range 50–51.5, the probability is about 0.0025, corresponding to a low bar, while for the range 74–75.5 it is about 0.0365, corresponding to a much higher bar.

The interpretation is the same for the curve of a continuous probability distribution. Its height also indicates the probability of different values in random sampling from a population with that distribution. More precisely, the *areas* under the curve give such probabilities for ranges of values. Probabilities of all the possible values must add up to one, so the area under the whole curve is one — i.e. a randomly sampled unit must have *some* value of the variable in question. More generally, the area under the curve for a range of values gives the probability that the value of a randomly sampled observation is in that range. These are the same principles that we have already used to derive P -values for tests in Sections 4.3.5 and 5.5.3.

Figure 6.5 illustrates this further with some results which hold for any normal distribution, whatever its mean and variance. The grey area in the figure corresponds to values from $\mu - \sigma$ to $\mu + \sigma$, i.e. those values which are no further than one standard deviation from the mean. The area of the grey region is 0.68, so the probability that a randomly sampled value from a normal distribution is within one standard deviation of the mean is 0.68. The two shaded regions either side of the grey area extend the area to 1.96 standard deviations below and above the mean. The probability of this region (the grey and shaded areas together) is 0.95. Rounding the 1.96 to 2, we can thus say that approximately 95% of observations drawn from a normal distribution tend to be within two standard deviations of the mean. This leaves the remaining 5% in the two tails of the distribution, further than 1.96 standard deviations from the mean (the two white areas in Figure 6.5). Because the normal distribution is symmetric, these two areas are of equal size and each thus has the probability 0.025 (i.e. $0.05/2$).

Such calculations can also be used to determine probabilities in particular examples. Returning to the blood pressure data, we might for example be interested in

- the proportion of people in some population whose diastolic blood pressure is higher than 90 (one possible cut-off point for high blood pressure or hypertension)
- the proportion of people with diastolic blood pressure below 60 (possibly indicating un-




Figure 6.5: Illustration of probabilities for the normal distribution. The probability of an observation being within one standard deviation of the mean (the grey area) is 0.68, and the probability of it being within 1.96 standard deviations of the mean (grey and shaded areas together) is 0.95.

usually low blood pressure or hypotension)

- the proportion of people in the normal pressure range of 60–90

Such figures might be of interest for example for predicting health service needs for treating hypertension. Suppose that we were reasonably confident (perhaps from surveys like the one described above) that the distribution of diastolic blood pressure in the population of interest was approximately normally distributed with mean 74.2 and variance 127.87 (and thus standard deviation 11.3). The probabilities of interest are then the areas of the regions shown in Figure 6.6.

The remaining question is how to calculate such probabilities. The short answer is “with a computer”. However, to explain an approach which is required for this in some computer packages and also to provide an alternative method which does not require a computer, we need to introduce one more new quantity. This is the **Z score**, which is defined as

$$Z = \frac{Y - \mu}{\sigma} \quad (6.8)$$

where Y can be any value of the variable of interest. For example, in the blood pressure example the Z scores corresponding to values 60 and 90 are $Z = (60 - 74.2)/11.3 = -1.26$ and $Z = (90 - 74.2)/11.3 = 1.40$ respectively. The Z score can be interpreted as the distance of the value Y from the mean μ , measured in standard deviations σ . Thus the blood pressure 60, with a Z score of -1.26 , is 1.26 standard deviations *below* (hence the negative sign) the mean, while 90 (with Z score 1.40) is 1.40 standard deviations *above* the mean.



Figure 6.6: Illustration of probabilities for a normal distribution in the blood pressure example, where $\mu = 74.2$ and $\sigma = 11.3$. The plot shows probabilities for the ranges of values at most 60 (“Low”), between 60 and 90 (“Mid”) and over 90 (“High”).

Table 6.2: Extract from the table of right-hand tail probabilities for normal Z scores. Here “Tail Prob.” is the probability that a value from the standard normal distribution is at least the value in the column labelled “ z ”. The full table is shown in the Appendix.

z	Tail Prob.
\vdots	\vdots
1.24	0.1075
1.25	0.1056
1.26	0.1038
1.27	0.1020
\vdots	\vdots
1.38	0.0838
1.39	0.0823
1.40	0.0808
1.41	0.0793
\vdots	\vdots

The probability distribution of the Z scores is a normal distribution with mean 0 and variance 1, i.e. $Z \sim N(0, 1)$. This is known as the **standard normal distribution**. The usefulness of Z scores lies in the fact that by transforming the original variable Y from the $N(\mu, \sigma^2)$ distribution into the standard normal distribution they remove the specific values of μ and σ from the calculation. With this trick, probabilities for any normal distribution can be calculated using a single table for Z scores. Such a table is given in the Appendix, and an extract from it is shown in Table 6.2 (note that it is not always presented exactly like this, as different books may use slightly different format or notation). The first column lists values of the Z score (a full table would typically give all values from 0.00 to about 3.50). The second column, labelled “Tail Prob.”, gives the probability that a Z score for a normal distribution is *larger than* the value given by z , i.e. the area of the region to the right of z .

0.1038 (rounded to 0.10 in Figure 6.6).

Finally, the probability of the “Mid” range of blood pressure is the remaining probability not in the two other regions. Because the whole area under the curve (the total probability) is 1, the required probability is obtained by subtraction as $1 - (0.0808 + 0.1038) = 0.8154$. In this example these values obtained from the normal approximation of the population distribution are very accurate. The exact proportions of the 4489 respondents who had diastolic blood pressure at most 60 or greater than 90 were 0.0996 and 0.0793 respectively, so rounded to two decimal places they were the same as the 0.10 and 0.08 obtained from the normal approximation.

These days we can use statistical computer programs to calculate such probabilities directly for a normal distribution with any mean and standard deviation. For example, SPSS has a function called `CDF.NORMAL(quant,mean,stddev)` for this purpose. It calculates the probability that the value from a normal distribution with mean *mean* and standard deviation *stddev* is **at most** *quant*.

In practice we do not usually know the population mean and variance, so their sample estimates will be used in such calculations. For example, for the sample in Table 6.1 we had $\bar{Y} = 72.6$ and $s = 12.7$. Using these values in a similar calculation as above gives the estimated proportion of people in the population with diastolic blood pressures over 90 as 8.5%. Even with a sample of only 50 observations, the estimate is reasonably close to the true population proportion of about 8.1%.

6.4 The normal distribution as a sampling distribution

We have already encountered the normal distribution in Section 5.5.3, in the role of the *sampling distribution* of a test statistic rather than as a model for the population distribution of a variable. In fact, the most important use of the normal distribution is as a sampling distribution, because in this role it often cannot be replaced by any other simple distributions. The reasons for this claim are explained in this section. We begin with the case of the distribution of the sample mean in samples from a normal population, before extending it with a result which provides the justification for the standard normal sampling distributions used for inference on proportions in Chapter 5, and even for the χ^2 sampling distribution of the χ^2 test in Chapter 4.

Recall from Section 4.3.4 that the sampling distribution of a statistic is its distribution across all possible random samples of a given size from a population. The statistic we focus on here is the sample mean \bar{Y} . If we assume that the population distribution is exactly normal, we have the following result:

- If the population distribution of a variable Y is normal with mean μ and variance σ^2 , the sampling distribution of the sample mean \bar{Y} for a random sample of size n is also a normal distribution, with mean μ and variance σ^2/n .

The mean and variance of this sampling distribution are worth discussing separately:

- The mean of the sampling distribution of \bar{Y} is equal to the population mean μ of Y . This means that while \bar{Y} from a single sample may be below or above the true μ , in repeated samples it would on average estimate the correct parameter. In statistical language, \bar{Y} is then an *unbiased estimate* of μ . More generally, most possible samples would give values of \bar{Y} not very far from μ , where the scale for “far” is provided by the standard deviation discussed below.
- The variance of the sampling distribution of \bar{Y} is σ^2/n or, equivalently, its standard deviation is σ/\sqrt{n} . This standard deviation is also known as the **standard error of**

the mean, and is often denoted by something like $\sigma_{\bar{Y}}$. It describes the variability of the sampling distribution. Its magnitude depends on σ , i.e. on the variability of Y in the population. More interestingly, it also depends on the sample size n , which appears in the denominator in σ/\sqrt{n} . This means that the standard error of the mean is smaller for large samples than for small ones. This is illustrated in Figure 6.7. It shows the sampling distribution of \bar{Y} for samples of sizes $n = 50$ and $n = 1000$ from a normal population with $\mu = 74.2$ and $\sigma = 11.3$, i.e. the population mean and standard deviation in the blood pressure example. It can be seen that while both sampling distributions are centered around the true mean $\mu = 74.2$, the distribution for the smaller sample is more spread out than that for the larger sample: more precisely, the standard error of the mean is $\sigma/\sqrt{n} = 11.3/\sqrt{50} = 1.60$ when $n = 50$ and $11.3/\sqrt{1000} = 0.36$ when $n = 1000$. Recalling from Section 6.3.2 that approximately 95% of the probability in a normal distribution is within two standard deviations of the mean, this means that about 95% of samples of size 50 in this case would give a value of \bar{Y} between $\mu - 2 * 1.60 = 74.2 - 3.2 = 71.0$ and $74.2 + 3.2 = 77.4$. For samples of size $n = 1000$, on the other hand, 95% of samples would yield \bar{Y} in the much narrower range of $74.2 - 2 * 0.36 = 73.5$ to $74.2 + 2 * 0.36 = 74.9$.




Figure 6.7: Illustration of the sampling distribution of the sample mean for two sample sizes. In both cases the population distribution is normal with $\mu = 74.2$ and $\sigma = 11.3$.

The connection between sample size and the variability of a sampling distribution applies not only to the sample mean but to (almost) all estimates of population parameters. In general, (i) the task of statistical inference is to use information in a sample to draw conclusions about population parameters; (ii) the expected magnitude of the sampling error, i.e. the remaining uncertainty about population parameters resulting from having information only on a sample, is characterised by the variability of the sampling distributions of estimates of the parameters; and (iii) other things being equal, the variability of a sampling distribution decreases when the sample size increases. Thus data really are the currency of statistics and more data are better than less data. In practice data collection of course costs time and money, so we cannot always obtain samples which are as large as we might otherwise want. Apart from resource constraints, the choice of sample size depends also on such things as the aims of the analysis, the level of

precision required, and guesses about the variability of variables in the population. Statistical considerations of the trade-offs between them in order to make decisions about sample sizes are known as *power* calculations. They will be discussed very briefly later, in Section 7.6.2.

In Figure 6.8 we use a computer simulation rather than a mathematical theorem to examine the sampling distribution of a sample mean. Here 100,000 simple random samples of size $n = 50$ were drawn from the $N = 4489$ values of blood pressure that we are treating as the finite population in this illustration. The sample mean \bar{Y} of blood pressure was calculated for each of these samples, and the histogram of these 100,000 values of \bar{Y} is shown in Figure 6.8. Also shown is the curve of the normal distribution with the mean μ and standard deviation $\sigma/\sqrt{50}$ determined by the theoretical result given above.




Figure 6.8: Example of the sampling distribution of the sample mean. The plot shows a histogram of the values of the sample mean in 100,000 samples of size $n = 50$ drawn from the 4489 values of diastolic blood pressure shown in Figure 6.1, for which the mean is $\mu = 74.2$ and standard deviation is $\sigma = 11.3$. Superimposed on the histogram is the curve of the approximate sampling distribution, which is normal with mean μ and standard deviation σ/\sqrt{n} .

The match between the curve and the histogram in Figure 6.8 is clearly very close. This is actually a nontrivial finding which illustrates a result which is of crucial importance for statistical inference. Recall that the normal curve shown in Figure 6.8 is derived from the mathematical result stated above, which assumed that the population distribution of Y is *exactly* normal. The histogram in Figure 6.8, on the other hand, is based on repeated samples from the actual population distribution of blood pressure, which, while quite close to a normal distribution as shown in Figure 6.2, is certainly not exactly normal. Despite this, it is clear that the normal curve describes the histogram essentially exactly.

If this was not true, that is if the sampling distribution that applies for the normal distribution was inadequate when the true population distribution was even slightly different from normal, the result would be of little practical use. No population distribution is ever exactly normal, and many are very far from normality. Fortunately, however, it turns out that quite the opposite is true, and that the sampling distribution of the mean is approximately the same for

nearly *all* population distributions. This is the conclusion from the **Central Limit Theorem** (CLT), one of the most remarkable results in all of mathematics. Establishing the CLT with increasing levels of generality has been the work of many mathematicians over several centuries, as different versions of it have been proved by, among others, de Moivre, Laplace, Cauchy, Chebyshev, Markov, Liapounov, Lindeberg, Feller, Lévy, Hoeffding, Robbins, and Rebolledo between about 1730 and 1980. One version of the CLT can be stated as

The (Lindeberg-Feller) Central Limit Theorem: For each $n = 1, 2, \dots$, let Y_{nj} , for $j = 1, 2, \dots, n$, be independent random variables with $E(Y_{nj}) = 0$ and $\text{var}(Y_{nj}) = \sigma_{nj}^2$. Let $Z_n = \sum_{j=1}^n Y_{nj}$, and let $B_n^2 = \text{var}(Z_n) = \sum_{j=1}^n \sigma_{nj}^2$. Suppose also that the following condition holds: for every $\epsilon > 0$,

$$\frac{1}{B_n^2} \sum_{j=1}^n E\{Y_{nj}^2 I(|Y_{nj}| \geq \epsilon B_n)\} \rightarrow 0 \text{ as } n \rightarrow \infty. \quad (6.9)$$

Then $Z_n/B_n \xrightarrow{\mathcal{L}} N(0, 1)$.

No, that will not come up in the examination. The theorem is given here just as a glimpse of how this topic would be introduced in a very different kind of text book³, and because it pleases the author of this coursepack to note that Jarl Lindeberg was Finnish. For our purposes, it is better to state the same result in English:

- If Y_1, Y_2, \dots, Y_n are a random sample of observations from (almost⁴) any distribution with a population mean μ and variance σ^2 , and if n is reasonably large, the sampling distribution of their sample mean \bar{Y} is approximately a normal distribution with mean μ and variance σ^2/n .

Thus the sampling distribution of the mean from practically any population distribution is approximately the same as when the population distribution is normal, as long as the sample size is “reasonably large”. The larger the sample size is, the closer the sampling distribution is to the normal distribution, and it becomes exactly normal when the sample size is infinitely large (i.e. “asymptotically”). What is large enough depends particularly on the nature of the population distribution. For continuous variables, the CLT approximation is typically adequate even for sample sizes as small as $n = 30$, so we can make use of the approximate normal sampling distribution when n is 30 or larger. This is, of course, simply a pragmatic rule of thumb which does not mean that the normal approximation is completely appropriate for $n = 30$ but entirely inappropriate for $n = 29$; rather, the approximation becomes better and better as the sample size increases, while below about 30 the chance of incorrect conclusions from using it becomes large enough for us not to usually want to take that risk.

We have seen in Figure 6.7 that in the blood pressure example the sampling distribution given by the Central Limit Theorem is essentially exact for samples of size $n = 50$. In this case this is hardly surprising, as the population distribution itself is already quite close to a normal distribution. The theorem is not, however, limited to such easy cases but works quite generally. To demonstrate this with a more severe test, let us consider a population distribution that is as far as possible from normal. This is the binomial distribution of a binary variable that was introduced in Section 5.3. If the probability parameter of this distribution is π , its mean and variance are $\mu = \pi$ and $\sigma^2 = \pi(1 - \pi)$, and the sample mean \bar{Y} of observations from the

³The data can be obtained from <http://www3.norc.ox.ac.uk/gss+website/>, which gives further information on the survey, including the full text of the questionnaires.

⁴ESS Round 5: European Social Survey Round 5 Data (2010). Data file edition 2.0. Norwegian Social Science Data Services, Norway. Data Archive and distributor of ESS data. The full data can be obtained from <http://ess.nsd.uib.no/ess/round5/>.

distribution is the sample proportion $\hat{\pi}$ (see the equation at the end of Section 5.4). The CLT then tells us that

- When n is large enough, the sampling distribution of the sample proportion $\hat{\pi}$ of a dichotomous variable Y with population proportion π is approximately a normal distribution with mean π and variance $\pi(1 - \pi)/n$.

This powerful result is illustrated in Figure 6.9. It is similar to Figure 6.8 in that it shows sampling distributions obtained from a computer simulation, together with the normal curve suggested by the CLT. For each plot, 5000 samples of size n were simulated from a population where π was 0.2. The sample proportion $\hat{\pi}$ was then calculated for each simulated sample, and the histogram of these 5000 values drawn. Four different sample sizes were used: $n = 10, 30, 100$, and 1000. It can be seen that the normal distribution is not a very good approximation of the sampling distribution of $\hat{\pi}$ when n is as small as 10 or even 30. For the larger values of 100 and 1000, however, the normal approximation is already quite good, as expected from the CLT.

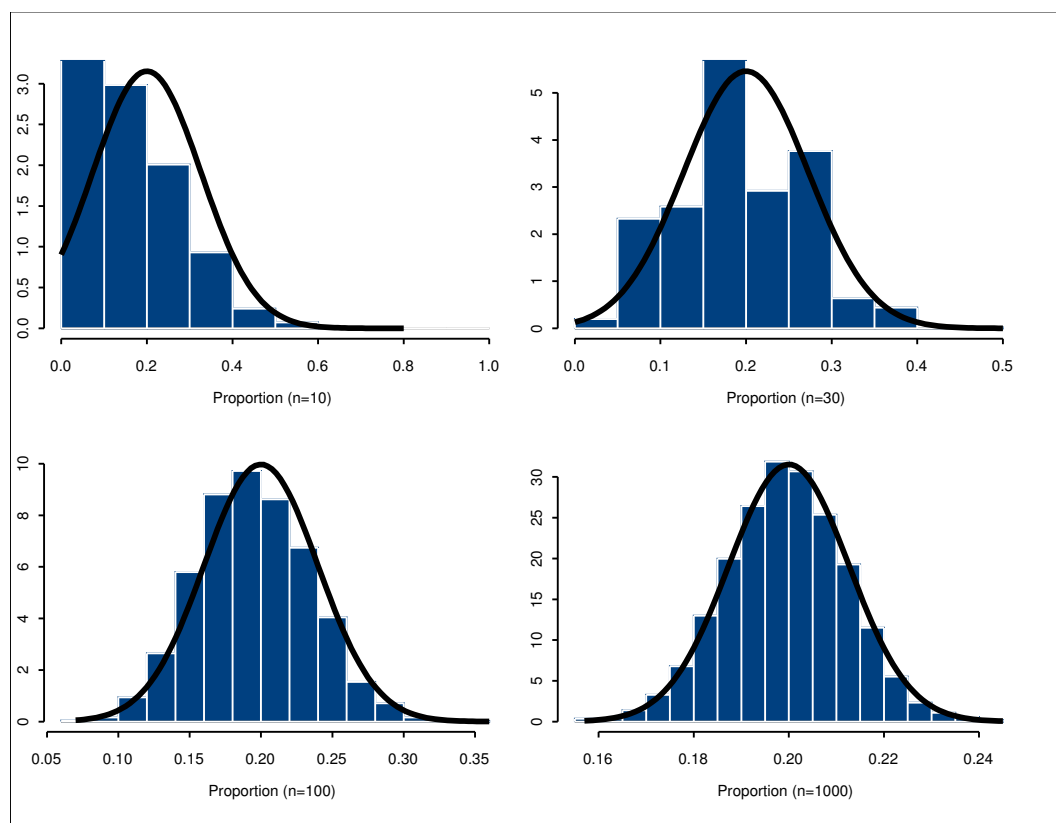


Figure 6.9: Illustration of the Central Limit Theorem for the sample proportion of a dichotomous variable. Each plot shows the histogram of the sample proportions $\hat{\pi}$ calculated for 5000 samples simulated from a population distribution with proportion $\pi = 0.2$, together with the normal curve with mean π and variance $\pi(1 - \pi)/n$. The sample sizes n are 10, 30, 100 and 1000.

The variability of the sampling distribution will again depend on n . In Figure 6.9, the observed range of values of $\hat{\pi}$ decreases substantially as n increases. When $n = 10$, values of between about 0 and 0.4 are quite common, whereas with $n = 1000$, essentially all of the samples give $\hat{\pi}$ between about 0.16 and 0.24, and a large majority are between 0.18 and 0.22. Thus increasing the sample size will again increase the precision with which we can estimate π , and decrease the uncertainty in inference about its true value.

The Central Limit Theorem is, with some additional results, the justification for the standard normal sampling distribution used for tests and confidence intervals for proportions in Chapter 5. The conditions for sample sizes mentioned there (at the beginning of Section 5.5.3 and 5.7) again derive from conditions for the CLT to be adequate. The same is also ultimately true for the χ^2 distribution and conditions for the χ^2 test in Chapter 4. Results like these, and many others, explain the central importance of the CLT in statistical methodology.

Chapter 7

Analysis of population means

7.1 Introduction and examples

This chapter introduces some basic methods of analysis for continuous, interval-level variables. The main focus is on statistical inference on population *means* of such variables, but some new methods of descriptive statistics are also described. The discussion draws on the general ideas that have already been explained for inference in Chapters 4 and 5, and for continuous distributions in Chapter 6. Few if any new concepts thus need to be introduced here. Instead, this chapter can focus on describing the specifics of these very commonly used methods for continuous variables.

As in Chapter 5, questions on both a single group and on comparisons between two groups are discussed here. Now, however, the main focus is on the two-group case. There we treat the group as the explanatory variable X and the continuous variable of interest as the response variable Y , and assess the possible associations between X and Y by comparing the distributions (and especially the means) of Y in the two groups.

The following five examples will be used for illustration throughout this chapter. Summary statistics for them are shown in Table 7.1.

Example 7.1: Survey data on diet

The National Diet and Nutrition Survey of adults aged 19–64 living in private households in Great Britain was carried out in 2000–01¹. One part of the survey was a food diary where the respondents recorded all food and drink they consumed in a seven-day period. We consider two variables derived from the diary: the consumption of fruit and vegetables in portions (of 400g) per day (with mean in the sample of size $n = 1724$ of $\bar{Y} = 2.8$, and standard deviation $s = 2.15$), and the percentage of daily food energy intake obtained from fat and fatty acids ($n = 1724$, $\bar{Y} = 35.3$, and $s = 6.11$).

¹ESS Round 5: European Social Survey Round 5 Data (2010). Data file edition 2.0. Norwegian Social Science Data Services, Norway. Data Archive and distributor of ESS data. The full data can be obtained from <http://ess.nsd.uib.no/ess/round5/>.

Table 7.1: Examples of analyses of population means used in Chapter 7. Here n and \bar{Y} denote the sample size and sample mean respectively, in the two-group examples 7.2–7.5 separately for the two groups. “Diff.” denotes the between-group difference of means, and s is the sample standard deviation of the response variable Y for the whole sample (Example 7.1), of the response variable within each group (Examples 7.2 and 7.3), or of the within-pair differences (Examples 7.4 and 7.5).

One sample	n	\bar{Y}	s	Diff.
<i>Example 7.1: Variables from the National Diet and Nutrition Survey</i>				
Fruit and vegetable consumption (400g portions)	1724	2.8	2.15	
Total energy intake from fat (%)	1724	35.3	6.11	
Two independent samples				
<i>Example 7.2: Average weekly hours spent on housework</i>				
Men	635	7.33	5.53	
Women	469	8.49	6.14	1.16
<i>Example 7.3: Perceived friendliness of a police officer</i>				
No sunglasses	67	8.23	2.39	
Sunglasses	66	6.49	2.01	-1.74
Two dependent samples				
<i>Example 7.4: Father’s personal well-being</i>				
Sixth month of wife’s pregnancy	109	30.69		
One month after the birth	109	30.77	2.58	0.08
<i>Example 7.5: Traffic flows on successive Fridays</i>				
Friday the 6th	10	128,385		
Friday the 13th	10	126,550	1176	-1835

Example 7.2: Housework by men and women

This example uses data from the 12th wave of the British Household Panel Survey (BHPS), collected in 2002. BHPS is an ongoing survey of UK households, measuring a range of socio-economic variables. One of the questions in 2002 was

“About how many hours do you spend on housework in an average week, such as time spent cooking, cleaning and doing the laundry?”

The response to this question (recorded in whole hours) will be the response variable Y , and the respondent’s sex will be the explanatory variable X . We consider only those respondents who were less than 65 years old at the time of the interview and who lived in single-person households (thus the comparisons considered here will not involve questions of the division of domestic work within families)².

²The data were obtained from <http://data.london.gov.uk/datastore/package/london-borough-profiles>.

We can indicate summary statistics separately for the two groups by using subscripts 1 for men and 2 for women (for example). The sample sizes are $n_1 = 635$ for men and $n_2 = 469$ for women, and the sample means of Y are $\bar{Y}_1 = 7.33$ and $\bar{Y}_2 = 8.49$. These and the sample standard deviations s_1 and s_2 are also shown in Table 7.1.

Example 7.3: Eye contact and perceived friendliness of police officers

This example is based on an experiment conducted to examine the effects of some aspects of the appearance and behaviour of police officers on how members of the public perceive their encounters with the police³. The subjects of the study were 133 people stopped by the Traffic Patrol Division of a detachment of the Royal Canadian Mounted Police. When talking to the driver who had been stopped, the police officer either wore reflective sunglasses which hid his eyes, or wore no glasses at all, thus permitting eye contact with the respondent. These two conditions define the explanatory variable X , coded 1 if the officer wore no glasses and 2 if he wore sunglasses. The choice of whether sunglasses were worn was made at random before a driver was stopped.

While the police officer went back to his car to write out a report, a researcher asked the respondent some further questions, one of which is used here as the response variable Y . It is a measure of the respondent's perception of the friendliness of the police officer, measured on a 10-point scale where large values indicate high levels of friendliness.

The article describing the experiment does not report all the summary statistics needed for our purposes. The statistics shown in Table 7.1 have thus been partially made up for use here. They are, however, consistent with the real results from the study. In particular, the direction and statistical significance of the difference between \bar{Y}_2 and \bar{Y}_1 are the same as those in the published report.

Example 7.4: Transition to parenthood

In a study of the stresses and feelings associated with parenthood, 109 couples expecting their first child were interviewed before and after the birth of the baby.⁴ Here we consider only data for the fathers, and only one of the variables measured in the study. This variable is a measure of personal well-being, obtained from a seven-item attitude scale, where larger values indicate higher levels of well-being. Measurements of it were obtained for each father at three time points: when the mother was six months pregnant, one month after the birth of the baby, and six months after the birth. Here we will use only the first two of the measurements. The response variable Y will thus be the measure of personal well-being, and the explanatory variable X will be the time of measurement (sixth month of the pregnancy or one month after the birth). The means of Y at the two times are shown in Table 7.1. As in Example 7.3, not all of the numbers needed here were given in the original article. Specifically, the standard error of the difference in Table 7.1 has been made up in such a way that the results of a significance test for the mean difference agree with those in the article.

If you download the “Profiles in Excel” workbook, you will find that one of the pages contains a map of the boroughs, and a tool for visualising the data on that map. A regular map of the boroughs can be found at for example at

<http://www.londoncouncils.gov.uk/londonfacts/londonlocalgovernment/londonmapandlinks/default.htm>.

³The data can be obtained from <http://www3.norc.ox.ac.uk/gss+website/>, which gives further information on the survey, including the full text of the questionnaires.

⁴ESS Round 5: European Social Survey Round 5 Data (2010). Data file edition 2.0. Norwegian Social Science Data Services, Norway Data Archive and distributor of ESS data. The full data can be obtained from <http://ess.nsd.uib.no/ess/round5/>.

Example 7.5: Traffic patterns on Friday the 13th

A common superstition regards the 13th day of any month falling on a Friday as a particularly unlucky day. In a study examining the possible effects of this belief on people's behaviour⁵, data were obtained on the numbers of vehicles travelling between junctions 7 and 8 and junctions 9 and 10 on the M25 motorway around London during every Friday the 13th in 1990–92. For comparison, the same numbers were also recorded during the previous Friday (i.e. the 6th) in each case. There are only ten such pairs here, and the full data set is shown in Table 7.2. Here the explanatory variable X indicates whether a day is Friday the 6th (coded as 1) or Friday the 13th (coded as 2), and the response variable is the number of vehicles travelling between two junctions.

Table 7.2: Data for Example 7.5: Traffic flows between junctions of the M25 on each Friday the 6th and Friday the 13th in 1990–92.

Date	Junctions	Friday the 6th	Friday the 13th	Difference
July 1990	7 to 8	139246	138548	-698
July 1990	9 to 10	134012	132908	-1104
September 1991	7 to 8	137055	136018	-1037
September 1991	9 to 10	133732	131843	-1889
December 1991	7 to 8	123552	121641	-1911
December 1991	9 to 10	121139	118723	-2416
March 1992	7 to 8	128293	125532	-2761
March 1992	9 to 10	124631	120249	-4382
November 1992	7 to 8	124609	122770	-1839
November 1992	9 to 10	117584	117263	-321

In each of these cases, we will regard the variable of interest Y as a continuous, interval-level variable. The five examples illustrate three different situations considered in this chapter. Example 7.1 includes two separate Y -variables (consumption of fruit and vegetables, and fat intake), each of which is considered for a single population. Questions of interest are about the mean of the variable in the population. This is analogous to the one-group questions on proportions in Sections 5.5 and 5.6. In this chapter the one-group case is discussed only relatively briefly, in Section 7.4.

The main focus here is on the case illustrated by Examples 7.2 and 7.3. These involve samples of a response variable (hours of housework, or preceived friendliness) from two groups (men and women, or police with or without sunglasses). We are then interested in comparing the distributions, and especially the means, of the response variable between the groups. This case will be discussed first. Descriptive statistics for it are described in Section 7.2, and statistical inference in Section 7.3.

Finally, examples 7.4 and 7.5 also involve comparisons between two groups, but of a slightly different kind than examples 7.2 and 7.3. The two types of cases differ in the nature of the two samples (groups) being compared. In Examples 7.2 and 7.3, the samples can be considered to be **independent**. What this claim means will be discussed briefly later; informally, it is justified in these examples because the subjects in the two groups are separate and unrelated individuals. In Examples 7.4 and 7.5, in contrast, the samples (before and after the birth of

⁵Strictly speaking, the analysis should incorporate sampling weights (variable *DWEIGHT*) to adjust for different sampling probabilities for different types of respondents. Here the weights are ignored. Using them would not change the main conclusions for these variables.

a child, or two successive Fridays) must be considered **dependent**, essentially because they concern measurements on the same units at two distinct times. This case is discussed in Section 7.5.

In each of the four two-group examples we are primarily interested in questions about possible association between the group variable X and the response variable Y . As before, this is the question of whether the conditional distributions of Y are different at the two levels of X . There is thus an association between X and Y if

- Example 7.2: The distribution of hours of housework is different for men than for women.
- Example 7.3: The distribution of perceptions of a police officer's friendliness is different when he is wearing mirrored sunglasses than when he is not.
- Example 7.4: The distribution of measurements of personal well-being is different at the sixth month of the pregnancy than one month after the birth.
- Example 7.5: The distributions of the numbers of cars on the motorway differ between Friday the 6th and the following Friday the 13th.

We denote the two values of X , i.e. the two groups, by 1 and 2. The mean of the population distribution of Y given $X = 1$ will be denoted μ_1 and the standard deviation σ_1 , and the mean and standard deviation of the population distribution given $X = 2$ are denoted μ_2 and σ_2 similarly. The corresponding sample quantities are the conditional sample means \bar{Y}_1 and \bar{Y}_2 and sample standard deviations s_1 and s_2 . For inference, we will focus on the population difference $\Delta = \mu_2 - \mu_1$ which is estimated by the sample difference $\hat{\Delta} = \bar{Y}_2 - \bar{Y}_1$. Some of the descriptive methods described in Section 7.2, on the other hand, also aim to summarise and compare other aspects of the two conditional sample distributions.

7.2 Descriptive statistics for comparisons of groups

7.2.1 Graphical methods of comparing sample distributions

There is an association between the group variable X and the response variable Y if the distributions of Y in the two groups are not the same. To determine the extent and nature of any such association, we need to compare the two distributions. This section describes methods of doing so for observed data, i.e. for examining associations in a sample. We begin with graphical methods which can be used to detect differences in any aspects of the two distributions. We then discuss some non-graphical summaries which compare specific aspects of the sample distributions, especially their means.

Although the methods of *inference* described later in this chapter will be limited to the case where the group variable X is dichotomous, many of the descriptive methods discussed below can just as easily be applied when more than two groups are being compared. This will be mentioned wherever appropriate. For inference in the multiple-group case some of the methods discussed in Chapter 8 are applicable.

In Section 2.5.2 we described four graphical methods of summarizing the sample distribution of one continuous variable Y : the histogram, the stem and leaf plot, the frequency polygon and the box plot. Each of these can be adapted for comparisons of two or more distributions, although some more conveniently than others. We illustrate the use three of the plots for this purpose, using the comparison of housework hours in Example 7.2 for illustration. Stem and leaf plots will not be shown, because they are less appropriate when the sample sizes are as large as they are in this example.

Two sample distributions can be compared by displaying histograms of them side by side, as shown in Figure 7.1. This is not a very common type of graph, and not ideal for visually comparing the two distributions, because the bars to be compared (here for men vs. women) end at opposite ends of the plot. A better alternative is to use frequency polygons. Since these represent a sample distribution by a single line, it is easy to include two of them in the same plot, as shown in Figure 7.2. Finally, Figure 7.3 shows two boxplots of reported housework hours, one for men and one for women.

The plots suggest that the distributions are quite similar for men and women. In both groups, the largest proportion of respondents stated that they do between 4 and 7 hours of housework a week. The distributions are clearly positively skewed, since the reported number of hours was much higher than average for a number of people (whereas less than zero hours were of course not recorded for anyone). The proportions of observations in categories including values 5, 10, 15, 20, 25 and 30 tend to be relatively high, suggesting that many respondents chose to report their answers in such round numbers. The box plots show that the median number of hours is higher for women than for men (7 vs. 6 hours), and women's responses have slightly less variation, as measured by both the IQR and the range of the whiskers. Both distributions have several larger, outlying observations (note that SPSS, which was used to produce Figure 7.3, divides outliers into moderate and “extreme” ones; the latter are observations more than 3 IQR from the end of the box, and are plotted with asterisks).




Figure 7.1: Histograms of the sample distributions of reported weekly hours of housework in Example 7.2, separately for men ($n = 635$) and women ($n = 469$).

Figures 7.1–7.3 also illustrate an important general point about such comparisons. Typically we focus on comparing *means* of the conditional distributions. Here the difference between the sample means is 1.16, i.e. women in the sample spend, on average, over an hour longer on housework per week than men. The direction of the difference could also be guessed from Figure 7.2, which shows that somewhat smaller proportions of women than of men report small numbers of hours, and larger proportions of women report large numbers. This difference




Figure 7.2: Frequency polygons of the sample distributions of reported weekly hours of housework in Example 7.2, separately for men and women. The points show the percentages of observations in the intervals of 0–3, 4–7, ..., 32–35 hours (plus zero percentages at each end of the curve).



Figure 7.3: Box plots of the sample distributions of reported weekly hours of housework in Example 7.2, separately for men and women.

will later be shown to be statistically significant, and it is also arguably relatively large in a substantive sense.

However, it is equally important to note that the two distributions summarized by the graphs are nevertheless largely similar. For example, even though the mean is higher for women, there are clearly many women who report spending hardly any time on housework, and many men who spend a lot of time on it. In other words, the two distributions overlap to a large extent. This obvious point is often somewhat neglected in public discussions of differences between groups such as men and women or different ethnic groups. It is not uncommon to see reports of research indicating that (say) men have higher or lower values of something or other than women. Such statements usually refer to differences of averages, and are often clearly important and interesting. Less helpful, however, is the tendency to discuss the differences almost as if the corresponding distributions had no overlap at all, i.e. as if *all* men were higher or lower in some characteristic than all women. This is obviously hardly ever the case.

Box plots and frequency polygons can also be used to compare more than two sample distributions. For example, the experimental conditions in the study behind Example 7.3 actually involved not only whether or not a police officer wore sunglasses, but also whether or not he wore a gun. Distributions of perceived friendliness given all four combinations of these two conditions could easily be summarized by drawing four box plots or frequency polygons in the same plot, one for each experimental condition.

7.2.2 Comparing summary statistics

Main features of sample distributions, such as their central tendencies and variations, are described using the summary statistics introduced in Section 2.6. These too can be compared between groups. Table 7.1 shows such statistics for the examples of this chapter. Tables like these are routinely reported for initial description of data, even if more elaborate statistical methods are later used.

Sometimes the association between two variables in a sample is summarized in a single *measure of association* calculated from the data. This is especially convenient when both of the variables are continuous (in which case the most common measure of association is known as the *correlation* coefficient). In this section we consider as such a summary the difference $\hat{\Delta} = \bar{Y}_2 - \bar{Y}_1$ of the sample means of Y in the two groups. These differences are also shown in Table 7.1.

The difference of means is important because it is also the focus of the most common methods of inference for two-group comparisons. For purely descriptive purposes it may be as or more convenient to report some other statistic. For example, the difference of means of 1.16 hours in Example 7.2 could also be described in *relative* terms by saying that the women's average is about 16 per cent higher than the men's average (because $1.16/7.33 = 0.158$, i.e. the difference represents 15.8 % of the men's average).

7.3 Inference for two means from independent samples

7.3.1 Aims of the analysis

Formulated as a statistical model in the sense discussed on page in Section 6.3.1, the assumptions of the analyses considered in this section are as follows:

1. We have a sample of n_1 independent observations of a variable Y in group 1, which have a population distribution with mean μ_1 and standard deviation σ_1 .

2. We have a sample of n_2 independent observations of Y in group 2, which have a population distribution with mean μ_2 and standard deviation σ_2 .
3. The two samples are independent, in the sense discussed following Example 7.5.
4. For now, we further assume that the population standard deviations σ_1 and σ_2 are equal, with a common value denoted by σ . This relatively minor assumption will be discussed further in Section 7.3.4.

We could have stated the starting points of the analyses in Chapters 4 and 5 also in such formal terms. It is not absolutely necessary to always do so, but we should at least remember that any statistical analysis is based on some such model. In particular, this helps to make it clear what our methods of analysis do and do not assume, so that we may critically examine whether these assumptions appear to be justified for the data at hand.

The model stated above does not require that the population distributions of Y should have the form of any particular probability distribution. It is often further assumed that these distributions are normal distributions, but this is not essential. Discussion of this question is postponed until Section 7.3.4.

The only new term in this model statement was the “independent” under assumptions 1 and 2. This statistical term can be roughly translated as “unrelated”. The condition can usually be regarded as satisfied when the units of analysis are different entities, as in Examples 7.2 and 7.3 where the units within each group are distinct individual people. In these examples the individuals in the two groups are also distinct, from which it follows that the two *samples* are independent as required by assumption 3. The same assumption of independent observations is also required by all of the methods described in Chapters 4 and 5, although we did not state this explicitly there.

This situation is illustrated by Example 7.2, where Y is the number of hours a person spends doing housework in a week, and the two groups are men (group 1) and women (group 2).

The quantity of main interest is here the difference of population means

$$\Delta = \mu_2 - \mu_1. \quad (7.1)$$

In particular, if $\Delta = 0$, the population means in the two groups are the same. If $\Delta \neq 0$, they are not the same, which implies that there is an association between Y and the group in the population.

Inference on Δ can be carried out using methods which are straightforward modifications of the ones introduced first in Chapter 5. For significance testing, the null hypothesis of interest is

$$H_0 : \Delta = 0, \quad (7.2)$$

to be tested against a two-sided ($H_a : \Delta \neq 0$) or one-sided ($H_a : \Delta > 0$ or $H_a : \Delta < 0$) alternative hypothesis. The test statistic used to test (7.2) is again of the form

$$t = \frac{\hat{\Delta}}{\hat{\sigma}_{\hat{\Delta}}} \quad (7.3)$$

where $\hat{\Delta}$ is a sample estimate of Δ , and $\hat{\sigma}_{\hat{\Delta}}$ its estimated standard error. Here the statistic is conventionally labelled t rather than z and called the *t-test statistic* because sometimes the

t -distribution rather than the normal is used as its sampling distribution. This possibility is discussed in Section 7.3.4, and we can ignore it until then.

Confidence intervals for the differences Δ are also of the familiar form

$$\hat{\Delta} \pm z_{\alpha/2} \hat{\sigma}_{\hat{\Delta}} \quad (7.4)$$

where $z_{\alpha/2}$ is the appropriate multiplier from the standard normal distribution to obtain the required confidence level, e.g. $z_{0.025} = 1.96$ for 95% confidence intervals. The multiplier is replaced with a slightly different one if the t -distribution is used as the sampling distribution, as discussed in Section 7.3.4.

The details of these formulas in the case of two-sample inference on means are described next, in Section 7.3.2 for the significance test and in Section 7.3.3 for the confidence interval.

7.3.2 Significance testing: The two-sample t -test

For tests of the difference of means $\Delta = \mu_2 - \mu_1$ between two population distributions, we consider the null hypothesis of no difference

$$H_0 : \Delta = 0. \quad (7.5)$$

In the housework example, this is the hypothesis that average weekly hours of housework in the population are the same for men and women. It is tested against an alternative hypothesis, either the two-sided alternative hypotheses

$$H_a : \Delta \neq 0 \quad (7.6)$$

or one of the one-sided alternative hypotheses

$$H_a : \Delta > 0 \text{ or } H_a : \Delta < 0$$

In the discussion below, we concentrate on the more common two-sided alternative.

The test statistic for testing (7.5) is of the general form (7.3). Here it depends on the data only through the sample means \bar{Y}_1 and \bar{Y}_2 and sample variances s_1^2 and s_2^2 of Y in the two groups. A point estimate of Δ is

$$\hat{\Delta} = \bar{Y}_2 - \bar{Y}_1. \quad (7.7)$$

In terms of the population parameters, the standard error of $\hat{\Delta}$ is

$$\sigma_{\hat{\Delta}} = \sqrt{\sigma_{\bar{Y}_2}^2 + \sigma_{\bar{Y}_1}^2} = \sqrt{\frac{\sigma_2^2}{n_2} + \frac{\sigma_1^2}{n_1}}. \quad (7.8)$$

When we assume that the population standard deviations σ_1 and σ_2 are equal, with a common value σ , (7.8) simplifies to

$$\sigma_{\hat{\Delta}} = \sigma \sqrt{\frac{1}{n_2} + \frac{1}{n_1}}. \quad (7.9)$$

The formula of the test statistic uses an estimate of this standard error, given by

$$\hat{\sigma}_{\hat{\Delta}} = \hat{\sigma} \sqrt{\frac{1}{n_2} + \frac{1}{n_1}} \quad (7.10)$$

where $\hat{\sigma}$ is an estimate of σ , calculated from

$$\hat{\sigma} = \sqrt{\frac{(n_2 - 1)s_2^2 + (n_1 - 1)s_1^2}{n_1 + n_2 - 2}}. \quad (7.11)$$

Substituting (7.7) and (7.10) into the general formula (7.3) gives the **two-sample t-test statistic for means**

$$t = \frac{\bar{Y}_2 - \bar{Y}_1}{\hat{\sigma} \sqrt{1/n_2 + 1/n_1}} \quad (7.12)$$

where $\hat{\sigma}$ is given by (7.11).

For an illustration of the calculations, consider again the housework Example 7.2. Here, denoting men by 1 and women by 2, $n_1 = 635$, $n_2 = 469$, $\bar{Y}_1 = 7.33$, $\bar{Y}_2 = 8.49$, $s_1 = 5.53$ and $s_2 = 6.14$. The estimated mean difference is thus

$$\hat{\Delta} = \bar{Y}_2 - \bar{Y}_1 = 8.49 - 7.33 = 1.16.$$

The common value of the population standard deviation σ is estimated from (7.11) as

$$\begin{aligned} \hat{\sigma} &= \sqrt{\frac{(n_2 - 1)s_2^2 + (n_1 - 1)s_1^2}{n_1 + n_2 - 2}} = \sqrt{\frac{(469 - 1)6.14^2 + (635 - 1)5.53^2}{635 + 469 - 2}} \\ &= \sqrt{33.604} = 5.797 \end{aligned}$$

and the estimated standard error of $\hat{\Delta}$ is given by (7.10) as

$$\hat{\sigma}_{\hat{\Delta}} = \hat{\sigma} \sqrt{\frac{1}{n_2} + \frac{1}{n_1}} = 5.797 \sqrt{\frac{1}{469} + \frac{1}{635}} = 0.353.$$

The value of the t-test statistic (7.12) is then obtained as

$$t = \frac{1.16}{0.353} = 3.29.$$

These values and other quantities explained later, as well as similar results for Example 7.3, are also shown in Table 7.3.

Table 7.3: Results of tests and confidence intervals for comparing means for two independent samples. For Example 7.2, the difference of means is between women and men, and for Example 7.3, it is between wearing and not wearing sunglasses. The test statistics and confidence intervals are obtained under the assumption of equal population standard deviations, and the P -values are for a test with a two-sided alternative hypothesis. See the text for the definitions of the statistics.

	$\hat{\Delta}$	$\hat{\sigma}_{\hat{\Delta}}$	t	P -value	95 % C.I.
Example 7.2: Average weekly hours spent on housework	1.16	0.353	3.29	0.001	(0.47; 1.85)
Example 7.3: Perceived friendliness of a police officer	-1.74	0.383	-4.55	< 0.001	(-2.49; -0.99)

If necessary, calculations like these can be carried out even with a pocket calculator. It is, however, much more convenient to leave them to statistical software. Figure 7.4 shows SPSS output for the two-sample t-test for the housework data. The first part of the table, labelled “Group Statistics”, shows the sample sizes n , means \bar{Y} and standard deviations s separately for the two groups. The quantity labelled “Std. Error Mean” is s/\sqrt{n} . This is an estimate of the standard error of the sample mean, which is the quantity σ/\sqrt{n} discussed in Section 6.4.

The second part of the table in Figure 7.4, labelled “Independent Samples Test”, gives results for the t-test itself. The test considered here, which assumes a common population standard deviation σ (and thus also variance σ^2), is found on the row labelled “Equal variances assumed”. The test statistic is shown in the column labelled “ t ”, and the difference $\hat{\Delta} = \bar{Y}_2 - \bar{Y}_1$ and its standard error $\hat{\sigma}_{\hat{\Delta}}$ are shown in the “Mean Difference” and “Std. Error Difference” columns respectively. Note that the difference (-1.16) has been calculated in SPSS between men and women rather than vice versa as in Table 7.3, but this will make no difference to the conclusions from the test.

In the two-sample situation with assumptions 1–4 at the beginning of Section 7.3.1, the sampling distribution of the t-test statistic (7.12) is approximately a standard normal distribution when the null hypothesis $H_0 : \Delta = \mu_2 - \mu_1 = 0$ is true in the population and the sample sizes are large enough. This is again a consequence of the Central Limit Theorem. The requirement for “large enough” sample sizes is fairly easy to satisfy. A good rule of thumb is that the sample sizes n_1 and n_2 in the two groups should both be at least 20 for the sampling distribution of the test statistic to be well enough approximated by the standard normal distribution. In the housework example we have data on 635 men and 469 women, so the sample sizes are clearly large enough. A variant of the test which relaxes the condition on the sample sizes is discussed in Section 7.3.4 below.

The P -value of the test is calculated from this sampling distribution in exactly the same way as for the tests of proportions in Section 5.5.3. In the housework example the value of the t -test statistic is $t = 3.29$. The P -value for testing the null hypothesis against the two-sided alternative (7.6) is then the probability, calculated from the standard normal distribution, of values that are at least 3.29 or at most -3.29. Each of these two probabilities is about 0.0005,

Group Statistics									
Sex - HH grid		N	Mean	Std. Deviation	Std. Error Mean				
Hours per week on housework	Male	635	7.33	5.528	.219				
	Female	469	8.49	6.141	.284				

Independent Samples Test										
		Levene's Test for Equality of Variances		t-test for Equality of Means						
		F	Sig.	t	df	Sig. (2-tailed)	Mean Difference	Std. Error Difference	95% Confidence Interval of the Difference	
									Lower	Upper
Hours per week on housework	Equal variances assumed	5.957	.015	-3.287	1102	.001	-1.160	.353	-1.853	-.468
	Equal variances not assumed			-3.236	945.777	.001	-1.160	.359	-1.864	-.457

Figure 7.4: SPSS output for a two-sample t -test in Example 7.2, comparing average weekly hours spent on housework between men and women.

so the P -value is $0.0005 + 0.0005 = 0.001$. In the SPSS output of Figure 7.4 it is given in the column labelled “Sig. (2-tailed)”, where “Sig.” is short for “significance” and “2-tailed” is a synonym for “2-sided”.

The P -value can also be calculated approximately using the table of the standard normal distribution (see Table 5.2, as explained in Section 5.5.3. Here the test statistic $t = 3.29$, which is larger than the critical values 1.65, 1.96 and 2.58 for the 0.10, 0.05 and 0.01 significance levels for a two-sided test, so we can report that $P < 0.01$. Here t is by chance actually equal (to two decimal places) to the critical value for the 0.001 significance level, so we could also report $P = 0.001$. These findings agree, as they should, with the exact P -value of 0.001 shown in the SPSS output.

In conclusion, the two-sample t -test in Example 7.2 indicates that there is very strong evidence (with $P = 0.001$ for the two-sided test) against the claim that the hours of weekly housework are on average the same for men and women in the population.

Here we showed raw SPSS output in Figure 7.4 because we wanted to explain its contents and format. Note, however, that such unedited computer output is rarely if ever appropriate in research reports. Instead, results of statistical analyses should be given in text or tables formatted in appropriate ways for presentation. See Table 7.3 and various other examples in this coursepack and textbooks on statistics.

To summarise the elements of the test again, we repeat them briefly, now for Example 7.3, the experiment on the effect of eye contact on the perceived friendliness of police officers (c.f. Table 7.1 for the summary statistics):

1. Data: samples from two groups, one with the experimental condition where the officer wore no sunglasses, with sample size $n_1 = 67$, mean $\bar{Y}_1 = 8.23$ and standard deviation $s_1 = 2.39$, and the second with the experimental condition where the officer did wear sunglasses, with $n_2 = 66$, $\bar{Y}_2 = 6.49$ and $s_2 = 2.01$.
2. Assumptions: the observations are random samples of statistically independent observations from two populations, one with mean μ_1 and standard deviation σ_1 , and the other with mean μ_2 and the same standard deviation σ_2 , where the standard deviations are equal, with value $\sigma = \sigma_1 = \sigma_2$. The sample sizes n_1 and n_2 are sufficiently large, say both at least 20, for the sampling distribution of the test statistic under the null hypothesis to be approximately standard normal.
3. Hypotheses: These are about the difference of the population means $\Delta = \mu_2 - \mu_1$, with null hypothesis $H_0 : \Delta = 0$. The two-sided alternative hypothesis $H_a : \Delta \neq 0$ is considered in this example.
4. The test statistic: the two-sample t -statistic

$$t = \frac{\hat{\Delta}}{\hat{\sigma}_{\hat{\Delta}}} = \frac{-1.74}{0.383} = -4.55$$

where

$$\hat{\Delta} = \bar{Y}_2 - \bar{Y}_1 = 6.49 - 8.23 = -1.74$$

and

$$\hat{\sigma}_{\hat{\Delta}} = \hat{\sigma} \sqrt{\frac{1}{n_2} + \frac{1}{n_1}} = 2.210 \times \sqrt{\frac{1}{66} + \frac{1}{67}} = 0.383$$

with

$$\hat{\sigma} = \sqrt{\frac{(n_2 - 1)s_2^2 + (n_1 - 1)s_1^2}{n_1 + n_2 - 2}} = \sqrt{\frac{65 \times 2.01^2 + 66 \times 2.39^2}{131}} = 2.210$$

5. The sampling distribution of the test statistic when H_0 is true: approximately the standard normal distribution.
6. The P -value: the probability that a randomly selected value from the standard normal distribution is at most -4.55 or at least 4.55 , which is about 0.000005 (reported as $P < 0.001$).
7. Conclusion: A two-sample t -test indicates very strong evidence that the average perceived level of the friendliness of a police officer is different when the officer is wearing reflective sunglasses than when the officer is not wearing such glasses ($P < 0.001$).

7.3.3 Confidence intervals for a difference of two means

A confidence interval for the mean difference $\Delta = \mu_1 - \mu_2$ is obtained by substituting appropriate expressions into the general formula (7.4). Specifically, here $\hat{\Delta} = \bar{Y}_2 - \bar{Y}_1$ and a 95% confidence interval for Δ is

$$(\bar{Y}_2 - \bar{Y}_1) \pm 1.96 \hat{\sigma} \sqrt{\frac{1}{n_2} + \frac{1}{n_1}} \quad (7.13)$$

where $\hat{\sigma}$ is obtained from equation 7.11. The validity of this again requires that the sample sizes n_1 and n_2 from both groups are reasonably large, say both at least 20. For the housework Example 7.2, the 95% confidence interval is

$$1.16 \pm 1.96 \times 0.353 = 1.16 \pm 0.69 = (0.47; 1.85)$$

using the values of $\bar{Y}_2 - \bar{Y}_1$ and its standard error calculated earlier. This interval is also shown in Table 7.3 and in the SPSS output in Figure 7.4. In the latter, the interval is given as $(-1.85; -0.47)$ because it is expressed for the difference defined in the opposite direction (men – women instead of vice versa). For Example 7.3, the 95% confidence interval is $-1.74 \pm 1.96 \times 0.383 = (-2.49; -0.99)$.

Based on the data in Example 7.2 we are thus 95 % confident that the difference between women's and men's average hours of reported weekly housework in the population is between 0.47 and 1.85 hours. In substantive terms this interval, from just under half an hour to nearly two hours, is arguably fairly wide in that its two end points might well be regarded as substantially different from each other. The difference between women's and men's average housework hours is thus estimated fairly imprecisely from this survey.

7.3.4 Variants of the test and confidence interval

Allowing unequal population variances

The two-sample t -test and confidence interval for the difference of means were stated above under the assumption that the standard deviations σ_1 and σ_2 of the variable of interest Y are the same in both of the two groups being compared. This assumption is not in fact essential. If it is omitted, we obtain formulas which differ from the ones discussed above only in one part of the calculations.

Suppose that we do allow the unknown values of σ_1 and σ_2 to be different from each other. In other words, we consider the model stated at the beginning of Section 7.3.1, without assumption 4 that $\sigma_1 = \sigma_2$. The test statistic is then still of the same form as before, i.e. $t = \hat{\Delta} / \hat{\sigma}_{\hat{\Delta}}$, with

$\hat{\Delta} = \bar{Y}_2 - \bar{Y}_1$. The only change in the calculations is that the estimate of the standard error of $\hat{\Delta}$, the formula of which is given by equation (7.8), now uses separate estimates of σ_1 and σ_2 . The obvious choices for these are the corresponding sample standard deviations, s_1 for σ_1 and s_2 for σ_2 . This gives the estimated standard error as

$$\hat{\sigma}_{\hat{\Delta}} = \sqrt{\frac{s_2^2}{n_2} + \frac{s_1^2}{n_1}}. \quad (7.14)$$

Substituting this to the formula of the test statistic yields the two-sample t -test statistic without the assumption of equal population standard deviations,

$$t = \frac{\bar{Y}_2 - \bar{Y}_1}{\sqrt{s_2^2/n_2 + s_1^2/n_1}}. \quad (7.15)$$

The sampling distribution of this under the null hypothesis is again approximately a standard normal distribution when the sample sizes n_1 and n_2 are both at least 20. The P -value for the test is obtained in exactly the same way as before, and the principles of interpreting the result of the test are also unchanged.

For the confidence interval, the only change from Section 7.3.3 is again that the estimated standard error is changed, so for a 95% confidence interval we use

$$(\bar{Y}_2 - \bar{Y}_1) \pm 1.96 \sqrt{\frac{s_2^2}{n_2} + \frac{s_1^2}{n_1}}. \quad (7.16)$$

In the housework example 7.2, the estimated standard error (7.14) is

$$\hat{\sigma}_{\hat{\Delta}} = \sqrt{\frac{6.14^2}{469} + \frac{5.53^2}{635}} = \sqrt{0.1285} = 0.359,$$

the value of the test statistic is

$$t = \frac{1.16}{0.359} = 3.23,$$

and the two-sided P -value is now $P = 0.001$. Recall that when the population standard deviations were assumed to be equal, we obtained $\hat{\sigma}_{\hat{\Delta}} = 0.353$, $t = 3.29$ and again $P = 0.001$. The two sets of results are thus very similar, and the conclusions from the test are the same in both cases. The differences between the two variants of the test are even smaller in Example 7.3, where the estimated standard error $\hat{\sigma}_{\hat{\Delta}} = 0.383$ is the same (to three decimal places) in both cases, and the results are thus identical⁶. In both examples the confidence intervals obtained from (7.13) and (7.16) are also very similar. Both variants of the two-sample analyses are shown in SPSS output (c.f. Figure 7.4), the ones assuming equal population standard deviations on the row labelled “Equal variances assumed” and the one without this assumption on the “Equal variances not assumed” row⁷.

Which methods should we then use, the ones with or without the assumption of equal population variances? In practice the choice rarely makes much difference, and the P -values and conclusions

⁶The data can be obtained from <http://bes2009-10.org/>, which gives further information on the survey, including the full text of the questionnaires. The data analysed in this class and homework are from the BES Campaign Internet Panel Survey, which has been divided into two data sets corresponding to two time periods leading up to the General Election.

⁷Official results obtained from www.olympic.org/london-2012-summer-olympics.

from the two versions of the test are typically very similar⁸. Not assuming the variances to be equal has the advantage of making fewer restrictive assumptions about the population. For this reason it should be used in the rare cases where the P -values obtained under the different assumptions are substantially different. This version of the test statistic is also slightly easier to calculate by hand, since (7.14) is a slightly simpler formula than (7.10)–(7.11). On the other hand, the test statistic which does assume equal standard deviations has the advantage that it is more closely related to analogous tests used in more general contexts (especially the method of linear regression modelling, discussed in Chapter 8). It is also preferable when the sample sizes are very small, as discussed below.

Using the t distribution

As discussed in Section 6.3, it is often assumed that the population distributions of the variables under consideration are described by particular probability distributions. In this chapter, however, such assumptions have so far been avoided. This is a consequence of the Central Limit Theorem, which ensures that as long as the sample sizes are large enough, the sampling distribution of the two-sample t -test statistic is approximately the standard normal distribution, irrespective of the forms of the population distributions of Y in the two groups. In this section we briefly describe variants of the test and confidence interval which *do* assume that the population distributions are of a particular form, specifically that they are normal distributions. This changes the sampling distribution that is used for the test statistic and for the multiplier of the confidence interval, but the analyses are otherwise unchanged.

For the significance test, there are again two variants depending on the assumptions about the the population standard deviations σ_1 and σ_2 . Consider first the case where these are assumed to be equal. The sampling distribution is then given by the following result, which now holds for *any* sample sizes n_1 and n_2 :

- In the two-sample situation specified by assumptions 1–4 at the beginning of Section 7.3.1 (including the assumption of equal population standard deviations, $\sigma_1 = \sigma_2 = \sigma$), and if also the distribution of Y is a normal distribution in both groups, the sampling distribution of the t -test statistic (7.12) is a t distribution with $n_1 + n_2 - 2$ degrees of freedom when the null hypothesis $H_0 : \Delta = \mu_2 - \mu_1 = 0$ is true in the population.

The **t distributions** mentioned in this result are a family of distributions with different degrees of freedom, in a similar way as the χ^2 distributions discussed in Section 4.3.4. All t distributions are symmetric around 0. Their shape is quite similar to that of the standard normal distribution, except that the variance of a t distribution is somewhat larger and its tails thus heavier. The difference is noticeable only when the degrees of freedom are small, as seen in Figure 7.5. This shows the curves for the t distributions with 6 and 30 degrees of freedom, compared to the standard normal distribution. It can be seen that the t_{30} distribution is already very similar to the $N(0, 1)$ distribution. With degrees of freedom larger than about 30, the difference becomes almost indistinguishable.

If we use this result for the test, the P -value is obtained from the t distribution with $n_1 + n_2 - 2$ degrees of freedom (often denoted $t_{n_1+n_2-2}$). The principles of doing this are exactly the same as those described in Section 5.5.3, and can be graphically illustrated by plots similar to those in Figure 5.1. Precise P -values are again obtained using a computer. In fact, P -values in SPSS output for the two-sample t -test (c.f. Figure 7.4) are actually those obtained from the t distribution (with the degrees of freedom shown in the column labelled “df”) rather than

⁸The data can be obtained from www3.norc.ox.ac.uk/GSS+Website/, which gives further information on the survey, including the full text of the questionnaires.




Figure 7.5: Curves of two t distributions with small degrees of freedom, compared to the standard normal distribution.

the standard normal distribution. Differences between the two are, however, very small if the sample sizes are even moderately large, because then the degrees of freedom $df = n_1 + n_2 - 2$ are large enough for the two distributions to be virtually identical. This is the case, for instance, in both of the examples considered so far in this chapter, where $df = 1102$ in Example 7.2 and $df = 131$ in Example 7.3.

If precise P -values from the t distribution are not available, upper bounds for them can again be obtained using appropriate tables, in the same way as in Section 5.5.3. Now, however, the critical values depend also on the degrees of freedom. Because of this, introductory text books on statistics typically include a table of critical values for t distributions for a selection of degrees of freedom. A table of this kind is shown in the Appendix at the end of this course pack. Each row of the table corresponds to a t distribution with the degrees of freedom given in the column labelled “df”. As here, such tables typically include all degrees of freedom between 1 and 30, plus a selection of larger values, here 40, 60 and 120.

The last row is labelled “ ∞ ”, the mathematical symbol for infinity. This corresponds to the standard normal distribution, as a t distribution with infinite degrees of freedom is equal to the standard normal. The practical implication of this is that the standard normal distribution is a good enough approximation for any t distribution with reasonably large degrees of freedom. The table thus lists individual degrees of freedom only up to some point, and the last row will be used for any values larger than this. For degrees of freedom between two values shown in the table (e.g. 50 when only 40 and 60 are given), it is best to use the values for the nearest available degrees of freedom *below* the required ones (e.g. use 40 for 50). This will give a “conservative” approximate P -value which may be slightly larger than the exact value.

As for the standard normal distribution, the table is used to identify critical values for different significance levels (c.f. the information in Table 5.2). For example, if the degrees of freedom are

20, the critical value for two-sided tests at the significance level 0.05 in the “0.025” column on the row labelled “20”. This is 2.086. In general, critical values for t distributions are somewhat larger than corresponding values for the standard normal distribution, but the difference between the two is quite small when the degrees of freedom are reasonably large.

The t -test and the t distribution are among the oldest tools of statistical inference. They were introduced in 1908 by W. S. Gosset⁹, initially for the one-sample case discussed in Section 7.4. Gosset was working as a chemist at the Guinness brewery at St. James’ Gate, Dublin. He published his findings under the pseudonym “Student”, and the distribution is often known as *Student’s t distribution*.

These results for the sampling distribution hold when the population standard deviations σ_1 and σ_2 are assumed to be equal. If this assumption is not made, the test statistic is again calculated using formulas (7.14) and (7.15). This case is mathematically more difficult than the previous one, because the sampling distribution of the test statistic under the null hypothesis is then not exactly a t distribution even when the population distributions are normal. One way of dealing with this complication (which is known as the Behrens–Fisher problem) is to find a t distribution which is a good approximation of the true sampling distribution. The degrees of freedom of this approximating distribution are given by

$$df = \frac{\left(\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}\right)^2}{\left(\frac{s_1^2}{n_1}\right)^2 \left(\frac{1}{n_1-1}\right) + \left(\frac{s_2^2}{n_2}\right)^2 \left(\frac{1}{n_2-1}\right)}. \quad (7.17)$$

This formula, which is known as the Welch–Satterthwaite approximation, is not particularly interesting or worth learning in itself. It is presented here purely for completeness, and to give an idea of how the degrees of freedom given in the SPSS output are obtained. In Example 7.2 (see Figure 7.4) these degrees of freedom are 945.777, showing that the approximate degrees of freedom from (7.17) are often not whole numbers. If approximate P -values are then obtained from a t -table, we need to use values for the nearest whole-number degrees of freedom shown in the table. This problem does not arise if the calculations are done with a computer.

Two sample t -test statistics (in two variants, under equal and unequal population standard deviations) have now been defined under two different sets of assumptions about the population distributions. In each case, the formula of the test statistic is the same, so the only difference is in the form of its sampling distribution under the null hypothesis. If the population distributions of Y in the two groups are assumed to be normal, the sampling distribution of the t -statistic is a t distribution with appropriate degrees of freedom. If the sample sizes are reasonably large, the sampling distribution is approximately standard normal, whatever the shape of the population distribution. Which set of assumptions should we then use? The following guidelines can be used to make the choice:

- The easiest and arguably most common case is the one where both sample sizes n_1 and n_2 are large enough (both greater than 20, say) for the standard normal approximation of the sampling distribution to be reasonably accurate. Because the degrees of freedom of the appropriate t distribution are then also large, the two sampling distributions are very similar, and conclusions from the test will be similar in either case. It is then purely a matter of convenience which sampling distribution is used:

⁹United Nations Development Programme *International Human Development Indicators*, <http://hdr.undp.org/en/data/>; World Bank *Worldwide Governance Indicators*, <http://info.worldbank.org/governance/wgi/pdf/wgidataset.xlsx>; World Bank *World Development Indicators*, <http://data.worldbank.org/indicator/SP.DYN.IMRT.IN>.

- If you use a computer (e.g. SPSS) to carry out the test or you are (e.g. in an exam) given computer output, use the P -value in the output. This will be from the t distribution.
- If you need to calculate the test statistic by hand and thus need to use tables of critical values to draw the conclusion, use the critical values for the standard normal distribution (see Table 5.2).
- When the sample sizes are small (e.g. if one or both of them are less than 20), only the t distribution can be used, and even then only if Y is approximately normally distributed in both groups in the population. For some variables (say weight or blood pressure) we might have some confidence that this is the case, perhaps from previous, larger studies. In other cases the normality of Y can only be assessed based on its sample distribution, which of course is not very informative when the sample is small. In most cases, some doubt will remain, so the results of a t -test from small samples should be treated with caution. An alternative is then to use *nonparametric* tests which avoid the assumption of normality, for example the so-called Wilcoxon–Mann–Whitney test. These, however, are not covered on this course.

There are also situations where the population distribution of Y cannot possibly be normal, so the possibility of referring to a t distribution does not arise. One example are the tests on population proportions that were discussed in Chapter 5. There the only possibility we discussed was to use the approximate standard normal sampling distribution, as long as the sample sizes were large enough. Because the t -distribution is never relevant there, the test statistic is conventionally called the z -test statistic rather than t . Sometimes the label z instead of t is used also for two-sample t -statistics described in this chapter. This does not change the test itself.

It is also possible to obtain a confidence interval for Δ which is valid for even very small sample sizes n_1 and n_2 , but only under the further assumption that the population distribution of Y in both groups is normal. This affects only the multiplier of the standard errors, which is now based on a t distribution. The appropriate degrees of freedom are again $df = n_1 + n_2 - 2$ when the population standard deviations are assumed equal, and approximately given by equation (7.17) if not. In this case the multiplier in (7.4) may be labelled $t_{\alpha/2}^{(df)}$ instead of $z_{\alpha/2}$ to draw attention to the fact that it comes from a t -distribution and depends on the degrees of freedom df as well as the significance level $1 - \alpha$.

Any multiplier $t_{\alpha/2}^{(df)}$ is obtained from the relevant t distribution using exactly the same logic as the one explained for the normal distribution in the previous section, using a computer or a table of t distributions. For example, in the t table in the Appendix, multipliers for a 95% confidence interval are the numbers given in the column labelled “0.025”. Suppose, for instance, that the sample sizes n_1 and n_2 are both 10 and population standard deviations are assumed equal, so that $df = 10 + 10 - 2 = 18$. The table shows that a t -based 95% confidence interval would then use the multiplier 2.101. This is somewhat larger than the corresponding multiplier 1.96 from the normal distribution, and the t -based interval is somewhat wider than one based on the normal distribution. The difference between the two becomes very small when the sample sizes are even moderately large, because then df is large and $t_{\alpha/2}^{(df)}$ is very close to 1.96.

The choice between confidence intervals based on the normal or a t distribution involves the same considerations as for the significance test. In short, if the sample sizes are not very small, the choice makes little difference and can be based on convenience. If you are calculating an interval by hand, a normal-based one is easier to use because the multiplier (e.g. 1.96 for 95% intervals) does not depend on the sample sizes. If, instead, a computer is used, it typically gives confidence intervals for differences of means based on the t distribution, so these are easier to use.

Finally, if one or both of the sample sizes are small, only t -based intervals can safely be used, and then only if you are confident that the population distributions of Y are approximately normal.

7.4 Tests and confidence intervals for a single mean

The task considered in this section is inference on the population mean of a continuous, interval-level variable Y in a single population. This is thus analogous to the analysis of a single proportion in Sections 5.5–5.6, but with a continuous variable of interest.

We use Example 7.1 on survey data on diet for illustration. We will consider two variables, daily consumption of portions of fruit and vegetables, and the percentage of total daily energy intake obtained from fat and fatty acids. These will be analysed separately, each in turn in the role of the variable of interest Y . Summary statistics for the variables are shown in Table 7.4

Table 7.4: Summary statistics, t -tests and confidence intervals for the mean for the two variables in Example 7.1 (variables from the Diet and Nutrition Survey). n =sample size; \bar{Y} =sample mean; s =sample standard deviation; μ_0 =null hypothesis about the population mean; t = t -test statistic; *: Alternative hypothesis $H_a : \mu \neq \mu_0$; †: Alternative hypotheses $H_a : \mu < 5$ and $\mu > 35$ respectively.

Variable	n	\bar{Y}	s	μ_0	t	P -value	P -value	95% CI for μ
						Two- sided*	One-sided†	
Fruit and vegetable consumption (400g portions)	1724	2.8	2.15	5	- 49.49	< 0.001	< 0.001	(2.70; 2.90)
Total energy intake from fat (%)	1724	35.3	6.11	35	2.04	0.042	0.021	(35.01; 35.59)

The setting for the analysis of this section is summarised as a statistical model for observations of a variable Y as follows:

1. The population distribution of Y has some unknown mean μ and unknown standard deviation σ .
2. The observations Y_1, Y_2, \dots, Y_n in the sample are a random sample from the population.
3. The observations are statistically independent, as discussed at the beginning of Section 7.3.1.

It is not necessary to assume that the population distribution has a particular form. However, this is again sometimes assumed to be a normal distribution, in which case the analyses may be modified in ways discussed below.

The only quantity of interest considered here is μ , the population mean of Y . In the diet examples this is the mean number of portions of fruit and vegetables, or mean percentage of energy derived from fat (both on an average day for an individual) among the members of the population (which for this survey is British adults aged 19–64).

Because no separate groups are being compared, questions of interest are now not about differences between different group means, but about the value of μ itself. The best single estimate (*point estimate*) of μ is the sample mean \bar{Y} . More information is provided by a confidence interval which shows which values of μ are plausible given the observed data.

Significance testing focuses on the question of whether it is plausible that the true value of μ is equal to a particular value μ_0 specified by the researcher. The specific value of μ_0 to be tested is suggested by the research questions. For example, we will consider $\mu_0 = 5$ for portions of fruit and vegetables and $\mu_0 = 35$ for the percentage of energy from fat. These values are chosen because they correspond to recommendations by the Department of Health that we should consume at least 5 portions of fruit and vegetables a day, and that fat should contribute no more than 35% of total energy intake. The statistical question is thus whether the average level of consumption in the population is at the recommended level.

In this setting, the null hypothesis for a significance test will be of the form

$$H_0 : \mu = \mu_0, \quad (7.18)$$

i.e. it claims that the unknown population mean μ is equal to the value μ_0 specified by the null hypothesis. This will be tested against the two-sided alternative hypothesis

$$H_a : \mu \neq \mu_0 \quad (7.19)$$

or one of the one-sided alternative hypotheses

$$H_a : \mu > \mu_0 \quad (7.20)$$

or

$$H_a : \mu < \mu_0. \quad (7.21)$$

For example, we might consider the one-sided alternative hypotheses $H_a : \mu < 5$ for portions of fruit and vegetables and $H_a : \mu > 35$ for the percentage of energy from fat. For both of these, the alternative corresponds to a difference from μ_0 in the unhealthy direction, i.e. less fruit and vegetables and more fat than are recommended.

To establish a connection to the general formulas that have been stated previously, it is again useful to express these hypotheses in terms of

$$\Delta = \mu - \mu_0, \quad (7.22)$$

i.e. the difference between the unknown true mean μ and the value μ_0 claimed by the null hypothesis. Because this is 0 if and only if μ and μ_0 are equal, the null hypothesis (7.18) can also be expressed as

$$H_0 : \Delta = 0, \quad (7.23)$$

and possible alternative hypotheses as

$$H_a : \Delta \neq 0, \quad (7.24)$$

$$H_0 : \Delta > 0 \quad (7.25)$$

and

$$H_0 : \Delta < 0, \quad (7.26)$$

corresponding to (7.19), (7.20) and (7.21) respectively.

The general formulas summarised in Section 7.3.1 can again be used, as long as their details are modified to apply to Δ defined as $\mu - \mu_0$. The resulting formulas are listed briefly below, and then illustrated using the data from the diet survey:

- The point estimate of the difference $\Delta = \mu - \mu_0$ is

$$\hat{\Delta} = \bar{Y} - \mu_0. \quad (7.27)$$

- The standard error of $\hat{\Delta}$, i.e. the standard deviation of its sampling distribution, is $\sigma_{\hat{\Delta}} = \sigma/\sqrt{n}$ (note that this is equal to the standard error $\sigma_{\bar{Y}}$ of the sample mean \bar{Y} itself¹⁰). This is estimated by

$$\hat{\sigma}_{\hat{\Delta}} = \frac{s}{\sqrt{n}}. \quad (7.28)$$

- The t -test statistic for testing the null hypothesis (7.23) is

$$t = \frac{\hat{\Delta}}{\hat{\sigma}_{\hat{\Delta}}} = \frac{\bar{Y} - \mu_0}{s/\sqrt{n}}. \quad (7.29)$$

- The sampling distribution of the t -statistic, when the null hypothesis is true, is approximately a standard normal distribution, when the sample size n is reasonably large. A common rule of thumb is that this sampling distribution is adequate when n is at least 30.
 - Alternatively, we may make the further assumption that the population distribution of Y is normal, in which case no conditions on n are required. The sampling distribution of t is then a t distribution with $n - 1$ degrees of freedom. The choice of which sampling distribution to refer to is based on the considerations outlined in Section 7.3.4. When n is 30 or larger, the two approaches give very similar results.
- P -values are obtained and the conclusions drawn in the same way as for two-sample tests, with appropriate modifications to the wording of the conclusions.
- A confidence interval for Δ , with confidence level $1 - \alpha$ and based on the approximate normal sampling distribution, is given by

$$\hat{\Delta} \pm z_{\alpha/2} \hat{\sigma}_{\hat{\Delta}} = (\bar{Y} - \mu_0) \pm z_{\alpha/2} \frac{s}{\sqrt{n}} \quad (7.30)$$

¹⁰ESS Round 5: European Social Survey Round 5 Data (2010). Data file edition 2.0. Norwegian Social Science Data Services, Norway. Data Archive and distributor of ESS data. The full data can be obtained from <http://ess.nsd.uib.no/ess/round5/>.

where $z_{\alpha/2}$ is the multiplier from the standard normal distribution for the required significance level (see Table 5.4), most often 1.96 for a 95% confidence interval. If an interval based on the t distribution is wanted instead, $z_{\alpha/2}$ is replaced by the corresponding multiplier $t_{\alpha/2}^{(n-1)}$ from the t_{n-1} distribution.

Instead of the interval (7.30) for the difference $\Delta = \mu - \mu_0$, it is usually more sensible to report a confidence interval for μ itself. This is given by

$$\bar{Y} \pm z_{\alpha/2} \frac{s}{\sqrt{n}}, \quad (7.31)$$

which is obtained by adding μ_0 to both end points of (7.30).

For the fruit and vegetable variable in the diet example, the mean under the null hypothesis is the dietary recommendation $\mu_0 = 5$. The estimated difference (7.27) is

$$\hat{\Delta} = 2.8 - 5 = -2.2$$

and its estimated standard error (7.28) is

$$\hat{\sigma}_{\hat{\Delta}} = \frac{2.15}{\sqrt{1724}} = 0.05178,$$

so the t -test statistic (7.29) is

$$t = \frac{-2.2}{0.05178} = -42.49.$$

To obtain the P -value for the test, $t = -42.49$ is referred to the sampling distribution under the null hypothesis, which can here be taken to be the standard normal distribution, as the sample size $n = 1723$ is large. If we consider the two-sided alternative hypothesis $H_a : \Delta \neq 0$ (i.e. $H_a : \mu \neq 5$), the P -value is the probability that a randomly selected value from the standard normal distribution is at most -42.49 or at least 42.49 . This is a very small probability, approximately $0.00 \dots 019$, with 268 zeroes between the decimal point and the 1. This is, of course, to all practical purposes zero, and can be reported as $P < 0.001$. The null hypothesis $H_0 : \mu = 5$ is rejected at any conventional level of significance. A t -test for the mean indicates very strong evidence that the average daily number of portions of fruit and vegetables consumed by members of the population differs from the recommended minimum of five.

If we considered instead the one-sided alternative hypothesis $H_a : \Delta < 0$ (i.e. $H_a : \mu < 5$), the observed sample mean $\bar{Y} = 2.8 < 5$ is in the direction of this alternative. The P -value is then the one-sided P -value divided by 2, which is here a small value reported as $P < 0.001$ again. The null hypothesis $H_0 : \mu = 5$ (and by implication also the one-sided null hypothesis $H_0 : \mu \geq 5$, as discussed at the end of Section 5.5.1) is thus also rejected in favour of this one-sided alternative, at any conventional significance level.

A 95% confidence interval for μ is obtained from (7.31) as

$$2.8 \pm 1.96 \times \frac{2.15}{\sqrt{1724}} = 2.8 \pm 1.96 \times 0.05178 = 2.8 \pm 0.10 = (2.70; 2.90).$$

We are thus 95% confident that the average daily number of portions of fruit and vegetables consumed by members of the population is between 2.70 and 2.90.

Figure 7.6 shows how these results for the fruit and vegetable variable are displayed in SPSS output. The label “portions” refers to the name given to the variable in the SPSS data file, and “Test Value = 5” indicates the null hypothesis value μ_0 being tested. Other parts of the SPSS output correspond to the information in Table 7.4 in fairly obvious ways, so “N” indicates

the sample size n (and not a population size, which is denoted by N in our notation), “Mean” the sample mean \bar{Y} , “Std. Deviation” the sample standard deviation s , “Std. Error Mean” the estimate of the standard error of the mean given by $s/\sqrt{n} = 2.15/\sqrt{1724} = 0.05178$, “Mean Difference” the difference $\hat{\Delta} = \bar{Y} - \mu_0 = 2.8 - 5 = -2.2$, and “t” the t -test statistic (7.29). The P -value against the two-sided alternative hypothesis is shown as “Sig. (2-tailed)” (reported in the somewhat sloppy SPSS manner as “.000”). This is actually obtained from the t distribution, the degrees of freedom of which ($n - 1 = 1723$) are given under “df”. Finally, the output also contains a 95% confidence interval for the difference $\Delta = \mu - \mu_0$, i.e. the interval (7.30).¹¹ This is given as $(-2.30; -2.10)$. To obtain the more convenient confidence interval (7.31) for μ itself, we only need to add $\mu_0 = 5$ to both end points of the interval shown by SPSS, to obtain $(-2.30 + 5; -2.10 + 5) = (2.70; 2.90)$ as before.

Similar results for the variable on the percentage of dietary energy obtained from fat are also shown in Table 7.4. Here $\mu_0 = 35$, $\hat{\Delta} = 35.3 - 35 = 0.3$, $\hat{\sigma}_{\hat{\Delta}} = 6.11/\sqrt{1724} = 0.147$, $t = 0.3/0.147$, and the two-sided P -value is $P = 0.042$. Here $P < 0.05$, so null hypothesis that the population average of the percentage of energy obtained from fat is 35 is rejected at the 5% level of significance. However, because $P > 0.01$, the hypothesis would not be rejected at the next conventional significance level of 1%. The conclusions are the same if we considered the one-sided alternative hypothesis $H_a : \mu > 35$, for which $P = 0.042/2 = 0.021$ (as the observed sample mean $\bar{Y} = 35.3$ is in the direction of H_a). In this case the evidence against the null hypothesis is thus somewhat less strong than for the fruit and vegetable variable, for which the P -value was extremely small. The 95% confidence interval for the population average of the fat variable is $35.3 \pm 1.96 \times 0.147 = (35.01; 35.59)$.

Analysis of a single population mean is a good illustration of some of the advantages of confidence intervals over significance tests. First, a confidence interval provides a summary of all the plausible values of μ even when, as is very often the case, there is no obvious single value μ_0 to be considered as the null hypothesis of the one-sample t -test. Second, even when such a significance test is sensible, the conclusion can also be obtained from the confidence interval, as discussed at the end of Section 5.6.4. In other words, $H_0 : \mu = \mu_0$ is rejected at a given significance level against a two-sided alternative hypothesis, if the confidence interval for μ at the corresponding confidence level does not contain μ_0 , and not rejected if the interval contains μ_0 . Here the 95% confidence interval $(2.70; 2.90)$ does not contain 5 for the fruit and vegetable variable, and the interval $(35.01; 35.59)$ does not contain 35 for the fat variable, so the null hypotheses with these values as μ_0 are rejected at the 5% level of significance.

The width of a confidence interval also gives information on how precise the results of the statistical analysis are. Here the intervals seem quite narrow for both variables, in that it seems that their end points (e.g. 2.7 and 2.9 for portions of fruit and vegetables) would imply qualitatively similar conclusions about the level of consumption in the population. Analysis of the sample of 1724 respondents in the National Diet and Nutrition Survey thus appears to have given us quite precise information on the population averages for most practical purposes. Of course, what is precise enough ultimately depends on what those purposes are. If much higher precision was required, the sample size in the survey would have to be correspondingly larger.

Finally, in cases where a null hypothesis is rejected by a significance test, a confidence interval has the additional advantage of providing a way to assess whether the observed deviation from the null hypothesis seems large in some *substantive* sense. For example, the confidence interval for the fat variable draws attention to the fact that the evidence against a population mean of 35 is not very strong. The lower bound of the interval is only 0.01 units above 35, which is very

¹¹This is another old idea. Different approaches to the problem of fitting curves to observations were gradually developed by Tobias Mayer, Rudjer Bošković and Pierre Simon Laplace from the 1750s onwards, and the method of least squares itself was presented by Adrien Marie Legendre in 1805.

One-Sample Statistics						
	N	Mean	Std. Deviation	Std. Error Mean		
portions	1724	2.8000	2.15000	.05178		

One-Sample Test						
	Test Value = 5					
	t	df	Sig. (2-tailed)	Mean Difference	95% Confidence Interval of the Difference	
					Lower	Upper
portions	-42.487	1723	.000	-2.20000	-2.3016	-2.0984

Figure 7.6: SPSS output for a t -test of a single mean. The output is for the variable on fruit and vegetable consumption in Table 7.4, with the null hypothesis $H - 0 : \mu = 5$.

little relative to the overall width (about 0.60) of the interval. The P -value (0.041) of the test, which is not much below the reference level of 0.05, also suggests this, but in a less obvious way. Even the upper limit (35.59) of the interval is arguably not very far from 35, so it suggests that we can be fairly confident that the population mean does not differ from 35 by very much in the substantive sense. This contrasts with the results for the fruit and vegetable variable, where all the values covered by the confidence interval (2.70; 2.90) are much more obviously far from the recommended value of 5.

7.5 Inference for dependent samples

In the two-sample cases considered in Section 7.3, the two groups being compared consisted of separate and presumably unrelated units (people, in all of these cases). It thus seemed justified to treat the groups as statistically independent. The third and last general case considered in this chapter is one where this assumption cannot be made, because there are some obvious connections between the groups. Examples 7.4 and 7.5 illustrate this situation. Specifically, in both cases we can find for each observation in one group a natural *pair* in the other group. In Example 7.4, the data consist of observations of a variable for a group of fathers at two time points, so the pairs of observations are clearly formed by the two measurements for each father. In Example 7.5 the basic observations are for separate days, but these are paired (*matched*) in that for each Friday the 13th in one group, the preceding Friday the 6th is included in the other. In both cases the existence of the pairings implies that we must treat the two groups as statistically *dependent*.

Data with dependent samples are quite common, largely because they are often very informative. Principles of good research design suggest that one key condition for being able to make valid and powerful comparisons between two groups is that the groups should be as similar as possible, apart from differing in the characteristic being considered. Dependent samples represent an attempt to achieve this through intelligent data collection. In Example 7.4, the comparison of interest is between a man's sense of well-being before and after the birth of his first child. It is likely that there are also other factors which affect well-being, such as personality and life circumstances unrelated to the birth of a child. Here, however, we can compare the well-being for the *same* men before and after the birth, which should mean that many of those other characteristics remain approximately unchanged between the two measurements. Information on the effects of the birth of a child will then mostly come not from overall levels of well-being but *changes* in it for each man.

In Example 7.5, time of the year and day of the week are likely to have a very strong effect on traffic levels. Comparing, say, Friday, November 13th to Friday, July 6th, let alone to Sunday, November 15th, would thus not provide much information about possible additional differences which were due specifically to a Friday being the 13th. To keep these other characteristics approximately constant and thus to focus on the effects of Friday the 13th, each such Friday has here been matched with the nearest preceding Friday. With this design, data on just ten matched pairs will (as seen below) allow us to conclude that the differences are statistically significant.

Generalisations of the research designs illustrated by Examples 7.4 and 7.5 allow for measurements at more than two occasions for each subject (so-called longitudinal or panel studies) and groups of more than two matched units (clustered designs). Most of these are analysed using statistical methods which are beyond the scope of this course. The paired case is an exception, for which the analysis is in fact easier than for two independent samples. This is because the pairing of observations allows us to reduce the analysis into a one-sample problem, simply by

considering within-pair *differences* in the response variable Y . Only the case where Y is a continuous variable is considered here. There are also methods of inference for comparing two (or more) dependent samples of response variables of other types, but they are not covered here.

The quantity of interest is again a population difference. This time it can be formulated as $\Delta = \mu_2 - \mu_1$, where μ_1 is the mean of Y for the first group (e.g. the first time point in Example 7.4) and μ_2 its mean for the second group. Methods of inference for Δ will again be obtained using the same general results which were previously applied to one-sample analyses and comparisons of two independent samples. The easiest way to do this is now to consider a new variable D , defined for each pair i as $D_i = Y_{2i} - Y_{1i}$, where Y_{1i} denotes the value of the first measurement of Y for pair i , and Y_{2i} is the second measurement of Y for the same pair. In Example 7.4 this is thus the difference between a man's well-being after the birth of his first baby, and the same man's well-being before the birth. In Example 7.5, D is the difference in traffic flows on a stretch of motorway between a Friday the 13th and the Friday a week earlier (these values are shown in the last column of Table 7.2). The number of observations of D is the number of pairs, which is equal to the sample sizes n_1 and n_2 in each of the two groups (the case where one of the two measurements might be missing for some pairs is not considered here). We will denote it by n .

The population mean of the differences D is also $\Delta = \mu_2 - \mu_1$, so the observed values D_i can be used for inference on Δ . An estimate of Δ is the sample average of D_i , i.e.

$$\hat{\Delta} = \bar{D} = \frac{1}{n} \sum_{i=1}^n D_i. \quad (7.32)$$

In other words, this is the average of the within-pair differences between the two measurements of Y . Its standard error is estimated by

$$\hat{\sigma}_{\hat{\Delta}} = \frac{s_D}{\sqrt{n}} \quad (7.33)$$

where s_D is the sample standard deviation of D , i.e.

$$s_D = \sqrt{\frac{\sum (D_i - \bar{D})^2}{n-1}}. \quad (7.34)$$

A test statistic for the null hypothesis $H_0 : \Delta = 0$ is given by

$$t = \frac{\hat{\Delta}}{\hat{\sigma}_{\hat{\Delta}}} = \frac{\bar{D}}{s_D/\sqrt{n}} \quad (7.35)$$

and its P -value is obtained either from the standard normal distribution or the t_{n-1} distribution. A confidence interval for Δ with confidence level $1 - \alpha$ is given by

$$\hat{\Delta} \pm q_{\alpha/2} \times \hat{\sigma}_{\hat{\Delta}} = \bar{D} \pm q_{\alpha/2} \times \frac{s_D}{\sqrt{n}} \quad (7.36)$$

where the multiplier $q_{\alpha/2}$ is either $z_{\alpha/2}$ or $t_{\alpha/2}^{(n-1)}$. These formulas are obtained by noting that this is simply a one-sample analysis with the differences D in place of the variable Y , and applying the formulas of Section 7.4 to the observed values of D .

Table 7.5: Results of tests and confidence intervals for comparing means of two dependent samples. For Example 7.4, the difference is between after and before the birth of the child, and for Example 7.5 it is between Friday the 13th and the preceding Friday the 6th. See the text for the definitions of the statistics. (* Obtained from the t_9 distribution; † Obtained from the standard normal distribution.)

	strut	strut	Test of $H_0 : \Delta = 0$	Test of $H_0 : \Delta = 0$	strut
	$\hat{\Delta}$	$\hat{\sigma}_{\hat{\Delta}}$	t	P -value	95 % C.I. for Δ
Example 7.4: Father's personal well-being	0.08	0.247	0.324	0.75†	(-0.40; 0.56)
Example 7.5: Traffic flows on successive Friday	-1835	372	-4.93	0.001*	(-2676; -994)

Results for Examples 7.4 and 7.5 are shown in Table 7.5. To illustrate the calculations, consider Example 7.5. The $n = 10$ values of D_i for it are shown in Table 7.2, and the summary statistics $\bar{D} = -1835$ and $s_D = 1176$ in Table 7.1. The standard error of \bar{D} is thus $s_D/\sqrt{n} = 1176/\sqrt{10} = 372$ and the value of the test statistic (7.35) is

$$z = \frac{-1835}{1176/\sqrt{10}} = \frac{-1835}{372} = -4.93.$$

This example differs from others we have considered so far in that the sample size of $n = 10$ is clearly too small for us to rely on large-sample results. It is thus not appropriate to refer the test statistic to a standard normal distribution. Instead, P -values can be obtained from a t distribution, but only if the population distribution of D itself can be assumed to be approximately normal. Here we have only the ten observed values of D to use for a rather informal assessment of whether this assumption appears to be reasonable. One value of D is smaller than -4000, and 2, 5, 2 of them are in the ranges -3000 to -2001, -2000 to -1001, and -1000 to -1 respectively. Apart from the smallest observation, the sample distribution of D is thus at least approximately symmetric. While this definitely does not prove that D is normally distributed, it is at least not obviously inconsistent with such a claim. We thus feel moderately confident that we can apply here tests and confidence intervals based on the t distribution.

The P -value, obtained from a t distribution with $n - 1 = 9$ degrees of freedom, for the test statistic -4.93 is approximately 0.001. Even with only ten pairs of observations, there is significant evidence that the volume of traffic on a Friday the 13th differs from that of the preceding Friday. A confidence interval for the difference is obtained from (7.36) as

$$-1835 \pm 2.26 \times 372 = (-2676; -994)$$

where the multiplier 2.26 is the quantity $t_{\alpha/2}^{(n-1)} = t_{0.975}^{(9)}$, obtained from a computer or a table of the t_9 -distribution. The interval shows that we are 95% confident that the average reduction in traffic on Friday the 13th on the stretches of motorway considered here is between 994 and 2676 vehicles. This seems like a substantial systematic difference, although not particularly large as

a proportion of the total volume of traffic on those roads. In the absence of other information we are tempted to associate the reduction with some people avoiding driving on a day they consider to be unlucky.

In Example 7.4 the P -value is 0.75, so we cannot reject the null hypothesis that $\Delta = 0$. There is thus no evidence that there was a difference in first-time fathers' self-assessed level of well-being between the time their wives were six months pregnant, and a month after the birth of the baby. This is also indicated by the 95% confidence interval $(-0.40; 0.56)$ for the difference, which clearly covers the value 0 of no difference.

7.6 Further comments on significance tests

Some further aspects of significance testing are discussed here. These are not practical issues that need to be actively considered every time you carry out a test. Instead, they provide context and motivation for the principles behind significance tests.

7.6.1 Different types of error

Consider for the moment the approach to significance testing where the outcome is presented in the form of a discrete claim or decision about the hypotheses, stating that the null hypothesis was either rejected or not rejected. This claim can either be correct or incorrect, depending on whether the null hypothesis is true in the population. There are four possibilities, summarized in Table 7.6. Two of these are correct decisions and two are incorrect. The two kinds of incorrect decisions are traditionally called

- **Type I error:** rejecting the null hypothesis when it is true
- **Type II error:** not rejecting the null hypothesis when it is false

The terms are unmemorably bland, but they do at least suggest an order of importance. Type I error is conventionally considered more serious than Type II, so what we most want to avoid is rejecting the null hypothesis unnecessarily. This implies that we will maintain the null hypothesis unless data provide strong enough evidence to justify rejecting it, a principle which is somewhat analogous to the “keep a theory until falsified” thinking of Popperian philosophy of science, or even the “innocent until proven guilty” principle of jurisprudence.

Table 7.6: The four possible combinations of the truth of a null hypothesis H_0 in a population and decision about it from a significance test.

		H_0	H_0
		Not Rejected	Rejected
H_0 is	True	Correct decision	Type I error
	False	Type II error	Correct decision

Despite our dislike of Type I errors, we will not try to avoid them completely. The only way to guarantee that the null hypothesis is never incorrectly rejected is never to reject it at all, whatever the evidence. This is not a useful decision rule for empirical research. Instead, we will decide in advance how high a probability of Type I error we are willing to tolerate, and then use a test procedure with that probability. Suppose that we use a 5% level of significance to make decisions from a test. The null hypothesis is then rejected if the sample yields a test statistic for

which the P -value is less than 0.05. If the null hypothesis is actually true, such values are, by the definition of the P -value, obtained with probability 0.05. Thus the significance level (α -level) of a test is the probability of making a Type I error. If we use a large α -level (say $\alpha = 0.10$), the null hypothesis is rejected relatively easily (whenever P -value is less than 0.10), but the chances of committing a Type I error are correspondingly high (also 0.10); with a smaller value like $\alpha = 0.01$, the error probability is lower because H_0 is rejected only when evidence against it is quite strong.

This description assumes that the true probability of Type I error for a test *is* equal to its stated α -level. This is true when the assumptions of the test (about the population distribution, sample size etc.) are satisfied. If the assumptions fail, the true significance level will differ from the stated one, i.e. the P -value calculated from the standard sampling distribution for that particular test will differ from the true P -value which would be obtained from the exact sampling distribution from the population in question. Sometimes the difference is minor and can be ignored for most practical purposes (the test is then said to be *robust* to violations of some of its assumptions). In many situations, however, using an inappropriate test may lead to incorrect conclusions: for example, a test which claims that the P -value is 0.02 when it is really 0.35 will clearly give a misleading picture of the strength of evidence against the null hypothesis. To avoid this, the task of statisticians is to develop valid (and preferably robust) tests for many different kinds of hypotheses and data. The task of the empirical researcher is to choose a test which is appropriate for his or her data.

In the spirit of regarding Type I errors as the most serious, the worst kind of incorrect test is one which gives too low a P -value, i.e. exaggerates the strength of evidence against the null hypothesis. Sometimes it is known that this is impossible or unlikely, so that the P -value is either correct or too high. The significance test is then said to be *conservative*, because its true rate of Type I errors will be the same or lower than the stated α -level. A conservative procedure of statistical inference is regarded as the next best thing to one which has the correct level of significance. For example, when the sample size is relatively large, P -values for all of the tests discussed in this chapter may be calculated from a standard normal or from a t distribution. P -values from a t distribution are then always somewhat larger. This means that using the t distribution is (very slightly) conservative when the population distributions are not normal, so that we can safely use the P -values from SPSS output of a t -test even in that case (this argument does not, however, justify using the t -test when Y is not normally distributed and the sample size is small, because the sampling distribution of the t -test statistic may then be very far from normal).

7.6.2 Power of significance tests

After addressing the question of Type I error by selecting an appropriate test and deciding on the significance level to be used, we turn our attention to Type II errors. The probability that a significance test will reject the null hypothesis when it is in fact not true, i.e. the probability of *avoiding* a Type II error, is known as the **power** of the test. It depends, in particular, on

- The nature of the test. If several valid tests are available for a particular analysis, we would naturally prefer one which tends to have the highest power. One aim of theoretical statistics is to identify the most powerful test procedures for different problems.
- The sample size: other things being equal, larger samples mean higher power.
- The true value of the population parameter to be tested, here the population mean or proportion. The power of any test will be highest when the true value is very different from the value specified by the null hypothesis. For example, it will obviously be easier

to detect that a population mean differs from a null value of $\mu_0 = 5$ when the true mean is 25 than when it is 5.1.

- The population variability of the variable. Since large population variance translates into large sampling variability and hence high levels of uncertainty, the power will be low when population variability is large, and high if the population variability is low.

The last three of these considerations are often used at the design stage of a study to get an idea of the sample size required for a certain level of power, or of the power achievable with a given sample size. Since data collection costs time and money, we would not want to collect a much larger sample than is required for a level of certainty sufficient for the purposes of a study. On the other hand, if a preliminary calculation reveals that the largest sample we can afford would still be unlikely to give enough information to detect interesting effects, the study might be best abandoned.

A power calculation requires the researcher to specify the kinds of differences from a null hypothesis which are large enough to be of practical or theoretical interest, so that she or he would want to be able to detect them with high probability (it must always be accepted that the power will be lower for smaller differences). For example, suppose that we are planning a study to compare the effects of two alternative teaching methods on the performance of students in an examination where possible scores are between 0 and 100. The null hypothesis is that average results are the same for students taught with each method. It is decided that we want enough data to be able to reject this with high probability if the true difference Δ of the average exam scores between the two groups is larger than 5 points, i.e. $\Delta < -5$ or $\Delta > 5$. The power calculation might then answer questions like

- What is the smallest sample size for which the probability of rejecting $H_0 : \Delta = 0$ is at least 0.9, when the true value of Δ is smaller than -5 or larger than 5 ?
- The largest sample sizes we can afford are 1000 in both groups, i.e. $n_1 = n_2 = 1000$. What is the probability this gives us of rejecting $H_0 : \Delta = 0$ when the true value of Δ is smaller than -5 or larger than 5 ?

To answer these questions, we would also need a rough guess of the population standard deviations σ_1 and σ_2 , perhaps obtained from previous studies. Such calculations employ further mathematical results for test statistics, essentially using their sampling distributions under specific alternative hypotheses. The details are, however, beyond the scope of this course.

7.6.3 Significance vs. importance

The P -value is a measure of the strength of evidence the data provide against the null hypothesis. This is not the same as the magnitude of the difference between sample estimates and the null hypothesis, or the practical importance of such differences. As noted above, the power of a test increases with increasing sampling size. One implication of this is that when n is large, even quite small observed deviations from the values that correspond exactly to the null hypothesis will be judged to be statistically significant. Consider, for example, the two dietary variables in Table 7.4. The sample mean of the fat variable is 35.3, which is significantly different (at the 5% level of significance) from μ_0 of 35. It is possible, however, that a difference of 0.3 might be considered unimportant in practice. In contrast, the sample mean of the fruit and vegetable variable is 2.8, and the difference from μ_0 of 5 seems not only strongly significant but also large for most practical purposes.

In contrast to the large-sample case, in small samples even quite large apparent deviations from the null hypothesis might still result in a large P -value. For example, in a very small study a

sample mean of the fat variable of, say, 30 or even 50 might not be interpreted as sufficiently strong evidence against a population mean of 35. This is obviously related to the discussion of statistical power in the previous section, in that it illustrates what happens when the sample is too small to provide enough information for useful conclusions.

In these and all other cases, decisions about what is or is not of practical importance are subject-matter questions rather than statistical ones, and would have to be based on information about the nature and implications of the variables in question. In our dietary examples this would involve at least medical considerations, and perhaps also financial implications of the public health costs of the observed situation or of possible efforts of trying to change it.

Chapter 8

Linear regression models

8.1 Introduction

This chapter continues the theme of analysing statistical associations between variables. The methods described here are appropriate when the response variable Y is a continuous, interval level variable. We will begin by considering bivariate situations where the only explanatory variable X is also a continuous variable. Section 8.2 first discusses graphical and numerical descriptive techniques for this case, focusing on two very commonly used tools: a *scatterplot* of two variables, and a measure of association known as the *correlation* coefficient. Section 8.3 then describes methods of statistical inference for associations between two continuous variables. This is done in the context of a statistical model known as the *simple linear regression model*.

The ideas of simple linear regression modelling can be extended to a much more general and powerful set of methods known as *multiple linear regression models*. These can have several explanatory variables, which makes it possible to examine associations between any explanatory variable and the response variable, while controlling for other explanatory variables. An important reason for the usefulness of these models is that they play a key role in statistical analyses which correspond to research questions that are causal in nature. As an interlude, we discuss issues of causality in research design and analysis briefly in Section 8.4. Multiple linear models are then introduced in Section 8.5. The models can also include categorical explanatory variables with any number of categories, as explained in Section 8.6.

The following example will be used for illustration throughout this chapter:

Example 8.1: Indicators of Global Civil Society

The *Global Civil Society 2004/5* yearbook gives tables of a range of characteristics of the countries of the world¹. The following measures will be considered in this chapter:

- **Gross Domestic Product (GDP)** per capita in 2001 (in current international dollars, adjusted for purchasing power parity)
- **Income level** of the country in three groups used by the Yearbook, as Low income, Middle income or High income
- **Income inequality** measured by the Gini index (with 0 representing perfect equality and 100 perfect inequality)

¹ESS Round 5: European Social Survey Round 5 Data (2010). Data file edition 2.0. Norwegian Social Science Data Services, Norway. Data Archive and distributor of ESS data. The full data can be obtained from <http://ess.nsd.uib.no/ess/round5/>.

- A measure of **political rights and civil liberties** in 2004, obtained as the average of two indices for these characteristics produced by the Freedom House organisation (1 to 7, with higher values indicating more rights and liberties)
- World Bank Institute's measure of control of **corruption** for 2002 (with high values indicating low levels of corruption)
- Net **primary school enrolment** ratio 2000-01 (%)
- **Infant mortality rate** 2001 (% of live births)

We will discuss various associations between these variables. It should be noted that the analyses are mainly illustrative examples, and the choices of explanatory and response variables do not imply any strong claims about causal connections between them. Also, the fact that different measures refer to slightly different years is ignored; in effect, we treat each variable as a measure of “recent” situation in the countries. The full data set used here includes 165 countries. Many of the variables are not available for all of them, so most of the analyses below use a smaller number of countries.

8.2 Describing association between two continuous variables

8.2.1 Introduction

Suppose for now that we are considering data on two continuous variables. The descriptive techniques discussed in this section do not strictly speaking require a distinction between an explanatory variable and a response variable, but it is nevertheless useful in many if not most applications. We will reflect this in the notation by denoting the variables X (for the explanatory variable) and Y (for the response variable). The observed data consist of the pairs of observations $(X_1, Y_1), (X_2, Y_2), \dots, (X_n, Y_n)$ of X and Y for each of the n subjects in a sample, or, with more concise notation, (X_i, Y_i) for $i = 1, 2, \dots, n$.

We are interested in analysing the association between X and Y . Methods for *describing* this association in the sample are first described in this section, initially with some standard graphical methods in Section 8.2.2. This leads to a discussion in Section 8.2.3 of what we actually mean by associations in this context, and then to a definition of numerical summary measures for such associations in Section 8.2.4. Statistical *inference* for the associations will be considered in Section 8.3.

8.2.2 Graphical methods

Scatterplots

The standard statistical graphic for summarising the association between two continuous variables is a **scatterplot**. An example of it is given in Figure 8.1, which shows a scatterplot of Control of corruption against GDP per capita for 61 countries for which the corruption variable is at least 60 (the motivation of this restriction will be discussed later). The two axes of the plot show possible values of the two variables. The horizontal axis, here corresponding to Control of corruption, is conventionally used for the explanatory variable X , and is often referred to as the **X-axis**. The vertical axis, here used for GDP per capita, then corresponds to the response variable Y , and is known as the **Y-axis**.




Figure 8.1: A scatterplot of Control of corruption vs. GDP per capita in the Global Civil Society data set, for 61 countries with Control of corruption at least 60. The dotted lines are drawn to the point corresponding to the United Kingdom.

The observed data are shown as points in the scatterplot, one for each of the n units. The location of each point is determined by its values of X and Y . For example, Figure 8.1 highlights the observation for the United Kingdom, for which the corruption measure (X) is 94.3 and GDP per capita (Y) is \$24160. The point for UK is thus placed at the intersection of a vertical line drawn from 94.3 on the X -axis and a horizontal line from 24160 on the Y -axis, as shown in the plot.

The principles of good graphical presentation on clear labelling, avoidance of spurious decoration and so on (c.f. Section 2.8) are the same for scatterplots as for any statistical graphics. Because the crucial visual information in a scatterplot is the shape of the cloud of the points, it is now often not necessary for the scales of the axes to begin at zero, especially if this is well outside the ranges of the observed values of the variables (as it is for the X -axis of Figure 8.1). Instead, the scales are typically selected so that the points cover most of the plotting surface. This is done by statistical software, but there are many situations where it is advisable to overrule the automatic selection (e.g. for making scatterplots of the same variables in two different samples directly comparable).

The main purpose of a scatterplot is to examine possible associations between X and Y . Loosely speaking, this means considering the shape and orientation of the cloud of points in the graph. In Figure 8.1, for example, it seems that most of the points are in a cluster sloping from lower left to upper right. This indicates that countries with low levels of Control of corruption (i.e. high levels of corruption itself) tend to have low GDP per capita, and those with little corruption tend to have high levels of GDP. A more careful discussion of such associations again relates them to the formal definition in terms of conditional distributions, and also provides a basis for the methods of inference introduced later in this chapter. We will resume the discussion of these issues in Section 8.2.3 below. Before that, however, we will digress briefly from the main thrust of this chapter in order to describe a slightly different kind of scatterplot.

Line plots for time series

A very common special case of a scatterplot is one where the observations correspond to measurements of a variable for the same unit at several occasions over time. This is illustrated by the following example (another one is Figure 2.9):

Example: Changes in temperature, 1903–2004

Figure 8.2 summarises data on average annual temperatures over the past century in five locations. The data were obtained from the GISS Surface Temperature (GISTEMP) database maintained by the NASA Goddard Institute for Space Studies². The database contains time series of average monthly surface temperatures from several hundred meteorological stations across the world. The five sites considered here are Haparanda in Northern Sweden, Independence, Kansas in the USA, Choshi on the east coast of Japan, Kimberley in South Africa, and the Base Orcadas Station on Laurie Island, off the coast of Antarctica. These were chosen rather haphazardly for this illustration, with the aim of obtaining a geographically scattered set of rural or small urban locations (to avoid issues with the heating effects of large urban areas). The temperature for each year at each location is here recorded as the difference from the temperature at that location in 1903³.

²The data were obtained from <http://data.london.gov.uk/datastore/package/london-borough-profiles>.

If you download the “Profiles in Excel” workbook, you will find that one of the pages contains a map of the boroughs, and a tool for visualising the data on that map. A regular map of the boroughs can be found at for example at

<http://www.londoncouncils.gov.uk/londonfacts/londonlocalgovernment/londonmapandlinks/default.htm>.

³The data can be obtained from <http://www3.norc.org/gss+website/>, which gives further information on




Figure 8.2: Changes of average annual temperature (11-year moving averages) from 1903 in five locations. See the text for further details. Source: The GISTEMP database <data.giss.nasa.gov/gistemp/>

Consider first the data for Haparanda only. Here we have two variables, year and temperature, and 102 pairs of observations of them, one for each year between 1903 and 2004. These pairs could now be plotted in a scatterplot as described above. Here, however, we can go further to enhance the visual effect of the plot. This is because the observations represent measurements of a variable (temperature difference) for the same unit (the town of Haparanda) at several successive times (years). These 102 measurements form a *time series* of temperature differences for Haparanda over 1903–2004. A standard graphical trick for such series is to connect the points for successive times by lines, making it easy for the eye to follow the changes over time in the variable on the Y-axis. In Figure 8.2 this is done for Haparanda using a solid line. Note that doing this would make no sense for scatter plots like the one in Figure 8.1, because all the points there represent different subjects, in that case countries.

We can easily include several such series in the same graph. In Figure 8.2 this is done by plotting the temperature differences for each of the five locations using different line styles. The graph now summarises data on three variables, year, temperature and location. We can then examine changes over time for any one location, but also compare patterns of changes between them. Here there is clearly much variation within and between locations, but also some common features. Most importantly, the temperatures have all increased over the past century. In all five locations the average annual temperatures at the end of the period were around 1–2°C higher than in 1903.

A set of time series like this is an example of dependent data in the sense discussed in Section 7.5. There we considered cases with pairs of observations, where the two observations in each pair had to be treated as statistically dependent. Here all of the temperature measurements for

one location are dependent, probably with strongest dependence between adjacent years and less dependence between ones further apart. This means that we will not be able to analyse these data with the methods described later in this chapter, because these assume statistically independent observations. Methods of statistical modelling and inference for dependent data of the kind illustrated by the temperature example are beyond the scope of this course. This, however, does not prevent us from using a plot like Figure 8.2 to *describe* such data.

8.2.3 Linear associations

Consider again statistically independent observations of (X_i, Y_i) , such as those displayed in Figure 8.1. Recall the definition that two variables are associated if the conditional distribution of Y given X is different for different values of X . In the two-sample examples of Chapter 7 this could be examined by comparing two conditional distributions, since X had only two possible values. Now, however, X has many (in principle, infinitely many) possible values, so we will need to somehow define and compare conditional distributions given each of them. We will begin with a rather informal discussion of how this might be done. This will lead directly to a more precise and formal definition introduced in Section 8.3.




Figure 8.3: The same scatterplot of Control of corruption vs. GDP per capita as in Figure 8.1, augmented by the best-fitting (least squares) straight line (solid line) and reference lines for two example values of Control of corruption (dotted lines).

Figure 8.3 shows the same scatterplot as Figure 8.1. Consider first one value of X (Control of corruption), say 65. To get a rough idea of the conditional distribution of Y (GDP per capita) given this value of X , we could examine the sample distribution of the values of Y for the units for which the value of X is close to 65. These correspond to the points near the vertical line drawn at $X = 65$ in Figure 8.3. This can be repeated for any value of X ; for example, Figure 8.3 also includes a vertical reference line at $X = 95$, for examining the conditional distribution

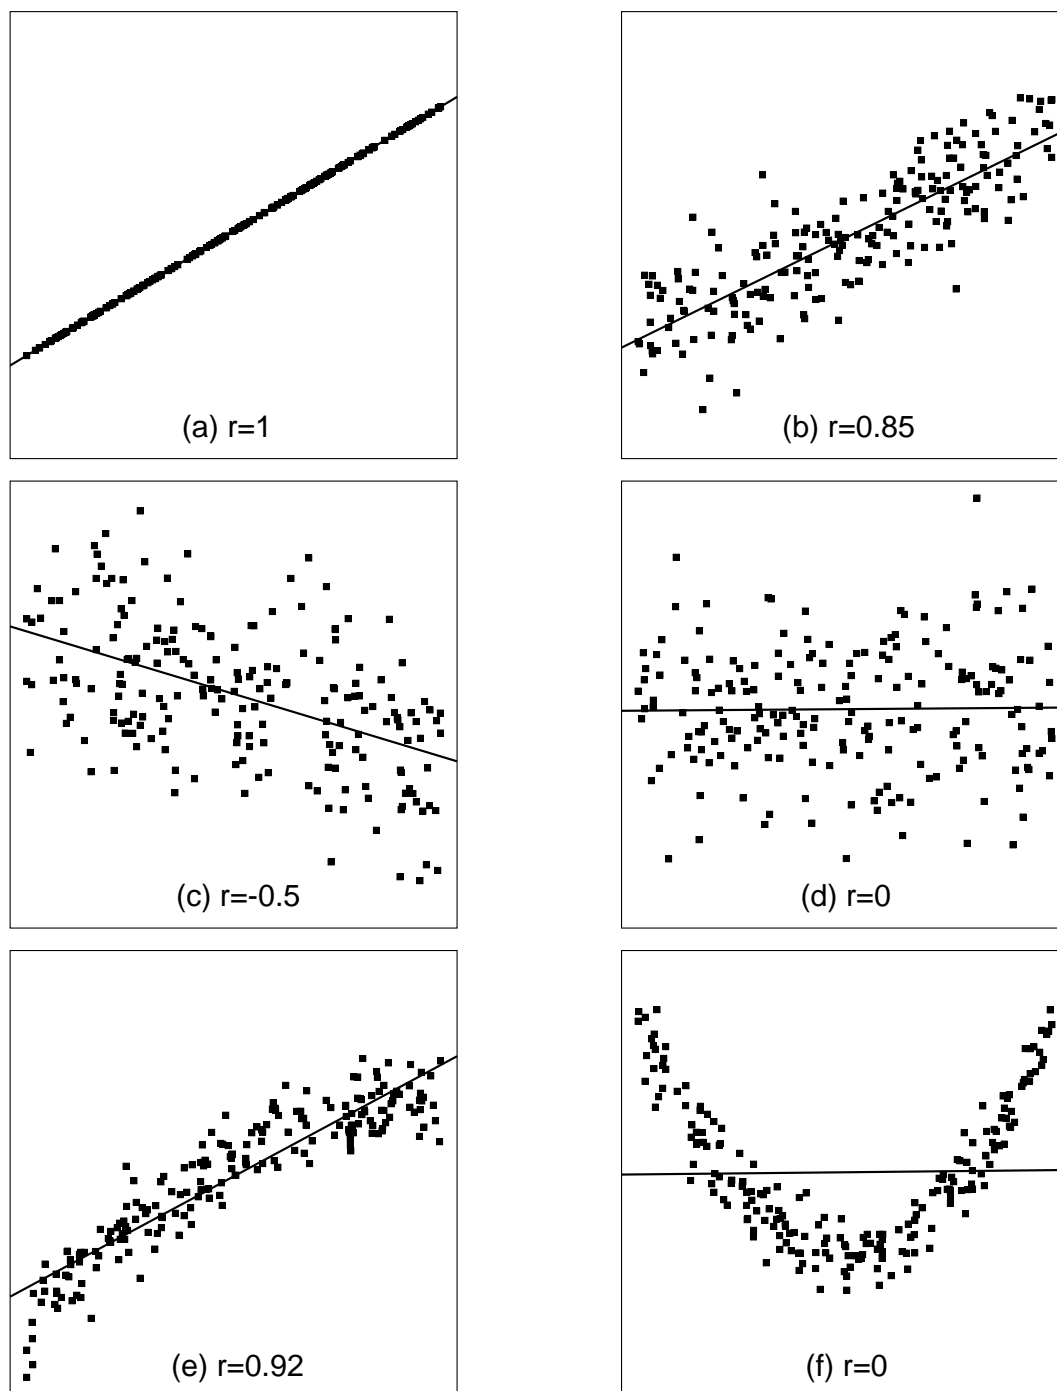
of Y given $X = 95$ ⁴.

As in Chapter 7, associations between variables will here be considered almost solely in terms of differences in the *means* of the conditional distributions of Y at different values of X . For example, Figure 8.3 suggests that the conditional mean of Y when X is 65 is around or just under 10000. At $X = 95$, on the other hand, the conditional mean seems to be between 20000 and 25000. The mean of Y is thus higher at the larger value of X . More generally, this finding is consistent across the scatterplot, in that the conditional mean of Y appears to increase when we consider increasingly large values of X , indicating that higher levels of Control of corruption are associated with higher average levels of GDP. This is often expressed by saying that the conditional mean of Y increases when we “increase” X ⁵. This is the sense in which we will examine associations between continuous variables: does the conditional mean of Y change (increase or decrease) when we increase X ? If it does, the two variables are associated; if it does not, there is no association of this kind. This definition also agrees with the one linking association with prediction: if the mean of Y is different for different values of X , knowing the value of X will clearly help us in making predictions about likely values of Y . Based on the information in Figure 8.3, for example, our best guesses of the GDPs of two countries would clearly be different if we were told that the control of corruption measure was 65 for one country and 95 for the other.

The *nature* of the association between X and Y is characterised by *how* the values of Y change when X increases. First, it is almost always reasonable to conceive these changes as reasonably smooth and gradual. In other words, if two values of X are close to each other, the conditional means of Y will be similar too; for example, if the mean of Y is 5 when $X = 10$, its mean when $X = 10.01$ is likely to be quite close to 5 rather than, say, 405. In technical terms, this means that the conditional mean of Y will be described by a smooth mathematical function of X . Graphically, the means of Y as X increases will then trace a smooth curve in the scatterplot. The simplest possibility for such a curve is a straight line. This possibility is illustrated by plot (a) of Figure ?? (this and the other five plots in the figure display artificial data, generated for this illustration). Here all of the points fall on a line, so that when X increases, the values of Y increase at a constant rate. A relationship like this is known as a **linear association** between X and Y . Linear associations are the starting point for examining associations between continuous variables, and often the only ones considered. In this chapter we too will focus almost completely on them.

⁴ESS Round 5: European Social Survey Round 5 Data (2010). Data file edition 2.0. Norwegian Social Science Data Services, Norway Data Archive and distributor of ESS data. The full data can be obtained from <http://ess.nsd.uib.no/ess/round5/>.

⁵Strictly speaking, the analysis should incorporate sampling weights (variable *DWEIGHT*) to adjust for different sampling probabilities for different types of respondents. Here the weights are ignored. Using them would not change the main conclusions for these variables.



In plot (a) of Figure ?? all the points are exactly on the straight line. This indicates a *perfect* linear association, where Y can be predicted exactly if X is known, so that the association is *deterministic*. Such a situation is neither realistic in practice, nor necessary for the association to be described as linear. All that is required for the latter is that the conditional *means* of Y given different values of X fall (approximately) on a straight line. This is illustrated by plot (b) of Figure ??, which shows a scatterplot of individual observations together with an approximation of the line of the means of Y given X (how the line was drawn will be explained later). Here the linear association is not perfect, as the individual points are not all on the same line but scattered around it. Nevertheless, the line seems to capture an important systematic feature of the data, which is that the *average* values of Y increase at an approximately constant rate as X increases. This combination of systematic and random elements is characteristic of all statistical associations, and it is also central to the formal setting for statistical inference

for linear associations described in Section 8.3 below.

The **direction** of a linear association can be either **positive** or **negative**. Plots (a) and (b) of Figure ?? show a positive association, because increasing X is associated with increasing average values of Y . This is indicated by the upward slope of the line describing the association. Plot (c) shows an example of a negative association, where the line slopes downwards and increasing values of X are associated with decreasing values of Y . The third possibility, illustrated by plot (d), is that the line slopes neither up nor down, so that the mean of Y is the same for all values of X . In this case there is no (linear) association between the variables.

Not all associations between continuous variables are linear, as shown by the remaining two plots of Figure ?. These illustrate two kinds of **nonlinear** associations. In plot (e), the association is still clearly *monotonic*, meaning that average values of Y change in the same direction — here increase — when X increases. The rate of this increase, however, is not constant, as indicated by the slightly curved shape of the cloud of points. The values of Y seem to increase faster for small values of X than for large ones. A straight line drawn through the scatterplot captures the general direction of the increase, but misses its nonlinearity. One practical example of such a relationship is the one between years of job experience and salary: it is often found that salary increases fastest early on in a person’s career and more slowly later on.

Plot (f) shows a nonlinear and nonmonotonic relationship: as X increases, average values of Y first decrease to a minimum, and then increase again, resulting in a U-shaped scatterplot. A straight line is clearly an entirely inadequate description of such a relationship. A nonmonotonic association of this kind might be seen, for example, when considering the dependence of the failure rates of some electrical components (Y) on their age (X). It might then be that the failure rates were high early (from quick failures of flawed components) and late on (from inevitable wear and tear) and lowest in between for “middle-aged but healthy” components.

Returning to real data, recall that we have so far considered control of corruption and GDP per capita only among countries with a Control of corruption score of at least 60. The scatterplot for these, shown in Figure 8.3, also includes a best-fitting straight line. The observed relationship is clearly positive, and seems to be fairly well described by a straight line. For countries with relatively low levels of corruption, the association between control of corruption and GDP can be reasonably well characterised as linear.

Consider now the set of all countries, including also those with high levels of corruption (scores of less than 60). In a scatterplot for them, shown in Figure 8.4, the points with at least 60 on the X -axis are the same as those in Figure 8.3, and the new points are to the left of them. The plot now shows a nonlinear relationship comparable to the one in plot (e) of Figure ?. The linear relationship which was a good description for the countries considered above is thus not adequate for the full set of countries. Instead, it seems that the association is much weaker for the countries with high levels of corruption, essentially all of which have fairly low values of GDP per capita. The straight line fitted to the plot identifies the overall positive association, but cannot describe its nonlinearity. This example further illustrates how scatterplots can be used to examine relationships between variables and to assess whether they can be best described as linear or nonlinear associations⁶.

So far we have said nothing about how the exact location and direction of the straight lines shown in the figures have been selected. These are determined so that the fitted line is in a certain sense the best possible one for describing the data in the scatterplot. Because the

⁶The data can be obtained from <http://bes2009-10.org/>, which gives further information on the survey, including the full text of the questionnaires. The data analysed in this class and homework are from the BES Campaign Internet Panel Survey, which has been divided into two data sets corresponding to two time periods leading up to the General Election.




Figure 8.4: A scatterplot of Control of corruption vs. GDP per capita for 163 countries in the Global Civil Society data set. The solid line is the best-fitting (least squares) straight line for the points.

calculations needed for this are also (and more importantly) used in the context of statistical inference for such data, we will postpone a description of them until Section 8.3.4. For now we can treat the line simply as a visual summary of the linear association in a scatterplot.

8.2.4 Measures of association: covariance and correlation

A scatterplot is a very powerful tool for examining sample associations of pairs of variables in detail. Sometimes, however, this is more than we really need for an initial summary of a data set, especially if there are many variables and thus many possible pairs of them. It is then convenient also to be able to summarise each pairwise association using a single-number measure of association. This section introduces the correlation coefficient, the most common such measure for continuous variables. It is a measure of the strength of *linear* associations of the kind defined above.

Suppose that we consider two variables, denoted X and Y . This again implies a distinction between an explanatory and a response variable, to maintain continuity of notation between different parts of this chapter. The correlation coefficient itself, however, is completely symmetric, so that its value for a pair of variables will be the same whether or not we treat one or the other of them as explanatory for the other. First, recall from equation of standard deviation towards the end of Section 2.6.2 that the sample standard deviations of the two variables are calculated as

$$s_x = \sqrt{\frac{\sum (X_i - \bar{X})^2}{n - 1}} \text{ and } s_y = \sqrt{\frac{\sum (Y_i - \bar{Y})^2}{n - 1}} \quad (8.1)$$

where the subscripts x and y identify the two variables, and \bar{X} and \bar{Y} are their sample means. A new statistic is the **sample covariance** between X and Y , defined as

$$s_{xy} = \frac{\sum (X_i - \bar{X})(Y_i - \bar{Y})}{n - 1}. \quad (8.2)$$

This is a measure of linear association between X and Y . It is positive if the sample association is positive and negative if the association is negative.

In theoretical statistics, covariance is the fundamental summary of sample and population associations between two continuous variables. For descriptive purposes, however, it has the inconvenient feature that its magnitude depends on the units in which X and Y are measured. This makes it difficult to judge whether a value of the covariance for particular variables should be regarded as large or small. To remove this complication, we can standardise the sample covariance by dividing it by the standard deviations, to obtain the statistic

$$r = \frac{s_{xy}}{s_x s_y} = \frac{\sum (X_i - \bar{X})(Y_i - \bar{Y})}{\sqrt{\sum (X_i - \bar{X})^2 \sum (Y_i - \bar{Y})^2}}. \quad (8.3)$$

This is the (sample) **correlation** coefficient, or correlation for short, between X and Y . It is also often (e.g. in SPSS) known as *Pearson's* correlation coefficient after Karl Pearson (of the χ^2 test, see first footnote in Chapter 4), although both the word and the statistic are really due to Sir Francis Galton⁷.

The properties of the correlation coefficient can be described by going through the same list as for the γ coefficient in Section 2.4.5. While doing so, it is useful to refer to the examples in Figure ??, where the correlations are also shown.

- **Sign:** Correlation is positive if the *linear* association between the variables is positive, i.e. if the best-fitting straight line slopes upwards (as in plots a, b and e) and negative if the association is negative (c). A zero correlation indicates complete lack of linear association (d and f).
- **Extreme values:** The largest possible correlation is +1 (plot a) and the smallest -1 , indicating perfect positive and negative linear associations respectively. More generally, the magnitude of the correlation indicates the strength of the association, so that the closer to +1 or -1 the correlation is, the stronger the association (e.g. compare plots a–d). It should again be noted that the correlation captures only the linear aspect of the association, as illustrated by the two nonlinear cases in Figure ?. In plot (e), there is curvature but also a strong positive trend, and the latter is reflected in a fairly high correlation. In plot (f), the trend is absent and the correlation is 0, even though there is an obvious nonlinear relationship. Thus the correlation coefficient is a reasonable initial summary of the strength of association in (e), but completely misleading in (f).
- **Formal interpretation:** The correlation coefficient cannot be interpreted as a Proportional Reduction in Error (PRE) measure, but its square can. The latter statistic, so-called coefficient of determination or R^2 , is described in Section 8.3.3.
- **Substantive interpretation:** As with any measure of association, the question of whether a particular sample correlation is high or low is not a purely statistical question, but depends on the nature of the variables. This can be judged properly only with the help of experience of correlations between similar variables in different contexts. As

⁷Official results obtained from www.olympic.org/london-2012-summer-olympics.

one very rough rule thumb it might be said that in many social science contexts correlations greater than 0.4 (or smaller than -0.4) would typically be considered noteworthy and ones greater than 0.7 quite strong.

Returning to real data, Table 8.1 shows the correlation coefficients for all fifteen distinct pairs of the six continuous variables in the Global Civil Society data set mentioned in Example 8.1. This is an example of a **correlation matrix**, which is simply a table with the variables as both its rows and columns, and the correlation between each pair of variables given at the intersection of corresponding row and column. For example, the correlation of GDP per capita and School enrolment is here 0.42. This is shown at the intersection of the first row (GDP) and fifth column (School enrolment), and also of the fifth row and first column. In general, every correlation is shown twice in the matrix, once in its upper triangle and once in the lower. The triangles are separated by a list of ones on the diagonal of the matrix. This simply indicates that the correlation of any variable with itself is 1, which is true by definition and thus of no real interest.

Table 8.1: Correlation matrix of six continuous variables in the Global Civil Society data set. See Example 8.1 for more information on the variables.

Variable	GDP	Gini	Pol.	Corrupt.	School	IMR
GDP per capita [GDP]	1	-0.39	0.51	0.77	0.42	-0.62
Income inequality [Gini]	-0.39	1	-0.15	-0.27	-0.27	0.42
Political rights [Pol.]	0.51	-0.15	1	0.59	0.40	-0.44
Control of corruption [Corrupt.]	0.77	-0.27	0.59	1	0.41	-0.64
School enrolment [School]	0.42	-0.27	0.40	0.41	1	-0.73
Infant mortality [IMR]	-0.62	0.42	-0.44	-0.64	-0.73	1

All of the observed associations in this example are in unsurprising directions. For example, School enrolment is positively correlated with GDP, Political rights and Control of corruption, and negatively correlated with Income inequality and Infant mortality. In other words, countries with large percentages of children enrolled in primary school tend to have high levels of GDP per capita and of political rights and civil liberties, and low levels of corruption, income inequality and infant mortality. The strongest associations in these data are between GDP per capita and Control of corruption ($r = 0.77$) and School enrolment and Infant mortality rate ($r = -0.73$), and the weakest between Income inequality on the one hand and Political rights, Control of corruption and School enrolment on the other (correlations of -0.15 , -0.27 and -0.27 respectively).

These correlations describe only the linear element of sample associations, but give no hint of any nonlinear ones. For example, the correlation of 0.77 between GDP and Control of corruption summarises the way the observations cluster around the straight line shown in Figure 8.4. The correlation is high because this increase in GDP as Control of corruption increases is quite strong, but it gives no indication of the nonlinearity of the association. A scatterplot is needed for revealing this feature of the data. The correlation for the restricted set of countries shown in Figure 8.3 is 0.82.

A correlation coefficient can also be defined for the joint population distribution of two variables. The sample correlation r can then be treated as an estimate of the population correlation, which is often denoted by ρ (the lower-case Greek “rho”). Statistical inference for the population correlation can also be derived. For example, SPSS automatically outputs significance tests for the null hypothesis that ρ is 0, i.e. that there is no linear association between X and Y in the population. Here, however, we will not discuss this, choosing to treat r purely as a descriptive

sample statistic. The next section provides a different set of tools for inference on population associations.

8.3 Simple linear regression models

8.3.1 Introduction

The rest of this course is devoted to the method of linear regression modelling. Its purpose is the analysis of associations in cases where the response variable is a continuous, interval level variable, and the possibly several explanatory variables can be of any type. We begin in this section with *simple* linear regression, where there is only one explanatory variable. We will further assume that this is also continuous. The situation considered here is thus the same as in the previous section, but here the focus will be on statistical inference rather than description. Most of the main concepts of linear regression can be introduced in this context. Those that go beyond it are described in subsequent sections. Section 8.5 introduces *multiple* regression involving more than one explanatory variable. The use of categorical explanatory variables in such models is explained in Section 8.6. Finally, Section 8.7 gives a brief review of some further aspects of linear regression modelling which are not covered on this course.

Example: Predictors of Infant Mortality Rate

The concepts of linear regression models will be illustrated as they are introduced with a second example from the Global Civil Society data set. The response variable will now be Infant Mortality Rate (IMR). This is an illuminating outcome variable, because it is a sensitive and unquestionably important reflection of a country's wellbeing; whatever we mean by "development", it is difficult to disagree that high levels of it should coincide with low levels of infant mortality. We will initially consider only one explanatory variable, Net primary school enrolment ratio, referred to as "School enrolment" for short. This is defined as the percentage of all children of primary school age who are enrolled in school. Enrolment numbers and the population size are often obtained from different official sources, which sometimes leads to discrepancies. In particular, School enrolment for several countries is recorded as over 100, which is logically impossible. This is an illustration of the kinds of measurement errors often affecting variables in the social sciences. We will use the School enrolment values as recorded, even though they are known to contain some error.

A scatterplot of IMR vs. School enrolment is shown in Figure 8.5, together with the best-fitting straight line. Later we will also consider three additional explanatory variables: Control of corruption, Income inequality and Income level of the country in three categories (c.f. Example 8.1). For further reference, Table 8.2 shows various summary statistics for these variables. Throughout, the analyses are restricted to those 111 countries for which all of the five variables are recorded. For this reason the correlations in Table 8.2 differ slightly from those in Table 8.1, where each correlation was calculated for all the countries with non-missing values of that pair of variables.

Table 8.2: Summary statistics for Infant Mortality Rate (IMR) and explanatory variables for it considered in the examples of Sections 8.3 and 8.5 ($n = 111$). See Example 8.1 for further information on the variables.

IMR	School	Control	Income
	enrolment	of	inequality
		corruption	

<i>Summary statistics</i>				
Mean	4.3	86.1	50.1	40.5
std. deviation	4.0	16.7	28.4	10.2
Minimum	0.3	30.0	3.6	24.4
Maximum	15.6	109.0	100.0	70.7
<i>Correlation matrix</i>				
IMR	1	-0.75	-0.60	0.39
School enrolment	-0.75	1	0.39	-0.27
Control of corruption	-0.60	0.39	1	-0.27
Income inequality	0.39	-0.27	-0.27	1
<i>Means for countries in different income categories</i>				
Low income ($n = 41$)	8.2	72.1	27.5	41.7
Middle income ($n = 48$)	2.8	92.5	50.8	43.3
High income ($n = 22$)	0.5	98.4	90.7	32.0

8.3.2 Definition of the model

The simple linear regression model defined in this section is a statistical model for a continuous, interval level response variable Y given a single explanatory variable X , such as IMR given School enrolment. The model will be used to carry out statistical inference on the association between the variables in a population (which in the IMR example is clearly again of the conceptual variety).

For motivation, recall first the situation considered in Section 7.3. There the data consisted of observations (Y_i, X_i) for $i = 1, 2, \dots, n$, which were assumed to be statistically independent. The response variable Y was continuous but X had only two possible values, coded 1 and 2. A model was then set up where the population distribution of Y had mean μ_1 and variance σ_1^2 for units with $X = 1$, and mean μ_2 and variance σ_2^2 when $X = 2$. In some cases it was further assumed that the population distributions were both normal, and that the population variances were equal, i.e. that $\sigma_1^2 = \sigma_2^2$, with their common value denoted σ^2 . With these further assumptions, which will also be used here, the model for Y given a dichotomous X stated that (1) observations for different units i were statistically independent; (2) each Y_i was sampled at random from a population distribution which was normal with mean μ_i and variance σ^2 ; and (3) μ_i depended on X_i so that it was equal to μ_1 if X_i was 1 and μ_2 if X_i was 2.

The situation in this section is exactly the same, except that X is now continuous instead of dichotomous. We will use the same basic model, but will change the specification of the conditional mean μ_i appropriately. In the light of the discussion in previous sections of this chapter, it is no surprise that this will be defined in such a way that it describes a linear association between X and Y . This is done by setting $\mu_i = \alpha + \beta X_i$, where α and β are unknown population parameters. This is the equation of straight line (we will return to it in the next section). With this specification, the model for observations $(Y_1, X_1), (Y_2, X_2), \dots, (Y_n, X_n)$ becomes

1. Observations for different units i ($= 1, 2, \dots, n$) are statistically independent.
2. Each Y_i is normally distributed with mean μ_i and variance σ^2 .
3. The means μ_i depend on X_i through $\mu_i = \alpha + \beta X_i$.

Often the model is expressed in an equivalent form where 2. and 3. are combined as




Figure 8.5: A scatterplot of net primary school enrolment ratio vs. Infant mortality rate for countries in the Global Civil Society data set ($n = 111$). The solid line is the best-fitting (least squares) straight line for the points.

$$Y_i = \alpha + \beta X_i + \epsilon_i \quad (8.4)$$

where each ϵ_i is normally distributed with mean 0 and variance σ^2 . The ϵ_i are known as **error terms** or **population residuals** (and the letter ϵ is the lower-case Greek “epsilon”). This formulation of the model clearly separates the mean of Y_i , which traces the straight line $\alpha + \beta X_i$ as X_i changes, from the variation around that line, which is described by the variability of ϵ_i .

The model defined above is known as the **simple linear regression model**:

- **Simple** because it has only one explanatory variable, as opposed to *multiple* linear regression models which will have more than one.
- **Linear** because it specifies a linear association between X and Y .⁸
- **Regression**: This is now an established part of the name of the model, although the origins of the word are not central to the use of the model⁹.
- **Model**, because this is a statistical model in the sense discussed in the middle of Section 6.3.1. In other words, the model is always only a simplified abstraction of the true,

⁸The data can be obtained from www3.norc.ox.ac.uk/GSS+Website/, which gives further information on the survey, including the full text of the questionnaires.

⁹United Nations Development Programme *International Human Development Indicators*, <http://hdr.undp.org/en/data/>; World Bank *Worldwide Governance Indicators*, <http://info.worldbank.org/governance/wgi/pdf/wgidataset.xlsx>; World Bank *World Development Indicators*, <http://data.worldbank.org/indicator/SP.DYN.IMRT.IN>.

immeasurably complex processes which determine the values of Y . Nevertheless, it is believed that a well-chosen model can be useful for explaining and predicting observed values of Y . This spirit is captured by the well-known statement by the statistician George Box¹⁰:

All models are wrong, but some are useful.

A model like this has the advantage that it reduces the examination of associations in the population to estimation and inference on a small number of model parameters, in the case of the simple linear regression model just α , β and σ^2 .

Of course, not all models are equally appropriate for given data, and some will be both wrong and useless. The results from a model should thus be seriously presented and interpreted only if the model is deemed to be reasonably adequate. For the simple linear regression model, this can be partly done by examining whether the scatterplot between X and Y appears to be reasonably consistent with a linear relationship. Some further comments on the assessment of model adequacy will be given in Section 8.7.

8.3.3 Interpretation of the model parameters

The simple linear regression model (8.4) has three parameters, α , β and σ^2 . Each of these has its own interpretation, which are explained in this section. Sometimes it will be useful to illustrate the definition with specific numerical values, for which we will use ones for the model for IMR given School enrolment in our example. SPSS output for this model is shown in Figure 8.6. Note that although these values are first used here to illustrate the interpretation of the *population* parameters in the model, they are of course only estimates (of a kind explained in the next section) of those parameters. Other parts of the SPSS output will be explained later in this chapter.

According to the model, the conditional mean (also often known as the conditional **expected value**) of Y given X in the population is (dropping the subscript i for now for notational simplicity) $\mu = \alpha + \beta X$. The two parameters α and β in this formula are known as **regression coefficients**. They are interpreted as follows:

- α is the expected value of Y when X is equal to 0. It is known as the **intercept** or **constant** term of the model.
- β is the change in the expected value of Y when X increases by 1 unit. It is known as the **slope** term or the **coefficient of X** .

Just to include one mathematical proof in this coursepack, these results can be derived as follows:

- When $X = 0$, the mean of Y is $\mu = \alpha + \beta X = \alpha + \beta \times 0 = \alpha + 0 = \alpha$.
- Compare two observations, one with value X of the explanatory variable, and the other with one unit more, i.e. $X + 1$. The corresponding means of Y are

with $X + 1$:	μ	$= \alpha + \beta \times (X + 1)$	$= \alpha + \beta X + \beta$
with X :	μ		$= \alpha + \beta X$
Difference:			β

¹⁰ESS Round 5: European Social Survey Round 5 Data (2010). Data file edition 2.0. Norwegian Social Science Data Services, Norway Data Archive and distributor of ESS data. The full data can be obtained from <http://ess.nsd.uib.no/ess/round5/>.

Model Summary				
Model	R	RSquare	Adjusted RSquare	Std. Error of the Estimate
1	.753 ^a	.567	.563	2.6173

a. Predictors: (Constant), Netprimaryschoolenrolment ratio2000-2001 (%)

ANOVA ^a				
Model	Sum of Squares	df	Mean Square	F
1 Regression	976.960	1	976.960	142.621
Residual	746.653	109	6.850	
Total	1723.613	110		

a. Predictors: (Constant), Netprimaryschoolenrolment ratio2000-2001 (%)

b. Dependent Variable: InfantMortalityRate2001(%)

Coefficients ^a				
Model		Unstandardized Coefficients		Standardized Coefficients
		B	Std. Error	
1	(Constant)	19.736	1.313	
	Netprimaryschool enrolmentratio 2000-2001 (%)	-.179	.015	

a. Dependent Variable: InfantMortalityRate2001(%)

which completes the proof of the claims above — Q.E.D. In case you prefer a graphical summary, this is given in Figure 8.7.




Figure 8.7: Illustration of the interpretation of the regression coefficients of a simple linear regression model.

The most important parameter of the model, and usually the only one really discussed in interpreting the results, is β , the regression coefficient of X . It is also called the slope because it is literally the slope of the regression line, as shown in Figure 8.7. It is the only parameter in the model which describes the association between X and Y , and it does so in the above terms of expected changes in Y corresponding to changes in X (β is also related to the correlation between X and Y , in a way explained in the next section). The sign of β indicates the direction of the association. When β is positive (greater than 0), the regression line slopes upwards and increasing X thus also increases the expected value of Y — in other words, the association between X and Y is positive. This is the case illustrated in Figure 8.7. If β is negative, the regression line slopes downwards and the association is also negative. Finally, if β is zero, the line is parallel with the X -axis, so that changing X does not change the expected value of Y . Thus $\beta = 0$ corresponds to no (linear) association between X and Y .

In the real example shown in Figure 8.6, X is School enrolment and Y is IMR. In SPSS output, the estimated regression coefficients are given in the “**Coefficients**” table in the column labelled “B” under “Unstandardized coefficients”. The estimated constant term α is given in the row labelled “(Constant)”, and the slope term on the next row, labelled with the name or label of the explanatory variable as specified in the SPSS data file — here “Net primary school enrolment ratio 2000-2001 (%)”. The value of the intercept is here 19.736 and the slope coefficient is -0.179 . The estimated regression line for expected IMR is thus $19.736 - 0.179X$, where X denotes School enrolment. This is the line shown in Figure 8.5.

Because the slope coefficient in the example is negative, the association between the variables is also negative, i.e. higher levels of school enrolment are associated with lower levels of infant mortality. More specifically, every increase of one unit (here one percentage point) in School enrolment is associated with a decrease of 0.179 units (here percentage points) in expected IMR.

Since the meaning of β is related to a unit increase of the explanatory variable, the interpretation of its magnitude depends on what those units are. In many cases one unit of X is too small or too large for convenient interpretation. For example, a change of one percentage point in School enrolment is rather small, given that the range of this variable in our data is 79 percentage points (c.f. Table 8.2). In such cases the results can easily be reexpressed by using multiples of β : specifically, the effect on expected value of Y of changing X by A units is obtained by multiplying β by A . For instance, in our example the estimated effect of increasing School enrolment by 10 percentage points is to decrease expected IMR by $10 \times 0.179 = 1.79$ percentage points.

The constant term α is a necessary part of the model, but it is almost never of interest in itself. This is because the expected value of Y at $X = 0$ is rarely specifically interesting. Very often $X = 0$ is also unrealistic, as in our example where it corresponds to a country with zero primary school enrolment. There are fortunately no such countries in the data, where the lowest School enrolment is 30%. It is then of no interest to discuss expected IMR for a hypothetical country where no children went to school. Doing so would also represent unwarranted *extrapolation* of the model beyond the range of the observed data. Even though the estimated linear model seems to fit reasonably well for these data, this is no guarantee that it would do so also for countries with much lower school enrolment, even if they existed.

The third parameter of the simple regression model is σ^2 . This is the variance of the conditional distribution of Y given X . It is also known as the **conditional variance** of Y , the **error variance** or the **residual variance**. Similarly, its square root σ is known as the conditional, error or **residual standard deviation**. To understand σ , let us consider a single value of X , such as one corresponding to one of the vertical dashed lines in Figure 8.7 or, say, school enrolment of 85 in Figure 8.5. The model specifies a distribution for Y given any such value of X . If we were to (hypothetically) collect a large number of observations, all with this same value of X , the distribution of Y for them would describe the conditional distribution of Y given that value of X . The model states that the average of these values, i.e. the conditional mean of Y , is $\alpha + \beta X$, which is the point on the regression line corresponding to X . The individual values of Y , however, would of course not all be on the line but somewhere around it, some above and some below.

The linear regression model further specifies that the form of the conditional distribution of Y is approximately normal. You can try to visualise this by imagining a normal probability curve (c.f. Figure 6.5) on the vertical line from X , centered on the regression line and sticking up from the page. The bell shape of the curve indicates that most of the values of Y for a given X will be close to the regression line, and only small proportions of them far from it. The residual standard deviation σ is the standard deviation of this conditional normal distribution, in essence describing how tightly concentrated values of Y tend to be around the regression line. The model assumes, mainly for simplicity, that the same value of σ applies to the conditional distributions at all values of X ; this is known as the assumption of *homoscedasticity*.

In SPSS output, an estimate of σ is given in the “**Model Summary**” table under the misleading label “Std. Error of the Estimate”. An estimate of the residual variance σ^2 is found also in the “**ANOVA**” table under “Mean Square” for “Residual”. In our example the estimate of σ is 2.6173 (and that of σ^2 is 6.85). This is usually not of direct interest for interpretation, but it will be a necessary component of some parts of the analysis discussed below.

8.3.4 Estimation of the parameters

Since the regression coefficients α and β and the residual standard deviation σ are unknown population parameters, we will need to use the observed data to obtain sensible estimates for


them. How to do so is now less obvious than in the cases of simple means and proportions considered before. This section explains the standard method of estimation for the parameters of linear regression models.

We will denote estimates of α and β by $\hat{\alpha}$ and $\hat{\beta}$ (“alpha-hat” and “beta-hat”) respectively (other notations are also often used, e.g. a and b). Similarly, we can define

$$\hat{Y} = \hat{\alpha} + \hat{\beta}X$$

for Y given any value of X . These are the values on the estimated regression line. They are known as **fitted values** for Y , and estimating the parameters of the regression model is often referred to as “fitting the model” to the observed data. The fitted values represent our predictions of expected values of Y given X , so they are also known as **predicted values** of Y .

In particular, fitted values $\hat{Y}_i = \hat{\alpha} + \hat{\beta}X_i$ can be calculated at the values X_i of the explanatory variable X for each unit i in the observed sample. These can then be compared to the corresponding values Y_i of the response variable. Their differences $Y_i - \hat{Y}_i$ are known as the (sample) **residuals**. These quantities are illustrated in Figure ???. This shows a fitted regression line, which is in fact the one for IMR given School enrolment also shown in Figure 8.5. Also shown are two points (X_i, Y_i) . These are also from Figure 8.5; the rest have been omitted to simplify the plot. The point further to the left is the one for Mali, which has School enrolment $X_i = 43.0$ and IMR $Y_i = 14.1$. Using the estimated coefficients $\hat{\alpha} = 19.736$ and $\hat{\beta} = -0.179$ in Figure 8.6, the fitted value for Mali is $\hat{Y}_i = 19.736 - 0.179 \times 43.0 = 12.0$. Their difference is the residual $Y_i - \hat{Y}_i = 14.1 - 12.0 = 2.1$. Because the observed value is here larger than the fitted value, the residual is positive and the observed value is above the fitted line, as shown in Figure ??.



The second point shown in Figure ?? corresponds to the observation for Ghana, for which $X_i = 58.0$ and $Y_i = 5.7$. The fitted value is then $\hat{Y}_i = 19.736 - 0.179 \times 58.0 = 9.4$ and the residual $Y_i - \hat{Y}_i = 5.7 - 9.4 = -3.7$. Because the observed value is now smaller than the fitted value, the residual is negative and the observed Y_i is below the fitted regression line.

So far we have still not explained how the specific values of the parameter estimates in Figure

8.6 were obtained. In doing so, we are faced with the task of identifying a regression line which provides the best fit to the observed points in a scatterplot like Figure 8.5. Each possible choice of $\hat{\alpha}$ and $\hat{\beta}$ corresponds to a different regression line, and some choices are clearly better than others. For example, it seems intuitively obvious that it would be better for the line to go through the cloud of points rather than stay completely outside it. To make such considerations explicit, the residuals can be used as a criterion of model fit. The aim will then be to make the total magnitude of the residuals as small as possible, so that the fitted line is as close as possible to the observed points Y_i in some overall sense. This cannot be done simply by adding up the residuals, because they can have different signs, and positive and negative residuals could thus cancel out each other in the addition. As often before, the way around this is to remove the signs by considering the squares of the residuals. Summing these over all units i in the sample leads to the sum of squared residuals

$$SSE = \sum (Y_i - \hat{Y}_i)^2.$$

Here SSE is short for Sum of Squares of Errors (it is also often called the Residual Sum of Squares or RSS). This is the quantity used as the criterion in estimating regression coefficients for a linear model. Different candidate values for $\hat{\alpha}$ and $\hat{\beta}$ lead to different values of \hat{Y}_i and thus of SSE . The final estimates are the ones which give the smallest value of SSE . Their formulas are

$$\hat{\beta} = \frac{\sum (X_i - \bar{X})(Y_i - \bar{Y})}{\sum (X_i - \bar{X})^2} = \frac{s_{xy}}{s_x^2} \quad (8.5)$$

and

$$\hat{\alpha} = \bar{Y} - \hat{\beta}\bar{X} \quad (8.6)$$

where \bar{Y} , \bar{X} , s_x and s_{xy} are the usual sample means, standard deviations and covariances for Y and X . These are known as the **least squares estimates** of the regression coefficients (or as Ordinary Least Squares or OLS estimates), and the reasoning used to obtain them is the **method of least squares**¹¹. Least squares estimates are almost always used for linear regression models, and they are the ones displayed by SPSS and other software. For our model for IMR given School enrolment, the estimates are the $\hat{\alpha} = 19.736$ and $\hat{\beta} = -0.179$ shown in Figure 8.6.

The estimated coefficients can be used to calculate predicted values for Y at any values of X , not just those included in the observed sample. For instance, in the infant mortality example the predicted IMR for a country with School enrolment of 80% would be $\hat{Y} = 19.736 - 0.179 \times 80 = 5.4$. Such predictions should usually be limited to the range of values of X actually observed in the data, and extrapolation beyond these values should be avoided.

The most common estimate of the remaining parameter of the model, the residual standard deviation σ , is

$$\hat{\sigma} = \sqrt{\frac{\sum (Y_i - \hat{Y}_i)^2}{n - (k + 1)}} = \sqrt{\frac{SSE}{n - (k + 1)}} \quad (8.7)$$

¹¹This is another old idea. Different approaches to the problem of fitting curves to observations were gradually developed by Tobias Mayer, Rudjer Bošković and Pierre Simon Laplace from the 1750s onwards, and the method of least squares itself was presented by Adrien Marie Legendre in 1805.

where k is here set equal to 1. This bears an obvious resemblance to the formula for the basic sample standard deviation, shown for Y_i in (8.1). One difference to that formula is that the denominator of (8.7) is shown as $n - (k + 1)$ rather than $n - 1$. Here $k = 1$ is the number of explanatory variables in the model, and $k + 1 = 2$ is the number of regression coefficients (α and β) including the constant term α . The quantity $n - (k + 1)$, i.e. here $n - 2$, is the **degrees of freedom** (df) of the parameter estimates. We will need it again in the next section. It is here given in the general form involving the symbol k , so that we can later refer to the same formula for models with more explanatory variables and thus k greater than 1. In SPSS output, the degrees of freedom are shown in the “ANOVA” table under “df” for “Residual”. In the infant mortality example $n = 111$, $k = 1$ and $df = 111 - 2 = 109$, as shown in Figure 8.6.

Finally, two connections between previous topics and the parameters $\hat{\alpha}$, $\hat{\beta}$ and $\hat{\sigma}$ are worth highlighting:

- The estimated slope $\hat{\beta}$ from (8.5) is related to the sample correlation r from (8.3) by $r = (s_x/s_y) \hat{\beta}$. In both of these it is $\hat{\beta}$ which carries information about the association between X and Y . The ratio s_x/s_y serves only to standardise the correlation coefficient so that it is always between -1 and $+1$. The slope coefficient $\hat{\beta}$ is not standardised, and the interpretation of its magnitude depends on the units of measurement of X and Y in the way defined in Section 8.3.3.
- Suppose we simplify the simple linear regression model (8.4) further by setting $\beta = 0$, thus removing βX from the model. The new model states that all Y_i are normally distributed with the same mean α and standard deviation σ . Apart from the purely notational difference of using α instead of μ , this is exactly the single-sample model considered in Section 7.4. Using the methods of this section to obtain estimates of the two parameters of this model also leads to exactly the same results as before. The least squares estimate of α is then $\hat{\alpha} = \bar{Y}$, obtained by setting $\hat{\beta} = 0$ in (8.6). Since there is no $\hat{\beta}$ in this case, $\hat{Y}_i = \bar{Y}$ for all observations, $k = 0$ and $df = n - (k + 1) = n - 1$. Substituting these into (8.7) shows that $\hat{\sigma}$ is then equal to the usual sample standard deviation s_y of Y_i .

Coefficient of determination (R^2)

The **coefficient of determination**, more commonly known as **R^2** (“R-squared”), is a measure of association very often used to describe the results of linear regression models. It is based on the same idea of sums of squared errors as least squares estimation, and on comparison of them between two models for Y . The first of these models is the very simple one where the explanatory variable X is not included at all. As discussed above, the estimate of the expected value of Y is then the sample mean \bar{Y} . This is the best prediction of Y we can make, if the same predicted value is to be used for all observations. The error in the prediction of each value Y_i in the observed data is then $Y_i - \bar{Y}$ (c.f. Figure ?? for an illustration of this for one observation). The sum of squares of these errors is $TSS = \sum (Y_i - \bar{Y})^2$, where TSS is short for “Total Sum of Squares”. This can also be regarded as a measure of the **total variation** in Y_i in the sample (note that $TSS/(n - 1)$ is the usual sample variance s_y^2).

When an explanatory variable X is included in the model, the predicted value for each Y_i is $\hat{Y}_i = \hat{\alpha} + \hat{\beta}X_i$, the error in this prediction is $Y_i - \hat{Y}_i$, and the error sum of squares is $SSE = \sum (Y_i - \hat{Y}_i)^2$. The two sums of squares are related by

$$\sum (Y_i - \bar{Y})^2 = \sum (Y_i - \hat{Y}_i)^2 + \sum (\hat{Y}_i - \bar{Y})^2. \quad (8.8)$$

Here $SSM = \sum (\hat{Y}_i - \bar{Y})^2 = TSS - SSE$ is the “Model sum of squares”. It is the reduction

in squared prediction errors achieved when we make use of X_i to predict values of Y_i with the regression model, instead of predicting \bar{Y} for all observations. In slightly informal language, SSM is the part of the total variation TSS “explained” by the fitted regression model. In this language, (8.8) can be stated as

Total variation of Y	=	Variation explained by regression	+	Unexplained variation
TSS	=	SSM	+	SSE

The R^2 statistic is defined as

$$R^2 = \frac{TSS - SSE}{TSS} = 1 - \frac{SSE}{TSS} = 1 - \frac{\sum(Y_i - \hat{Y}_i)^2}{\sum(Y_i - \bar{Y})^2}. \quad (8.9)$$

This is the *proportion* of the total variation of Y in the sample explained by the fitted regression model. Its smallest possible value is 0, which is obtained when $\hat{\beta} = 0$, so that X and Y are completely unassociated, X provides no help for predicting Y , and thus $SSE = TSS$. The largest possible value of R^2 is 1, obtained when $\hat{\sigma} = 0$, so that the observed Y can be predicted perfectly from the corresponding X and thus $SSE = 0$. More generally, R^2 is somewhere between 0 and 1, with large values indicating strong linear association between X and Y .

R^2 is clearly a Proportional Reduction of Error (PRE) measure of association of the kind discussed in Section 2.4.5, with $E_1 = TSS$ and $E_2 = SSE$ in the notation of equation for the PRE measure of association in Section 2.4.5. It is also related to the correlation coefficient. In simple linear regression, R^2 is the square of the correlation r between X_i and Y_i . Furthermore, the square root of R^2 is the correlation between Y_i and the fitted values \hat{Y}_i . This quantity, known as the **multiple correlation coefficient** and typically denoted R , is always between 0 and 1. It is equal to the correlation r between X_i and Y_i when r is positive, and the absolute value (removing the $-$ sign) of r when r is negative. For example, for our infant mortality model $r = -0.753$, $R^2 = r^2 = 0.567$ and $R = \sqrt{R^2} = 0.753$.

In SPSS output, the “**ANOVA**” table shows the model, error and total sums of squares SSM , SSE and TSS in the “Sum of Squares column”, on the “Regression”, “Residual” and “Total” rows respectively. R^2 is shown in “**Model summary**” under “R Square” and multiple correlation R next to it as “R”. Figure 8.6 shows these results for the model for IMR given School enrolment. Here $R^2 = 0.567$. Using each country’s level of school enrolment to predict its IMR thus reduces the prediction errors by 56.7% compared to the situation where the predicted IMR is the overall sample mean (here 4.34) for every country. Another conventional way of describing this R^2 result is to say that the variation in rates of School enrolment explains 56.7% of the observed variation in Infant mortality rates.

R^2 is a useful statistic with a convenient interpretation. However, its importance should not be exaggerated. R^2 is rarely the only or the most important part of the model results. This may be the case if the regression model is fitted solely for the purpose of *predicting* future observations of the response variable. More often, however, we are at least or more interested in examining the nature and strength of the associations between the response variable and the explanatory variable (later, variables), in which case the regression coefficients are the main parameters of interest. This point is worth emphasising because in our experience many users of linear regression models tend to place far too much importance on R^2 , often hoping to treat it as the ultimate measure of the goodness of the model. We are frequently asked questions along the lines of “My model has R^2 of 0.42 — is that good?”. The answer tends to be “I have no idea” or, at best, “It depends”. This not a sign of ignorance, because it really does depend:

- Which values of R^2 are large or small or “good” is not a statistical question but a substantive one, to which the answer depends on the nature of the variables under consideration. For example, most associations between variables in the social sciences involve much unexplained variation, so their R^2 values tend to be smaller than for quantities in, say, physics. Similarly, even in social sciences models for aggregates such as countries often have higher values of R^2 than ones for characteristics of individual people. For example, the $R^2 = 0.567$ in our infant mortality example (let alone the $R^2 = 0.753$ we will achieve for a multiple linear model for IMR in Section 8.6) would be unachievably high for many types of individual-level data.
- In any case, achieving large R^2 is usually not the ultimate criterion for selecting a model, and a model can be very useful without having a large R^2 . The R^2 statistic reflects the magnitude of the variation around the fitted regression line, corresponding to the residual standard deviation $\hat{\sigma}$. Because this is an accepted part of the model, R^2 is not a measure of how well the model fits: we can have a model which is essentially true (in that X is linearly associated with Y) but has large residual standard error and thus small R^2 .

8.3.5 Statistical inference for the regression coefficients

The only parameter of the simple linear regression model for which we will describe methods of statistical inference is the slope coefficient β . Tests and confidence intervals for population values of the intercept α are rarely and ones about the residual standard deviation σ almost never substantively interesting, so they will not be considered. Similarly, the only null hypothesis on β discussed here is that its value is zero, i.e.

$$H_0 : \beta = 0. \quad (8.10)$$

Recall that when β is 0, there is no linear association between the explanatory variable X and the response variable Y . Graphically, this corresponds to a regression line in the population which is parallel to the X -axis (see plot (d) of Figure ?? for an illustration of such a line in a sample). The hypothesis (8.10) can thus be expressed in words as

$$H_0 : \text{“There is no linear association between } X \text{ and } Y \text{ in the population”}. \quad (8.11)$$

Tests of this are usually carried out against a two-sided alternative hypothesis $H_a : \beta \neq 0$, and we will also concentrate on this case.

Formulation (8.11) implies that the hypothesis that $\beta = 0$ is equivalent to one that the population correlation ρ between X and Y is also 0. The test statistic presented below for testing (8.10) is also identical to a common test statistic for $\rho = 0$. A test of $\beta = 0$ can thus be interpreted also as a test of no correlation in the population.

The tests and confidence intervals involve both the estimate $\hat{\beta}$ and its estimated standard error, which we will here denote $\widehat{\text{se}}(\hat{\beta})$.¹² It is calculated as

$$\widehat{\text{se}}(\hat{\beta}) = \frac{\hat{\sigma}}{\sqrt{\sum (X_i - \bar{X})^2}} = \frac{\hat{\sigma}}{s_x \sqrt{n-1}} \quad (8.12)$$

¹²It would have been more consistent with related notation used in Chapter 7 to denote it something like $\hat{\sigma}_{\hat{\beta}}$, but this would later become somewhat cumbersome.

where $\hat{\sigma}$ is the estimated residual standard deviation given by (8.7), and s_x is the sample standard deviation of X . The standard error indicates the level of precision with which $\hat{\beta}$ estimates the population parameter β . The last expression in (8.12) shows that the sample size n appears in the denominator of the standard error formula. This means that the standard error becomes smaller as the sample size increases. In other words, the precision of estimation increases when the sample size increases, as with all the other estimates of population parameters we have considered before. In SPSS output, the estimated standard error is given under “Std. Error” in the “**Coefficients**” table. Figure 8.6 shows that $\widehat{\text{se}}(\hat{\beta}) = 0.015$ for the estimated coefficient $\hat{\beta}$ of School enrolment.

The test statistic for the null hypothesis (8.10) is once again of the general form (see the beginning of Section 5.5.2), i.e. a point estimate divided by its standard error. Here this gives

$$t = \frac{\hat{\beta}}{\widehat{\text{se}}(\hat{\beta})}. \quad (8.13)$$

The logic of this is the same as in previous applications of the same idea. Since the null hypothesis (8.10) claims that the population β is zero, values of its estimate $\hat{\beta}$ far from zero will be treated as evidence against the null hypothesis. What counts as “far from zero” depends on how precisely β is estimated from the observed data by $\hat{\beta}$ (i.e. how much uncertainty there is in $\hat{\beta}$), so $\hat{\beta}$ is standardised by dividing by its standard error to obtain the test statistic.

When the null hypothesis (8.10) is true, the sampling distribution of the test statistic (8.13) is a t distribution with $n - 2$ degrees of freedom (i.e. $n - (k + 1)$ where $k = 1$ is the number of explanatory variables in the model). The P -value for the test against a two-sided alternative hypothesis $\beta \neq 0$ is then the probability that a value from a t_{n-2} distribution is at least as far from zero as the value of the observed test statistic. As for the tests of one and two means discussed in Chapter 7, it would again be possible to consider a large-sample version of the test which relaxes the assumption that Y_i given X_i are normally distributed, and uses (thanks to the Central Limit Theorem again) the standard normal distribution to obtain the P -value. With linear regression models, however, the t distribution version of the test is usually used and included in standard computer output, so only it will be discussed here. The difference between P -values from the t_{n-2} and standard normal distributions is in any case minimal when the sample size is reasonably large (at least 30, say).

In the infant mortality example shown in Figure 8.6, the estimated coefficient of School enrolment is $\hat{\beta} = -0.179$, and its estimated standard error is $\widehat{\text{se}}(\hat{\beta}) = 0.015$, so the test statistic is

$$t = \frac{-0.179}{0.015} = -11.94$$

(up to some rounding error). This is shown in the “ t ” column of the “**Coefficients**” table. The P -value, obtained from the t distribution with $n - 2 = 109$ degrees of freedom, is shown in the “Sig.” column. Here $P < 0.001$, so the null hypothesis is clearly rejected. The data thus provide very strong evidence that primary school enrolment is associated with infant mortality rate in the population.

In many analyses, rejecting the null hypothesis of no association will be entirely unsurprising. The question of interest is then not *whether* there is an association in the population, but *how strong* it is. This question is addressed with the point estimate $\hat{\beta}$, combined with a confidence interval which reflects the level of uncertainty in $\hat{\beta}$ as an estimate of the population parameter β . A confidence interval for β with the confidence level $1 - \alpha$ is given by

$$\hat{\beta} \pm t_{\alpha/2}^{(n-2)} \widehat{\text{se}}(\hat{\beta}) \quad (8.14)$$

where the multiplier $t_{\alpha/2}^{(n-2)}$ is obtained from the t_{n-2} distribution as in previous applications of t -based confidence intervals (c.f. the description in Section 7.3.4). For a 95% confidence interval (i.e. one with $\alpha = 0.05$) in the infant mortality example, the multiplier is $t_{0.025}^{(109)} = 1.98$, and the endpoints of the interval are

$$-0.179 - 1.98 \times 0.015 = -0.209 \text{ and } -0.179 + 1.98 \times 0.015 = -0.149.$$

These are also shown in the last two columns of the “**Coefficients**” table of SPSS output. In this example we are thus 95% confident that the expected change in IMR associated with an increase of one percentage point in School enrolment is a decrease of between 0.149 and 0.209 percentage points. If you are calculating this confidence interval by hand, it is (if the sample size is at least 30) again acceptable to use the multiplier 1.96 from the standard normal distribution instead of the t -based multiplier. Here this would give the confidence interval $(-0.208; -0.150)$.

It is often more convenient to interpret the slope coefficient in terms of larger or smaller increments in X than one unit. As noted earlier, a point estimate for the effect of this is obtained by multiplying $\hat{\beta}$ by the appropriate constant. A confidence interval for it is calculated similarly, by multiplying the end points of an interval for $\hat{\beta}$ by the same constant. For example, the estimated effect of a 10-unit increase in School enrolment is $10 \times \hat{\beta} = -1.79$, and a 95% confidence interval for this is $10 \times (-0.209; -0.149) = (-2.09; -1.49)$. In other words, we are 95% confident that the effect is a decrease of between 2.09 and 1.49 percentage points.

8.4 Interlude: Association and causality

Felix, qui potuit rerum cognoscere causas,
atque metus omnis et inexorabile fatum
subiecit pedibus strepitumque Acherontis avari

Blessed is he whose mind had power to probe
The causes of things and trample underfoot
All terrors and inexorable fate
And the clamour of devouring Acheron

(Publius Vergilius Maro: *Georgica* (37-30 BCE), 2.490-492;
translation by L. P. Wilkinson)

These verses from Virgil's *Georgics* are the source of the LSE motto — “Rerum cognoscere causas”, or “To know the causes of things” — which you can see on the School's coat of arms on the cover of this coursepack. As the choice of the motto suggests, questions on *causes* and *effects* are of great importance in social and all other sciences. Causal connections are the mechanisms through which we try to understand and predict what we observe in the world, and the most interesting and important research questions thus tend to involve claims about causes and effects.

We have already discussed several examples of statistical analyses of *associations* between variables. Association is not the same as causation, as two variables can be statistically associated without being in any way directly causally related. Finding an association is thus not *sufficient* for establishing a causal link. It is, however, *necessary* for such a claim: if two variables are not associated, they will not be causally connected either. This implies that examination of associations must be a part of any analysis aiming to obtain conclusions about causal effects.

Definition and analysis of causal effects are considered in more detail on the course MY400 and in much greater depth still on MY457. Here we will discuss only the following simplified empirical version of the question¹³. Suppose we are considering two variables X and Y , and suspect that X is a cause of Y . To support such a claim, we must be able to show that the following three conditions are satisfied:

1. There is a statistical association between X and Y .
2. An appropriate time order: X comes before Y .
3. All alternative explanations for the association are ruled out.

The first two conditions are relatively straightforward, at least in principle. Statistical associations are examined using the kinds of techniques covered on this course, and decisions about whether or not there is an association are usually made largely with the help of statistical inference. Note also that making *statistical* associations one of the conditions implies that this empirical definition of causal effects is not limited to *deterministic* effects, where a particular value of X always leads to exactly the same value of Y . Instead, we consider *probabilistic* causal effects, where changes in X make different values of Y more or less likely. This is clearly crucial in the social sciences, where hardly any effects are even close to deterministic.

The second condition is trivial in many cases where X must logically precede Y in time: for example, a person's sex is clearly determined before his or her income at age 20. In other cases the order is less obvious: for example, if we consider the relationship between political attitudes and readership of different newspapers, it may not be clear whether attitude came before choice

¹³Here adapted from a discussion in Agresti and Finlay, *Statistical Methods for the Social Sciences* (1997).

of paper of vice versa. Clarifying the time order in such cases requires careful research design, often involving measurements taken at several different times.

The really difficult condition is the third one. The list of “all alternative explanations” is essentially endless, and we can hardly ever be sure that all of them have been “ruled out”. Most of the effort and ingenuity in research design and analysis in a study of any causal hypothesis usually goes into finding reasonably convincing ways of eliminating even the most important alternative explanations. Here we will discuss only one general class of such explanations, that of spurious associations due to common causes of X and Y . Suppose that we observe an association, here denoted symbolically by $X - Y$, and would like to claim that this implies a causal connection $X \rightarrow Y$. One situation where such a claim is *not* justified is when both X and Y are caused by a third variable Z , as in the graph in Figure 8.8. If we here consider only X and Y , they will appear to be associated, but the connection is not a causal one. Instead, it is a **spurious association** induced by the dependence on both variables on the common cause Z .



Figure 8.8: A graphical representation of a spurious association between X and Y , explained by dependence on a common cause Z .

To illustrate a spurious association with a silly but memorable teaching example, suppose that we examine a sample of house fires in London, and record the number of fire engines sent to each incident (X) and the amount of damage caused by the fire (Y). There will be a strong association between X and Y , with large numbers of fire engines associated with large amounts of damage. It is also reasonably clear that the number of fire engines is determined before the final extent of damage. The first two conditions discussed above are thus satisfied. We would, however, be unlikely to infer from this that the relationship is causal and conclude that we should try to reduce the cost of fires by dispatching fewer fire engines to them. This is because the association between the number of fire engines and the amount of damages is due to both of them being influenced by the size of the fire (Z). Here this is of course obvious, but in most real research questions possible spurious associations are less easy to spot.

How can we then rule out spurious associations due to some background variables Z ? The usual approach is to try to remove the association between X and Z . This means in effect setting up comparisons between units which have different values of X but the same or similar values of Z . Any differences in Y can then more confidently be attributed to X , because they cannot be due to differences in Z . This approach is known as **controlling** for other variables Z in examining the association between X and Y .

The most powerful way of controlling for background variables is to conduct a **randomized experiment**, where the values of the explanatory variable X can be set by the researcher, and are assigned at random to the units in the study. For instance, of the examples considered in Chapters 5 and 7, Examples 5.3, 5.4 and 7.3 are randomized experiments, each with an intervention variable X with two possible values (placebo or real vaccine, one of two forms of a survey question, and police officer wearing or not wearing sunglasses, respectively). The randomization assures that units with different values of X are on average similar in *all* variables Z which precede X and Y , thus ruling out the possibility of spurious associations due to such variables.

Randomized experiments are for practical or ethical reasons infeasible in many contexts, especially in the social sciences. We may then resort to other, less powerful research designs which help to control for some background variables. This, however, is usually only partially successful, so we may also need methods of control applied at the analysis stage of the research process. This is known as **statistical control**. The aim of statistical control is to estimate and test associations between X and Y while effectively holding the values of some other variables Z constant in the analysis. When the response variable is continuous, the most common way of doing this is the method of multiple linear regression which is described in the next section. When all the variables are categorical, one simple way of achieving the control is analysis of multiway contingency tables, which is described in Chapter 9.

8.5 Multiple linear regression models

8.5.1 Introduction

Simple linear regression becomes multiple linear regression when more than one explanatory variable is included in the model. How this is done is explained in Section 8.5.2 below. The definition of the model builds directly on that of the simple linear model, and most of the elements of the multiple model are either unchanged or minimally modified from the simple one. As a result, we can in Section 8.5.3 cover much of the multiple linear model simply by referring back to the descriptions in Section 8.3. One aspect of the model is, however, conceptually expanded when there are multiple explanatory variables, and requires a careful discussion. This is the meaning of the regression coefficients of the explanatory variables. The interpretation of and inference for these parameters are discussed in Section 8.5.4. The crucial part of this interpretation, and the main motivation for considering multiple regression models, is that it is one way of implementing the ideas of statistical control in analyses for continuous response variables.

The concepts are be illustrated with a further example from the Global Civil Society data set. The response variable will still be the Infant mortality rate of a country, and there will be three explanatory variables: School enrolment, Control of corruption and Income inequality as measured by the Gini index (see Example 8.1). Results for this model are shown in Table 8.5, to which we will refer throughout this section. The table is also an example of the kind of format in which results of regression models are typically reported. Presenting raw computer output such as that in Figure 8.6 is normally not appropriate in formal research reports.

Table 8.5: Response variable: Infant Mortality Rate (%). Results for a linear regression model for Infant mortality rate given three explanatory variables in the Global Civil Society data. $\hat{\sigma} = 2.23$; $R^2 = 0.692$; $n = 111$; $df = 107$

Explanatory variable	Coefficient	Standard error	t	P -value	95 % Confidence interval
Constant	16.40				
School enrolment (%)	-0.139	0.014	-9.87	< 0.001	(-0.167; -0.111)
Control of corruption	-0.046	0.008	-5.53	< 0.001	(-0.062; -0.029)

Income inequality	0.055	0.022	2.50	0.014	(0.011; 0.098)
-------------------	-------	-------	------	-------	-------------------

8.5.2 Definition of the model

Having multiple explanatory variables requires a slight extension of notation. Let us denote the number of explanatory variables in the model by k ; in our example $k = 3$. Individual explanatory variables are then denoted with subscripts as X_1, X_2, \dots, X_k , in the example for instance as $X_1 = \text{School enrolment}$, $X_2 = \text{Control of corruption}$ and $X_3 = \text{Income inequality}$. Observations for individual units i (with values $i = 1, 2, \dots, n$) in the sample are indicated by a further subscript, so that $X_{1i}, X_{2i}, \dots, X_{ki}$ denote the observed values of the k explanatory variables for unit i .

The multiple linear regression model is essentially the same as the simple linear model. The values Y_i of the response variable in the sample are again assumed to be statistically independent, and each of them is regarded as an observation sampled from a normal distribution with mean μ_i and variance σ^2 . The crucial change is that the expected values μ_i now depend on the multiple explanatory variables through

$$\mu_i = \alpha + \beta_1 X_{1i} + \beta_2 X_{2i} + \dots + \beta_k X_{ki} \quad (8.15)$$

where the coefficients $\beta_1, \beta_2, \dots, \beta_k$ of individual explanatory variables are now also identified with subscripts. As in (8.4) for the simple linear model, the multiple model can also be expressed in the concise form

$$Y_i = \alpha + \beta_1 X_{1i} + \beta_2 X_{2i} + \dots + \beta_k X_{ki} + \epsilon_i \quad (8.16)$$

where the error term (population residual) ϵ_i is normally distributed with mean 0 and variance σ^2 .

The expected value of Y as defined in (8.15) is a linear function of X_1, X_2, \dots, X_k . If there are two explanatory variables ($k = 2$), μ is described by a flat *plane* as X_1 and X_2 take different values (think of a flat sheet of paper, at an angle and extended indefinitely in all directions, cutting across a room in the air). A plane is the two-dimensional generalisation of a one-dimensional straight line. The actual observations of Y_i now correspond to points in a three-dimensional space, and they are generally not on the regression plane (think of them as a swarm of bees in the air, some perhaps sitting on that sheet of paper but most hovering above or below it). When k is larger than 2, the regression surface is a higher-dimensional linear object known as a hyperplane. This is impossible to visualise in our three-dimensional world, but mathematically the idea of the model remains unchanged. In each case, the observed values of Y exist in a yet higher dimension, so they cannot in general be predicted exactly even with multiple explanatory variables. A regression plane defined by several X -variables does nevertheless allow for more flexibility for μ_i than a straight line, so it is in general possible to predict Y_i more accurately with a multiple regression model than a simple one. This, however, is not usually the only or main criterion for selecting a good regression model, for reasons discussed in Section 8.5.4 below.

8.5.3 Unchanged elements from simple linear models

As mentioned at the beginning of this section, most elements of the multiple linear regression model are the same or very similar as for the simple model, and require little further explanation:

- The **constant term** (intercept) α is interpreted as the expected value of Y when all of the explanatory variables have the value 0. This can be seen by setting $X_{1i}, X_{2i}, \dots, X_{ki}$ all to 0 in (8.15). As before, α is rarely of any substantive interest. In the example in Table 8.5, the estimated value of α is 16.40.
- The **residual standard deviation** σ is the standard deviation of the conditional distribution of Y given the values of all of X_1, X_2, \dots, X_k . It thus describes the magnitude of the variability of Y around the regression plane. The model assumes that σ is the same at all values of the explanatory variables. In Table 8.5, the estimate of σ is 2.23.
- **Estimates of the regression coefficients** are here denoted with hats as $\hat{\alpha}$ and $\hat{\beta}_1, \hat{\beta}_2, \dots, \hat{\beta}_k$, and fitted (predicted) values for Y_i are given by

$$\hat{Y}_i = \hat{\alpha} + \hat{\beta}_1 X_{1i} + \hat{\beta}_2 X_{2i} + \dots + \hat{\beta}_k X_{ki}. \quad (8.17)$$

The estimated regression coefficients are again obtained with the method of **least squares** by finding the values for $\hat{\alpha}, \hat{\beta}_1, \hat{\beta}_2, \dots, \hat{\beta}_k$ which make the error sum of squares $SSE = \sum (Y_i - \hat{Y}_i)^2$ as small as possible. This is both mathematically and intuitively the same exercise as least squares estimation for a simple linear model, except with more dimensions: here we are finding the best-fitting hyperplane through a high-dimensional cloud of points rather than the best-fitting straight line through a two-dimensional scatterplot.

With more than one explanatory variable, the computational formulas for the estimates become difficult to write down¹⁴ and essentially impossible to calculate by hand. This is not a problem in practice, as they are easily computed by statistical software such as SPSS. In Table 8.5, the least squares estimates of the regression coefficients are shown in the “Coefficient” column. Each row of the table gives the coefficient for one explanatory variable, identified in the first column. A similar format is adopted in SPSS output, where the “**Coefficients**” table looks very similar to the main part of Table 8.5. The arrangement of other parts of SPSS output for multiple linear regression is essentially unchanged from the format shown in Figure 8.6.

- Predicted values for Y can be calculated from (8.17) for any set of values of the explanatory variables (whether those observed in the data or not, as long as extrapolation outside their observed ranges is avoided). This is often very useful for illustrating the implications of a fitted model. For example, Table 8.2 shows that the sample averages of the explanatory variables in Table 8.5 are approximately 86 for School enrolment (X_1), 50 for Control of corruption (X_2) and 40 for Income inequality (X_3). The predicted IMR for a hypothetical “average” country with all these values would be

$$\hat{Y} = 16.4 - 0.139 \times 86 - 0.046 \times 50 + 0.055 \times 40 = 4.35$$

using the estimated intercept $\hat{\alpha} = 16.4$, and the estimated regression coefficients $\hat{\beta}_1 = -0.139$, $\hat{\beta}_2 = -0.046$ and $\hat{\beta}_3 = 0.055$ for X_1 , X_2 and X_3 . For further illustration, we might compare this to other predicted values calculated for, say, different combinations of large and/or small values of the explanatory variables.

¹⁴At least until we adopt extended, so-called matrix notation. In this, the least squares estimates are expressible simply as $\hat{\beta} = (\mathbf{X}'\mathbf{X})^{-1}(\mathbf{X}'\mathbf{Y})$.

- The estimated residual standard error $\hat{\sigma}$ is again calculated from (8.7), using the appropriate value of k . Here $n = 111$ and $k = 3$, and so the degrees of freedom are $df = n - (k + 1) = n - 4 = 107$. The estimate is $\hat{\sigma} = 2.23$.
- The explanation of the coefficient of determination R^2 is entirely unchanged from the one given under “Coefficient of determination (R^2)” in Section 8.3.4. It is still calculated with the formula (8.9), and its interpretation is also the same. The R^2 statistic thus describes the proportion of the sample variation in Y explained by the regression model, i.e. by the variation in the explanatory variables. Similarly, the multiple correlation coefficient $R = \sqrt{R^2}$ is again the correlation between the observed Y_i and the fitted values \hat{Y}_i . In our example, $R^2 = 0.692$ (and $R = \sqrt{0.692} = 0.832$), i.e. about 69.2% of the observed variation in IMR between countries is explained by the variation in levels of School enrolment, Control of corruption and Income inequality between them. Compared to the $R^2 = 0.567$ for the simple regression model in Figure 8.6, adding the two new explanatory variables has increased R^2 by 12.5 percentage points, which seems like a reasonably large increase.

8.5.4 Interpretation and inference for the regression coefficients

Interpretation

The concept of statistical control was outlined in Section 8.4 above. In essence, its idea is to examine the association between a response variable and a particular explanatory variable, while holding all other explanatory variables at constant values. This is useful for assessing claims about causal effects, but also more broadly whenever we want to analyse associations free of the confounding effects of other variables.

When all of the variables were categorical, statistical control could be carried out obviously and transparently by considering partial tables, where the control variables are literally held constant. This is not possible when some of the control variables are continuous, because they then have too many different values for it to be feasible to consider each one separately. Instead, statistical control is implemented with the help of a multiple regression model, and interpreted in terms of the regression coefficients.

Consider, for example, a linear regression model with three explanatory variables X_1 , X_2 and X_3 . This specifies the expected value of Y as

$$\mu = \alpha + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_3 \quad (8.18)$$

for any values of X_1 , X_2 and X_3 . Suppose now that we consider a second observation, which has the same values of X_1 and X_2 as before, but the value of X_3 larger by one unit, i.e. $X_3 + 1$. The expected value of Y is now

$$\mu = \alpha + \beta_1 X_1 + \beta_2 X_2 + \beta_3 (X_3 + 1) = \alpha + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_3 + \beta_3. \quad (8.19)$$

Subtracting (8.18) from (8.19) leaves us with β_3 . In other words, β_3 is the change in expected value of Y when X_3 is increased by one unit, while keeping the values of X_1 and X_2 unchanged. The same result would obviously be obtained for X_1 and X_2 , and for models with any number of explanatory variables. Thus in general

- The regression coefficient of any explanatory variable in a multiple linear regression model shows the change in expected value of the response variable Y when that explanatory variable is increased by one unit, while holding all other explanatory variables constant.

When there is only one explanatory variable, the “while holding...” part is omitted and the interpretation becomes the one for simple linear regression in Section 8.3.3.

This interpretation of the regression coefficients is obtained by “increasing by one unit” and “holding constant” values of explanatory variables by mathematical manipulations alone. It is thus true within the model even when the values of the explanatory variables are not and cannot actually be controlled and set at different values by the researcher. This, however, also implies that this appealing interpretation is a mathematical construction which does not automatically correspond to reality. In short, the interpretation of the regression coefficients is always mathematically true, but whether it is also an approximately correct description of an association in the real world depends on the appropriateness of the model for the data and study at hand. In some studies it is indeed possible to manipulate at least some explanatory variables, and corresponding regression models can then help to draw reasonably strong conclusions about associations between variables. Useful results can also be obtained in studies where no variables are in our control (so-called *observational studies*), as long as the model is selected carefully. This requires, in particular, that a linear model specification is adequate for the data, and that no crucially important explanatory variables have been omitted from the model.

In the IMR example, the estimated coefficients in Table 8.5 are interpreted as follows:

- Holding levels of Control of corruption and Income inequality constant, increasing School enrolment by one percentage point decreases expected IMR by 0.139 percentage points.
- Holding levels of School enrolment and Income inequality constant, increasing Control of corruption by one unit decreases expected IMR by 0.046 percentage points.
- Holding levels of School enrolment and Control of corruption constant, increasing Income inequality by one unit increases expected IMR by 0.055 percentage points.

Instead of “holding constant”, we often talk about “controlling for” other variables in such statements. As before, it may be more convenient to express the interpretations in terms of other increments than one unit (e.g. ten units of the measure of Income inequality) by multiplying the coefficient by the corresponding value.

The association between the response variable Y and a particular explanatory variable X described by the coefficient of X in a multiple regression model is known as a **partial association** between X and Y , controlling for the other explanatory variables in the model. This will often differ from the association estimated from a simple regression model for Y given X , because of the correlations between the control variables and X and Y . In the infant mortality example, the estimated effect of School enrolment was qualitatively unaffected by controlling for the other variables, and decreased in magnitude from -0.179 to -0.139.

Inference

Inference for the regression coefficients in a multiple linear model differs from that for the simple model in interpretation but not in execution. Let $\hat{\beta}_j$ denote the estimated coefficient of an explanatory variable X_j (where j may be any of $1, 2, \dots, k$), and let $\widehat{\text{se}}(\hat{\beta}_j)$ denote its estimated standard error. The standard errors cannot now be calculated by hand, but they are routinely produced by computer packages and displayed as in Table 8.5. A t -test statistic for the null hypothesis discussed below is given by

$$t = \frac{\hat{\beta}_j}{\widehat{\text{se}}(\hat{\beta}_j)}. \quad (8.20)$$

This is identical in form to statistic (8.13) for the simple regression model. The corresponding null hypotheses are, however, subtly but crucially different in the two cases. In a multiple model, (8.20) is a test statistic for the null hypothesis

$$H_0 : \beta_j = 0, \text{ other regression coefficients are unrestricted} \quad (8.21)$$

against the alternative hypothesis

$$H_a : \beta_j \neq 0, \text{ other regression coefficients are unrestricted.}$$

Here the statement about “unrestricted” other parameters implies that neither hypothesis makes any claims about the values of other coefficients than β_j , and these are allowed to have any values. The null hypothesis is a claim about the association between X_j and Y when the other explanatory variables are already included in the model. In other words, (8.20) tests

$$H_0 : \text{ There is no partial association between } X_j \text{ and } Y, \\ \text{controlling for the other explanatory variables.}$$

The sampling distribution of (8.20) when the null hypothesis (8.21) holds is a t distribution with $n - (k + 1)$ degrees of freedom, where k is again the number of explanatory variables in the model. The test statistic and its P -value from the $t_{n-(k+1)}$ distribution are shown in standard computer output, in a form similar to Table 8.5.

It is important to note two things about test results for multiple regression models, such as those in Table 8.5. First, (8.21) implies that if the null hypothesis is not rejected, X_j is not associated with Y , *if* the other explanatory variables are already included in the model. We would typically react to this by removing X_j from the model, while keeping the other variables in it. This is because of a general principle that models should usually be as simple (*parsimonious*) as possible, and not include variables which have no partial effect on the response variable. Second, the k tests and P -values actually refer to k *different* hypotheses of the form (8.21), one for each explanatory variable. This raises the question of what to do if, say, tests for two variables have large P -values, suggesting that either of them could be removed from the model. The appropriate reaction is to remove one of the variables (perhaps the one with the larger P -value) rather than both at once, and then see whether the other still remains nonsignificant (if so, it can then also be removed). This is part of the general area of **model selection**, principles and practice of which are mostly beyond the scope of this course; some further comments on it are given in Section 8.7.

In the example shown in Table 8.5, the P -values are small for the tests for all of the coefficients. Each of the three explanatory variables thus has a significant effect on the response even after controlling for the other two, so none of the variables should be removed from the model.

A confidence interval with confidence level $1 - \alpha$ for any β_j is given by

$$\hat{\beta}_j \pm t_{\alpha/2}^{(n-(k+1))} \widehat{\text{se}}(\hat{\beta}_j). \quad (8.22)$$

This is identical in form and interpretation to the interval (8.14) for simple regression (except that the degrees of freedom are now $df = n - (k + 1)$), so no new issues arise. The confidence intervals for the coefficients in our example (where $df = n - 4 = 107$ and $t_{0.025}^{(107)} = 1.98$) are shown in Table 8.5.

8.6 Including categorical explanatory variables

8.6.1 Dummy variables

Our models for Infant mortality rate so far did not include some more basic characteristics of the countries than school enrolment, corruption and income inequality. In particular, it seems desirable to control for the wealth of a country, which is likely to be correlated with both a health outcome like infant mortality and the other measures of development used as explanatory variables. We will do this by adding to the model the income level of the country, classified in the Global Civil Society Yearbook into three groups as Low, Middle or High income. Here one reason for considering income as a categorical variable is obviously to obtain an illustrative example for this section. However, using a variable like income in a grouped form is also more generally common practice. It also has the advantage that it is one way of dealing with cases where the effects on Y of the corresponding continuous explanatory variable may be nonlinear.

Summary statistics in Table 8.2 show that income group is associated with both IMR and the explanatory variables considered so far: countries with higher income tend to have lower levels of infant mortality, and higher school enrolment, less corruption and less income inequality than lower-income countries. It thus seems that controlling for income is potentially necessary, and may change the conclusions from the model.

Trying to add income level to the model confronts us with a new problem: how can a categorical explanatory variable like this be used in linear regression? This question is not limited to the current example, but is unavoidable in the social sciences. Even just the standard background variables such as sex, marital status, education level, party preference, employment status and region of residence considered in most individual-level studies are mostly categorical. Similarly, most survey data on attitudes and behaviour are collected in a categorical form, and even variables such as age or income which are originally continuous are often used in a grouped form. Categorical variables are thus ubiquitous in the social sciences, and it is essential to be able to use them also in regression models. How this is done is explained in this section, again illustrated with the infant mortality example. Section 8.6.2 then describes a different example for further illustration of the techniques.

The key to handling categorical explanatory variables is the use of dummy variables. A **dummy variable** (or **indicator variable**) is a variable with only two values, 0 and 1. Its value is 1 if a unit is in a particular category of a categorical variable, and 0 if it is not in that category. For example, we can define for each country the variable

$$D_m = \begin{cases} 1 & \text{if Income level is "Middle"} \\ 0 & \text{otherwise.} \end{cases}$$

This would typically be referred to as something like “dummy for middle income level”. Note that the label D_m used here has no special significance, and was chosen simply to make it easy to remember. Dummy variables will be treated below as regular explanatory variables, and we could denote them as X s just like all the others. A dummy for high income level is defined similarly as

$$D_h = \begin{cases} 1 & \text{if Income level is "High"} \\ 0 & \text{otherwise.} \end{cases}$$

The two variables D_m and D_h are enough to identify the income level of any country. If a country is in the middle-income group, the two dummies have the values $D_m = 1$ and $D_h = 0$ (as no country can be in two groups), and if it has high income, the dummies are $D_m = 0$ and $D_h = 1$. For low-income countries, both $D_m = 0$ and $D_h = 0$. There is thus no need to define

a dummy for low income, because this category is identified by the two other dummies being both zero. The same is true in general: if a categorical variable has K categories, only $K - 1$ dummy variables are needed to identify the category of every unit. Note, in particular, that *dichotomous* variables with only two categories ($K = 2$) are fully identified by just one dummy variable. The category which is not given a dummy of its own is known as the **reference category** or **baseline category**. Any category can be the baseline, and this is usually chosen in a way to make interpretation (discussed below) convenient. The results of the model will be the same, whichever baseline is used.

Categorical variables are used as explanatory variables in regression models by including the dummy variables for them in the model. The results for this in our example are shown in Table 8.6. This requires no changes in the definition or estimation of the model, and the parameter estimates, standard errors and quantities for statistical inference are obtained exactly as before even when some of the explanatory variables are dummy variables. The only aspect which requires some further explanation is the interpretation of the coefficients of the dummy variables.

Table 8.6: Response variable: Infant Mortality Rate (%). Results for a linear regression model for Infant mortality rate in the Global Civil Society data, given the three explanatory variables in Table 8.5 and Income level in three groups. $\hat{\sigma} = 2.01$; $R^2 = 0.753$; $n = 111$; $df = 105$.

Explanatory variable	Coefficient	Std. error	t	P -value	95 % Conf. interval
Constant	12.00				
School enrolment (%)	-0.091	0.016	-5.69	< 0.001	(-0.123; -0.059)
Control of corruption	-0.020	0.011	-1.75	0.083	(-0.043; 0.003)
Income inequality	0.080	0.021	3.75	< 0.001	(0.038; 0.122)
Income level (Reference group: Low)					
Middle	-3.210	0.631	-5.09	< 0.001	(-4.461; -1.958)
High	-3.296	1.039	-3.17	0.002	(-5.357; -1.235)

Recall that the regression coefficient of a continuous explanatory variable X is the expected change in the response variable when X is increased by one unit, holding all other explanatory variables constant. Exactly the same interpretation works for dummy variables, except that it is limited to the only one-unit increase possible for them, i.e. from 0 to 1. For example, consider two (hypothetical) countries with values 0 and 1 for the dummy D_m for middle income, but with the same values for the three continuous explanatory variables. How about the other dummy variable D_h , for high income? The interpretation requires that this too is held constant in the comparison. If this constant value was 1, it would not be possible for D_m to be 1 because every country is in only one income group. Thus the only value at which D_h can be held constant while D_m is changed is 0, so that the comparison will be between a country with $D_m = 1$, $D_h = 0$ and one with $D_m = 0$, $D_h = 0$, both with the same values of the other explanatory variables. In other words, the interpretation of the coefficient of D_m refers to a comparison in expected value of Y between a middle-income country and a country in the baseline category of low income, controlling for the other explanatory variables. The same applies to the coefficient of D_h , and of dummy variables in general:

- The coefficient of a dummy variable for a particular level of a categorical explanatory

variable is interpreted as the *difference* in the expected value of the response variable Y between a unit with that level of the categorical variable and a unit in the baseline category, holding all other explanatory variables constant.

Here the estimated coefficient of D_m is -3.21 . In other words, comparing a middle-income country and a low-income country, both with the same levels of School enrolment, Control of corruption and Income inequality, the expected IMR is 3.21 percentage points lower in the middle-income country than in the low-income one. Similarly, a high-income country has an expected IMR 3.296 percentage points lower than a low-income one, other things being equal. The expected difference between the two non-reference levels is obtained as the difference of their coefficients (or by making one of them the reference level, as discussed below); here $-3.296 - (-3.210) = -0.086$, so a high-income country has an expected IMR 0.086 percentage points lower than a middle-income one, controlling for the other explanatory variables.

Predicted values are again obtained by substituting values for the explanatory variables, including appropriate zeros and ones for the dummy variables, into the estimated regression equation. For example, the predicted IMR for a country with School enrolment of 86 %, Control of corruption score of 50 and Income inequality of 40 is

$$\begin{aligned}\hat{Y} &= 12.0 - 0.091 \times 86 - 0.020 \times 50 + 0.080 \times 40 - 3.210 \times 0 - 3.296 \times 0 \\ &= 6.37 \text{ for a low-income country, and} \\ \hat{Y} &= 12.0 - 0.091 \times 86 - 0.020 \times 50 + 0.080 \times 40 - 3.210 \times 1 - 3.296 \times 0 \\ &= 6.37 - 3.21 = 3.16 \text{ for a middle-income country,}\end{aligned}$$

with a difference of 3.21, as expected. Note how the constant term 12.0 sets the level for the baseline (low-income) group, and the coefficient -3.21 shows the change from that level when considering a middle-income country instead. Note also that we should again avoid unrealistic combinations of the variables in such predictions. For example, the above values would not be appropriate for high-income countries, because there are no such countries in these data with Control of corruption as low as 50.

The choice of the reference category does not affect the fitted model, and exactly the same results are obtained with any choice. For example, if high income is used as the reference category instead, the coefficients of the three continuous variables are unchanged from Table 8.6, and the coefficients of the dummy variables for low and middle incomes are 3.296 and 0.086 respectively. The conclusions from these are the same as above: controlling for the other explanatory variables, the difference in expected IMR is 3.296 between low and high-income, 0.086 between middle and high-income and $3.296 - 0.086 = 3.210$ between low and middle-income countries. Because the choice is arbitrary, the baseline level should be selected in whichever way is most convenient for stating the interpretation. If the categorical variable is ordinal (as it is here), it makes most sense for the baseline to be the first or last category. In other words the dummy-variable treatment makes no distinction between nominal and ordinal categorical variables. Both are treated effectively as nominal in fitting the model, and information on any ordering of the categories is ignored.

Significance tests and confidence intervals are obtained for coefficients of dummy variables exactly as for any regression coefficients. Since the coefficient is in this case interpreted as an expected difference between a level of a categorical variable and the reference level, the null hypothesis of a zero coefficient is the hypothesis that there is no such difference. For example, Table 8.6 shows that the coefficients of both the middle income and high income dummies are clearly significantly different from zero. This shows that, controlling for the other explanatory variables, expected infant mortality for both middle and high-income countries is different from that in low-income countries. The 95% confidence intervals in Table 8.6 are intervals for this difference.

On the other hand, the coefficients of the two higher groups are very similar, which suggests that they may not be different from each other. This can be confirmed by fitting the same model with high income as the reference level, including dummies for low and middle groups. In this model (not shown here), the coefficient of the middle income dummy corresponds to the difference of the Middle and High groups. Its P -value is 0.907 and 95% confidence interval $(-1.37; 1.54)$, so the difference is clearly not significant. This suggests that we could simplify the model further by combining the two higher groups and considering only two income groups, low vs. middle/high.

In cases like this where a categorical explanatory variable has more than two categories, t -tests of individual coefficients are tests of hypotheses about no differences between individual categories, not the hypothesis that the variable has no effect overall. This is the hypothesis that the coefficients of the dummies for *all* of the categories are zero. This requires a slightly different test, which will not be considered here. In our example the low income category is so obviously different from the other two that it is clear that the hypothesis of no overall income effect would be rejected.

The main reason for including income group in the example was not to study income effects themselves (it is after all not all that surprising that infant mortality is highest in poor countries), but to control for them when examining partial associations between IMR and the other explanatory variables. These describe the estimated effects of these continuous variables when comparing countries with similar income levels. Comparing the results in Tables 8.5 and 8.6, it can be seen that the effect of School enrolment remains significant and negative (with higher enrolment associated with lower mortality), although its magnitude decreases somewhat after controlling for income group. Some but not all of its estimated effect in the first model is thus explained away by the fact that income is associated with both primary school enrolment and infant mortality, with richer countries having both higher enrolment and lower mortality.

The effect of Income inequality also remains significant in the larger model, even with a slightly increased coefficient. Countries with larger income inequality tend to have higher levels of infant mortality, even when we compare countries with similar levels of income. The effect of Control of corruption, on the other hand, is no longer significant in Table 8.6. This variable is strongly associated with income (as seen in Table 8.2), with the more corrupt countries typically being poor. Controlling for income, however, level of corruption appears to have little further effect on infant mortality. This also suggests that we might simplify the model by omitting the corruption variable.

One final remark on dummy variables establishes a connection to the techniques discussed in Chapter 7. There we described statistical inference for comparisons of the population means of a continuous response variable Y between two groups, denoted 1 and 2. Suppose now that we fit a simple linear regression model for Y , with a dummy variable for group 2 as the only explanatory variable. This gives exactly the same results as the two-sample t -tests and confidence intervals (under the assumption of equal variances in the groups) in Section 7.3. Related to the notation of that section, the coefficients from the model are $\hat{\alpha} = \bar{Y}_1$, $\hat{\beta} = \bar{Y}_2 - \bar{Y}_1$, and $\hat{\sigma}$ from (8.7) is equal to (see equation 11 in Section 7.3.2). Similarly, the standard error (8.12) is the same as $\hat{\sigma}_{\hat{\Delta}}$ in the standard error equation in Section 7.3.2, and the test statistic (8.13) and confidence interval (8.14) are identical with the t -test statistic in Section 7.3.2 and the t distribution version of the equation in Section 7.3.3 respectively.

The connection between linear regression and the two-sample t -test is an illustration of how statistical methods are not a series of separate tricks for different situations, but a set of connected procedures unified by common principles. Whenever possible, methods for more complex problems have been created by extending those for simpler ones, and simple analyses are in turn special cases of more general ones. Although these connections are unfortunately often some-

what obscured by changes in language and notation, trying to understand them is very useful for effective learning of statistics.

8.6.2 A second example

Because the analysis of the models for infant mortality was presented piecemeal to accompany the introduction of different elements of linear regression, an overall picture of that example may have been difficult to discern. This section describes a different analysis in a more concise manner. It is particularly an illustration of the use of dummy variables, as most of the explanatory variables are categorical. The example concerns the relationship between minimum wage and employment, and uses data originally collected and analysed by David Card and Alan Krueger.¹⁵ Most of the choices of analyses and variables considered here are based on those of Card and Krueger. Their article should be consulted for discussion and more detailed analyses.

A minimum wage of \$5.05 per hour came into effect in the U.S. state of New Jersey on April 1 1992. This represented an increase from the previous, federally mandated minimum wage of \$4.25 per hour. Conventional economic theory predicts that employers will react to such an increase by cutting their work force. Here the research hypothesis is thus that employment will be reduced among businesses affected by the new minimum wage. This can be addressed using suitable data, examining a statistical hypothesis of no association between measures of minimum wage increase and change of employment, controlling for other relevant explanatory variables.

Card and Krueger collected data for 410 fast food restaurants at two times, about one month before and eight months after the minimum wage increase came into effect in New Jersey. Only the 368 restaurants with known values of all the variables used in the analyses are included here. Of them, 268 were in New Jersey and had starting wages below \$5.05 at the time of the first interview, so that these had to be increased to meet the new minimum wage. The remaining 100 restaurants provide a control group which was not affected by the change: 75 of them were in neighbouring eastern Pennsylvania where the minimum wage remained at \$4.25, and 25 were in New Jersey but had starting wages of at least \$5.05 even before the increase. The theoretical prediction is that the control group should experience a smaller negative employment change than the restaurants affected by the wage increase, i.e. employment in the control restaurants should not decrease or at least decrease less than in the affected restaurants. Card and Krueger argue that fast food restaurants provide a good population for examining the research question, because they employ many low-wage workers, generally comply with minimum wage legislation, do not receive tips which would complicate wage calculations, and are relatively easy to sample and interview.

The response variable considered here is the change between the two interviews in full-time equivalent employment at the restaurant, defined as the number of full-time workers (including managers) plus half the number of part-time workers. This will be called “Employment change” below. We consider two variables which indicate how the restaurant was affected by the minimum wage increase. The first is simply a dummy variable which is 1 for those New Jersey restaurants where wages needed to be raised because of the increase, and 0 for the other restaurants. These will be referred to as “Affected” and “Unaffected” restaurants respectively. The second variable is also 0 for the unaffected restaurants; for the affected ones, it is the proportion by which their previous starting wage had to be increased to meet the new minimum wage. For example, if the previous starting wage was the old minimum of \$4.25, this “Wage gap” is $(5.05 - 4.25)/4.25 = 0.188$. Finally, we will also use information on the chain the restaurant

¹⁵Card, D. and Krueger, A. B. (1994). Minimum wages and employment: A case study of the fast-food industry in New Jersey and Pennsylvania. *The American Economic Review* **84**, 772–793.

belongs to (Burger King, Roy Rogers, Wendy's or KFC) and whether it is owned by the parent company or the franchisee. These will be included in the analyses as partial control for other factors affecting the labour market, which might have had a differential impact on different types of restaurants over the study period. Summary statistics for the variables are shown in Table 8.7.

Table 8.7: Summary statistics for the variables considered in the minimum wage example of Section 8.6.2. Mean employment change: Among unaffected restaurants: -2.93 ; Among affected restaurants: $+0.68$.

Group	%	(<i>n</i>)	Minimum Affected % (<i>n</i>)	wage variable Wage gap (mean for affected restaurants)	Response vari- able: Em- ploy- ment change (mean)
Overall	100	(368)	72.8 (268)	0.115	-0.30
Ownership					
Franchisee	64.7	(238)	71.8 (171)	0.122	-0.17
Company	35.3	(130)	74.6 (97)	0.103	-0.52
Chain					
Burger King	41.0	(151)	73.5 (111)	0.129	$+0.02$
Roy Rogers	24.7	(91)	72.5 (66)	0.104	-1.89
Wendy's	13.0	(48)	60.4 (29)	0.086	-0.28
KFC	21.2	(78)	79.5 (62)	0.117	$+0.94$

Table 8.8: Response variable: Change in Full-time equivalent employment. Two fitted models for Employment change given exposure to minimum wage increase and control variables. See the text for further details.

Variable	Model (1) Coefficient (std error)	(<i>t</i>) <i>P</i> -value	Model (2) Coeffi- cient (std error)	(<i>t</i>) <i>P</i> - value
Constant	-2.63		-1.54	
Affected by the increase	3.56 (1.02)	(3.50) 0.001	— —	— —
Wage gap	— —	— —	15.88 (6.04)	(2.63) 0.009
Ownership (vs. Franchisee)				
Company	0.22 (1.06)	(0.20) 0.84	0.43 (1.07)	(0.40) 0.69

Chain (vs. Burger King)				
Roy Rogers	-2.00	(-1.56)	-1.84	(-1.43)
	(1.28)	0.12	(1.29)	0.15
Wendy's	0.15	(0.11)	0.36	(0.24)
	(1.44)	0.92	(1.46)	0.81
KFC	0.64	(0.51)	0.81	(0.64)
	(1.25)	0.61	(1.26)	0.52
R^2	0.046		0.032	

Table 8.8 shows results for two linear regression models for Employment change, one using the dummy for affected restaurants and one using Wage gap as an explanatory variable. Both include the same dummy variables for ownership and chain of the restaurant. Consider first the model in column (1), which includes the dummy variable for affected restaurants. The estimated coefficient of this is 3.56, which is statistically significant (with $t = 3.50$ and $P = 0.001$). This means that the estimated expected Employment change for the restaurants affected by the minimum wage increase was 3.56 full-time equivalent employees larger (in the positive direction) than for unaffected restaurants, controlling for the chain and type of ownership of the restaurant. This is the opposite of the theoretical prediction that the difference would be negative, due to the minimum wage increase leading to reductions of work force among affected restaurants but little change for the unaffected ones. In fact, the summary statistics in Table 8.7 show (although without controlling for chain and ownership) that average employment actually *increased* in absolute terms among the affected restaurants, but decreased among the unaffected ones.

The coefficients of the control variables in Table 8.8 describe estimated differences between company-owned and franchisee-owned restaurants, and between Burger Kings and restaurants of other chains, controlling for the other variables in the model. All of these coefficients have high P -values for both models, suggesting that the differences are small. In fact, the only one which is borderline significant, after all the other control dummies are successively removed from the model (not shown here), is that Employment change seems to have been more negative for Roy Rogers restaurants than for the rest. This side issue is not investigated in detail here. In any case, the control variables have little influence on the effect of the variable of main interest: if all the control dummies are removed from Model (1), the coefficient of the dummy variable for affected restaurants becomes 3.61 with a standard error of 1.01, little changed from the estimates in Table 8.8. This is not entirely surprising, as the control variables are weakly associated with the variable of interest: as seen in Table 8.7, the proportions of affected restaurants are mostly fairly similar among restaurants of all chains and types of ownership.

In their article, Card and Krueger carefully explore (and confirm) the robustness of their findings by considering a series of variants of the analysis, with different choices of variables and sets of observations. This is done to try to rule out the possibility that the main conclusions are reliant on, and possibly biased by, some specific features of the data and variables in the initial analysis, such as missing data or measurement error. Such sensitivity analyses would be desirable in most social science contexts, where single definitely best form of analysis is rarely obvious. Here we will carry out a modest version of such an assessment by considering the Wage gap variable as an alternative measure of the impact of minimum wage increase, instead of a dummy for affected restaurants. This is a continuous variable, but one whose values are 0 for all unaffected restaurants and vary only among the affected ones. The logic of using Wage gap as an explanatory variable here is that Employment change could be expected to depend not only on *whether* a restaurant had to increase its starting wage to meet the new minimum wage, but also on *how large* that compulsory increase was.

The results for the second analysis are shown as Model (2) in Table 8.8. The results are qualitatively the same as for Model (1), in that the coefficients of the control dummies are not significant, and that of Wage gap (which is 15.88) is significant and positive. The estimated employment change is thus again larger for affected than for unaffected restaurants, and their difference now even increases when the wage rise required from an affected restaurant increases. To compare these results more directly to Model (1), we can consider a comparison between an unaffected restaurant (with Wage gap 0) and an affected one with Wage gap equal to its mean among the affected restaurants, which is 0.115 (c.f. Table 8.7). The estimated difference in Employment change between them, controlling for ownership and chain, is $0.115 \times 15.88 = 1.83$, which is somewhat lower than the 3.56 estimated from model (1).

This example is also a good illustration of the limitations of the R^2 statistic. The R^2 values of 0.046 and 0.032 are very small, so over 95% of the observed variation in employment changes remains unexplained by the variation in the three explanatory variables. In other words, there are large differences in Employment change experienced even by affected or unaffected restaurants of the same chain and type of ownership. This would make *predicting* the employment change for a particular restaurant a fairly hopeless task with these explanatory variables. However, prediction is not the point here. The research question focuses on possible differences in *average* changes in employment, and finding such differences is of interest even if variation around the averages is large.

In summary, the analysis provides no support for the theoretical prediction that the restaurants affected by a minimum wage increase would experience a larger negative job change than control restaurants. In fact, there was a small but significant difference in the opposite direction in both models described here, and in all of the analyses considered by Card and Krueger. The authors propose a number of tentative explanations for this finding, but do not present any of them as definitive.

8.7 Other issues in linear regression modelling

The material in this chapter provides a reasonably self-contained introduction to linear regression models. However, it is not possible for a course like this to cover comprehensively all aspects of the models, so some topics have been described rather superficially and several have been omitted altogether. In this section we briefly discuss some of them. First, three previously unexplained small items in standard SPSS output are described. Second, a list of further topics in linear regression is given.

An example of SPSS output for linear regression models was given in Figure 8.6. Most parts of it have been explained above, but three have not been mentioned. These can be safely ignored, because each is of minor importance in most analyses. However, it is worth giving a brief explanation so as not to leave these three as mysteries:

- “Adjusted R Square” in the “**Model Summary**” table is a statistic defined as $R_{adj}^2 = [(n-1)R^2 - k]/(n-k-1)$. This is most relevant in situations where the main purpose of the model is prediction of future observations of Y . The population value of the R^2 statistic is then a key criterion of model selection. R_{adj}^2 is a better estimate of it than standard R^2 . Unlike R^2 , R_{adj}^2 does not always increase when new explanatory variables are added to the model. As a sample statistic, R_{adj}^2 does not have the same interpretation as the proportion of variation of Y explained as standard R^2 .
- The last two columns of the “**ANOVA**” (Analysis of Variance) table show the test statistic

and P -value for the so-called F -test¹⁶. The null hypothesis for this is that the regression coefficients of *all* the explanatory variables are zero, i.e. $\beta_1 = \beta_2 = \dots = \beta_k = 0$. In the case of simple regression ($k = 1$), this is equivalent to the t -test for $\beta = 0$. In multiple regression, it implies that none of the explanatory variables have an effect on the response variable. In practice, this is rejected in most applications. Rejecting the hypothesis implies that at least one of the explanatory variables is associated with the response, but the test provides no help for identifying *which* of the individual partial effects are significant. The F -test is thus usually largely irrelevant. More useful is an extended version of it (which is not included in the default output), which is used for hypotheses that a set of several regression coefficients (but not all of them) is zero. For example, this could be used in the example of Table 8.6 to test if income level had no effect on IMR, i.e. if the coefficients of the dummies for *both* middle and high income were zero.

- The “Standardized Coefficients/Beta” in the “**Coefficients**” table are defined as $(s_{xj}/s_y)\hat{\beta}_j$, where $\hat{\beta}_j$ is the estimated coefficient of X_j , and s_{xj} and s_y are sample standard deviations of X_j and Y respectively. This is equal to the correlation of Y and X_j when X_j is the only explanatory variable in the model, but not otherwise. The standardized coefficient describes the expected change in Y in units of its sample standard error, when X_j is increased by one standard error s_{xj} , holding other explanatory variables constant. The aim of this exercise is to obtain coefficients which are more directly comparable between different explanatory variables. Ultimately it refers to the question of **relative importance** of the explanatory variables, i.e. “Which of the explanatory variables in the model is the most important?” This is understandably of interest in many cases, often perhaps more so than any other aspect of the model. Unfortunately, however, relative importance is also one of the hardest questions in modelling, and one without a purely statistical solution. Despite their appealing title, standardized coefficients have problems of their own and do not provide a simple tool for judging relative importance. For example, their values depend not only on the strength of association described by $\hat{\beta}_j$ but also on the standard deviation s_{xj} , which can be different in different samples.

Sensible comparisons of the magnitudes of expected changes in Y in response to changes in individual explanatory variables can usually be presented even without reference to standardized coefficients, simply by using the usual coefficients $\hat{\beta}_j$ and carefully considering the effects of suitably chosen increments of the explanatory variables. In general, it is also worth bearing in mind that questions of relative importance are often conceptually troublesome, for example between explanatory variables with very different practical implications. For instance, suppose that we have fitted a model for a measure of the health status of a person, given the amount of physical exercise the person takes (which can be changed by him/herself), investment in preventive healthcare in the area where the person lives (which can be changed, but with more effort and not by the individual) and the person’s age (which cannot be manipulated at all). The values of the unstandardized or standardized coefficients of these explanatory variables can certainly be compared, but it is not clear what statements about the relative sizes of the effects of “increasing” them would really mean.

A further course on linear regression (e.g. first half of MY452) will typically examine the topics covered on this course in more detail, and then go on to discuss further issues. Here we will give just a list of some such topics, in no particular order:

- Model **diagnostics** to examine whether a particular model appears to be adequate for the data. The residuals $Y_i - \hat{Y}_i$ are a key tool in this, and the most important graphical

¹⁶The sampling distribution of this test is the F distribution. The letter in both refers to Sir Ronald Fisher, the founder of modern statistics.

diagnostic is simply a scatterplot of the residuals against the fitted values \hat{Y}_i . One task of diagnostics is to identify individual **outliers** and **influential observations** which have a substantial impact on the fitted model.

- Modelling **nonlinear effects** of the explanatory variables. This is mostly done simply by including transformed values like squares X^2 or logarithms $\log(X)$ as explanatory variables in the model. It is sometimes also useful to transform the response variable, e.g. using $\log(Y)$ as the response instead of Y .
- Including **interactions** between explanatory variables in the model. This is achieved simply by including products of them as explanatory variables.
- Identifying and dealing with problems caused by extremely high correlations between the explanatory variables, known as problems of **multicollinearity**.
- **Model selection** to identify the best sets of explanatory variables to be used in the model. This may employ both significance tests and other approaches.
- Analysis of Variance (**ANOVA**) and Analysis of Covariance (**ANCOVA**), which are terms used for models involving only or mostly categorical explanatory variables, particularly in the context of randomized experiments. Many of these models can be fitted as standard linear regression models with appropriate use of dummy variables, but the conventional terminology and notation for ANOVA and ANCOVA are somewhat different from the ones used here.

Chapter 9

Analysis of 3-way contingency tables

In Section 2.4 and Chapter 4 we discussed the analysis of two-way contingency tables (crosstabulations) for examining the associations between two categorical variables. In this section we extend this by introducing the basic ideas of **multiway contingency tables** which include more than two categorical variables. We focus solely on the simplest instance of them, a **three-way table** of three variables.

This topic is thematically related also to some of Chapter 8, in that a multiway contingency table can be seen as a way of implementing for categorical variables the ideas of statistical control that were also a feature of the multiple linear regression model of Section 8.5. Here, however, we will not consider formal regression models for categorical variables (these are mentioned only briefly at the end of the chapter). Instead, we give examples of analyses which simply apply familiar methods for two-way tables repeatedly for tables of two variables at fixed values of a third variable.

The discussion is organised around three examples. In each case we start with a two-way table, and then introduce a third variable which we want to control for. This reveals various features in the examples, to illustrate the types of findings that may be uncovered by statistical control.

Example 9.1: Berkeley admissions

Table 9.1 summarises data on applications for admission to graduate study at the University of California, Berkeley, for the fall quarter 1973¹. The data are for five of the six departments with the largest number of applications, labelled below Departments 2–5 (Department 1 will be discussed at the end of this section). Table 9.1 shows the two-way contingency table of the sex of the applicant and whether he or she was admitted to the university.

Table 9.1: Table of sex of applicant vs. admission in the Berkeley admissions data. The column labelled ‘% Yes’ is the percentage of applicants admitted within each row. $\chi^2 = 38.4$, $df = 1$, $P < 0.001$.

	Admitted	Admitted		
tabularnewline Sex	No	Yes	% Yes	Total
Male	1180	686	36.8	1866
Female	1259	468	27.1	1727

¹ESS Round 5: European Social Survey Round 5 Data (2010). Data file edition 2.0. Norwegian Social Science Data Services, Norway Data Archive and distributor of ESS data. The full data can be obtained from <http://ess.nsd.uib.no/ess/round5/>.

Total	2439	1154	32.1	3593
-------	------	------	------	------

The percentages in Table 9.1 show that men were more likely to be admitted, with a 36.8% success rate compared to 27.1% for women. The difference is strongly significant, with $P < 0.001$ for the χ^2 test of independence. If this association was interpreted causally, it might be regarded as evidence of sex bias in the admissions process. However, other important variables may also need to be considered in the analysis. One of them is the academic department to which an applicant had applied. Information on the department as well as sex and admission is shown in Table 9.2.

Table 9.2: Sex of applicant vs. admission by academic department in the Berkeley admissions data.

		Admitted		% Yes	Total
tabularnewline Department	Sex	No	Yes		
2	Male	207	353	63.0	560
	Female	8	17	68.0	25
	Total	215	370	63.2	585
$\chi^2 = 0.25, P = 0.61$					
3	Male	205	120	36.9	325
	Female	391	202	34.1	593
	Total	596	322	35.1	918
$\chi^2 = 0.75, P = 0.39$					
4	Male	279	138	33.1	417
	Female	244	131	34.9	375
	Total	523	269	34.0	792
$\chi^2 = 0.30, P = 0.59$					
5	Male	138	53	27.7	191
	Female	299	94	23.9	393
	Total	437	147	25.2	584
$\chi^2 = 1.00, P = 0.32$					
6	Male	351	22	5.9	373
	Female	317	24	7.0	341
	Total	668	46	6.4	714
$\chi^2 = 0.38, P = 0.54$					
Total		2439	1154	32.1	3593

Table 9.2 is a *three-way* contingency table, because each of its internal cells shows the number of applicants with a particular combination of three variables: department, sex and admission status. For example, the frequency 207 in the top left corner indicates that there were 207 male applicants to department 2 who were not admitted. Table 9.2 is presented in the form of a series of tables of sex vs. admission, just like in the original two-way table 9.1, but now with one table for each department. These are known as **partial tables** of sex vs. admission, **controlling for** department. The word “control” is used here in the same sense as before: each partial table summarises the data for the applicants to a single department, so the variable “department” is literally held constant within the partial tables.

Table 9.2 also contains the marginal distributions of sex and admission status within each department. They can be used to construct the other two possible two-way tables for these variables, for department vs. sex of applicant and department vs. admission status. This information, summarised in Table 9.3, is discussed further below.

The association between sex and admission within each partial table can be examined using methods for two-way tables. For every one of them, the χ^2 test shows that the hypothesis of independence cannot be rejected, so there is no evidence of sex bias within any department. The apparent association in Table 9.1 is thus spurious, and disappears when we control for department. Why this happens can be understood by considering the distributions of sex and admissions across departments, as shown in Table 9.3. Department is clearly associated with sex of the applicant: for example, almost all of the applicants to department 2, but only a third of the applicants to department 5 are men. Similarly, there is an association between department and admission: for example, nearly two thirds of the applicants to department 2 but only a quarter of the applicants to department 5 were admitted. It is the combination of these associations which induces the spurious association between sex and admission in Table 9.1. In essence, women had a lower admission rate overall because relatively more of them applied to the more selective departments and fewer to the easy ones.

Table 9.3: Percentages of male applicants and applicants admitted by department in the Berkeley admissions data.

	Department				
tabularnewline Of all applicants	2	3	4	5	6
% Male	96	35	53	33	52
% Admitted	63	35	34	25	6
Number of applicants	585	918	792	584	714

One possible set of causal connections leading to a spurious association between X and Y was represented graphically by Figure 8.8. There are, however, other possibilities which may be more appropriate in particular cases. In the admissions example, department (corresponding to the control variable Z) cannot be regarded as the cause of the sex of the applicant. Instead, we may consider the causal chain $\text{Sex} \longrightarrow \text{Department} \longrightarrow \text{Admission}$. Here department is an *intervening variable* between sex and admission rather than a common cause of them. We can still argue that sex has an effect on admission, but it is an *indirect effect* operating through the effect of sex on choice of department. The distinction is important for the original research question behind these data, that of possible sex bias in admissions. A direct effect of sex on likelihood of admission might be evidence of such bias, because it might indicate that departments are treating male and female candidates differently. An indirect effect of the kind found here does not suggest bias, because it results from the applicants' own choices of which department to apply to.

In the admissions example a strong association in the initial two-way table was “explained away” when we controlled for a third variable. The next example is one where controlling leaves the initial association unchanged.

Example 9.2: Importance of short-term gains for investors (continued)

Table 2.7 showed a relatively strong association between a person's age group and his or her attitude towards short-term gains as an investment goal. This association is also strongly significant, with $P < 0.001$ for the χ^2 test of independence. Table 9.5 shows the crosstabulations of these variables, now controlling also for the respondent's sex. The association is now still significant in both partial tables. An investigation of the row proportions suggests that the pattern of association is very similar in both tables, as is its strength as measured by the γ statistic ($\gamma = -0.376$ among men and $\gamma = -0.395$ among women). The conclusions obtained from the original two-way table are thus unchanged after controlling for sex.

Table 9.4: Frequencies of respondents by age group and attitude towards short-term gains in Example 9.2, controlling for sex of respondent. The numbers below the frequencies are proportions within rows. $\chi^2 = 82.4$, $df = 9$, $P < 0.001$. $\gamma = -0.376$.

MEN Age group	Irrelevant	Slightly Important	Very important	Total	
Under 45	29 0.250	35 0.302	30 0.259	22 0.190	116 1.000
45–54	83 0.371	60 0.268	52 0.232	29 0.129	224 1.000
55–64	116 0.580	40 0.200	28 0.140	16 0.080	200 1.000
65 and over	150 0.649	53 0.229	16 0.069	12 0.052	231 1.000
Total	378 0.490	188 0.244	126 0.163	79 0.102	771 1.000

Table 9.5: Frequencies of respondents by age group and attitude towards short-term gains in Example 9.2, controlling for sex of respondent. The numbers below the frequencies are proportions within rows. $\chi^2 = 27.6$, $df = 9$, $P = 0.001$. $\gamma = -0.395$.

WOMEN	Age group	Irrelevant	Slightly Important	Very important	Total	
			important	important		
	Under 45	8 0.267	10 0.333	8 0.267	4 0.133	30 1.000
	45–54	28 0.483	17 0.293	5 0.086	8 0.138	58 1.000
	55–64	37 0.698	9 0.170	3 0.057	4 0.075	53 1.000
	65 and over	43 0.717	11 0.183	3 0.050	3 0.050	60 1.000
	Total	116 0.577	47 0.234	19 0.095	19 0.095	201 1.000

Example 9.3: *The Titanic*

The passenger liner RMS *Titanic* hit an iceberg and sank in the North Atlantic on 14 April 1912, with heavy loss of life. Table 9.6 shows a crosstabulation of the people on board the *Titanic*, classified according to their status (as male passenger, female or child passenger, or member of the ship's crew) and whether they survived the sinking². The χ^2 test of independence has $P < 0.001$ for this table, so there are statistically significant differences in probabilities of

²The data were obtained from <http://data.london.gov.uk/datastore/package/london-borough-profiles>.

If you download the “Profiles in Excel” workbook, you will find that one of the pages contains a map of the boroughs, and a tool for visualising the data on that map. A regular map of the boroughs can be found at for example at

<http://www.londoncouncils.gov.uk/londonfacts/londonlocalgovernment/londonmapandlinks/default.htm>.

survival between the groups. The table suggests, in particular, that women and children among the passengers were more likely to survive than male passengers or the ship's crew.

Table 9.6: Survival status of the people aboard the *Titanic*, divided into three groups. The numbers in brackets are proportions of survivors and non-survivors within each group. $\chi^2 = 418$, $df = 2$, $P < 0.001$.

Group	Survivor		Total
	Yes	No	
Male passenger	146 (0.181)	659 (0.819)	805 (1.000)
Female or child passenger	353 (0.691)	158 (0.309)	511 (1.000)
Crew member	212 (0.240)	673 (0.760)	885 (1.000)
Total	711 (0.323)	1490 (0.677)	2201 (1.000)

We next control also for the class in which a person was travelling, classified as first, second or third class. Since class does not apply to the ship's crew, this analysis is limited to the passengers, classified as men vs. women and children. The two-way table of sex by survival status for them is given by Table 9.6, ignoring the row for crew members. This association is strongly significant, with $\chi^2 = 344$ and $P < 0.001$.

Table 9.7: Survival status of the passengers of the *Titanic*, classified by class and sex. The numbers below the frequencies are proportions within rows.

Class	Group	Survivor Yes	No	Total
First	Man	57	118	175
		0.326	0.674	1.000
	Woman or child	146	4	150
		0.973	0.027	1.000
	Total	203	122	325
Second	Man	14	154	168
		0.083	0.917	1.000
	Woman or child	104	13	117
		0.889	0.111	1.000
	Total	118	167	285
Third	Man	75	387	462
		0.162	0.838	1.000
	Woman or child	103	141	244
		0.422	0.578	1.000
	Total	178	528	706
		0.252	0.748	1.000

Total	499	817	1316
	0.379	0.621	1.000

Two-way tables involving class (not shown here) suggest that it is mildly associated with sex (with percentages of men 54%, 59% and 65% in first, second and third class respectively) and strongly associated with survival (with 63%, 41% and 25% of the passengers surviving). It is thus possible that class might influence the association between sex and survival. This is investigated in Table 9.7, which shows the partial associations between sex and survival status, controlling for class. This association is strongly significant (with $P < 0.001$ for the χ^2 test) in every partial table, so it is clearly not explained away by associations involving class. The direction of the association is also the same in each table, with women and children more likely to survive than men among passengers of every class.

The presence and direction of the association in the two-way Table 9.6 are thus preserved and similar in every partial table controlling for class. However, there appear to be differences in the *strength* of the association between the partial tables. Considering, for example, the ratios of the proportions in each class, women and children were about 3.0 ($= 0.973/0.326$) times more likely to survive than men in first class, while the ratio was about 10.7 in second class and 2.6 in the third. The contrast of men vs. women and children was thus strongest among second-class passengers. This example differs in this respect from the previous ones, where the associations were similar in each partial table, either because they were all essentially zero (Example 9.1) or non-zero but similar in both direction and strength (Example 9.2).

We are now considering three variables, class, sex and survival. Although it is not necessary for this analysis to divide them into explanatory and response variables, introducing such a distinction is useful for discussion of the results. Here it is most natural to treat survival as the response variable, and both class and sex as explanatory variables for survival. The associations in the partial tables in Table 9.7 are then partial associations between the response variable and one of the explanatory variables (sex), controlling for the other explanatory variable (class). As discussed above, the strength of this partial association is different for different values of class. This is an example of a statistical **interaction**. In general, there is an interaction between two explanatory variables if the strength and nature of the partial association of (either) one of them on a response variable depends on the value at which the other explanatory variable is controlled. Here there is an interaction between class and sex, because the association between sex and survival is different at different levels of class.

Interactions are an important but challenging element of many statistical analyses. Important, because they often correspond to interesting and subtle features of associations in the data. Challenging, because understanding and interpreting them involves talking about (at least) three variables at once. This can seem rather complicated, at least initially. It adds to the difficulty that interactions can take many forms. In the *Titanic* example, for instance, the nature of the class-sex interaction was that the association between sex and survival was in the same direction but of different strengths at different levels of class. In other cases associations may disappear in some but not all of the partial tables, or remain strong but in different directions in different ones. They may even all or nearly all be in a different direction from the association in the original two-way table, as in the next example.

Example 9.4: Smoking and mortality

A health survey was carried out in Whickham near Newcastle upon Tyne in 1972–74, and a follow-up survey of the same respondents twenty years later³. Here we consider only the

³The data can be obtained from <http://www3.norc.org/gss+website/>, which gives further information on the survey, including the full text of the questionnaires.

$n = 1314$ female respondents who were classified by the first survey either as current smokers or as never having smoked. Table 9.8 shows the crossclassification of these women according to their smoking status in 1972–74 and whether they were still alive twenty years later. The χ^2 test indicates a strongly significant association (with $P = 0.003$), and the numbers suggest that a smaller proportion of smokers than of nonsmokers had died between the surveys. Should we thus conclude that smoking helps to keep you alive? Probably not, given that it is known beyond reasonable doubt that the causal relationship between smoking and mortality is in the opposite direction. Clearly the picture has been distorted by failure to control for some relevant further variables. One such variable is the age of the respondents.

Table 9.8: Table of smoking status in 1972–74 vs. twenty-year survival among the respondents in Example 9.4. The numbers below the frequencies are proportions within rows.

Smoker	Dead	Alive	Total
Yes	139	443	582
	0.239	0.761	1.000
No	230	502	732
	0.314	0.686	1.000
Total	369	945	1314
	0.281	0.719	1.000

Table 9.9 shows the partial tables of smoking vs. survival controlling for age at the time of the first survey, classified into seven categories. Note first that this three-way table appears somewhat different from those shown in Tables 9.2, 9.5 and 9.7. This is because one variable, survival status, is summarised only by the percentage of survivors within each combination of age group and smoking status. This is a common trick to save space in three-way tables involving dichotomous variables like survival here. The full table can easily be constructed from these numbers if needed. For example, 98.4% of the nonsmokers aged 18–24 were alive at the time of the second survey. Since there were a total of 62 respondents in this group (as shown in the last column), this means that 61 of them (i.e. 98.4%) were alive and 1 (or 1.6%) was not.

The percentages in Table 9.9 show that in five of the seven age groups the proportion of survivors is higher among nonsmokers than smokers, i.e. these partial associations in the sample are in the opposite direction from the association in Table 9.8. This reversal is known as **Simpson's paradox**. The term is somewhat misleading, as the finding is not really paradoxical in any logical sense. Instead, it is again a consequence of a particular pattern of associations between the control variable and the other two variables.

Table 9.9: Percentage of respondents in Example 9.4 surviving at the time of the second survey, by smoking status and age group in 1972–74.

Age group	% Alive after 20 years:		Number (in 1972–74):	
	Smokers	Nonsmokers	Smokers	Nonsmokers
18–24	96.4	98.4	55	62
25–34	97.6	96.8	124	157
35–44	87.2	94.2	109	121
45–54	79.2	84.6	130	78
55–64	55.7	66.9	115	121
65–74	19.4	21.7	36	129

75–	0.0	0.0	12	64
All age groups	76.1	68.6	582	732

The two-way tables of age by survival and age by smoking are shown side by side in Table 9.10. The table is somewhat elaborate and unconventional, so it requires some explanation. The rows of the table correspond to the age groups, identified by the second column, and the frequencies of respondents in each age group are given in the last column. The left-hand column shows the percentages of survivors within each age group. The right-hand side of the table gives the two-way table of age group and smoking status. It contains percentages calculated both within the rows (without parentheses) and columns (in parentheses) of the table. As an example, consider numbers for the age group 18–24. There were 117 respondents in this age group at the time of the first survey. Of them, 47% were then smokers and 53% were nonsmokers, and 97% were still alive at the time of the second survey. Furthermore, 10% of all the 582 smokers, 9% of all the 732 nonsmokers and 9% of the 1314 respondents overall were in this age group.

Table 9.10: Two-way contingency tables of age group vs. survival (on the left) and age group vs. smoking (on the right) in Example 6.4. The percentages in parentheses are column percentages (within smokers or nonsmokers) and the ones without parentheses are row percentages (within age groups).

		Row and column %		Total	Count
tabularnewline % Alive	Age group	Smokers	Nonsmokers	%	
97	18–24	47 (10)	53 (9)	100 (9)	117
97	25–34	44 (21)	56 (21)	100 (21)	281
91	35–44	47 (19)	53 (17)	100 (18)	230
81	45–54	63 (22)	38 (11)	100 (16)	208
61	55–64	49 (20)	51 (17)	100 (13)	236
21	65–74	22 (6)	78 (18)	100 (13)	165
0	75–	17 (2)	83 (9)	100 (6)	77
72	Total %	44 (100)	56 (100)	100 (100)	
945	Total count	582	732		1314

Table 9.10 shows a clear association between age and survival, for understandable reasons mostly unconnected with smoking. The youngest respondents of the first survey were highly likely and the oldest unlikely to be alive twenty years later. There is also an association between age and smoking: in particular, the proportion of smokers was lowest among the oldest respondents. The implications of this are seen perhaps more clearly by considering the column proportions, i.e. the age distributions of smokers and nonsmokers in the original sample. These show that the group of nonsmokers was substantially older than that of the smokers; for example, 27% of the nonsmokers but only 8% of the smokers belonged to the two oldest age groups. It is this imbalance which explains why nonsmokers, more of whom are old, appear to have lower chances

of survival until we control for age.

The discussion above refers to the *sample* associations between smoking and survival in the partial tables, which suggest that mortality is higher among smokers than nonsmokers. In terms of statistical significance, however, the evidence is fairly weak: the association is even borderline significant only in the 35–44 and 55–64 age groups, with P -values of 0.063 and 0.075 respectively for the χ^2 test. This is an indication not so much of lack of a real association but of the fact that these data do not provide much power for detecting it. Overall twenty-year mortality is a fairly rough measure of health status for comparisons of smokers and nonsmokers, especially in the youngest and oldest age groups where mortality is either very low or very high for everyone. Differences are likely to have been further diluted by many of the original smokers having stopped smoking between the surveys. This example should thus not be regarded as a serious examination of the health effects of smoking, for which much more specific data and more careful analyses are required⁴.

The Berkeley admissions data discussed earlier provide another example of a partial Simpson's paradox. Previously we considered only departments 2–6, for none of which there was a significant partial association between sex and admission. For department 1, the partial table indicates a strongly significant difference in favour of women, with 82% of the 108 female applicants admitted, compared to 62% of the 825 male applicants. However, the two-way association between sex and admission for departments 1–6 combined remains strongly significant and shows an even larger difference in favour of men than before. This result can now be readily explained as a result of imbalances in sex ratios and rates of admission between departments. Department 1 is both easy to get into (with 64% admitted) and heavily favoured by men (88% of the applicants). These features combine to contribute to higher admissions percentages for men overall, even though within department 1 itself women are more likely to be admitted.

In summary, the examples discussed above demonstrate that many things can happen to an association between two variables when we control for a third one. The association may disappear, indicating that it was spurious, or it may remain similar and unchanged in all of the partial tables. It may also become different in different partial tables, indicating an interaction. Which of these occurs depends on the patterns of associations between the control variable and the other two variables. The crucial point is that the two-way table alone cannot reveal which of these cases we are dealing with, because the counts in the two-way table could split into three-way tables in many different ways. The only way to determine how controlling for other variables will affect an association is to actually do so. This is the case not only for multiway contingency tables, but for all methods of statistical control, in particular multiple linear regression and other regression models.

Two final remarks round off our discussion of multiway contingency tables:

- Extension of the ideas of three-way tables to four-way and larger contingency tables is obvious and straightforward. In such tables, every cell corresponds to the subjects with a particular combination of the values of four or more variables. This involves no new conceptual difficulties, and the only challenge is how to arrange the table for convenient presentation. When the main interest is in associations between a particular pair of two variables, the usual solution is to present a set of partial two-way tables for them, one for each combination of the other (control) variables. Suppose, for instance, that in the university admissions example we had obtained similar data at two different years, say 1973 and 2003. We would then have four variables: year, department, sex and admission status. These could be summarised as in Table 9.2, except that each partial table for sex

⁴ESS Round 5: European Social Survey Round 5 Data (2010). Data file edition 2.0. Norwegian Social Science Data Services, Norway. Data Archive and distributor of ESS data. The full data can be obtained from <http://ess.nsd.uib.no/ess/round5/>.

vs. admission would now be conditional on the values of both year and department. The full four-way table would then consist of ten 2×2 partial tables, one for each of the ten combinations of two years and five departments, (i.e. applicants to Department 2 in 1973, Department 2 in 2003, and so on to Department 6 in 2003).

- The only inferential tool for multiway contingency tables discussed here was to arrange the table as a set of two-way partial tables and to apply the χ^2 test of independence to each of them. This is a perfectly sensible approach and a great improvement over just analysing two-way tables. There are, however, questions which cannot easily be answered with this method. For example, when can we say that associations in different partial tables are different enough for us to declare that there is evidence of an interaction? Or what if we want to consider many different partial associations, either for a response variable with each of the other variables in turn, or because there is no single response variable? More powerful methods are required for such analyses. They are multiple regression models like the multiple linear regression of Section 8.5, but modified so that they become appropriate for categorical response variables. Some of these models are introduced on the course MY452.

Chapter 10

More statistics...

You will no doubt be pleased to learn that the topics covered on this course have not quite exhausted the list of available statistical methods. In this chapter we outline some of the most important further areas of statistics, so that you are at least aware of their existence and titles. For some of them, codes of LSE courses which cover these methods are given in parentheses.

A very large part of advanced statistics is devoted to further types of **regression models**. The basic idea of them is the same as for multiple linear regression, i.e. modelling expected values of response variables given several explanatory variables. The issues involved in the form and treatment of explanatory variables are usually almost exactly the same as for linear models. Different classes of regression models are needed mainly to accommodate different types of response variables:

- Models for **categorical response variables**. These exist for situations where the response variable is dichotomous (**binary regression**, especially **logistic models**), has more than two unordered (**multinomial logistic models**) or ordered (**ordinal regression models**) categories, or is a count, for example in a contingency table (**Poisson regression**, **loglinear models**). Despite the many different titles, all of these models are closely connected (MY452)
- Models for cases where the response is a length of time to some event, such as a spell of unemployment, interval between births of children or survival of a patient in a medical study. These techniques are known as **event history analysis**, **survival analysis** or **lifetime data analysis**. Despite the different terms, all refer to the same statistical models.

Techniques for the analysis of **dependent data**, which do not require the assumption of statistically independent observations used by almost all the methods on this course:

- **Time series analysis** for one or more long sequence of observations of the same quantity over time. For example, each of the five temperature sequences in Figure 8.2 is a time series of this kind.
- Regression models for **hierarchical data**, where some sets of observations are not independent of each other. There are two main types of such data: **longitudinal** or **panel data** which consist of short time series for many units (e.g. answers by respondents in successive waves of a panel survey), and **nested** or **multilevel data** where basic units are grouped in natural groups or clusters (e.g. pupils in classes and schools in an educational study). Both of these can be analysed using the same general classes of models, which in turn are generalisations of linear and other regression models used for independent data (ST416 for models for multilevel data and ST442 for models for longitudinal data).

Methods for **multivariate data**. Roughly speaking, this means data with several variables for comparable quantities treated on an equal footing, so that none of them is obviously a response to the others. For example, results for the ten events in the decathlon data of the week 7 computer class or, more seriously, the responses to a series of related attitude items in a survey are multivariate data of this kind.

- Various methods of **descriptive multivariate analysis** for jointly summarising and presenting information on the many variables, e.g. **cluster analysis**, **multidimensional scaling** and **principal component analysis** (MY455 for principal components analysis).
- Model-based methods for multivariate data. These are typically **latent variable models**, which also involve variables which can never be directly observed. The simplest latent variable technique is **exploratory factor analysis**, and others include **confirmatory factor analysis**, **structural equation models**, and **latent trait** and **latent class** models (MY455).

Some types of **research design** may also involve particular statistical considerations:

- **Sampling theory** for the design of probability samples, e.g. for surveys (part of MY456, which also covers methodology of surveys in general).
- **Design of experiments** for more complex randomized experiments.

Finally, some areas of statistics are concerned with broader and more fundamental aspects of statistical analysis, such as alternative forms of model specification and inference (e.g. **non-parametric methods**) or the basic ideas of inference itself (e.g. **Bayesian statistics**). These and the more specific tools further build on the foundations of all statistical methods, which are the subject of **probability theory** and **mathematical statistics**. However, you are welcome, if you wish, to leave the details of these fields to professional statisticians, if only to keep them too in employment.

Appendix

Computer classes

Some general instructions on computing and the SPSS package are given first below. It makes most sense to read these together with the instructions for individual computer classes.

General instructions

Using the networked computers at LSE

- To access IT facilities at LSE you need an IT account with its **Username** and **Password**. Please see <http://www.lse.ac.uk/intranet/LSEServices/IMT/guides/accounts/activateAccount.aspx> for instructions on how to activate your account. In case of any problems, please ask for assistance at the IT help desk (Library 1st floor).
- Various introductory **documents** can be accessed through the IMT services web pages at <http://www.lse.ac.uk/intranet/LSEServices/IMT/home.aspx>.
- **Logging in** to use Windows: When you arrive at a networked computer, wait for Windows to start up (if the machine is not already on). Type in **CTRL + ALT + Delete** and the **Enter Network Password** screen will appear. Type in your username and your password and press **Enter** or click on the **OK** button. This will log you on to the computer.

Data downloading

The instructions for each class will give the name of a file or files which will be used for that exercise. In order to do the class, you will need to download the file to your H: space (i.e. your personal file storage space on the LSE network, shown as disk drive H: on a networked computer once you have logged on). You can download all the data files for the course, as well as other course-related material, from the web-based **Moodle** system. See instructions in the beginning of this book for how to register for MY451 on Moodle.

Introduction to SPSS

General information and documentation

SPSS (formerly Statistical Package for the Social Sciences) is a widely used general-purpose statistical software package. It will be used for all the computer classes on this course. The

current version on the LSE network is SPSS 21. This section gives some general information on the structure and use of SPSS. The discussion is brief and not meant to be comprehensive. The instructions given here and in the descriptions of individual computer classes below will be sufficient for the purposes of this course. If, however, you wish to find out more about SPSS, more information and examples can be found in the SPSS help files and tutorials found under the **Help** menu of the program, and in introductory guide books such as

Field, A. (2013). *Discovering Statistics using IBM SPSS Statistics* (4th ed). Sage. Kinnear, P. R. and Gray, C. D. (2012). *SPSS 19 Made Simple*. Psychology Press. Pallant, J. (2013). *SPSS Survival Manual* (5th ed). Open University Press.

These are given here purely as examples (there are many others) and not as recommendations. We have not reviewed any of these books in detail and so cannot make any comparisons between them.

Starting SPSS

To start SPSS, double-click on the SPSS icon on the Windows desktop. Alternatively, click on the **Start** button at the bottom left corner, and select **All Programs**, then **Specialist and teaching software**, **Statistics**, **SPSS**, and finally **SPSS 21** (or some obvious variant of these, in case the exact wording on your desktop is slightly different).

An initial screen for opening data files appears. Click on **Cancel** to get rid of this and to enter the data editor (which will be discussed further below).

Exiting from SPSS

Select **Exit** from the **File** menu or click on the X at the upper right corner of the SPSS data editor window. You may then be prompted to save the information in the open windows; in particular, you should save the contents of the data editor in a file (see below) if you have made any changes to it.

SPSS windows

There are several different types of windows in SPSS. The two most important are

- **Data editor:** A data set is displayed in the Data Editor window. Several of these can be open at a time. The data editor which you have selected (clicked on) most recently defines the active data set, and the procedures you request from the menus are applied to this data set until you select a different active data set. The data editor window has two parts, accessed by clicking on the two tabs at the bottom of the window:
 - **Data view**, which shows the data matrix in the spreadsheet-like form discussed in Section 1.2.1, with units in the rows and variables in the columns.
 - **Variable view**, which shows information about the variables.

Working with the data editor will be practised in the first computer class. The contents of the data editor, i.e. the data matrix and associated information, can be saved in an SPSS data file. Such files have names with the extension **.sav**.

- **Output viewer:** Output from statistical analyses carried out on the data will appear here. The output can be printed directly from the viewer or copied and pasted to other

programs. The contents of the viewer can also be saved in a file, with a name with the extension **.spv** (since version 17; in previous versions of SPSS the extension was **.spo**).

There are also other windows, for example for editing SPSS graphs. They will be discussed in the instructions to individual computer classes where necessary.

Menus

SPSS has a menu-based interface, which can be used to access most of its features for statistical analysis, manipulation of data, loading, saving and printing files, and so on.

- The procedures for statistical analysis are found under the **Analyze** menu, which provides further drop-down menus for choosing different methods.
 - Similarly, procedures for various statistical graphics are found under **Graphs**. We will be using procedures found under **Graphs / Legacy Dialogs**. Here “legacy” means that these are the graphics menus which were included also in previous versions of SPSS. The current version also contains a second, new set of menus for the same graphs, under **Graphs / Chart Builder**. We do not regard these as an improvement in usability, so we will continue to use the old menus. You are welcome to explore the capabilities of the “Chart Builder” on your own.
- Eventually the menu choices lead to a **dialog box** with various boxes and buttons for specifying details of the required analysis. Most dialog boxes contain buttons which open new dialog boxes for further options. The details of these choices for the methods covered on this course are described in the instructions to individual computer classes.
- Almost all of the dialog boxes have options which are not needed for our classes and never mentioned in the instructions. Some of these simply modify the output, others request variants of the statistical methods which will not be used in these classes. All such options have default values which can be left untouched here. You are, however, welcome to experiment with these additional choices to see what they do. Further information on them can be accessed through the **Help** button in each dialog box.

Notational conventions for the instructions

Because analyses in SPSS are carried out by making choices from the menus, the instructions for the computer classes need to describe these choices somehow. To reduce the length and tedium of the instructions, we will throughout present them in a particular format explained below. Because this information is rather abstract if read in isolation, it is best to go through it while carrying out specific instructions for the first few computer classes.

- The appropriate menu choices for obtaining the dialog box for the required analysis are first given in bold, for example as follows:

Analyze/Descriptive statistics/Frequencies

This is short for “Click on the menu item **Analyze** at the top of the window; from the drop-down menu, select **Descriptive statistics** and then click on **Frequencies**.” This particular choice opens a dialog box for constructing various descriptive statistics and graphs (as discussed in Chapter 2).

Unless otherwise mentioned, subsequent instructions then refer to choices in the most recently opened dialog box, without repeating the full path to it.

- For all of the statistical analyses, we need first to specify which variables the analyses should be applied to. This is done by entering the names of those variables in appropriate boxes in the dialog boxes. For example, the dialog box opened above has a box labelled **Variable(s)** for this purpose. The dialog box also includes a separate box containing a list of all the variables in the data set. The required variables are selected from this list and moved to the choice boxes (and back again, when choices are changed) by clicking on an arrow button between the boxes. For example, suppose that a data set contains a grouped age variable called *AGEGROUP*, for which we want to construct a frequency table. The class instructions may then state in words “Place *AGEGROUP* in the **Variable(s)** box”, or sometimes just

Variable(s)/*AGEGROUP*

both of which are short for “In the dialog box opened above, click on the name *AGEGROUP* in the list of variables, and then click on the arrow button to move the name into the **Variable(s)** box”. Sometimes we may also use a generic instruction of the form

Variable(s)/< Variables>

where < Variables> indicates that this is where we would put the name of any variables for which we want to obtain a frequency table. Note that here and in many other procedures, it is possible to select several variables at once. For the Frequencies procedure used as an example here, this simply means that a separate frequency table is constructed for each selected variable.

- Other choices in a dialog box determine details of the analysis and its output. In most cases the selection is made from a fixed list of possibilities provided by SPSS, by clicking on the appropriate box or button. In the instructions, the choice is indicated by listing a path to it, for example as

Charts/Chart Type/Bar charts

in the above example (this requests the so-called bar chart). The items on such a list are labels for various items in the dialog boxes. For example, here **Charts** is a button which opens a new subsidiary dialog box, **Chart Type** is the title of a list of options in this new dialog box, and **Bar charts** is the choice we want to select. In other words, the above instruction is short for “In the dialog box opened above, click on the button **Charts** to open a new dialog box. Under **Chart type**, select **Bar charts** by clicking on a button next to it.”

- Some choices need to be made by typing in some information rather than selecting from a list of options. Specific instructions for this will be given when needed.
- After choices are made in subsidiary dialog boxes, we return to the main dialog box by clicking on **Continue**. Once all the required choices have been made, the analysis is executed by clicking on **OK** in the main dialog box. This should be reasonably obvious, so we will omit explicit instructions to do so.

A useful feature of SPSS is the **dialog recall** button, which is typically sixth from the left in the top row of buttons in the Output viewer window; the button shows a rectangle with a green arrow pointing down from it. Clicking on this gives a menu of recently used procedures, and choosing one of these brings up the relevant dialog box, with the previously used choices selected. This is useful when you want to rerun a procedure, e.g. to try different choices for its options. It is usually quicker to reopen a dialog box using the dialog recall button than through the menus.

SPSS session options

Various options for controlling the format of SPSS output and other features can be found under **Edit/Options**. For example, an often useful choice is **General/Variable Lists/Display names**. This instructs SPSS to display the names of variables in the variable lists of all procedures, instead of the (typically much longer) descriptive labels of the variables. In large data sets this may make it easier to find the right variables from the list. This may be further helped by selecting **General/Variable Lists/Alphabetical**, which causes the names to be listed in an alphabetical order rather than the order in which the variables are included in the data set.

Printing from SPSS

All the computers in the public rooms are connected to one of the laser printers. When you print a document or a part of it, you need to have credit on your printing account. See <http://www.lse.ac.uk/intranet/LSEServices/IMT/guides/printing.aspx> for further information.

- You can print your results from the Output Viewer either by selecting **File/Print** or by clicking on Print on the toolbar (the button with a little picture of a printer). Please note that SPSS output is often quite long, so this may result in much more printout than you really want.
- Alternatively, in the Output Viewer, select the objects to be printed, select **Edit / Copy**, open a Word or Excel document and **Paste**. You can make any changes or corrections in this document before printing it. This method gives you more control over what gets printed than printing directly from SPSS.
- At the printer terminal, type in your username and password. The files sent for printing are then listed. Select the appropriate file number and follow the instructions given by the computer.

SPSS control language

Early versions of SPSS had no menu-based interface. Instead, commands were executed by specifying them in SPSS command language. This language is still there, underlying the menus, and each choice of commands and options from the menus can also be specified in the control language. We will not use this approach on this course, so you can ignore this section if you wish. However, there are some very good reasons why you might want to learn about the control language if you need to work with SPSS for, say, analyses for your thesis or dissertation:

- Because the control language commands can be saved in a file, they preserve a record of how an analysis was done. This may be important for checking that there were no errors, and for rerunning the analyses later if needed.
- For repetitive analyses, modifying and rerunning commands in the control language is quicker and less tedious than using the menus repeatedly.
- Some advanced SPSS procedures are not included in the menus, and can only be accessed through the control language.

The main cost of using the control language is learning its syntax. This is initially much harder than using the menus, but becomes easier with experience. The easiest way to begin learning the syntax is to request SPSS to print out the commands corresponding to choices made from the menus. Two easy ways of doing so are

- Selecting the session option (i.e. under **Edit/Options**) **Viewer/Display commands in the log**. This causes the commands corresponding to the menu choices to be displayed in the output window.
- Clicking on the **Paste** button in a dialog box (instead of **OK**) after selecting an analysis. This opens a *Syntax window* where the corresponding commands are now displayed. The commands in a syntax window can be edited and executed, and also saved in a file (with the extension **.sps**) for future use.

WEEK 2 class: Descriptive statistics for categorical data, and entering data

Data set

The data file **ESS5_sample.sav** will be used today. It contains a simplified sample of data from UK respondents in the 2010 European Social Survey (Round 5). The questions in the survey that you see here were designed By Dr Jonathan Jackson and his team as part of a module investigating public trust in the criminal justice system. Further information about the study can be found at

www.lse.ac.uk/methodology/whosWho/Jackson/jackson_ESS.aspx¹

The main purpose of today's class is to introduce you to the layout of SPSS and to show you how to produce some basic tables and graphs for categorical variables. Additionally, we provide instructions on how to enter data into a new SPSS data file, using the Data Editor. This exercise is not strictly needed for the course, but we include it for two purposes. Firstly, students often find this a helpful way of learning how the software works. Secondly, this exercise may be a useful introduction for students who go on to collect or collate data for their own empirical research.

Classwork

Part 1: The layout of an SPSS data file

1. **Opening an SPSS data file:** this is done from **File/Open/Data**, selecting the required file from whichever folder it is saved in in the usual Windows way. Do this to open **ESS5_sample.sav**.
2. **Information in the Variable View window.** The data file is now displayed in the Data Editor. Its Data View window shows the data as a spreadsheet (i.e. a data matrix). We will first consider the information in the Variable View window, accessed by clicking on the Variable View tab at the bottom left corner of the window. The columns of this window show various pieces of information about the variables. Take a little while familiarising yourself with them. The most important of the columns in Variable View are
 - **Name** of the variable in the SPSS data file. The names in this column (also shown as the column headings in Data View) will be used to refer to specific variables in all of the instructions for these computer classes.
 - **Type** of the variable. Here most of the variables are *Numeric*, i.e. numbers, and a few are *String*, which means text. Clicking on the entry for a variable in this column and then on the button (with three dots on it) revealed by this shows a list of other possibilities.
 - **Width** and **Decimals** control the total number of digits and the number of decimal places displayed in Data View. Clicking on an entry in these columns reveals buttons which can be used to increase or decrease these values. Here all but two of the numeric variables are coded as whole numbers, so Decimals has been set to 0 for them.
 - **Label** is used to enter a longer description of the variable. Double-clicking on an entry allows you to edit the text.

¹ESS Round 5: European Social Survey Round 5 Data (2010). Data file edition 2.0. Norwegian Social Science Data Services, Norway Data Archive and distributor of ESS data. The full data can be obtained from <http://ess.nsd.uib.no/ess/round5/>.

- **Values** shows labels for individual values of a variable. This is mostly relevant for categorical variables, such as most of the ones in these data. Such variables are coded in the data set as numbers, and the Values entry maintains a record of the meanings of the categories the numbers correspond to. You can see examples of this by clicking on some of the entries in the Values column and then on the resulting button. The value labels can also be displayed for each observation in Data View by selecting **View/Value Labels** in that window.
 - **Missing** specifies *missing data codes*, i.e. values which are not actual measurements but indicators that an observation should be treated as missing. There may be several such codes. For example, variables in these data often have separate missing data codes for cases where a respondent was never asked a question (“Not applicable”, often abbreviated NAP), replied “Don’t know” (DK) or otherwise failed to provide an answer (“Refusal” or “No answer”; NA); the explanations of these values are found in the Values column. An alternative to using missing data codes (so-called *User missing* values) is to enter no value (a *System missing* value) for an observation in the data matrix. This is displayed as a full stop (.) in Data View. There are no such values in these data.
 - **Measure** indicates the measurement level of a variable, as *Nominal*, *Ordinal* or *Scale* (meaning interval). This is mostly for the user’s information, as SPSS makes little use of this specification.
3. Any changes made to the data file are preserved by saving it again from **File/Save** (or by clicking on the Save File button of the toolbar, which the one with the picture of a diskette). You will also be prompted to do so when exiting SPSS or when trying to open a new data file. Today you should not save any changes you may have made to ESS5_sample.sav, so click **No** if prompted to do so below.

Part 2: Descriptive statistics for categorical variables

Most of the statistics required for this class are found in SPSS under **Analyze/Descriptive Statistics/Frequencies** as follows:

- Names of the variables for which the statistics are requested are placed in the **Variable(s)** box. To make it easy to find variables in the list box on the left, you may find it convenient to change the way the variables are displayed in the list; see under “SPSS Session Options” for instructions.
- Tables of frequencies: select **Display frequency tables**
- Bar charts: **Charts/Chart Type/Bar charts**. Note that under **Chart Values** you can choose between frequencies or percentage labels on the vertical axis.
- Pie charts: **Charts/Chart Type/Pie charts**

In addition, we will construct some two-way tables or cross-tabulations, by selecting **Analyze/Descriptive Statistics/Crosstabs**. In the dialog box that opens, request a contingency table between two variables by entering

- The name of the row variable into the **Row(s)** box, and
- The name of the column variable into the **Column(s)** box.
- **Cells/Percentages** for percentages within the table: **Row** gives percentages within each row (i.e. frequencies divided by row totals), **Column** percentages within columns, and

Total percentages out of the total sample size.

The labels in the SPSS output should be self-explanatory. Note that in this and all subsequent classes, the output may also include some entries corresponding to methods and statistics not discussed on this course. They can be ignored here.

1. The first variable in the data set, GOODJOB, asks respondents whether they generally feel that the police are doing a good job in their country. There are three response categories for this item: “a good job”, “neither a good job nor a bad job”, or “a bad job”. Obtain a frequency table and bar chart to investigate the distribution of responses to this question.

Check that you understand how to interpret the output you obtain. In particular, make sure that you understand the information displayed in each of the columns in the main table, and that you can see the connection between the information in the table and the information represented in the bar chart.

2. The last variable in the set, AGE_GRP, records in which of the following age groups each respondent falls: up to 29 years of age, 30-49, or 50+ years. Let us consider the association between age group and opinions of the police. Obtain a two-way contingency table of GOODJOB by AGE_GRP. To make interpretation easier, request percentages within each of the age groups. If you use AGE_GRP as the row variable, then include row percentages in the output.

Interpret the resulting table. Are opinions of the police distributed differently among the three different age groups? Does there appear to be an association between age group and attitude?

3. If you have time after completing the data entry exercise (below), you may wish to return to this data set and explore frequencies and contingency tables for some of the other variables in the set.

Part 3: Entering data directly into Data Editor

This procedure may be useful to know if the data you are analysing are not in any electronic form at the beginning of the analysis, for example if you start with a pile of filled-in questionnaires from a survey. For practice, we will enter the following small, artificial data set:

Sex: Man; Age: 45; Weight: 11 st 3 lbs

Sex: Woman; Age: 68; Weight: 8 st 2 lbs

Sex: Woman; Age: 17; Weight: 9 st 6 lbs

Sex: Man; Age: 28; Weight: 13 st 8 lbs

Sex: Woman; Age: 16; Weight: 7 st 8lbs

1. Select **File/New/Data** to clear the Data Editor. Go to Variable View and enter into the first four rows of the Name column names for the variables, for example *sex*, *age*, *wstones* and *wpounds*.
2. Switch to Data View and type the data above into the appropriate columns, one unit (respondent) per row. Note that the person's weight is entered into two columns, one for stones and one for pounds. Enter sex using numerical codes, e.g. 1 for women and 2 for men.
3. Save the file as a new SPSS data file (**File/Save as**), giving it a name of your choice. You should also resave the file (from **File/Save** or by clicking the File Save button) after each of the changes made below.
4. Practise modifying the information in Variable View by adding the following information for the sex variable:
 - Enter the label *Sex of the respondent* into the Label column.
 - Click on the Values cell and then on the resulting button to open a dialog box for entering value labels. Enter **Value: 1; Value Label: Woman; Add**. Repeat for men, and click **OK** when finished.
5. **Transforming variables:** It is often necessary to derive new variables from existing ones. We will practise the two most common examples of this:
 1. **Creating a grouped variable:** Suppose, for example, that we want to define a grouped age variable with three categories: less than 25 years, 25–54 and 55 or over. This is done as follows:
 - Select **Transform/Recode into Different Variables**. This opens a dialog box which is used to define the rule for how values of the existing variable are to be grouped into categories of the new one.
 - Move the name of the age variable to the **Input Variable → Output Variable** box.
 - Under **Output Variable**, enter the **Name** of the new variable, for example *agegroup*, and click **Change**.
 - Click on **Old and New Values**. Enter **Old Value/Range: Lowest through 24** and **New Value/Value: 1**, and click **Add**.
 - Repeat for the other two categories, selecting **Range: 25 through 54** and **Range: 55 through highest** for **Old value**, and 2 and 3 respectively for **New value**.

- You should now see the correct grouping instructions in the **Old** → **New** box. Click **Continue** and **OK** to create the new variable.
 - Check the new variable in Data View. At this stage you should normally enter in Variable View the value labels of the age groups.
2. **Calculations on variables:** Some new variables are obtained through mathematical calculations on existing ones. For example, suppose we want to include weight in kilograms as well as stones and pounds. Using the information that one stone is equal to 6.35 kg and one pound is about 0.45 kg, the transformation is carried out as follows:
- Select **Transform/Compute Variable**. This opens a dialog box which is used to define the rule for calculating the values of the new variable.
 - Enter **Target variable:** *weightkg* (for example; this is the name of the new variable) and **Numeric Expression:** $6.35 * wstones + 0.45 * wpounds$; for the latter, you can either type in the formula or use the variable list and calculator buttons in a fairly obvious way.

WEEK 2 HOMEWORK

The homework exercise for this week is to complete the multiple choice quiz which you can find in the Moodle resource for MY451. Answers to the questions are also included there, including feedback on why the incorrect answer are incorrect. The first part of the quiz asks for answers to the class exercise, and the second part asks you to identify the level of measurement of some different variables.

WEEK 3 class

Descriptive statistics for continuous variables

Data set: The data file used today is **london-borough-profiles.sav**. It contains a selection of data on the 33 London boroughs obtained from the *London Datastore*, which publishes a range of statistical data about the city, collated by the Greater London Authority's *GLA Intelligence Unit* ².

Descriptive statistics in SPSS

This week you will produce and examine descriptive statistics for a number of individual variables. As for last week, almost all of the statistics required for this class can be obtained in SPSS under **Analyze/Descriptive Statistics/Frequencies**. Note that you will probably not find the tables of frequencies very useful, because continuous variables can take so many different values. So for this class, uncheck the **Display frequency tables** option in the dialog box.

- Measures of central tendency: **Mean, Median and Mode** under **Statistics / Central Tendency**
- Measures of variation: **Range, Std. deviation and Variance** under **Statistics/Dispersion**. For the Interquartile range, select **Statistics/ Percentile values/Quartiles** and calculate by hand the difference between the third and first quartiles given (as Percentiles 75 and 25 respectively) in the output.
- Histograms: **Charts/Chart Type/Histograms**

Two charts needed today are not found under the Frequencies procedure:

- **Stem and leaf plots**, which are obtained from **Analyze/Descriptive Statistics/Explore** by entering variable(s) under **Dependent list** and selecting **Display/Plots** and **Plots/Descriptive/Stem-and-leaf**. You can place more than one variable under the **Dependent list** in order to compare variables.
- **Box plots** are also automatically generated through this dialog box, regardless of whether you want to see them! So this is the simplest way to produce them.

Most of these statistics and charts can be obtained in other ways as well, for example from **Analyze/ Descriptive Statistics/Descriptives** or **Graphs/Legacy Dialogs/Histogram**, or **Graphs/Legacy Dialogs/Boxplot**, but we will not use these alternatives today. Feel free to investigate them in your own time if you wish.

The labels in the SPSS output should be self-explanatory. Note that in this and all subsequent classes, the output may also include some entries corresponding to methods and statistics not discussed on this course. They can be ignored here.

Classwork

1. The variable *YOUTH_DEPRIVATION* records for each borough the percentage of children who live in out-of-work families. This is an indicator of deprivation, with higher

²The data were obtained from <http://data.london.gov.uk/datastore/package/london-borough-profiles>.

If you download the "Profiles in Excel" workbook, you will find that one of the pages contains a map of the boroughs, and a tool for visualising the data on that map. A regular map of the boroughs can be found at for example at

<http://www.londoncouncils.gov.uk/londonfacts/londonlocalgovernment/londonmapandlinks/default.htm>.

values indicating a worse situation for each borough. Investigate the distribution of this variable across London boroughs by obtaining its mean, median, minimum and maximum, quartiles and standard deviation, and a histogram. Obtain also a stem and leaf plot and a box plot. Note that double-clicking on a histogram (or any other SPSS graph) opens it in a new window, where the graph can be further edited by changing titles, colours etc. The graph can also be exported from SPSS into other software. Check that you understand how to find the measures of central tendency and dispersion from the output. Does the distribution of *YOUTH_DEPRIVATION* appear to be symmetrically distributed or skewed?

2. Consider now the variable *CRIME*, which records the numbers of reported crimes for every 1000 inhabitants, over the years 2011-12. Obtain some summary descriptive statistics, a histogram and a box plot for this variable. Is the distribution of the variable symmetric or skewed to the left or right? *CRIME* is one of many variables in this data set which have outliers, i.e. boroughs with unusually large or small values of the variable. Normally statistical analysis focuses on the whole data rather than individual observations, but the identities of individual outliers are often also of interest. The outliers can be seen most easily in the box plots, where SPSS labels them with their case numbers, so that you can identify them easily in the data set. For example, 1 would indicate the 1st case in the data set. If you click on to the Data View you can see that this 1st case is the City of London. Which borough is the outlier for *CRIME*?

HOMEWORK

For the questions below, select the relevant SPSS output to include in your homework and write brief answers to the specific questions. Remember SPSS produces some outputs that you do not need. Feel free to transcribe tables or modify charts if you wish to improve their presentation.

1. The variable *VOTING* records voter turnout in a borough, specifically the percentage of eligible voters who voted in the local elections in 2010. Obtain descriptive statistics, a histogram and a box plot for this variable. What is the range of the variable, and what is its inter-quartile range? Are there any outliers? Is the distribution of voter turnout symmetrical or skewed? How can you tell?
2. In the data set employment rates are given overall, but also separately for males and females. The employment rate is the percentage of working age population who are in employment. Compare and contrast male and female employment rates across the boroughs, using the variables *MALE_EMPLOYMENT* and *FEMALE_EMPLOYMENT*. Comment on the differences and/or similarities in their descriptive statistics: minimum and maximum, mean, median and standard deviation. Obtain histograms for these two variables. Are the distributions of male employment and female employment symmetrical or skewed?

WEEK 4 class: Two-way contingency tables

Data set: The data file used today is **GSS2010.sav**. It contains a selection of variables on attitudes and demographic characteristics for 2044 respondents in the 2010 U.S. General Social Survey (GSS)³. The full data set contains 790 variables. For convenience the version you are analysing today contains just a selection of those items.

Analysing two-way contingency tables in SPSS

All of the analyses needed for this week's class are found under **Analyze/Descriptive Statistics/Crosstabs**. We will be obtaining contingency tables between two variables, as in Week 2 class, with the following commands:

- The name of the row variable into the **Row(s)** box, and
- The name of the column variable into the **Column(s)** box.
- **Cells/Percentages** for percentages within the table: **Row** gives percentages within each row (i.e. frequencies divided by row totals), **Column** percentages within columns, and **Total** percentages out of the total sample size.

The only additional output we will need today is obtained by selecting

- **Statistics/Chi-square** for the χ^2 test of independence
- (If you are interested in the γ measure of association for ordinal variables, outlined in the coursepack, you may obtain it using **Statistics/Ordinal/Gamma**. In the output the γ statistic is shown in the "Symmetric measures" table in the "Value" column for "Gamma". We will not use this measure today, but feel free to ask if you are interested in it.)

Classwork

Suppose we want to use the GSS data to investigate whether in the U.S. population sex and age are associated with attitudes towards women's roles. The respondent's sex is included in the data as the variable *SEX*, and age as *AGEGROUP* in three groups: 18-34, 35-54, and 55 or over. The three attitude variables we consider are

- *FEFAM*: Level of agreement with the following statement: "It is much better for everyone involved if the man is the achiever outside the home and the woman takes care of the home and family". Available response options are Strongly agree, Agree, Disagree, and Strongly disagree.
 - *FEPOL*: Level of agreement with the following statement: "Most men are better suited emotionally for politics than are most women". Available response options are: Agree and Disagree.
 - *FEPRES*: Response to the following statement: "If your party nominated a woman for President, would you vote for her if she were qualified for the job?" Available response options are Yes and No.
1. Consider first the association between sex and attitude towards male and female work roles, by constructing a contingency table between *SEX* and *FEFAM*. To make interpretation of the results easier, include also appropriate percentages. Here it makes most sense to

³The data can be obtained from <http://www3.norc.ox.ac.uk/gss+website/>, which gives further information on the survey, including the full text of the questionnaires.

treat sex as an explanatory variable for attitude, so we want to examine percentages of attitudes within categories of male and female. If you use *SEX* as the row variable, this means including the Row percentages in the output. Request also the χ^2 -test statistic. In SPSS output, results for the χ^2 test are given below the two-way table itself in a table labelled “Chi-Square Tests”, in the row “Pearson Chi-Square”. The test statistic itself is given under “Value” and its *P*-value under “Asymp. Sig. (2-sided)”. By considering the χ^2 test statistic and its *P*-value, do you think there is enough evidence to conclude that males and females differ in their views on male and female work roles? If there is, how would you describe the association?

2. Consider now the association between age and attitude towards male and female work roles, by constructing a table between *AGEGROUP* and *FEFAM*. Interpret the results, and compare them to your findings in Exercise 1.
3. Examine differences between men and women in their views about women’s suitability for politics, using a table between *SEX* and *FEPOL*. Interpret the results. (Note: ignore the last two columns of the χ^2 test output, labelled ‘Exact Sig. (2-sided)’ and ‘Exact Sig. (1-sided)’, and use the result under “Asymp. Sig. (2-sided)” as in the other tables.)

HOMEWORK

1. What is the null hypothesis for the χ^2 test that you carried out in analysis 2 in the class, for the table of *AGEGROUP* by *FEFAM*?
2. State the χ^2 test statistic, degrees of freedom and *P*-value for this table, and interpret these results.
3. Interpret the table of percentages to describe the nature of the association between *AGEGROUP* and *FEFAM*.
4. Consider now the association between age and attitude towards voting for a female President, by constructing a table between *AGEGROUP* and *FEPRES*. In the population, do people in different age groups differ in their willingness to vote for a female President? Interpret the results of the χ^2 test and illustrate your answer with one or two percentages from the two-way table.

WEEK 5 class: Inference for two population means

Data set: The data file used today is **ESS5_GBFR.sav**. It contains data for a selection of variables from the 2010 European Social Survey for respondents in Great Britain and France⁴. Only a few of the variables are used in the exercises; the rest are included in the data set as examples of the kinds of information obtained from this survey.

Two-sample inference for means in SPSS

- *t*-tests and confidence intervals for two independent samples for inference on the difference of the population means: **Analyze/Compare Means/Independent-Samples T Test**. The variable of interest Y is placed under **Test Variable(s)** and the explanatory variable X under **Grouping Variable**. The values of X identifying the two groups being compared are defined under **Define Groups**.
- *Box plots* for descriptive purposes are obtained from **Analyze/Descriptive Statistics/Explore**. Here we want to draw side-by-side box plots for values of a response variable Y , one plot for each distinct value of an explanatory variable X . The name of Y is placed under **Dependent List** and that of X under **Factor List**. Box plots are obtained by selecting **Plots/Boxplots/Factor levels together**.
- Tests and confidence intervals for single means (c.f. Section 7.4) are not considered today. These are obtained from **Analyze/Compare Means/One-Sample T Test**. They can also be used to carry out inference for comparisons of means between two *dependent* samples (c.f. Section 7.5).

Classwork Consider the survey data in the file *ESS5_GBFR.sav*. We will examine two variables, and carry out statistical inference to compare their means among the survey populations of adults in Great Britain and France⁵.

1. The variable *WKHTOT* shows the number of hours per week the respondent normally works in his or her main job. Obtain box plots and descriptive statistics for this variable separately for each country (identified by the variable *CNTRY*). Compare measures of central tendency and variation for *WKHTOT* between the two countries. What do you observe?
2. Obtain a *t*-test and confidence interval for the difference of weekly working hours between Britain and France (specify the values of the country variable as **Define Groups/Group 1: GB** and **Group 2: FR** as coded in the data). Details of SPSS output for this are explained in Chapter 7; you can use the results under the assumption of equal population variances. What do you conclude? Is there a statistically significant difference in the average values of *WKHTOT* between the two countries? What does the confidence interval suggest about the size of the difference?
3. The variable *STFJBOT* asks those in paid work, “How satisfied are you with the balance between the time you spend on your paid work and the time you spend on other aspects of your life?”. Respondents are asked to rate their level of satisfaction on a scale from 0-10, where 0 means “Extremely dissatisfied” and 10 means “Extremely satisfied”. Repeat exercises 1 and 2 for this variable, and compare also histograms of *STFJBOT* for each country. What do you observe?

⁴ESS Round 5: European Social Survey Round 5 Data (2010). Data file edition 2.0. Norwegian Social Science Data Services, Norway Data Archive and distributor of ESS data. The full data can be obtained from <http://ess.nsd.uib.no/ess/round5/>.

⁵Strictly speaking, the analysis should incorporate sampling weights (variable *DWEIGHT*) to adjust for different sampling probabilities for different types of respondents. Here the weights are ignored. Using them would not change the main conclusions for these variables.

HOMEWORK

1. Write up your answers to the second class exercise, answering these specific questions:
 1. What are the observed sample means for *WKHTOT* for French and British respondents?
 2. Is there a statistically significant difference in the average values of *WKHTOT* between the two countries? State the value of the test statistic and its corresponding *P*-value. You may assume equal population variances for this test.
 3. Interpret the 95% confidence interval for the difference.
2. The variable *WKHSCH* asks respondents, “How many hours a week, if any, would you choose to work, bearing in mind that your earnings would go up or down according to how many hours you work?”. Is there a statistically significant difference between ideal (rather than actual) work hours for French and British respondents? Carry out a t-test and report and interpret the results.
3. The variable *STFMJOB* asks respondents, “How satisfied are you in your main job?”. Respondents are asked to rate their level of satisfaction on a scale from 0-10, where 0 means “Extremely dissatisfied” and 10 means “Extremely satisfied”. Is there a statistically significant difference, at the 5% level of significance, between mean levels of job satisfaction for French and British respondents? Answer this question by using the 95% confidence interval for the difference in means (you need the full t-test output to obtain the confidence interval, but you need not report the results of the t-test itself for this question).

WEEK 7 class: Inference for population proportions

Data sets: Files `BES2010post_lastdebate.sav` and `BES2010pre_lastdebate.sav`.

Inference on proportions in SPSS

- SPSS menus do not provide procedures for calculating the tests and confidence intervals for proportions discussed in Chapter 5. This is not a serious limitation, as the calculations are quite simple.
- It is probably easiest to use a pocket calculator for the calculations, and this is the approach we recommend for this class. The only part of the analysis it cannot do is calculating the precise P -value for the tests, but even this can be avoided by using critical values from a statistical table such as the one at the end of this Coursepack to determine approximate P -values (or by using an online P -value calculator — see “About Week 4 class” on the Moodle page for suggested links).

Classwork

The survey data set `BES2010post_lastdebate.sav` contains part of the information collected by the British Election Study, an ongoing research programme designed to understand voter choices in the UK ⁶.

In the run-up to the UK General Election on 6 May 2010, opinion polls reported quite dramatic changes in popularity of the Liberal Democrat party. Key to their increasing popularity was the performance of their party leader, Nick Clegg, in a series of three televised debates between the leaders of the three main political parties (the other participants were Gordon Brown for Labour and David Cameron for the Conservative party). The debates were broadcast between 15 and 29 April 2010.

The data in `BES2010post_lastdebate.sav` contain information on respondents’ voting intentions, obtained after the debates had ended (i.e. between 30 April and 6 May).

1. `VOTE_LIBDEM` is a dichotomous variable indicating whether a respondent intended to vote for the Liberal Democrats (value 1) or some other party (0) in the 2010 General Election. The value of this variable is by definition missing for those who had not decided which way they would vote or who did not intend to vote at all, so they are automatically excluded from the analysis. The parameter of interest π is now the population proportion of those who *say* they would vote Liberal Democrat. We will compare it to 0.23, the proportion of the vote the party actually received in 2010. The analysis is thus one-sample inference on a population proportion, and the relevant formulas are (11) for the test statistic and (15) for the confidence interval that can be found in Section 5.5.2 and 5.6.2 respectively.
 - Begin by creating a frequency table of `VOTE_LIBDEM`. This should show that the sample estimate of π is 0.260, out of 3226 non-missing responses. Thus $n = 3226$ and $\hat{\pi} = 0.260$ in the notation of Chapter 5.
 - For the one-sample significance test, the value of π under the null hypothesis is $\pi_0 = 0.230$. Using the specific formula of the test statistic in Section 5.5.2, the value

⁶The data can be obtained from <http://bes2009-10.org/>, which gives further information on the survey, including the full text of the questionnaires. The data analysed in this class and homework are from the BES Campaign Internet Panel Survey, which has been divided into two data sets corresponding to two time periods leading up to the General Election.

of the test statistic z is thus given by the calculation

$$z = \frac{0.260 - 0.230}{\sqrt{0.230 \times (1 - 0.230)/3226}}$$

Calculate this using a calculator. The result should be $z = 4.049$.

- The (two-sided) P -value for this is the probability that a value from the standard normal distribution is at most -4.049 or at least 4.049 . Evaluate this approximately by comparing the value of z to critical values from the standard normal distribution (c.f. Table 5.2) as explained in Section 5.5.3. Here, for example, z is larger than 1.96 , so the two-sided P -value must be smaller than 0.05 . Convince yourself that you understand this statement.
- Calculate a 95% confidence interval for the population proportion of prospective Liberal Democrat voters, using equation (15) at the end of Section 5.6.2.

What do you conclude about the proportions of prospective and actual Liberal Democrat voters? Why might the two differ from each other?

2. The variable *TVDEBATE* indicates whether the respondent reports having watched any of the three televised debates (1 for Yes, at least one watched, 0 otherwise - this includes “no” and “don’t know” responses). We will compare the proportion of people intending to vote Liberal Democrat amongst those who watched some or all of the debates with those who did not, using the two-sample methods of analysis discussed in Section 5.7. The formula of the z -test statistic for testing the hypothesis of equal population proportions is thus the two-sample z -test statistic for proportions (see middle of Section 5.7), and a confidence interval for the difference of the proportions is (25) in Section 5.7.
 - Begin by calculating the relevant sample proportions. The easiest way to do this is by creating a two-way contingency table between *TVDEBATE* and *VOTE_LIBDEM* as you did in the Week 2 and 4 classes. The results required for the analysis considered here are all shown in the resulting table. Convince yourself that these show that, in the notation of Section 5.7,

$$- n_1 = 930 \text{ and } \hat{\pi}_1 = 0.218 (= 203/930),$$

$$- n_2 = 2296 \text{ and } \hat{\pi}_2 = 0.277 (= 636/2296),$$

where 1 denotes respondents who did not watch any of the debates and 2 those who watched at least some. The pooled estimated proportion $\hat{\pi}$ (formula 21 in Section 5.7) used in the test statistic (23) is here $\hat{\pi} = 0.260$, shown on the “Total” row.

- Calculate the test statistic, its P -value and a 95% confidence for the difference in population proportions, using the relevant formulas. For example, the test statistic is here given by

$$z = \frac{0.277 - 0.218}{\sqrt{0.260 \times (1 - 0.260) \times (1/2296 + 1/930)}}.$$

What do you conclude? Is there evidence that those who watched at least some of the leaders’ debates were more likely to declare an intention to vote Liberal Democrat? If there is, how big is the difference in proportions of prospective Liberal Democrat voters between the debate-watchers and debate-non-watchers?

HOMEWORK

Write up your answers to the second class exercise. In particular, answer the following specific questions:

1. What proportion of respondents say that they did watch at least some of the leaders' debates? And what proportion did not? Of those who watched at least some of the leaders' debates, what proportion said they intended to vote Liberal Democrat? And what proportion of those who did *not* watch any of the leaders' debates said they intended to vote Liberal Democrat?
2. Calculate the test statistic and find its corresponding approximate P -value for the difference in population proportions of prospective Liberal Democrat voters among those who did and did not watch the leaders' debates. Show your working. State the conclusion from the test.
3. Calculate a 95% confidence interval around this difference. State its lower and upper limits.
4. Write a brief substantive interpretation of your results.

The data set *BES2010pre_lastdebate.sav* contains responses to the same question - whether respondents intended to vote Liberal Democrat or not - but asked before the last of the party leaders' debates. Repeat the analysis you carried out for the first class exercise, but using this data set. In other words carry out a one-sample analysis, of the kind done in exercise 1 above, to compare the proportion of respondents who said they intended to vote Liberal Democrat with the proportion who actually did. Answer the following questions:

1. State the null hypothesis for the test.
2. Calculate the test statistic and find its corresponding approximate P -value. Show your workings.
3. Give a brief interpretation of the results. Do they differ from the other data set? Can you think of any reasons for this? (This last question invites some speculation - do not worry if you don't have any ideas! But see the sample answer if you are interested in our speculation.)

WEEK 7 class: Correlation and simple linear regression 1

Data set: Files `decathlon2012.sav`.

Scatterplots, correlation and simple linear regression in SPSS

- A scatterplot is obtained from **Graphs/Legacy Dialogs/“Scatter/Dot”/ Simple Scatter/Define**. The variables for the X-axis and Y-axis are placed in the **X Axis** and **Y Axis** boxes respectively. Double-clicking on the plot in the Output window opens it in a *Chart Editor*, where various additions to the graph can be made. A fitted straight line is added from **Elements/Fit Line at Total**. A least squares fitted line is the default under this option, so it is drawn immediately and you can just click **Close**. Closing the Chart Editor commits the changes to the Output window.
- A correlation matrix is obtained from **Analyze/Correlate/Bivariate**, when **Correlation Coefficients/Pearson** is selected (which is the default, so you should not need to change it). The variables included in the correlation matrix are placed into the **Variables** box. The output also includes a test for the hypothesis that the population correlation is 0, but we will ignore it.
- Linear regression models are obtained from **Analyze/Regression/Linear**. The response variable is placed under **Dependent** and the explanatory variable under **Independent(s)**. The dialog box has many options for various additional choices. Today you can leave all of them at their default values, except that you should select **Statistics/Regression Coefficients/Confidence intervals** to include also 95% confidence intervals for the regression coefficients in the output.

Classwork

Decathlon is a sport where the participants complete ten different athletics events over two days. Their results in each are then translated into points, and the winner is the competitor with the highest points total for the ten events. The file *decathlon2012.sav* contains the results of the decathlon competition at the 2012 Olympics in London for the 26 athletes who finished the competition.⁷ The results for each event are given both in their original units (variables with names beginning with “mark_”) and in decathlon points (names beginning with “points_”). The ten events are identified by the variable labels in Variable View. The variable *points_total* gives the final points total for each competitor.

1. Create a scatterplot between the result (X-axis) and points (Y-axis) for one event, the 100-metre sprint (variables *MARK_100M* and *POINTS_100M*), and add a fitted line. This simply provides information on the calculation used to transform the result into points. Clearly a linear calculation is used for this, at least over the range of results in these data. Notice the downward slope of the line: the faster the result, the higher the number of points. From now on, for simplicity we will consider only the points variables for each event.
2. Obtain the correlation matrix for all pairs of variables among the ten individual points scores and the total score. Consider first correlations between the individual events only. Which correlations tend to be high (say over 0.5), which ones close to zero and which ones even negative? Can you think of any reasons for this? Draw scatterplots and fitted

⁷Official results obtained from www.olympic.org/london-2012-summer-olympics.

lines for a few pairs of variables with different sizes of correlations (here the variables are treated symmetrically, so it does not matter which one is placed on the X -axis). Can these associations be reasonably described as linear?

3. Consider now the correlations between the ten event scores and the final score *POINTS_TOTAL*. Which of them is highest, and which one lowest? Examine the scatterplot and fitted line between points for 100 metres (*POINTS_100M*) and the total score (*POINTS_TOTAL*). Fit a line to this scatterplot variables, with *POINTS_100M* as the explanatory variable. Interpret the results. Does there appear to be an association between the points for 100 metres and the total score? What is the nature of the association?

Suppose you were told that a competitor received 800 points (a time of about 11.3 seconds) for 100 metres, the first event of the decathlon. Based on the fitted model, what final points score would you predict for him? You can calculate this fitted value with a pocket calculator. What would be the predicted value if the 100-metre score was 950 points (about 10.6 s) instead?

HOMEWORK

1. Briefly discuss the correlation matrix produced in the class. Pick out a few examples for illustration - which correlations are highest, and which ones lowest, and which ones negative? You may comment on correlations between individual events, as well as on correlations between the final score and individual events.
2. Obtain the scatterplot and linear regression model for the total score given points for the long jump, one of the field events (*POINTS_LONGJUMP*). Is the score for long jump strongly or weakly associated with the final score? Interpret the slope coefficient. Suppose you were told that a competitor received 900 points (a jump of about 7.4 metres) for the long jump. Based on the fitted model, what final points score would you predict for him?
3. Obtain the scatterplot and linear regression model for the total score given points for throwing the discus, another of the field events (*POINTS_DISCUS*). Interpret the slope coefficient. Is the score for discus strongly or weakly associated with the final score?

WEEK 8 class: Simple linear regression and 3-way tables

Data set: File **GSS2010.SAV**. This contains a selection of variables on attitudes and demographic characteristics for 2044 respondents in the 2010 U.S. General Social Survey (GSS)⁸. Only a few of the variables are used in the exercises.

Classwork - linear regression

Here we will focus on the variables *EDUC*, *PAEDUC*, *MAEDUC* and *SPEDUC*. These show the number of years of education completed by, respectively, the survey respondent him/herself, and the respondent's father, mother and spouse.

1. Obtain basic descriptive statistics for the variables. Here they can be compared directly, because the meaning of the variable is similar in each case. We can even draw side-by-side box plots for the variables (rather than for values of a single variable at different levels of another, as before). These can be obtained from **Analyze/Descriptive Statistics/Explore** by placing all the variables under **Dependent List** and selecting **Plots/Boxplots/Dependents together**. You should then also select **Options/Missing Values/Exclude cases pairwise** to include all non-missing values for each variable (here *SPEDUC* has for obvious reasons more missing values than the others).
2. Obtain the correlation matrix of the four variables. Which correlations are highest, and which ones lowest?
3. Draw a scatterplot with fitted line for *EDUC* given *PAEDUC*. Fit a linear regression model between these variables, regression of *EDUC* (response variable) on *PAEDUC* (explanatory variable). Interpret the results. Is there a statistically significant linear association between a person's years of schooling and those of his/her father? Interpret the estimated regression coefficient, *t*-statistic and *P*-value, and 95 per cent confidence interval.
4. Based on the fitted model, what is the predicted number of years of education for a respondent whose father completed 12 years of education?

HOMEWORK: Simple linear regression and three-way tables

The homework exercise uses the same data set for two different types of analysis.

Linear regression

Draw a scatterplot with fitted line for *EDUC* given *MAEDUC*. Fit a linear regression model between these variables, regression of *EDUC* (response variable) on *MAEDUC* (explanatory variable).

1. Interpret the results: Is there a statistically significant linear association between a person's years of schooling and those of his/her mother? Interpret the estimated regression coefficient, *t*-statistic and *P*-value, and 95 per cent confidence interval.
2. Based on the fitted model, what is the predicted number of years of education for a respondent whose mother completed 10 years of education?

⁸The data can be obtained from www3.norc.ox.ac.uk/GSS+Website/, which gives further information on the survey, including the full text of the questionnaires.

3. Interpret the R-squared statistic for the model.

Analysing multiway contingency tables in SPSS

Three-way contingency tables are again obtained from **Analyze/Descriptive Statistics/Crosstabs**. The only change from Week 4 class is that the conditioning variable is now placed in the **Layer 1 of 1** box. This produces a series of partial two-way tables between the row and column variables specified in the **Row(s)** and **Column(s)** boxes, one for each category of the **Layer** variable. Percentages and χ^2 test are similarly calculated separately for each partial table. For this example we elaborate on the first two exercises from Week 4 class. To remind you, the categorical variables we are analysing are these:

- The respondent's sex, recorded as the variable *SEX*.
 - age as *AGEGROUP* in three groups: 18-34, 35-54 and 55 or over.
 - *FEFAM*: Level of agreement with the following statement: "It is much better for everyone involved if the man is the achiever outside the home and the woman takes care of the home and family", with response options Strongly agree, Agree, Disagree, and Strongly disagree.
1. First remind yourself of the associations between *SEX* and *FEFAM* and between *AGEGROUP* and *FEFAM*. Obtain the two-way contingency table between *FEFAM* and *SEX*, including appropriate percentages and χ^2 test of independence. Repeat the procedure for *FEFAM* by *AGEGROUP*. What do you learn about the associations between attitude and sex, and between attitude and age?
 2. Sociologists would suggest that the relationship between sex and attitude towards male and female work roles might be different for different age groups. In other words, age might modify the association between sex and attitude. Investigate this possible interaction between the three variables. Create a three-way table where *FEFAM* is the column variable, *SEX* the row variable and *AGEGROUP* the layer (conditioning) variable. Study the SPSS output, and make sure you understand how this shows three partial tables of *FEFAM* vs. *SEX*, one for each possible value of *AGEGROUP*. Examine and interpret the associations in the three partial tables. State the results of the χ^2 test for each partial table, and illustrate your interpretations with some appropriate percentages. Finally, summarise your findings: are there differences in the nature, strength or significance of the association between sex and attitude, depending on the age group? Comment on how this interpretation differs from the initial two-way table of *FEFAM* and *SEX*.

WEEK 9 class: Multiple linear regression

Data set: File `humandevlopment2011.sav`.

Multiple linear regression in SPSS

- Multiple linear regression is obtained from **Analyze/Regression/Linear**, by placing all of the required explanatory variables in the **Independent(s)** box. No other changes from last week are required.
- To include categorical explanatory variables, the necessary dummy variables have to be created first. The ones for today's class are already included in the data set. If you need to create dummy variables for your own analyses in the future, it is usually easiest to do so from **Transform/Compute Variable**. Some of the buttons on the keypad shown in that dialog box are *logical operators* for defining conditions for which the outcome is either 1 (True) or 0 (False), as required by a dummy variable. For example, the categorical variable *INCOME_GROUP* in today's data set has the value 3 if the country is in the high income group. The dummy variable *HIGH_INCOME* was created from this by entering **Target Variable:** *HIGH_INCOME* and **Numeric Expression:** *INCOME_GROUP*=3. This means that the new variable *HIGH_INCOME* will have the value 1 for countries for which *INCOME_GROUP* is equal to 3, and will be 0 otherwise. Other logical operators may also be used: for example, *urban_pop*\$<\$50 would produce 1 if the variable *URBAN_POP* was less than 50 and 0 otherwise.

Classwork

The file *humandevlopment2011.sav* contains data on a number of indicators of what might broadly be called development, for 194 countries in 2011. These were collated from two international data agency sources⁹. The response variable considered today is *SCHOOL_YEARS*, which records for each country the mean number of years of schooling taken by the adult population. We treat it here as a general indicator of the educational situation in a country, which is an important aspect of development. We will consider the following explanatory variables for it:

- *URBAN_POP*: the degree of urbanisation of the country, specifically the percentage of the country's population living in urban areas variable
- *GOVERNANCE*, a continuous variable constructed from expert opinion surveys to reflect the perceived effectiveness of government in delivering services.
- *INFANT_MORTALITY*, number of infants dying before 1 year old, per 1,000 live births — a “proxy” indicator representing the health of the population
- *INCOME_GROUP*, classified as low, middle or high income economies. This is also provided in the form of three dummy variables: *LOW_INCOME*, *MIDDLE_INCOME* and *HIGH_INCOME*.

⁹United Nations Development Programme *International Human Development Indicators*, <http://hdr.undp.org/en/data/>; World Bank *Worldwide Governance Indicators*, <http://info.worldbank.org/governance/wgi/pdf/wgidataset.xlsx>; World Bank *World Development Indicators*, <http://data.worldbank.org/indicator/SP.DYN.IMRT.IN>.

1. Obtain some descriptive statistics for the continuous variables, to gain an impression of their ranges. A quick way of doing this is via **Analyze/Descriptive Statistics/Frequencies**, unchecking the “Display frequency tables” and requesting minimum and maximum values.
2. Investigate the idea that increased urbanisation is linked to greater availability of schooling for people. Obtain a scatterplot and a simple linear regression model for *SCHOOL_YEARS* given *URBAN_POP*. What do you observe in the scatterplot? Interpret the regression output.
3. Now consider the possibility that schooling may also be explained by the effectiveness of governments in providing public services (such as education). Fit a multiple linear regression model for *SCHOOL_YEARS* given both *URBAN_POP* and *GOVERNANCE*. Compare the estimated coefficient of *URBAN_POP* for this model with the coefficient of the same variable in the model in Question 2. What do you conclude? Does the association between schooling and urbanisation change when we control for government effectiveness? If so, in what way? Interpret the estimate coefficient of *GOVERNANCE* in the fitted model, the results of its *t*-test and its 95% confidence interval.
4. Next consider the possible explanatory value of the income wealth of a country for understanding variation in schooling years. Include income by entering two of the three dummy variables for income group. For the most convenient interpretation, we suggest that you leave “low income” as the reference group, and enter the dummies for *MIDDLE_INCOME* and *HIGH_INCOME* in the model. Interpret the values of the estimated regression coefficients for the two income dummy variables. In addition, for each one state the null hypothesis for its *t*-test, and interpret the result of the test and 95% confidence intervals.
5. Using this model, what level of schooling would you predict for a country with 70% urban population, a score of 1.5 on governance, and a high income economy?
6. Using this model, what level of schooling would you predict for a country with 30% urban population, a score of -0.2 on governance, and a low income economy?

HOMEWORK

1. Write up your answers to the last three questions in the class exercise.
2. Finally, consider one more possible explanatory variable: *INFANT_MORTALITY*. Add this variable to the multiple linear regression model fitted above. Is it statistically significant, at the 1% level of significance? Interpret the value of its estimated coefficient, and its 95% confidence interval. Take care to make sense of the sign (positive or negative) of the coefficient.
3. Has the inclusion of *INFANT_MORTALITY* modified the interpretation of any of the other explanatory variables in the model? Are they all statistically significant, at the 5% level of significance? Briefly outline the similarities and differences between the results for this final model and the model fitted in the class exercise.

WEEK 10 class: Review and Multiple linear regression

Data set: File **ESS5GB_trust.sav**.

This class is for you to revisit any topics of your choosing. Make the most of the opportunity to ask your class teachers any questions you have about any of the course material, and to practise any of the analyses you have learned during the course.

As an optional exercise, the data file **ESS5GB_trust.sav** is provided. This contains a selection of variables from the survey of British respondents that forms the 2010 wave of the European Social Survey¹⁰.

We suggest that you use the data to practise multiple linear regression modelling on one or more of the variables capturing people's levels of trust in institutions. For these questions, respondents were asked the following: "Using this card, please tell me on a score of 0-10 how much you personally trust each of the institutions I read out. 0 means you do not trust an institution at all, and 10 means you have complete trust." The institutions (and their variable names) are:

- *trstprl*: Trust in country's parliament
- *trstlgl*: Trust in the legal system
- *trstplc*: Trust in the police
- *trstplt*: Trust in politicians
- *trstprrt*: Trust in political parties
- *trstep*: Trust in the European Parliament
- *trstun*: Trust in the United Nations

After you choose a response variable that interests you, you will need to select some potential explanatory variables to test. The data set contains a number of variables. Some are socio-demographic, such as age and gender. Some are attitudinal or behavioural, such as amount of time spent reading newspapers. You will need to make a judgement about the levels of measurement of the variables, and how to enter them into the model. Use the "Values" column in the SPSS Variable View to check how each variable is coded. Note: we suggest that it is not too much of a compromise to treat the variables on television, radio and newspaper consumption as continuous, interval level variables. Note also: we have provided dummy variables for the categorical variables in the data set.

HOMEWORK

As this is the last week of the course, there is no homework. You can find further information on this and the other class exercises and homeworks in the model answers, which will be posted in the Moodle site.

¹⁰ESS Round 5: European Social Survey Round 5 Data (2010). Data file edition 2.0. Norwegian Social Science Data Services, Norway Data Archive and distributor of ESS data. The full data can be obtained from <http://ess.nsd.uib.no/ess/round5/>.

Statistical tables

Explanation of the “Table of standard normal tail probabilities” in Section @ref(s__disttables__Z):

- The table shows, for values of Z between 0 and 3.5, the probability that a value from the standard normal distribution is *larger than* Z (i.e. the “right-hand” tail probabilities).
 - For example, the probability of values larger than 0.50 is 0.3085.
- For negative values of Z , the probability of values *smaller than* Z (the “left-hand” tail probability) is equal to the right-hand tail probability for the corresponding positive value of Z .
 - For example, the probability of values smaller than -0.50 is also 0.3085.

Table of standard normal tail probabilities

z	Prob.	z	Prob.	z	Prob.	z	Prob.	z	Prob.	z	Prob.
0.00	0.5000	0.50	0.3085	1.00	0.1587	1.50	0.0668	2.00	0.0228	2.50	0.0062
0.01	0.4960	0.51	0.3050	1.01	0.1562	1.51	0.0655	2.01	0.0222	2.52	0.0059
0.02	0.4920	0.52	0.3015	1.02	0.1539	1.52	0.0643	2.02	0.0217	2.54	0.0055
0.03	0.4880	0.53	0.2981	1.03	0.1515	1.53	0.0630	2.03	0.0212	2.56	0.0052
0.04	0.4840	0.54	0.2946	1.04	0.1492	1.54	0.0618	2.04	0.0207	2.58	0.0049
0.05	0.4801	0.55	0.2912	1.05	0.1469	1.55	0.0606	2.05	0.0202	2.60	0.0047
0.06	0.4761	0.56	0.2877	1.06	0.1446	1.56	0.0594	2.06	0.0197	2.62	0.0044
0.07	0.4721	0.57	0.2843	1.07	0.1423	1.57	0.0582	2.07	0.0192	2.64	0.0041
0.08	0.4681	0.58	0.2810	1.08	0.1401	1.58	0.0571	2.08	0.0188	2.66	0.0039
0.09	0.4641	0.59	0.2776	1.09	0.1379	1.59	0.0559	2.09	0.0183	2.68	0.0037
0.10	0.4602	0.60	0.2743	1.10	0.1357	1.60	0.0548	2.10	0.0179	2.70	0.0035
0.11	0.4562	0.61	0.2709	1.11	0.1335	1.61	0.0537	2.11	0.0174	2.72	0.0033
0.12	0.4522	0.62	0.2676	1.12	0.1314	1.62	0.0526	2.12	0.0170	2.74	0.0031
0.13	0.4483	0.63	0.2643	1.13	0.1292	1.63	0.0516	2.13	0.0166	2.76	0.0029
0.14	0.4443	0.64	0.2611	1.14	0.1271	1.64	0.0505	2.14	0.0162	2.78	0.0027
0.15	0.4404	0.65	0.2578	1.15	0.1251	1.65	0.0495	2.15	0.0158	2.80	0.0026
0.16	0.4364	0.66	0.2546	1.16	0.1230	1.66	0.0485	2.16	0.0154	2.82	0.0024
0.17	0.4325	0.67	0.2514	1.17	0.1210	1.67	0.0475	2.17	0.0150	2.84	0.0023
0.18	0.4286	0.68	0.2483	1.18	0.1190	1.68	0.0465	2.18	0.0146	2.86	0.0021
0.19	0.4247	0.69	0.2451	1.19	0.1170	1.69	0.0455	2.19	0.0143	2.88	0.0020
0.20	0.4207	0.70	0.2420	1.20	0.1151	1.70	0.0446	2.20	0.0139	2.90	0.0019
0.21	0.4168	0.71	0.2389	1.21	0.1131	1.71	0.0436	2.21	0.0136	2.92	0.0018
0.22	0.4129	0.72	0.2358	1.22	0.1112	1.72	0.0427	2.22	0.0132	2.94	0.0016
0.23	0.4090	0.73	0.2327	1.23	0.1093	1.73	0.0418	2.23	0.0129	2.96	0.0015
0.24	0.4052	0.74	0.2296	1.24	0.1075	1.74	0.0409	2.24	0.0125	2.98	0.0014
0.25	0.4013	0.75	0.2266	1.25	0.1056	1.75	0.0401	2.25	0.0122	3.00	0.0013
0.26	0.3974	0.76	0.2236	1.26	0.1038	1.76	0.0392	2.26	0.0119	3.02	0.0013
0.27	0.3936	0.77	0.2206	1.27	0.1020	1.77	0.0384	2.27	0.0116	3.04	0.0012
0.28	0.3897	0.78	0.2177	1.28	0.1003	1.78	0.0375	2.28	0.0113	3.06	0.0011
0.29	0.3859	0.79	0.2148	1.29	0.0985	1.79	0.0367	2.29	0.0110	3.08	0.0010
0.30	0.3821	0.80	0.2119	1.30	0.0968	1.80	0.0359	2.30	0.0107	3.10	0.0010
0.31	0.3783	0.81	0.2090	1.31	0.0951	1.81	0.0351	2.31	0.0104	3.12	0.0009
0.32	0.3745	0.82	0.2061	1.32	0.0934	1.82	0.0344	2.32	0.0102	3.14	0.0008
0.33	0.3707	0.83	0.2033	1.33	0.0918	1.83	0.0336	2.33	0.0099	3.16	0.0008
0.34	0.3669	0.84	0.2005	1.34	0.0901	1.84	0.0329	2.34	0.0096	3.18	0.0007
0.35	0.3632	0.85	0.1977	1.35	0.0885	1.85	0.0322	2.35	0.0094	3.20	0.0007
0.36	0.3594	0.86	0.1949	1.36	0.0869	1.86	0.0314	2.36	0.0091	3.22	0.0006
0.37	0.3557	0.87	0.1922	1.37	0.0853	1.87	0.0307	2.37	0.0089	3.24	0.0006
0.38	0.3520	0.88	0.1894	1.38	0.0838	1.88	0.0301	2.38	0.0087	3.26	0.0006
0.39	0.3483	0.89	0.1867	1.39	0.0823	1.89	0.0294	2.39	0.0084	3.28	0.0005
0.40	0.3446	0.90	0.1841	1.40	0.0808	1.90	0.0287	2.40	0.0082	3.30	0.0005
0.41	0.3409	0.91	0.1814	1.41	0.0793	1.91	0.0281	2.41	0.0080	3.32	0.0005
0.42	0.3372	0.92	0.1788	1.42	0.0778	1.92	0.0274	2.42	0.0078	3.34	0.0004
0.43	0.3336	0.93	0.1762	1.43	0.0764	1.93	0.0268	2.43	0.0075	3.36	0.0004
0.44	0.3300	0.94	0.1736	1.44	0.0749	1.94	0.0262	2.44	0.0073	3.38	0.0004
0.45	0.3264	0.95	0.1711	1.45	0.0735	1.95	0.0256	2.45	0.0071	3.40	0.0003
0.46	0.3228	0.96	0.1685	1.46	0.0721	1.96	0.0250	2.46	0.0069	3.42	0.0003
0.47	0.3192	0.97	0.1660	1.47	0.0708	1.97	0.0244	2.47	0.0068	3.44	0.0003
0.48	0.3156	0.98	0.1635	1.48	0.0694	1.98	0.0239	2.48	0.0066	3.46	0.0003
0.49	0.3121	0.99	0.1611	1.49	0.0681	1.99	0.0233	2.49	0.0064	3.48	0.0003

Table of critical values for t -distributions

df	0.100	0.050	0.025	0.010	0.005	0.001	0.0005
1	3.078	6.314	12.706	31.821	63.657	318.309	636.619
2	1.886	2.920	4.303	6.965	9.925	22.327	31.599
3	1.638	2.353	3.182	4.541	5.841	10.215	12.924
4	1.533	2.132	2.776	3.747	4.604	7.173	8.610
5	1.476	2.015	2.571	3.365	4.032	5.893	6.869
6	1.440	1.943	2.447	3.143	3.707	5.208	5.959
7	1.415	1.895	2.365	2.998	3.499	4.785	5.408
8	1.397	1.860	2.306	2.896	3.355	4.501	5.041
9	1.383	1.833	2.262	2.821	3.250	4.297	4.781
10	1.372	1.812	2.228	2.764	3.169	4.144	4.587
11	1.363	1.796	2.201	2.718	3.106	4.025	4.437
12	1.356	1.782	2.179	2.681	3.055	3.930	4.318
13	1.350	1.771	2.160	2.650	3.012	3.852	4.221
14	1.345	1.761	2.145	2.624	2.977	3.787	4.140
15	1.341	1.753	2.131	2.602	2.947	3.733	4.073
16	1.337	1.746	2.120	2.583	2.921	3.686	4.015
17	1.333	1.740	2.110	2.567	2.898	3.646	3.965
18	1.330	1.734	2.101	2.552	2.878	3.610	3.922
19	1.328	1.729	2.093	2.539	2.861	3.579	3.883
20	1.325	1.725	2.086	2.528	2.845	3.552	3.850
21	1.323	1.721	2.080	2.518	2.831	3.527	3.819
22	1.321	1.717	2.074	2.508	2.819	3.505	3.792
23	1.319	1.714	2.069	2.500	2.807	3.485	3.768
24	1.318	1.711	2.064	2.492	2.797	3.467	3.745
25	1.316	1.708	2.060	2.485	2.787	3.450	3.725
26	1.315	1.706	2.056	2.479	2.779	3.435	3.707
27	1.314	1.703	2.052	2.473	2.771	3.421	3.690
28	1.313	1.701	2.048	2.467	2.763	3.408	3.674
29	1.311	1.699	2.045	2.462	2.756	3.396	3.659
30	1.310	1.697	2.042	2.457	2.750	3.385	3.646
40	1.303	1.684	2.021	2.423	2.704	3.307	3.551
60	1.296	1.671	2.000	2.390	2.660	3.232	3.460
120	1.289	1.658	1.980	2.358	2.617	3.160	3.373
∞	1.282	1.645	1.960	2.326	2.576	3.090	3.291

Explanation: For example, the value 3.078 in the top left corner indicates that for a t -distribution with 1 degree of freedom the probability of values greater than 3.078 is 0.100. The last row shows critical values for the standard normal distribution.

Table of critical values for χ^2 distributions

df	0.100	0.050	0.010	0.001
1	2.71	3.84	6.63	10.828
2	4.61	5.99	9.21	13.816
3	6.25	7.81	11.34	16.266
4	7.78	9.49	13.28	18.467
5	9.24	11.07	15.09	20.515
6	10.64	12.59	16.81	22.458
7	12.02	14.07	18.48	24.322
8	13.36	15.51	20.09	26.124
9	14.68	16.92	21.67	27.877
10	15.99	18.31	23.21	29.588
11	17.28	19.68	24.72	31.264
12	18.55	21.03	26.22	32.909
13	19.81	22.36	27.69	34.528
14	21.06	23.68	29.14	36.123
15	22.31	25.00	30.58	37.697
16	23.54	26.30	32.00	39.252
17	24.77	27.59	33.41	40.790
18	25.99	28.87	34.81	42.312
19	27.20	30.14	36.19	43.820
20	28.41	31.41	37.57	45.315
25	34.38	37.65	44.31	52.620
30	40.26	43.77	50.89	59.703
40	51.81	55.76	63.69	73.402
50	63.17	67.50	76.15	86.661
60	74.40	79.08	88.38	99.607
70	85.53	90.53	100.43	112.317
80	96.58	101.88	112.33	124.839
90	107.57	113.15	124.12	137.208
100	118.50	124.34	135.81	149.449

Explanation: For example, the value 2.71 in the top left corner indicates that for a χ^2 distribution with 1 degree of freedom the probability of values greater than 2.71 is 0.100.